

Federated Learning

From Theory to Practice

Alexander Jung*

January 23, 2025

*AJ is currently Associate Professor for Machine Learning at Aalto University (Finland).
This work has been partially funded by the Academy of Finland (decision numbers 331197, 349966, 363624) and the European Union (grant number 952410).

Preface

This book is about federated learning (FL), a paradigm for distributed machine learning that enables collaborative training across multiple decentralized datasets. FL leverages similarities between local datasets to collaboratively train models while preserving data privacy. The primary focus of this book is to explore the theoretical foundations and algorithmic techniques underpinning FL, emphasizing its mathematical models and optimization frameworks.

This book revolves around two central mathematical constructs: FL networks and generalized total variation (GTV) minimization. We use FL networks to model relationships between clients, their datasets, and the locally trained models. GTV minimization, which is defined for a given FL network, formulates FL as a regularized instance of empirical risk minimization. The regularization term captures variation across local models connected through the FL network, offering a principled way to balance personalization and collaboration. We obtain practical FL systems by solving GTV minimization with established distributed optimization techniques.

(Our Main Goal) *The book trains the reader in formulating FL applications as instances of GTV minimization over a FL network. It also demonstrates how to study and design FL algorithms by applying distributed optimization techniques to solve GTV minimization.*

FL has found applications in healthcare, finance, telecommunications, smart cities, and more. In these fields, FL enables organizations to build robust predictive models without ever centralizing sensitive data. The distributed

nature of FL also enables scalability, allowing large-scale systems to train collaboratively without a central data repository.

We believe that FL becomes a main paradigm for the development of trustworthy artificial intelligence. Indeed, FL merges established machine learning concepts with considerations for privacy, security, and robustness. Roughly speaking, FL is privacy-friendly by not sharing raw data but only updates of model parameters. The de-centralized nature of FL systems offers robustness against different forms of cyber-attacks such as data poisoning.

Audience. The only mathematical background required of the reader is a good knowledge of calculus and linear algebra. Any reader with a completed undergraduate course on basic mathematical analysis (covering concepts such as norms, convergence and derivatives) should be able to follow every argument and discussion in the book. Even readers who are not well-versed in calculus or linear algebra should still get all of the essential ideas and important points. Prior exposure to machine learning and optimization is helpful but not required as we develop most of the required concepts from first principles.

Acknowledgement. The development of this text has benefited from student feedback received for the course CS-E4740 Federated Learning which has been offered at Aalto University during 2023. The author is indebted to Olga Kuznetsova, Diana Pfau and Shamsiat Abdurakhmanova for providing feedback on early drafts.

Contents

1	Introduction	1
1.1	Main Tools	4
1.2	Main Goal of the Book	5
1.3	Outline	6
1.4	Problems	8
2	ML Basics	11
2.1	Learning Goals	11
2.2	Three Components and a Design Principle	11
2.3	Computational Aspects of ERM	15
2.4	Statistical Aspects of ERM	16
2.5	Validation and Diagnosis of ML	20
2.6	Regularization	25
2.7	Problems	28
3	A Design Principle for FL	30
3.1	Learning Goals	31
3.2	FL Networks	31
3.3	Generalized Total Variation	34
3.4	Generalized Total Variation Minimization	40
3.4.1	Computational Aspects of GTVMin	43
3.4.2	Statistical Aspects of GTVMin	45
3.5	How to Handle Non-Parametric Models	48
3.6	Interpretations	49
3.7	Problems	51

3.8	Proofs	56
3.8.1	Proof of Proposition 3.1	56
4	Gradient Methods	57
4.1	Learning Goals	57
4.2	Gradient Descent	58
4.3	Learning Rate	60
4.4	When to Stop?	62
4.5	Perturbed Gradient Step	66
4.6	Handling Constraints - Projected Gradient Descent	68
4.7	Generalizing the Gradient Step	70
4.8	Gradient Methods as Fixed-Point Iterations	73
4.9	Problems	75
5	FL Algorithms	77
5.1	Learning Goals	77
5.2	Gradient Descent for GTVMin	78
5.3	Message Passing Implementation	80
5.4	FedSGD	85
5.5	FedAvg	89
5.6	FedRelax	94
5.7	Asynchronous FL Algorithms	96
5.8	Problems	99
5.9	Proofs	100
5.9.1	Proof of Proposition 5.1	100
5.9.2	Proof of Proposition 5.2	101

6	Main Flavours of FL	104
6.1	Learning Goals	105
6.2	Single-Model FL	106
6.3	Clustered FL	108
6.4	Horizontal FL	112
6.5	Vertical FL	114
6.6	Personalized Federated Learning	115
6.7	Few-Shot Learning	118
6.8	Problems	118
6.9	Proofs	120
6.9.1	Proof of Proposition 6.1	120
7	Graph Learning	122
7.1	Learning Goals	123
7.2	Edges as Design Choice	123
7.3	Measuring (Dis-)Similarity Between Datasets	128
7.4	Graph Learning Methods	131
7.5	Problems	134
8	Trustworthy FL	135
8.1	Learning Goals	135
8.2	Seven Key Requirements by the EU	136
8.2.1	KR1 - Human Agency and Oversight.	136
8.2.2	KR2 - Technical Robustness and Safety.	137
8.2.3	KR3 - Privacy and Data Governance.	138
8.2.4	KR4 - Transparency.	139

8.2.5	KR5 - Diversity, Non-Discrimination and Fairness. . . .	140
8.2.6	KR6 - Societal and Environmental Well-Being.	141
8.2.7	KR7 - Accountability.	142
8.3	Technical Robustness of FL Systems	142
8.3.1	Sensitivity Analysis	143
8.3.2	Estimation Error Analysis	144
8.3.3	Network Resilience	147
8.3.4	Stragglers	148
8.4	Subjective Explainability of FL Systems	151
8.5	Overview of Coding Assignment	155
9	Privacy-Protection in FL	1
9.1	Learning Goals	1
9.2	Measuring Privacy Leakage	2
9.3	Ensuring Differential Privacy	9
9.4	Private Feature Learning	12
9.5	Problems	16
9.6	Overview of Coding Assignment	17
9.6.1	Where Are You?	17
9.6.2	Ensuring Privacy with Pre-Processing	18
9.6.3	Ensuring Privacy with Post-Processing	18
9.6.4	Private Feature Learning	18
10	Data and Model Poisoning in FL	1
10.1	Learning Goals	2
10.2	Attack Types	2

10.3 Data Poisoning	6
10.4 Model Poisoning	7
10.5 Assignment	7
10.5.1 Denial-of-Service Attack	8
10.5.2 Backdoor Attack	8
Glossary	1

Lists of Symbols

Sets and Functions

$a \in \mathcal{A}$ This statement indicates that the object a is an element of the set \mathcal{A} .

$a := b$ This statement defines a to be shorthand for b .

$|\mathcal{A}|$ The cardinality (number of elements) of a finite set \mathcal{A} .

$\mathcal{A} \subseteq \mathcal{B}$ \mathcal{A} is a subset of \mathcal{B} .

$\mathcal{A} \subset \mathcal{B}$ \mathcal{A} is a strict subset of \mathcal{B} .

\mathbb{N} The natural numbers $1, 2, \dots$

\mathbb{R} The real numbers x [1].

\mathbb{R}_+ The non-negative real numbers $x \geq 0$.

\mathbb{R}_{++} The positive real numbers $x > 0$.

$\{0, 1\}$	The set consisting of the two real numbers 0 and 1.
$[0, 1]$	The closed interval of real numbers x with $0 \leq x \leq 1$.
$\operatorname{argmin}_{\mathbf{w}} f(\mathbf{w})$	The set of minimizers for a real-valued function $f(\mathbf{w})$.
$\mathbb{S}^{(n)}$	The set of unit-norm vectors in \mathbb{R}^{n+1} .
$\log a$	The logarithm of the positive number $a \in \mathbb{R}_{++}$.
$h(\cdot) : \mathcal{A} \rightarrow \mathcal{B} : a \mapsto h(a)$	<p>A function (map) that accepts any element $a \in \mathcal{A}$ from a set \mathcal{A} as input and delivers a well-defined element $h(a) \in \mathcal{B}$ of a set \mathcal{B}. The set \mathcal{A} is the domain of the function h and the set \mathcal{B} is the codomain of h. ML aims at finding (or learning) a function h (<i>hypothesis</i>) that reads in the features \mathbf{x} of a data point and delivers a prediction $h(\mathbf{x})$ for its label y.</p>
$\nabla f(\mathbf{w})$	<p>The gradient of a differentiable real-valued function $f : \mathbb{R}^d \rightarrow \mathbb{R}$ is the vector $\nabla f(\mathbf{w}) = \left(\frac{\partial f}{\partial w_1}, \dots, \frac{\partial f}{\partial w_d} \right)^T \in \mathbb{R}^d$ [2, Ch. 9].</p>

Matrices and Vectors

$\mathbf{x} = (x_1, \dots, x_d)^T$	A vector of length d with its j -th entry being x_j .
\mathbb{R}^d	The set of vectors $\mathbf{x} = (x_1, \dots, x_d)^T$ consisting of d real-valued entries $x_1, \dots, x_d \in \mathbb{R}$.
$\mathbf{I}_{l \times d}$	A generalized identity matrix with l rows and d columns. The entries of $\mathbf{I}_{l \times d} \in \mathbb{R}^{l \times d}$ are equal to 1 along the main diagonal and equal to 0 otherwise.
\mathbf{I}_d, \mathbf{I}	A square identity matrix of size $d \times d$. If the size is clear from the context, we drop the subscript.
$\ \mathbf{x}\ _2$	The Euclidean (or ℓ_2) norm of the vector $\mathbf{x} = (x_1, \dots, x_d)^T \in \mathbb{R}^d$ defined as $\ \mathbf{x}\ _2 := \sqrt{\sum_{j=1}^d x_j^2}$.
$\ \mathbf{x}\ $	Some norm of the vector $\mathbf{x} \in \mathbb{R}^d$ [3]. Unless specified otherwise, we mean the Euclidean norm $\ \mathbf{x}\ _2$.
\mathbf{X}^T	The transpose of a matrix $\mathbf{X} \in \mathbb{R}^{m \times d}$. A square real-valued matrix $\mathbf{X} \in \mathbb{R}^{m \times m}$ is called symmetric if $\mathbf{X} = \mathbf{X}^T$.
\mathbf{x}^T	The transpose of a matrix that has the vector $\mathbf{x} \in \mathbb{R}^d$ as its single column.
$\mathbf{0} = (0, \dots, 0)^T$	The vector in \mathbb{R}^d with each entry equal to zero.
$\mathbf{1} = (1, \dots, 1)^T$	The vector in \mathbb{R}^d with each entry equal to one.

$(\mathbf{v}^T, \mathbf{w}^T)^T$	The vector of length $d + d'$ obtained by concatenating the entries of vector $\mathbf{v} \in \mathbb{R}^d$ with the entries of $\mathbf{w} \in \mathbb{R}^{d'}$.
$\text{span}\{\mathbf{B}\}$	The span of a matrix $\mathbf{B} \in \mathbb{R}^{a \times b}$, which is the subspace of all linear combinations of columns of \mathbf{B} , $\text{span}\{\mathbf{B}\} = \{\mathbf{B}\mathbf{a} : \mathbf{a} \in \mathbb{R}^b\} \subseteq \mathbb{R}^a$.
\mathbb{S}_+^d	The set of all positive semi-definite (psd) matrices of size $d \times d$.
$\det(\mathbf{C})$	The determinant of the matrix \mathbf{C} .
$\mathbf{A} \otimes \mathbf{B}$	The Kronecker product of \mathbf{A} and \mathbf{B} [4].

Probability Theory

$\mathbb{E}_p\{f(\mathbf{z})\}$ The expectation of a function $f(\mathbf{z})$ of a RV \mathbf{z} whose probability distribution is $p(\mathbf{z})$. If the probability distribution is clear from context we just write $\mathbb{E}\{f(\mathbf{z})\}$.

$p(\mathbf{x}, y)$ A (joint) probability distribution of a RV whose realizations are data points with features \mathbf{x} and label y .

$p(\mathbf{x}|y)$ A conditional probability distribution of a RV \mathbf{x} given the value of another RV y [5, Sec. 3.5].

$p(\mathbf{x}; \mathbf{w})$ A parametrized probability distribution of a RV \mathbf{x} . The probability distribution depends on a parameter vector \mathbf{w} . For example, $p(\mathbf{x}; \mathbf{w})$ could be a multivariate normal distribution with the parameter vector \mathbf{w} given by the entries of the mean vector $\mathbb{E}\{\mathbf{x}\}$ and the covariance matrix $\mathbb{E}\left\{(\mathbf{x} - \mathbb{E}\{\mathbf{x}\})(\mathbf{x} - \mathbb{E}\{\mathbf{x}\})^T\right\}$.

$\mathcal{N}(\mu, \sigma^2)$ The probability distribution of a Gaussian RV $x \in \mathbb{R}$ with mean (or expectation) $\mu = \mathbb{E}\{x\}$ and variance $\sigma^2 = \mathbb{E}\{(x - \mu)^2\}$.

$\mathcal{N}(\boldsymbol{\mu}, \mathbf{C})$ The multivariate normal distribution of a vector-valued Gaussian RV $\mathbf{x} \in \mathbb{R}^d$ with mean (or expectation) $\boldsymbol{\mu} = \mathbb{E}\{\mathbf{x}\}$ and covariance matrix $\mathbf{C} = \mathbb{E}\{(\mathbf{x} - \boldsymbol{\mu})(\mathbf{x} - \boldsymbol{\mu})^T\}$.

Machine Learning

r	An index $r = 1, 2, \dots$, that enumerates data points.
m	The number of data points in (the size of) a dataset.
\mathcal{D}	A dataset $\mathcal{D} = \{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$ is a list of individual data points $\mathbf{z}^{(r)}$, for $r = 1, \dots, m$.
d	Number of features that characterize a data point.
x_j	The j -th feature of a data point. The first feature of a given data point is denoted x_1 , the second feature x_2 and so on.
\mathbf{x}	The feature vector $\mathbf{x} = (x_1, \dots, x_d)^T$ of a data point whose entries are the individual features of a data point.
\mathcal{X}	The feature space \mathcal{X} is the set of all possible values that the features \mathbf{x} of a data point can take on.
\mathbf{z}	Beside the symbol \mathbf{x} , we sometimes use \mathbf{z} as another symbol to denote a vector whose entries are individual features of a data point. We need two different symbols to distinguish between <i>raw</i> and <i>learnt</i> features [6, Ch. 9].
$\mathbf{x}^{(r)}$	The feature vector of the r -th data point within a dataset.
$x_j^{(r)}$	The j -th feature of the r -th data point within a dataset.

\mathcal{B}	A mini-batch (subset) of randomly chosen data points.
B	The size of (the number of data points in) a mini-batch.
y	The label (quantity of interest) of a data point.
$y^{(r)}$	The label of the r -th data point.
$(\mathbf{x}^{(r)}, y^{(r)})$	The features and label of the r -th data point.
\mathcal{Y}	<p>The label space \mathcal{Y} of a ML method consists of all potential label values that a data point can carry. The nominal label space might be larger than the set of different label values arising in a given dataset (e.g., a training set). We refer to ML problems (methods) using a numeric label space, such as $\mathcal{Y} = \mathbb{R}$ or $\mathcal{Y} = \mathbb{R}^3$, as regression problems (methods). ML problems (methods) that use a discrete label space, such as $\mathcal{Y} = \{0, 1\}$ or $\mathcal{Y} = \{cat, dog, mouse\}$ are referred to as classification problems (methods).</p>
η	Learning rate (step-size) used by gradient-based methods.
$h(\cdot)$	A hypothesis map that reads in features \mathbf{x} of a data point and delivers a prediction $\hat{y} = h(\mathbf{x})$ for its label y .
$\mathcal{Y}^{\mathcal{X}}$	Given two sets \mathcal{X} and \mathcal{Y} , we denote by $\mathcal{Y}^{\mathcal{X}}$ the set of all possible hypothesis maps $h : \mathcal{X} \rightarrow \mathcal{Y}$.

	A hypothesis space or model used by a ML method. The hypothesis space consists of different hypothesis maps $h : \mathcal{X} \rightarrow \mathcal{Y}$ between which the ML method has to choose .
\mathcal{H}	
$d_{\text{eff}}(\mathcal{H})$	The effective dimension of a hypothesis space \mathcal{H} .
B^2	The squared bias of a learnt hypothesis \hat{h} delivered by a ML algorithm that is fed with data points which are modelled as realizations of RVs. If data is modelled as realizations of RVs, also the delivered hypothesis \hat{h} is the realization of a RV.
V	The variance of the (parameters of the) hypothesis delivered by a ML algorithm. If the input data for this algorithm is interpreted as realizations of RVs, so is the delivered hypothesis a realization of a RV.
$L((\mathbf{x}, y), h)$	The loss incurred by predicting the label y of a data point using the prediction $\hat{y} = h(\mathbf{x})$. The prediction \hat{y} is obtained from evaluating the hypothesis $h \in \mathcal{H}$ for the feature vector \mathbf{x} of the data point.
E_v	The validation error of a hypothesis h , which is its average loss incurred over a validation set.
$\hat{L}(h \mathcal{D})$	The empirical risk or average loss incurred by the hypothesis h on a dataset \mathcal{D} .

E_t	The training error of a hypothesis h , which is its average loss incurred over a training set.
t	A discrete-time index $t = 0, 1, \dots$ used to enumerate sequential events (or time instants).
t	An index that enumerates learning tasks within a multi-task learning problem.
α	A regularization parameter that controls the amount of regularization.
$\lambda_j(\mathbf{Q})$	The j -th eigenvalue (sorted either ascending or descending) of a psd matrix \mathbf{Q} . We also use the shorthand λ_j if the corresponding matrix is clear from context.
$\sigma(\cdot)$	The activation function used by an artificial neuron within an artificial neural network (ANN).
$\mathcal{R}_{\hat{y}}$	A decision region within a feature space.
\mathbf{w}	A parameter vector $\mathbf{w} = (w_1, \dots, w_d)^T$ of a model, e.g, the weights of a linear model or in a ANN.
$h^{(\mathbf{w})}(\cdot)$	A hypothesis map that involves tunable model parameters w_1, \dots, w_d , stacked into the vector $\mathbf{w} = (w_1, \dots, w_d)^T$.
$\phi(\cdot)$	A feature map $\phi : \mathcal{X} \rightarrow \mathcal{X}' : \mathbf{x} \mapsto \mathbf{x}' := \phi(\mathbf{x}) \in \mathcal{X}'$.

Federated Learning

$\mathcal{G} = (\mathcal{V}, \mathcal{E})$	An undirected graph whose nodes $i \in \mathcal{V}$ represent devices within a federated learning (FL) network. The undirected weighted edges \mathcal{E} represent similarities and connectivity between the devices.
$i \in \mathcal{V}$	A node in the FL network that represents some device which can access a local dataset and train a local model.
$\mathcal{G}^{(\mathcal{C})}$	The induced sub-graph of \mathcal{G} using the nodes in $\mathcal{C} \subseteq \mathcal{V}$.
$\mathbf{L}^{(\mathcal{G})}$	The Laplacian matrix of a graph \mathcal{G} .
$\mathbf{L}^{(\mathcal{C})}$	The Laplacian matrix of the induced graph $\mathcal{G}^{(\mathcal{C})}$.
$\mathcal{N}^{(i)}$	The neighbourhood of a node i in a graph \mathcal{G} .
$d^{(i)}$	The weighted degree $d^{(i)} := \sum_{i' \in \mathcal{N}^{(i)}} A_{i,i'}$ of node i in a graph \mathcal{G} .
$d_{\max}^{(\mathcal{G})}$	The maximum weighted node degree of a graph \mathcal{G} .

$\mathcal{D}^{(i)}$	A local dataset $\mathcal{D}^{(i)}$ at device $i \in \mathcal{V}$ of an FL network.
m_i	The number of data points (sample size) contained in the local dataset $\mathcal{D}^{(i)}$ at node $i \in \mathcal{V}$.
$\mathbf{x}^{(i,r)}$	The features of the r -th data point in the local dataset $\mathcal{D}^{(i)}$.
$y^{(i,r)}$	The label of the r -th data point in the local dataset $\mathcal{D}^{(i)}$.
$\mathbf{w}^{(i)}$	The local model parameters of device i within an FL network.
$L_i(\mathbf{w})$	The local loss function used by device i to measure the usefulness of some choice \mathbf{w} for the local model parameters .
$L^{(d)}(\mathbf{x}, h, h')$	The loss incurred by a hypothesis h' on a data point with features \mathbf{x} and label $h(\mathbf{x})$ that is obtained from another hypothesis.
$\text{stack}\{\mathbf{w}^{(i)}\}_{i=1}^n$	The vector $\left((\mathbf{w}^{(1)})^T, \dots, (\mathbf{w}^{(n)})^T \right)^T \in \mathbb{R}^{dn}$ that is obtained by vertically stacking the local model parameters $\mathbf{w}^{(i)} \in \mathbb{R}^d$.

1 Introduction

We are surrounded by devices, such as smartphones or wearables, generating decentralized collections of local datasets [7–11]. These local datasets typically have an intrinsic network structure that arises from functional constraints (connectivity between devices) or statistical similarities.

For example, the management of pandemics uses contact networks to relate local datasets generated by patients. Network medicine relates data about diseases via co-morbidity networks [12]. Social science uses notions of acquaintance to relate data collected from be-friended individuals [13]. Another example for network-structured data are the weather observations collected at Finnish Meteorological Institute (FMI) stations. The FMI stations generate local datasets which tend to have similar statistical properties for nearby stations.

Federated learning (FL) is an umbrella term for distributed optimization techniques to train machine learning (ML) models from decentralized collections of local datasets [14–18]. The idea is to carry out the model training directly at the location of data generation (such as your smartphone or a heart rate sensor). This is different from the basic ML workflow, which is to (i) collect data at a single location (computer) and (ii) then train a single ML model on this data.

It can be beneficial to train different ML models at the locations of actual data generation [19] for several reasons:

- **Privacy.** FL methods are appealing for applications involving sensitive data (such as healthcare) as they do not require the exchange of raw data



Fig. 1.1: Left: A basic ML method uses a single dataset to train a single model. Right: Decentralized collection of devices (or *clients*) with the ability to generate data and train models.

but only model parameters (or their updates) [16, 17]. By exchanging only (updates of) model parameters, FL methods are considered privacy-friendly in the sense of not leaking (too much) sensitive information that is contained in the local datasets (see Chapter 9).

- **Robustness.** By relying on decentralized data and computation, FL methods offer robustness (to some extent) against hardware failures (such as *stragglers*) and cyber attacks. One form of cyber attack on FL systems is data poisoning which we discuss in Chapter 10.
- **Parallel Computing.** We can interpret a collection of interconnected devices as a parallel computer. One example for such a parallel computer is a mobile network constituted by smartphones that can communicate via radio links. This parallel computer allows to speed up computations required for the training of ML models (see Chapter 4).

- **Democratization of ML.** FL allows to combine (or pool) the computational resources of many low-cost devices in order to train high-dimensional models such as a Large Language Model (LLM). Instead of using a few powerful computers we combine the contributions of many low-complexity devices [54, 55].
- **Trading Computation against Communication.** Consider a FL application where local datasets are generated by low-complexity devices at remote locations (think of a wildlife camera) that cannot be easily accessed. The cost of communicating raw local datasets to some central unit (which then trains a single global ML model) might be much higher than the computational cost incurred by using the low-complexity devices to (partially) train ML models [20].
- **Personalization.** FL trains personalized ML models for collections of devices (e.g., smartphones or wearables) with computational capabilities and can access local datasets [21]. A key challenge for ensuring personalization is the heterogeneity of local datasets [22, 23]. Indeed, the statistical properties of different local datasets might vary significantly such that they cannot be accurately modelled as independent and identically distributed (i.i.d.). Each local dataset induces a separate learning task that consists of learning useful parameter values for a local model. FL systems train personalized models for devices by combining the information carried in their local datasets (see Chapter 6).

1.1 Main Tools

Euclidean space. Our main mathematical structure for the study and design of FL systems is the Euclidean space \mathbb{R}^d . We expect familiarity with the basic properties of the algebraic and geometric structure of \mathbb{R}^d [24,25]. For example, we often use the spectral decomposition of psd matrices that naturally arise in the formulation of FL applications. We will also use the geometric structure of \mathbb{R}^d , which is defined by the inner-product $\mathbf{w}^T \mathbf{w}' := \sum_{j=1}^d w_j w'_j$ between two vectors $\mathbf{w}, \mathbf{w}' \in \mathbb{R}^d$ and the induced norm $\|\mathbf{w}\|_2 := \sqrt{\mathbf{w}^T \mathbf{w}} = \sqrt{\sum_{j=1}^d w_j^2}$.

Calculus. A main toolbox for the design the FL algorithms are variants of gradient descent (GD). The common idea of these gradient-based methods is to approximate a function $f(\mathbf{w})$ locally by a linear function. This local linear approximation is determined by the gradient $\nabla f(\mathbf{w})$. We, therefore, expect some familiarity with multivariable calculus [2].

Fixed-Point Iterations. Each algorithm that we discuss in this book can be interpreted as a fixed-point iteration of some (typically non-linear) operator $\mathcal{P} : \mathbb{R}^d \rightarrow \mathbb{R}^d$. These operators depend on the local datasets and personal models used within a FL system. A prime example of such a non-linear operator is the gradient step of gradient-based methods (see Chapter 4). The computational properties (such as convergence speed) of these FL algorithms are determined by the contraction properties of the underlying operator [26].

1.2 Main Goal of the Book

The overarching goal of the book is to study FL applications and algorithms using mathematical structures from network theory and optimization theory. In particular, we develop the concept of a FL network to represent a real-world FL application. We then formulate FL as an optimization problem over a given FL network. This lends naturally to a principled design of FL algorithms using distributed optimization methods.

The nodes of a FL network represent devices that train a local (personalized) model based on some local loss function. These loss functions might be obtained from local datasets but this is not relevant for our approach. Formally, each node in a FL network is assigned a model (or hypothesis space) and a local loss function.

Some devices of a FL network are connected by links which can be used to share messages (e.g., intermediate results of computations) during the FL process. We represent these links by the undirected weighted edges of an FL network. The edges of an FL network not only represent communication links but also some notion of statistical similarity between (the data generated by) different devices.

We will see later in the book (see Ch. 7.4) how to construct measures for statistical similarities between local datasets. However, we can view the edges also as a design choice that determines the behaviour of FL algorithms that built on top of a FL network.

From an engineering point of view, the main goal of the book is to establish a flexible design principle for FL systems (see Ch. 3). This approach is inspired by empirical risk minimization as a main tool for the analysis and design

of ML systems [6, 141]. Similar to empirical risk minimization, our design principle is formulated as an optimization problem, which we refer to as GTVMinimization (GTVMin).

GTVMin optimizes model parameters for the local models by balancing the sum of the incurred local loss with some measure for their variation across the edges of the FL network. Thus, GTVMin amounts to minimizing the local loss while also requiring the model parameters and connected nodes to be similar. We obtain different FL methods by using different measures for the variation of model parameters across edges.

Once we formulate a FL application as GTVMin, we can construct FL algorithms by applying distributed optimization methods to solve GTVMin. The optimization methods studies in this book are all instances of fixed-point iterations. We obtain different methods by choosing from different fixed-point operators that have the solutions of GTVMin as their fixed-points. One important example for such a fixed-point iteration is GD.

1.3 Outline

This book is roughly divided into three parts:

- **Part I: ML Refresher.** Chapter 2 introduces data, models and loss functions as three main components of ML. This chapter also explains how these components are combined within empirical risk minimization (ERM). We also discuss how to regularize ERM via manipulating its three main components. We then explain when and how to solve regularized ERM via simple GD methods in Chapter 4. Overall, this part serves two main purposes: (i) to briefly recap basic concepts of

ML in a simple centralized setting and (ii) to highlight ML techniques (such as regularization) that are particularly relevant for the design and analysis of FL methods.

- **Part II: FL Theory and Methods.** Chapter 3 introduces the FL network as a mathematical structure for representing collections of devices that generate local datasets and train local (personalized) models. An FL network also contains undirected weighted edges that connect some of the nodes. The edges represent statistical similarities between local datasets as well as communication links for the implementation of FL algorithms. Chapter 3 formulates FL as an instance of regularized empirical risk minimization (RERM) which we refer to as GTVMin. GTVMin uses the variation of personalized model parameters across edges in the FL network as regularizer. We will see that GTVMin couples the training of tailored (or *personalized*) ML models such that well-connected nodes (clusters) in the FL network obtain similar trained models. Chapter 4 discusses variations of gradient descent as our main algorithmic toolbox for solving GTVMin. Chapter 5 shows how FL algorithms can be obtained in a principled fashion by applying optimization methods, such as gradient-based methods, to GTVMin. We will obtain FL algorithms that can be implemented as iterative message passing methods for the distributed training of tailored (or *personalized*) models. Chapter 6 shows how some main flavours of FL can be interpreted as special cases of GTVMin. The usefulness of GTVMin crucially depends on the choice for the weighted edges in the FL network. Chapter 7 discusses basic principles of graph learning

methods. The common idea of these methods is to place edges between the nodes of a FL network that are most similar. Similarity measures can be obtained via statistical inference procedures or via learnt vector representations. For example, we can map a dataset to a vector by evaluating the gradient of the average loss incurred over this dataset.

- **Part III: Trustworthy AI.** Chapter 8 enumerates seven key requirements for trustworthy artificial intelligence (AI) that have been put forward by the European Union. These key requirements include the protection of privacy as well as robustness against (intentional) perturbations of data or computation. We then discuss how FL algorithms can ensure privacy protection in Chapter 9. Chapter 10 discusses how to evaluate and ensure the robustness of FL methods against intentional perturbations (poisoning) of local dataset.

1.4 Problems

1.1. Complexity of Matrix Inversion. Choose your favourite computer architecture (represented by a mathematical model) and think about how much computation is required - in the worst case - by the most efficient algorithm that can invert any given invertible matrix $\mathbf{Q} \in \mathbb{R}^{d \times d}$? Try also to reflect on how practical your chosen computer architecture is, i.e., is it possible to buy such a computer in your nearest electronics shop?

1.2. Vector Spaces and Euclidean Norm. Consider data points, each characterized by a feature vector $\mathbf{x} \in \mathbb{R}^d$ with entries x_1, x_2, \dots, x_d .

- Show that the set of all feature vectors forms a vector space under

standard addition and scalar multiplication.

- Calculate the Euclidean norm of the vector $\mathbf{x} = (1, -2, 3)^T$.
- If $\mathbf{x}^{(1)} = (1, 2, 3)^T$ and $\mathbf{x}^{(2)} = (-1, 0, 1)^T$, compute $3\mathbf{x}^{(1)} - 2\mathbf{x}^{(2)}$.

1.3. Matrix Operations in Linear Models. Linear regression methods learn model parameters $\hat{\mathbf{w}} \in \mathbb{R}^d$ via solving the optimization problem:

$$\hat{\mathbf{w}} = \arg \min_{\mathbf{w} \in \mathbb{R}^d} \|\mathbf{y} - \mathbf{X}\mathbf{w}\|_2^2,$$

with some matrix $\mathbf{X} \in \mathbb{R}^{m \times d}$, and some vector $\mathbf{y} \in \mathbb{R}^m$.

- Derive a closed-form expression for $\hat{\mathbf{w}}$ that is valid for *arbitrary* \mathbf{X}, \mathbf{y} .
- Discuss the conditions under which $\mathbf{X}^T \mathbf{X}$ is invertible.
- Compute $\hat{\mathbf{w}}$ for the following dataset:

$$\mathbf{X} = \begin{pmatrix} 1 & 2 \\ 3 & 4 \\ 5 & 6 \end{pmatrix}, \quad \mathbf{y} = \begin{pmatrix} 7 \\ 8 \\ 9 \end{pmatrix}.$$

- Compute $\hat{\mathbf{w}}$ for the following dataset: The r -th row of \mathbf{X} , for $r = 1, \dots, 28$, is given by the temperature recordings (with a 10-minute interval) during day $r/\text{Mar}/2023$ at FMI weather station *Kustavi Isokari*. The r -th row of \mathbf{y} is the maximum daytime temperature during day $r + 1/\text{Mar}/2023$ at the same weather station.

1.4. Eigenvalues and Positive Semi-Definiteness. The convergence properties of widely-used ML methods rely on the properties of psd matrices. Let $\mathbf{Q} = \mathbf{X}^T \mathbf{X}$, where $\mathbf{X} \in \mathbb{R}^{m \times d}$.

1. Prove that \mathbf{Q} is psd.
2. Compute the eigenvalues of \mathbf{Q} for $\mathbf{X} = \begin{pmatrix} 1 & 2 \\ 3 & 4 \end{pmatrix}$.
3. Compute the eigenvalues of \mathbf{Q} for the matrix \mathbf{X} used in Problem 1.3 that is constituted by FMI temperature recordings.

2 ML Basics

This chapter covers basic ML techniques instrumental for federated learning (FL). Content-wise, this chapter is more extensive compared to the following chapters. However, this chapter should be considerably easier to follow than the following chapters as it mainly refreshes pre-requisite knowledge.

2.1 Learning Goals

After completing this chapter, you will

- be familiar with the concept of data points (their features and labels), model and loss function,
- be familiar with empirical risk minimization (ERM) as a design principle for ML systems,
- know why and how validation is performed,
- be able to diagnose ML methods by comparing training error with validation error,
- be able to regularize ERM via modifying data, model and loss.

2.2 Three Components and a Design Principle

Machine Learning (ML) revolves around learning a hypothesis map h out of a hypothesis space \mathcal{H} that allows to accurately predict the label of a data point solely from its features. One of the most crucial steps in applying ML methods to a given application domain is the definition or choice of what

precisely a data point is. Coming up with a good choice or definition of data points is not trivial as it influences the overall performance of a ML method in many different ways.

As a recurring example for an FL application we use weather prediction. Here, data points represent the daily weather conditions around Finnish Meteorological Institute (FMI) weather stations. We denote a specific data point by \mathbf{z} . It is characterized by the following features:

- name of the FMI weather station, e.g., “TurkuRajakari”
- latitude lat and longitude lon of the weather station, e.g., $\text{lat} := 60.37788$, $\text{lon} := 22.0964$,
- timestamp of the measurement in the format YYYY-MM-DD HH:MM:SS, e.g., 2023-12-31 18:00:00

It is convenient to stack the features into a feature vector \mathbf{x} . The label $y \in \mathbb{R}$ of such a data point is the maximum daytime temperature in degree Celsius, e.g., -20 . We indicate the features \mathbf{x} and label y of a data point via the notation $\mathbf{z} = (\mathbf{x}, y)$.¹

We predict the label of a data point with features \mathbf{x} by the function value $h(\mathbf{x})$ of a hypothesis (map) $h(\cdot)$. The prediction will typically be not perfect, i.e., $h(\mathbf{x}) \neq y$. ML methods use a loss function $L((\mathbf{x}, y), h)$ to measure the error incurred by using the prediction $h(\mathbf{x})$ as a guess for the true label y . The

¹Strictly speaking, a data point \mathbf{z} might not be exactly the same as the pair of features \mathbf{x} and label y . Indeed, the data point might have additional properties that are neither used as features nor as label. A more precise notation would then be $\mathbf{x}(\mathbf{z})$ and $y(\mathbf{z})$, indicating that the features \mathbf{x} and label y are functions of the data point \mathbf{z} .

choice of loss function crucially influences the statistical and computational properties of the resulting ML method (see [6, Ch. 2]). In what follows, unless stated otherwise, we use the squared error loss $L(\mathbf{z}, h) := (y - h(\mathbf{x}))^2$ to measure the prediction error.

It seems natural to choose (or learn) a hypothesis that minimizes the average loss (or empirical risk) on a given set of data points

$$\mathcal{D} := \{(\mathbf{x}^{(1)}, y^{(1)}), \dots, (\mathbf{x}^{(m)}, y^{(m)})\}.$$

This is known as ERM,

$$\hat{h} \in \operatorname{argmin}_{h \in \mathcal{H}} (1/m) \sum_{r=1}^m (y^{(r)} - h(\mathbf{x}^{(r)}))^2. \quad (2.1)$$

As the notation in (2.1) indicates (using the symbol “ \in ” instead of “ $:=$ ”), there can be several different solutions to the optimization problem (2.1). Unless specified otherwise, \hat{h} can be used to denote any hypothesis in \mathcal{H} that has minimum average loss over \mathcal{D} .

Several important machine learning (ML) methods use a parametric model \mathcal{H} : Each hypothesis $h \in \mathcal{H}$ is defined by parameters $\mathbf{w} \in \mathbb{R}^d$, often indicated by the notation $h^{(\mathbf{w})}$. One important example for a parametrized model is the linear model [6, Sec. 3.1],

$$\mathcal{H}^{(d)} := \{h : \mathbb{R}^d \mapsto \mathbb{R} : h(\mathbf{x}) := \mathbf{w}^T \mathbf{x}\}. \quad (2.2)$$

Linear regression methods, for example, learn the parameters of a linear model by minimizing the average squared error loss. For linear regression, ERM becomes an optimization over the parameter space \mathbb{R}^d ,

$$\hat{\mathbf{w}}^{(\text{LR})} \in \operatorname{argmin}_{\mathbf{w} \in \mathbb{R}^d} \underbrace{(1/m) \sum_{r=1}^m (y^{(r)} - \mathbf{w}^T \mathbf{x}^{(r)})^2}_{:= f(\mathbf{w})}. \quad (2.3)$$

Note that (2.3) amounts to finding the minimum of a smooth and convex function

$$f(\mathbf{w}) = (1/m) \left[\mathbf{w}^T \mathbf{X}^T \mathbf{X} \mathbf{w} - 2 \mathbf{y}^T \mathbf{X} \mathbf{w} + \mathbf{y}^T \mathbf{y} \right] \quad (2.4)$$

$$\text{with the feature matrix } \mathbf{X} := (\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)})^T \quad (2.5)$$

$$\text{and the label vector } \mathbf{y} := (y^{(1)}, \dots, y^{(m)})^T \text{ of the training set } \mathcal{D}. \quad (2.6)$$

Inserting (2.4) into (2.3) allows to formulate linear regression as

$$\hat{\mathbf{w}}^{(\text{LR})} \in \underset{\mathbf{w} \in \mathbb{R}^d}{\operatorname{argmin}} \mathbf{w}^T \mathbf{Q} \mathbf{w} + \mathbf{w}^T \mathbf{q} \quad (2.7)$$

$$\text{with } \mathbf{Q} := (1/m) \mathbf{X}^T \mathbf{X}, \mathbf{q} := -(2/m) \mathbf{X}^T \mathbf{y}.$$

The matrix $\mathbf{Q} \in \mathbb{R}^{d \times d}$ is positive semi-definite (psd) with corresponding eigenvalue decomposition (EVD),

$$\mathbf{Q} = \sum_{j=1}^d \lambda_j(\mathbf{Q}) \mathbf{u}^{(j)} (\mathbf{u}^{(j)})^T. \quad (2.8)$$

The EVD (2.8) involves non-negative eigenvalues

$$0 \leq \lambda_1(\mathbf{Q}) \leq \dots \leq \lambda_d(\mathbf{Q}). \quad (2.9)$$

To train a ML model \mathcal{H} means to solve ERM (2.1) (or (2.3) for linear regression); the dataset \mathcal{D} is therefore referred to as a training set. The trained model results in the learnt hypothesis \hat{h} . We obtain practical ML methods by applying optimization algorithms to solve (2.1). Two key questions studied in ML are:

- **Computational aspects.** How much compute do we need to solve (2.1) ?




Fig. 2.1: ERM (2.1) for linear regression minimizes a convex quadratic function $\mathbf{w}^T \mathbf{Q} \mathbf{w} + \mathbf{w}^T \mathbf{q}$.

- **Statistical aspects.** How useful is the solution \hat{h} to (2.1) in general, i.e., how accurate is the prediction $\hat{h}(\mathbf{x})$ for the label y of an **arbitrary** data point with features \mathbf{x} ?

2.3 Computational Aspects of ERM

ML methods use optimization algorithms to solve (2.1), resulting in the learnt hypothesis \hat{h} . Within this book, we use optimization algorithms that are iterative methods: Starting from an initial choice $h^{(0)}$, they construct a sequence

$$h^{(0)}, h^{(1)}, h^{(2)}, \dots,$$

which are hopefully increasingly accurate approximations to a solution \hat{h} of (2.1). The computational complexity of such a ML method can be measured by the number of iterations required to guarantee some prescribed level of approximation.

For a parametric model and a smooth loss function, we can solve (2.3) by

gradient-based methods: Starting from an initial parameters $\mathbf{w}^{(0)}$, we iterate the gradient step:

$$\begin{aligned}\mathbf{w}^{(k)} &:= \mathbf{w}^{(k-1)} - \eta \nabla f(\mathbf{w}^{(k-1)}) \\ &= \mathbf{w}^{(k-1)} + (2\eta/m) \sum_{r=1}^m \mathbf{x}^{(r)} (y^{(r)} - (\mathbf{w}^{(k-1)})^T \mathbf{x}^{(r)}).\end{aligned}\quad (2.10)$$

How much computation do we need for one iteration of (2.10)? How many iterations do we need? We will try to answer the latter question in Chapter 4. The first question can be answered more easily for a typical computational infrastructure (e.g., “Python running on a commercial Laptop”). The evaluation of (2.10) then typically requires around m arithmetic operations (addition, multiplication).

It is instructive to consider the special case of a linear model that does not use any feature, i.e., $h(\mathbf{x}) = w$. For this extreme case, the ERM (2.3) has a simple closed-form solution:

$$\hat{w} = (1/m) \sum_{r=1}^m y^{(r)}.\quad (2.11)$$

Thus, for this special case of the linear model, solving (2.11) amounts to summing m numbers $y^{(1)}, \dots, y^{(m)}$. The amount of computation, measured by the number of elementary arithmetic operations, required by (2.11) is proportional to m .

2.4 Statistical Aspects of ERM

We can train a linear model on a given training set as ERM (2.3). But how useful is the solution $\hat{\mathbf{w}}$ of (2.3) for predicting the labels of data points outside

the training set? Consider applying the learnt hypothesis $h^{(\hat{\mathbf{w}})}$ to an arbitrary data point not contained in the training set. What can we say about the resulting prediction error $y - h^{(\hat{\mathbf{w}})}(\mathbf{x})$ in general? In other words, how well does $h^{(\hat{\mathbf{w}})}$ generalize beyond the training set.

The most widely used approach to study the generalization of ML methods is via probabilistic models. Here, we interpret each data point as a realization of an independent and identically distributed (i.i.d.) RV with probability distribution $p(\mathbf{x}, y)$. Under this i.i.d. assumption, we can evaluate the overall performance of a hypothesis $h \in \mathcal{H}$ via the expected loss (or risk)

$$\mathbb{E}\{L((\mathbf{x}, y), h)\}. \quad (2.12)$$

One example for a probability distribution $p(\mathbf{x}, y)$ relates the label y with the features \mathbf{x} of a data point as

$$y = \bar{\mathbf{w}}^T \mathbf{x} + \varepsilon \text{ with } \mathbf{x} \sim \mathcal{N}(\mathbf{0}, \mathbf{I}), \varepsilon \sim \mathcal{N}(0, \sigma^2), \mathbb{E}\{\varepsilon \mathbf{x}\} = \mathbf{0}. \quad (2.13)$$

A simple calculation reveals the expected squared error loss of a given linear hypothesis $h(\mathbf{x}) = \mathbf{x}^T \hat{\mathbf{w}}$ as²

$$\mathbb{E}\{(y - h(\mathbf{x}))^2\} = \|\bar{\mathbf{w}} - \hat{\mathbf{w}}\|^2 + \sigma^2. \quad (2.14)$$

The component σ^2 can be interpreted as the intrinsic noise level of the label y . We cannot hope to find a hypothesis with an expected loss below this level.

²Strictly speaking, the relation (2.14) only applies for constant (deterministic) model parameters $\hat{\mathbf{w}}$ that do not depend on the RVs whose realizations are the observed data points (see, e.g., (2.13)). However, the learnt model parameters $\hat{\mathbf{w}}$ might be the output of a ML method (such as (2.3)) that is applied to a dataset \mathcal{D} consisting of i.i.d. realizations from some underlying probability distribution. In this case, we need to replace the expectation on the LHS of (2.14) with a conditional expectation $\mathbb{E}\{(y - h(\mathbf{x}))^2 | \mathcal{D}\}$.

The first component of the RHS in (2.14) is the estimation error $\|\bar{\mathbf{w}} - \hat{\mathbf{w}}\|^2$ of a ML method that reads in the training set and delivers an estimate $\hat{\mathbf{w}}$ (e.g., via (2.3)) for the parameters of a linear hypothesis.

We next study the estimation error $\bar{\mathbf{w}} - \hat{\mathbf{w}}$ incurred by the specific estimate $\hat{\mathbf{w}} = \hat{\mathbf{w}}^{(\text{LR})}$ (2.7) delivered by linear regression methods. To this end, we first use the probabilistic model (2.13) to decompose the label vector \mathbf{y} in (2.6) as

$$\mathbf{y} = \mathbf{X}\bar{\mathbf{w}} + \mathbf{n}, \text{ with } \mathbf{n} := (\varepsilon^{(1)}, \dots, \varepsilon^{(m)})^T. \quad (2.15)$$

Inserting (2.15) into (2.7) yields

$$\hat{\mathbf{w}}^{(\text{LR})} \in \underset{\mathbf{w} \in \mathbb{R}^d}{\operatorname{argmin}} \mathbf{w}^T \mathbf{Q} \mathbf{w} + \mathbf{w}^T \mathbf{q}' + \mathbf{w}^T \mathbf{e} \quad (2.16)$$

$$\text{with } \mathbf{Q} := (1/m) \mathbf{X}^T \mathbf{X}, \mathbf{q}' := -(2/m) \mathbf{X}^T \mathbf{X} \bar{\mathbf{w}}, \text{ and } \mathbf{e} := -(2/m) \mathbf{X}^T \mathbf{n}. \quad (2.17)$$

Figure 2.2 depicts the objective function of (2.16). It is a perturbation of the convex quadratic function $\mathbf{w}^T \mathbf{Q} \mathbf{w} + \mathbf{w}^T \mathbf{q}'$, which is minimized at $\mathbf{w} = \bar{\mathbf{w}}$. In general, the minimizer $\hat{\mathbf{w}}^{(\text{LR})}$ delivered by linear regression is different from $\bar{\mathbf{w}}$ due the perturbation term $\mathbf{w}^T \mathbf{e}$ in (2.16).

The following result bounds the deviation between $\hat{\mathbf{w}}^{(\text{LR})}$ and $\bar{\mathbf{w}}$ under the assumption that the matrix $\mathbf{Q} = (1/m) \mathbf{X}^T \mathbf{X}$ is invertible.³

Proposition 2.1. *Consider a solution $\hat{\mathbf{w}}^{(\text{LR})}$ to the ERM instance (2.16) that is applied to the dataset (2.15). If the matrix $\mathbf{Q} = (1/m) \mathbf{X}^T \mathbf{X}$ is invertible, with minimum eigenvalue $\lambda_1(\mathbf{Q}) > 0$,*

$$\|\hat{\mathbf{w}}^{(\text{LR})} - \bar{\mathbf{w}}\|_2^2 \leq \frac{\|\mathbf{e}\|_2^2}{\lambda_1^2} \stackrel{(2.17)}{=} \frac{4}{m^2} \frac{\|\mathbf{X}^T \mathbf{n}\|_2^2}{\lambda_1^2}. \quad (2.18)$$

³Can you think of sufficient conditions on the feature matrix of the training set that ensure $\mathbf{Q} = (1/m) \mathbf{X}^T \mathbf{X}$ is invertible?



Fig. 2.2: The estimation error of linear regression is determined by the effect of the perturbation term $\mathbf{w}^T \mathbf{e}$ on the minimizer of the convex quadratic function $\mathbf{w}^T \mathbf{Q} \mathbf{w} + \mathbf{w}^T \mathbf{q}'$.

Proof. Let us rewrite (2.16) as

$$\hat{\mathbf{w}}^{(\text{LR})} \in \underset{\mathbf{w} \in \mathbb{R}^d}{\operatorname{argmin}} f(\mathbf{w}) \text{ with } f(\mathbf{w}) := (\mathbf{w} - \bar{\mathbf{w}})^T \mathbf{Q} (\mathbf{w} - \bar{\mathbf{w}}) + \mathbf{e}^T (\mathbf{w} - \bar{\mathbf{w}}). \quad (2.19)$$

Clearly $f(\bar{\mathbf{w}}) = 0$ and, in turn, $f(\hat{\mathbf{w}}) = \min_{\mathbf{w} \in \mathbb{R}^d} f(\mathbf{w}) \leq 0$. On the other hand,

$$\begin{aligned} f(\mathbf{w}) &\stackrel{(2.19)}{=} (\mathbf{w} - \bar{\mathbf{w}})^T \mathbf{Q} (\mathbf{w} - \bar{\mathbf{w}}) + \mathbf{e}^T (\mathbf{w} - \bar{\mathbf{w}}) \\ &\stackrel{(a)}{\geq} (\mathbf{w} - \bar{\mathbf{w}})^T \mathbf{Q} (\mathbf{w} - \bar{\mathbf{w}}) - \|\mathbf{e}\|_2 \|\mathbf{w} - \bar{\mathbf{w}}\|_2 \\ &\stackrel{(b)}{\geq} \lambda_1 \|\mathbf{w} - \bar{\mathbf{w}}\|_2^2 - \|\mathbf{e}\|_2 \|\mathbf{w} - \bar{\mathbf{w}}\|_2. \end{aligned} \quad (2.20)$$

Step (a) used Cauchy–Schwarz inequality and (b) used the EVD (2.8) of \mathbf{Q} . Evaluating (2.20) for $\mathbf{w} = \hat{\mathbf{w}}$ and combining with $f(\hat{\mathbf{w}}) \leq 0$ yields (2.18). \square

The bound (2.18) suggests that the estimation error $\hat{w}^{(\text{LR})} - \bar{w}$ is small if $\lambda_1(\mathbf{Q})$ is large. This smallest eigenvalue of the matrix $\mathbf{Q} = (1/m) \mathbf{X}^T \mathbf{X}$ could

be controlled by suitable choice (or transformation) of features \mathbf{x} of a data point. Trivially, we can increase $\lambda_1(\mathbf{Q})$ by a factor 100 if we scale each feature by a factor 10. However, this approach would also scale (by factor 100) the error term $\|\mathbf{X}^T \mathbf{n}\|_2^2$ in (2.18). For some applications, we can find feature transformations (“whitening”) that increases $\lambda_1(\mathbf{Q})$ but does not increase $\|\mathbf{X}^T \mathbf{n}\|_2^2$. We finally note that the error term $\|\mathbf{X}^T \mathbf{n}\|_2^2$ in (2.18) vanishes if the noise vector \mathbf{n} is orthogonal to the columns of the feature matrix \mathbf{X} .

It is instructive to evaluate the bound (2.18) for the special case where each data point has the same feature $x = 1$. Here, the probabilistic model (2.15) reduces to a *signal in noise* model,

$$y^{(r)} = x^{(r)}\bar{w} + \varepsilon^{(r)} \text{ with } x^{(r)} = 1, \quad (2.21)$$

with some true underlying parameter \bar{w} . The noise terms $\varepsilon^{(r)}$, for $r = 1, \dots, m$, are realizations of i.i.d. RVs with probability distribution $\mathcal{N}(0, \sigma^2)$. The feature matrix then becomes $\mathbf{X} = \mathbf{1}$ and, in turn, $\mathbf{Q} = \mathbf{1}$, $\lambda_1(\mathbf{Q}) = 1$. Inserting these values into (2.18), results in the bound

$$(\hat{w}^{(\text{LR})} - \bar{w})^2 \leq 4 \|\mathbf{n}\|_2^2 / m^2. \quad (2.22)$$

For the labels and features in (2.21), the solution of (2.16) is given by

$$\hat{w}^{(\text{LR})} = (1/m) \sum_{r=1}^m y^{(r)} \stackrel{(2.21)}{=} \bar{w} + (1/m) \sum_{r=1}^m \varepsilon^{(r)}. \quad (2.23)$$

2.5 Validation and Diagnosis of ML

The above analysis of the generalization error started from postulating a probabilistic model for the generation of data points. However, this probabilistic model might be wrong and the bound (2.18) does not apply. Thus, we might

want to use a more data-driven approach for assessing the usefulness of a learnt hypothesis \hat{h} obtained, e.g., from solving ERM (2.1).

Loosely speaking, validation tries to find out if a learnt hypothesis \hat{h} performs similar inside and outside the training set. In its most basic form, validation amounts to computing the average loss of a learnt hypothesis \hat{h} on some data points not included in the training set. We refer to these data points as the validation set.

Algorithm 2.1 summarizes a single iteration of a prototypical ML workflow that consists of model training and validation. The workflow starts with an initial choice for a dataset \mathcal{D} , model \mathcal{H} and loss function $L(\cdot, \cdot)$. We then repeat Algorithm 2.1 several times. After each repetition, based on the resulting training error and validation error, we modify (some of) the design choices for dataset, model and loss function. These design choices, including the decision of when to stop repeating Algorithm 2.1, are often guided by simple heuristics [6, Ch.6.6].

We can diagnose a ERM-based ML method, such as Algorithm 2.1, by comparing its training error with its validation error. This diagnosis is further enabled if we know a baseline $E^{(\text{ref})}$. One important source for a baseline $E^{(\text{ref})}$ are probabilistic models for the data points (see Section 2.4).

Given a probabilistic model $p(\mathbf{x}, y)$, we can compute the minimum achievable risk (2.12). Indeed, the minimum achievable risk is precisely the expected loss of the Bayes estimator $\hat{h}(\mathbf{x})$ of the label y , given the features \mathbf{x} of a data point. The Bayes estimator $\hat{h}(\mathbf{x})$ is fully determined by the probability distribution $p(\mathbf{x}, y)$ [27, Chapter 4].

A further potential source for a baseline $E^{(\text{ref})}$ is an existing, but for

Algorithm 2.1 One Iteration of ML Training and Validation

Input: dataset \mathcal{D} , model \mathcal{H} , loss function $L(\cdot, \cdot)$

- 1: split \mathcal{D} into a training set $\mathcal{D}^{(\text{train})}$ and a validation set $\mathcal{D}^{(\text{val})}$
- 2: learn a hypothesis via solving ERM

$$\hat{h} \in \operatorname{argmin}_{h \in \mathcal{H}} \sum_{(\mathbf{x}, y) \in \mathcal{D}^{(\text{train})}} L((\mathbf{x}, y), h) \quad (2.24)$$

- 3: compute resulting training error

$$E_t := (1/|\mathcal{D}^{(\text{train})}|) \sum_{(\mathbf{x}, y) \in \mathcal{D}^{(\text{train})}} L((\mathbf{x}, y), \hat{h})$$

- 4: compute validation error

$$E_v := (1/|\mathcal{D}^{(\text{val})}|) \sum_{(\mathbf{x}, y) \in \mathcal{D}^{(\text{val})}} L((\mathbf{x}, y), \hat{h})$$

Output: learnt hypothesis (or trained model) \hat{h} , training error E_t and validation error E_v

some reason unsuitable, ML method. This existing ML method might be computationally too expensive to be used for the ML application at hand. However, we might still use its statistical properties as a baseline.

We can also use the performance of human experts as a baseline. If we want to develop a ML method that detects skin cancer from images, a possible baseline is the classification accuracy achieved by experienced dermatologists [28].

We can diagnose a ML method by comparing the training error E_t with the validation error E_v and (if available) the baseline $E^{(\text{ref})}$.

- $E_t \approx E_v \approx E^{(\text{ref})}$: The training error is on the same level as the validation error and the baseline. There seems to be little point in trying to improve the method further since the validation error is already close to the baseline. Moreover, the training error is not much smaller than the validation error which indicates that there is no overfitting.
- $E_v \gg E_t$: The validation error is significantly larger than the training error, which hints at overfitting. We can address overfitting either by reducing the effective dimension of the hypothesis space or by increasing the size of the training set. To reduce the effective dimension of the hypothesis space, we can use fewer features (in a linear model), a smaller maximum depth of decision trees or fewer layers in an artificial neural network (ANN). Instead of this coarse-grained discrete model pruning, we can also reduce the effective dimension of a hypothesis space continuously via regularization (see [6, Ch. 7]).
- $E_t \approx E_v \gg E^{(\text{ref})}$: The training error is on the same level as the

validation error and both are significantly larger than the baseline. Thus, the learnt hypothesis seems to not overfit the training set. However, the training error achieved by the learnt hypothesis is significantly larger than the baseline. There can be several reasons for this to happen. First, it might be that the hypothesis space is too small, i.e., it does not include a hypothesis that provides a good approximation for the relation between features and label of a data point. One remedy to this situation is to use a larger hypothesis space, e.g., by including more features in a linear model, using higher polynomial degrees in polynomial regression, using deeper decision trees or ANNs with more hidden layers (deep net). Second, besides the model being too small, another reason for a large training error could be that the optimization algorithm used to solve ERM (2.24) is not working properly (see Chapter 4).

- $E_t \gg E_v$: The training error is significantly larger than the validation error. The idea of ERM (2.24) is to approximate the risk (2.12) of a hypothesis by its average loss on a training set $\mathcal{D} = \{(\mathbf{x}^{(r)}, y^{(r)})\}_{r=1}^m$. The mathematical underpinning for this approximation is the law of large numbers which characterizes the average of (realizations of) i.i.d. RVs. The accuracy of this approximation depends on the validity of two conditions: First, the data points used for computing the average loss “should behave” like realizations of i.i.d. RVs with a common probability distribution. Second, the number of data points used for computing the average loss must be sufficiently large.

Whenever the training set or validation set differs significantly from realizations of i.i.d. RVs, the interpretation (and comparison) of the

training error and the validation error of a learnt hypothesis becomes more difficult. As an extreme case, the validation set might consist of data points for which every hypothesis incurs small average loss (see Figure 2.3). Here, we might try to increase the size of the validation set by collecting more labeled data points or by using data augmentation (see Sec. 2.6). If the size of both training set and validation set is large but we still obtain $E_t \gg E_v$, one should verify if data points in these sets conform to the i.i.d. assumption. There are principled statistical tests for the validity of the i.i.d. assumption for a given dataset (see [29] and references therein).




Fig. 2.3: Example for an unlucky split into training set and validation set for the model $\mathcal{H} := \{h^{(1)}, h^{(2)}, h^{(3)}\}$.

2.6 Regularization

Consider a ERM-based ML method using a hypothesis space \mathcal{H} and dataset \mathcal{D} (we assume all data points are used for training). A key parameter for such

a ML method is the ratio $d_{\text{eff}}(\mathcal{H})/|\mathcal{D}|$ between the model size $d_{\text{eff}}(\mathcal{H})$ and the number $|\mathcal{D}|$ of data points. The tendency of the ML method to overfit increases with the ratio $d_{\text{eff}}(\mathcal{H})/|\mathcal{D}|$.

Regularization techniques decrease the ratio $d_{\text{eff}}(\mathcal{H})/|\mathcal{D}|$ via three (essentially equivalent) approaches:

- collect more data points, possibly via data augmentation (see Fig. 10.9),
- add penalty term $\alpha\mathcal{R}\{h\}$ to average loss in ERM (2.1) (see Fig. 10.9),
- shrink the hypothesis space, e.g., by adding constraints on the model parameters such as $\|\mathbf{w}\|_2 \leq 10$.

It can be shown that these three perspectives (corresponding to the three components data, model and loss) on regularization are closely related [6, Ch. 7]. For example, adding a penalty term $\alpha\mathcal{R}\{h\}$ in ERM (2.1) is equivalent to ERM (2.1) with a pruned hypothesis space $\mathcal{H}^{(\alpha)} \subseteq \mathcal{H}$. Using a larger α typically results in a smaller $\mathcal{H}^{(\alpha)}$.

One important example of regularization is to add a penalty term to the average loss is ridge regression. In particular, ridge regression uses the regularizer $\mathcal{R}\{h\} := \|\mathbf{w}\|_2^2$ for a linear hypothesis $h(\mathbf{x}) := \mathbf{w}^T \mathbf{x}$. Thus, ridge regression learns the parameters of a linear hypothesis via solving

$$\widehat{\mathbf{w}}^{(\alpha)} \in \operatorname{argmin}_{\mathbf{w} \in \mathbb{R}^d} \left[(1/m) \sum_{r=1}^m (y^{(r)} - \mathbf{w}^T \mathbf{x}^{(r)})^2 + \alpha \|\mathbf{w}\|_2^2 \right]. \quad (2.25)$$

The objective function in (2.25) can be interpreted as the objective function of linear regression applied to a modification of the training set \mathcal{D} : We replace each data point $(\mathbf{x}, y) \in \mathcal{D}$ by a sufficiently large number of i.i.d. realizations

of

$$(\mathbf{x} + \mathbf{n}, y) \text{ with } \mathbf{n} \sim \mathcal{N}(\mathbf{0}, \alpha \mathbf{I}). \quad (2.26)$$

Thus, ridge regression (2.25) is equivalent to linear regression applied to an augmented variant \mathcal{D}' of \mathcal{D} . The augmentation \mathcal{D}' is obtained by replacing each data point $(\mathbf{x}, y) \in \mathcal{D}$ with a sufficiently large number of noisy copies. Each copy of (\mathbf{x}, y) is obtained by adding a i.i.d. realization \mathbf{n} of a zero-mean Gaussian noise with covariance matrix $\alpha \mathbf{I}$ to the features \mathbf{x} (see (2.26)). The label of each copy of (\mathbf{x}, y) is equal to y , i.e., the label is not perturbed.




Fig. 2.4: Equivalence between data augmentation and loss penalization.

To study the computational aspects of ridge regression, let us rewrite (2.25) as

$$\begin{aligned} \hat{\mathbf{w}}^{(\alpha)} &\in \underset{\mathbf{w} \in \mathbb{R}^d}{\operatorname{argmin}} \mathbf{w}^T \mathbf{Q} \mathbf{w} + \mathbf{w}^T \mathbf{q} \\ \text{with } \mathbf{Q} &:= (1/m) \mathbf{X}^T \mathbf{X} + \alpha \mathbf{I}, \mathbf{q} := (-2/m) \mathbf{X}^T \mathbf{y}. \end{aligned} \quad (2.27)$$

Thus, like linear regression (2.7), also ridge regression minimizes a convex quadratic function. A main difference between linear regression (2.7) and

ridge regression (for $\alpha > 0$) is that the matrix \mathbf{Q} in (2.27) is guaranteed to be invertible for any training set \mathcal{D} . In contrast, the matrix \mathbf{Q} in (2.7) for linear regression might be singular for some training sets.⁴

The statistical aspects of the solutions to (2.27) (i.e., the parameters learnt by ridge regression) crucially depends on the value of α . This choice can be guided by an error analysis using a probabilistic model for the data (see Proposition 2.1). Instead of using a probabilistic model, we can also compare the training error and validation error of the hypothesis $h(\mathbf{x}) = (\widehat{\mathbf{w}}^{(\alpha)})^T \mathbf{x}$ learnt by ridge regression with different values of α .

2.7 Problems

2.1. Fundamental Limits for Linear Regression. Linear regression learns model parameters of a linear model in order to minimize the risk $\mathbb{E}\{(y - \mathbf{w}^T \mathbf{x})^2\}$ where (\mathbf{x}, y) is a RV. In practice we do not observe the RV (\mathbf{x}, y) itself but a (realization of a) sequence of i.i.d. samples $(\mathbf{x}^{(t)}, y^{(t)})$, for $t = 1, 2, \dots$. The minimax risk is a lower bound on the risk obtained by any learning method. Determine the minimax risk in terms of the probability distribution of (\mathbf{x}, y) .

2.2. Penalty Term as Data Augmentation. Consider a ML method that trains a model, with model parameters \mathbf{w} . The training uses ERM with squared error loss. Show that regularization of the model training via adding a penalty term $\alpha \|\mathbf{w}\|_2^2$ is equivalent to a specific form of data augmentation. What is the augmented training set ?

⁴Consider the extreme case where all features of each data point in the training set \mathcal{D} are zero.

2.3. Data Augmentation via Linear Interpolation. Consider a ML method that trains a model, with model parameters \mathbf{w} , from a training set \mathcal{D} . Each data point $\mathbf{z} \in \mathcal{D}$ is characterized by feature vector $\mathbf{x} \in \mathbb{R}^d$ and label $y \in \mathbb{R}$, i.e., $\mathbf{z} = (\mathbf{x}, y)$. We augment the training set by adding, for each pair of two different data points $\mathbf{z}, \mathbf{z}' \in \mathcal{D}$, the synthetic data points $\tilde{\mathbf{z}}^{(r)} := \mathbf{z} + (\mathbf{z}' - \mathbf{z})r/100$ and , for $r = 0, \dots, 99$. Does this augmentation typically increase the training error.

2.4. Ridge Regression via Deterministic Data Augmentation. Ridge regression is obtained from linear regression by adding the penalty term $\alpha \|\mathbf{w}\|_2^2$ to the average squared error loss incurred by the hypothesis $h^{(\mathbf{w})}$ on the training set \mathcal{D} ,

$$\min_{\mathbf{w}} (1/m) \sum_{r=1}^m (y^{(r)} - h(\mathbf{x}^{(r)}))^2 + \alpha \|\mathbf{w}\|_2^2. \quad (2.28)$$

Construct an augmented training set \mathcal{D}' such that the objective function of (2.28) coincides with the objective function of plain linear regression using \mathcal{D}' as training set. To construct \mathcal{D}' , add carefully chosen data points to the original training set $\mathcal{D} = \left\{ (y^{(1)}, \mathbf{x}^{(1)}), \dots, (y^{(m)}, \mathbf{x}^{(m)}) \right\}$. Generalize the construction of \mathcal{D}' to implement a generalized form of ridge regression,

$$\min_{\mathbf{w}} (1/m) \sum_{r=1}^m (y^{(r)} - h(\mathbf{x}^{(r)}))^2 + \alpha \|\mathbf{w} - \tilde{\mathbf{w}}\|_2^2. \quad (2.29)$$

Here, we used some prescribed reference model parameters $\tilde{\mathbf{w}}$. Note that (2.29) reduces to basic ridge regression (2.28) for the specific choice $\tilde{\mathbf{w}} = \mathbf{0}$.

3 A Design Principle for FL

Chapter 2 reviewed ML methods that use numeric arrays to store data points (their features and labels) and model parameters. We have also discussed ERM (and its regularization) as a main design principle for practical ML systems. This chapter extends the basic ML concepts from a centralized *single-dataset single-model* setting to FL applications involving distributed collections of data and models.

Section 3.2 introduces federated learning (FL) networks as a mathematical model for collections of devices that generate local datasets and train local models. Section 3.4 presents our main design principle for FL systems. This principle, referred to as GTV minimization (GTVMin), is a special case of ERM using a particular choice of loss function and model.

The model (or hypothesis space) used by GTVMin is a product-space generated by the local models at the nodes of an FL network. The loss function of GTVMin consists of two parts: the sum of loss functions at each node and a penalty term that measures the variation of local models across the edges of the FL network. The precise form of this penalty term is a design choice and depends on the nature of the local models.

For parametrized local models, we can construct a penalty by computing the differences of model parameters across the edges in the FL network. For non-parametric local models, we can measure the variation of two local models (connected by an edge) by comparing their predictions on a common test-dataset. We obtain the generalized total variation (GTV) by summing their variations over all edges in the FL network.

3.1 Learning Goals

After completing this chapter, you will

- be familiar with the concept of an FL network,
- know how to characterize connectivity of an FL network via the spectrum of the Laplacian matrix,
- know some measures for the variation of local models,
- be able to formulate FL as instances of GTVMin.

3.2 FL Networks

Consider a FL system that consists of devices, indexed by $i = 1, \dots, i$, each with the ability to generate a local dataset $\mathcal{D}^{(i)}$ and to train a personalized model $\mathcal{H}^{(i)}$. These devices collaborate with each other via some communication network in order to learn a local hypothesis $h^{(i)} \in \mathcal{H}^{(i)}$. We measure the quality of $h^{(i)} \in \mathcal{H}^{(i)}$ via some loss function $L_i(h^{(i)})$.

We now introduce the concept of a FL network as a mathematical model for FL applications. A FL network consists of an undirected weighted graph $\mathcal{G} = (\mathcal{V}, \mathcal{E})$ with nodes $\mathcal{V} := \{1, \dots, n\}$ and undirected edges \mathcal{E} between pairs of different nodes. Each node $\mathcal{V} \in \mathcal{V}$ represents some devices with different amounts of computational resources.

An undirected edge $\{i, i'\} \in \mathcal{E}$ in an FL network represents some form of similarity between device i and device i' . The amount of similarity is represented by an edge-weight $A_{i,i'}$. We can collect edge-weights into an

adjacency matrix $\mathbf{A} \in \mathbb{R}^{n \times n}$, with $A_{i,i'} = A_{i',i}$. Fig. 3.1 depicts an example for an FL network.




Fig. 3.1: Example of an FL network whose nodes $i \in \mathcal{V}$ represent devices. Each device i generates a local dataset $\mathcal{D}^{(i)}$ and trains a local model $\mathcal{H}^{(i)}$. Some of the devices i, i' are connected by an undirected edge $\{i, i'\}$ with a positive edge-weight $A_{i,i'}$.

Note that the undirected edges \mathcal{E} of an FL network encode a symmetric notion of similarity between devices: If the device i is similar to the device i' , i.e., $\{i, i'\} \in \mathcal{E}$, then also the device i' is similar to the device i . For some FL applications, an asymmetric notion of similarity, represented by directed edges, could be more accurate. However, the generalization of a FL network to directed graphs is beyond the scope of this book.

It can be convenient to replace a given FL network \mathcal{G} with an equivalent fully connected FL network \mathcal{G}' (see Figure 3.2). The fully connected graph \mathcal{G}' contains an edge between every pair of two different nodes i, i' ,

$$\mathcal{E}' = \{\{i, i'\} \text{ , with some } i, i' \in \mathcal{V}, i \neq i'\}.$$

The edge weights are chosen $A'_{i,i'} = A_{i,i'}$ for any edge $\{i, i'\} \in \mathcal{E}$ and $A'_{i,i'} = 0$ if the original FL network \mathcal{G} does not contain an edge between nodes i, i' .

A FL network is more than the undirected weighted graph \mathcal{G} : It also includes the local dataset $\mathcal{D}^{(i)}$ and local model $\mathcal{H}^{(i)}$ (or its model parameters



Fig. 3.2: Left: Some FL network \mathcal{G} consisting of $n = 4$ nodes. Right: Equivalent fully connected FL network \mathcal{G}' with the same nodes and non-zero edge weights $A'_{i,i'} = A_{i,i'}$ for $\{i, i'\} \in \mathcal{E}$ and $A'_{i,i'} = 0$ for $\{i, i'\} \notin \mathcal{E}$.

$\mathbf{w}^{(i)}$) for each device $i \in \mathcal{V}$. The details of the generation and the format of a local dataset will not be important in what follows. Indeed, the main purpose of a local dataset will be as a means to construct a loss function in order to evaluate a given choice of model parameters. However, to build intuition, we can think of a local dataset $\mathcal{D}^{(i)}$ as a labelled dataset

$$\mathcal{D}^{(i)} := \{(\mathbf{x}^{(i,1)}, y^{(i,1)}), \dots, (\mathbf{x}^{(i,m_i)}, y^{(i,m_i)})\}. \quad (3.1)$$

Here, $\mathbf{x}^{(i,r)}$ and $y^{(i,r)}$ denote, respectively, the features and the label of the r th data point in the local dataset $\mathcal{D}^{(i)}$. Note that the size m_i of the local dataset can vary between different nodes $i \in \mathcal{V}$.

It is convenient to collect the feature vectors $\mathbf{x}^{(i,r)}$ and labels $y^{(i,r)}$ into a feature matrix $\mathbf{X}^{(i)}$ and label vector $\mathbf{y}^{(i)}$, respectively,

$$\mathbf{X}^{(i)} := (\mathbf{x}^{(i,1)}, \dots, \mathbf{x}^{(i,m_i)})^T, \text{ and } \mathbf{y} := (y^{(1)}, \dots, y^{(m_i)})^T. \quad (3.2)$$

The local dataset $\mathcal{D}^{(i)}$ can then be represented compactly by the feature matrix $\mathbf{X}^{(i)} \in \mathbb{R}^{m_i \times d}$ and the vector $\mathbf{y}^{(i)} \in \mathbb{R}^{m_i}$.

Besides the local dataset $\mathcal{D}^{(i)}$, each node $i \in \mathcal{G}$ also carries a local model $\mathcal{H}^{(i)}$. Our focus is on parametric local models with by model parameters

$\mathbf{w}^{(i)} \in \mathbb{R}^d$, for $i = 1, \dots, n$. The usefulness of a specific choice for the local model parameter $\mathbf{w}^{(i)}$ is then measured by a local loss function $L_i(\mathbf{w}^{(i)})$, for $i = 1, \dots, n$. Note that we can use different local loss functions $L_i(\cdot) \neq L_{i'}(\cdot)$ at different nodes $i, i' \in \mathcal{V}$.

We now have introduced all the components of an FL network. Strictly speaking, a FL network is a tuple $(\mathcal{G}, \{\mathcal{H}^{(i)}\}_{i \in \mathcal{V}}, \{L_i(\cdot)\}_{i \in \mathcal{V}})$ consisting of an undirected weighted graph \mathcal{G} , a local model $\mathcal{H}^{(i)}$ and local loss function $L_i(\cdot)$ for each node $i \in \mathcal{V}$. In principle, all of these components are design choices that influence the computational and statistical properties of the FL algorithms presented in Chapter 5. The main focus of this book will be on the effect of the network structure, i.e., the edges \mathcal{E} of the graph \mathcal{G} , on the resulting FL systems.

The role (or meaning) of an edge $\{i, i'\}$ in a FL network is two-fold: First, it represents a communication link that allows to exchange messages between devices i, i' (see Sec. 5.3). Second, an edge $\{i, i'\}$ indicates similar statistical properties of local datasets generated by devices i, i' . It then seems natural to learn similar hypothesis maps $h^{(i)}, h^{(i')}$. This is actually the main idea behind all the FL algorithms that we will discuss in the rest of this book. To make this idea precise, we next discuss how to obtain a quantitative measures for how much local hypothesis maps $h^{(i)}$ vary across the edges $\{i, i'\} \in \mathcal{E}$ of an FL network.

3.3 Generalized Total Variation

Consider some FL network with nodes $i = 1, \dots, n$, undirected edges \mathcal{E} with edge weights $A_{i,i'} > 0$, for each $\{i, i'\} \in \mathcal{E}$. For each edge $\{i, i'\} \in \mathcal{E}$, we

want to couple the training of the corresponding local models $\mathcal{H}^{(i)}, \mathcal{H}^{(i')}$. The strength of this coupling is determined by $A_{i,i'}$. We implement the coupling by penalizing the discrepancy between (the model parameters) of local models (see Section 3.4).

We can measure the discrepancy between two local models $h^{(i)}, h^{(i')}$ across $\{i, i'\} \in \mathcal{E}$ in different ways. For example, we can compare their predictions on a common test set \mathcal{D} by computing $\sum_{\mathbf{x} \in \mathcal{D}} (h^{(i)}(\mathbf{x}) - h^{(i')}(\mathbf{x}))^2$. In what follows, our main focus will be on FL networks with parametrized local models $\mathcal{H}^{(i)}$, each having their own model parameters $\mathbf{w}^{(i)} \in \mathbb{R}^d$ with some fixed dimension d .

We measure the discrepancy between parametric local models by a function $\phi(\mathbf{w}^{(i)} - \mathbf{w}^{(i')})$ of the difference $\mathbf{w}^{(i)} - \mathbf{w}^{(i')}$. The function ϕ is arbitrary as long as it is monotonically increasing in some norm on the Euclidean space \mathbb{R}^d [18, 30]. This requirement ensures that $\phi(\mathbf{w}^{(i)} - \mathbf{w}^{(i')}) = \phi(\mathbf{w}^{(i')} - \mathbf{w}^{(i)})$ and, in turn, allows to use $\phi(\mathbf{w}^{(i)} - \mathbf{w}^{(i')})$ as a measure of variation across an undirected edge $\{i, i'\} \in \mathcal{E}$.

By summing up the edge-wise variations (weighted by the edge weights), we obtain the total variation of a collection of local model parameters,

$$\sum_{\{i, i'\} \in \mathcal{E}} A_{i, i'} \phi(\mathbf{w}^{(i)} - \mathbf{w}^{(i')}). \quad (3.3)$$

Our main focus will be on the special case of (3.3), obtained for $\phi(\cdot) := \|\cdot\|_2^2$,

$$\sum_{\{i, i'\} \in \mathcal{E}} A_{i, i'} \left\| \mathbf{w}^{(i)} - \mathbf{w}^{(i')} \right\|_2^2. \quad (3.4)$$

The choice of penalty $\phi(\cdot)$ has crucial impact on the computational and statistical properties of the FL algorithms presented in Ch. 5. Our main

choice during the rest of this book will be the penalty function $\phi(\cdot) := \|\cdot\|_2^2$. This choice often allows to formulate FL as the minimization of a smooth convex function, which can be done via simple gradient-based methods (see Ch. 7). On the other hand, choosing ϕ to be a norm results in FL algorithms that require more computation but less training data [30].

The connectivity of an FL network \mathcal{G} can be characterized locally - around a node $i \in \mathcal{V}$ - by its weighted node degree

$$d^{(i)} := \sum_{i' \in \mathcal{N}^{(i)}} A_{i,i'}. \quad (3.5)$$

Here, we used the neighbourhood $\mathcal{N}^{(i)} := \{i' \in \mathcal{V} : \{i, i'\} \in \mathcal{E}\}$ of node $i \in \mathcal{V}$. A global characterization for the connectivity of \mathcal{G} is the maximum weighted node degree

$$d_{\max}^{(\mathcal{G})} := \max_{i \in \mathcal{V}} d^{(i)} \stackrel{(3.5)}{=} \max_{i \in \mathcal{V}} \sum_{i' \in \mathcal{N}^{(i)}} A_{i,i'}. \quad (3.6)$$

Besides inspecting its (maximum) node degrees, we can study the connectivity of \mathcal{G} also via the eigenvalues and eigenvectors of its Laplacian matrix $\mathbf{L}^{(\mathcal{G})} \in \mathbb{R}^{n \times n}$.⁵ The Laplacian matrix of an undirected weighted graph \mathcal{G} is defined element-wise as (see Fig. 10.7)

$$L_{i,i'}^{(\mathcal{G})} := \begin{cases} -A_{i,i'} & \text{for } i \neq i', \{i, i'\} \in \mathcal{E} \\ \sum_{i'' \neq i} A_{i,i''} & \text{for } i = i' \\ 0 & \text{else.} \end{cases} \quad (3.7)$$

⁵The study of graphs via the eigenvalues and eigenvectors of associated matrices is the main subject of spectral graph theory [31, 32].



Fig. 3.3: Left: Example for an FL network \mathcal{G} with three nodes $i = 1, 2, 3$. Right: Laplacian matrix $\mathbf{L}^{(\mathcal{G})} \in \mathbb{R}^{3 \times 3}$ of \mathcal{G} .

The Laplacian matrix is symmetric and psd which follows from the identity

$$\begin{aligned} \mathbf{w}^T (\mathbf{L}^{(\mathcal{G})} \otimes \mathbf{I}) \mathbf{w} &= \sum_{\{i, i'\} \in \mathcal{E}} A_{i, i'} \left\| \mathbf{w}^{(i)} - \mathbf{w}^{(i')} \right\|_2^2 \\ \text{for any } d \in \mathbb{N}, \mathbf{w} &:= \underbrace{\left((\mathbf{w}^{(1)})^T, \dots, (\mathbf{w}^{(n)})^T \right)^T}_{=:\text{stack}\{\mathbf{w}^{(i)}\}_{i=1}^n} \in \mathbb{R}^{d \cdot n}. \end{aligned} \quad (3.8)$$

As a psd matrix, $\mathbf{L}^{(\mathcal{G})}$ possesses an EVD

$$\mathbf{L}^{(\mathcal{G})} = \sum_{i=1}^n \lambda_i \mathbf{u}^{(i)} (\mathbf{u}^{(i)})^T, \quad (3.9)$$

with orthonormal eigenvectors $\mathbf{u}^{(1)}, \dots, \mathbf{u}^{(n)}$ and corresponding list of eigenvalues

$$0 = \lambda_1(\mathbf{L}^{(\mathcal{G})}) \leq \lambda_2(\mathbf{L}^{(\mathcal{G})}) \leq \dots \leq \lambda_n(\mathbf{L}^{(\mathcal{G})}). \quad (3.10)$$

We just write λ_i instead of $\lambda_i(\mathbf{L}^{(\mathcal{G})})$ if the Laplacian matrix $\mathbf{L}^{(\mathcal{G})}$ is clear from context. The eigenvalue $\lambda_i(\mathbf{L}^{(\mathcal{G})})$ corresponds to the eigenvector $\mathbf{u}^{(i)}$, i.e., $\mathbf{L}^{(\mathcal{G})} \mathbf{u}^{(i)} = \lambda_i(\mathbf{L}^{(\mathcal{G})}) \mathbf{u}^{(i)}$ for $i = 1, \dots, n$.

It is important to note that the ordered list of eigenvalues (3.10) is uniquely determined for a given Laplacian matrix. In contrast, the eigenvectors $\mathbf{u}^{(i)}$ in

(3.9) are not unique in general.⁶

The ordered eigenvalues $\lambda_i(\mathbf{L}^{(\mathcal{G})})$ in (3.10) can be computed (or characterized) via the Courant–Fischer–Weyl min-max characterization (CFW) [3, Thm. 8.1.2.]. Two important special cases of this characterization are [31, 32]

$$\begin{aligned} \lambda_n(\mathbf{L}^{(\mathcal{G})}) &\stackrel{\text{CFW}}{=} \max_{\substack{\mathbf{v} \in \mathbb{R}^n \\ \|\mathbf{v}\|=1}} \mathbf{v}^T \mathbf{L}^{(\mathcal{G})} \mathbf{v} \\ &\stackrel{(3.8)}{=} \max_{\substack{\mathbf{v} \in \mathbb{R}^n \\ \|\mathbf{v}\|=1}} \sum_{\{i,i'\} \in \mathcal{E}} A_{i,i'} (v_i - v_{i'})^2 \end{aligned} \quad (3.11)$$

and

$$\begin{aligned} \lambda_2(\mathbf{L}^{(\mathcal{G})}) &\stackrel{\text{CFW}}{=} \min_{\substack{\mathbf{v} \in \mathbb{R}^n \\ \mathbf{v}^T \mathbf{1} = 0 \\ \|\mathbf{v}\|=1}} \mathbf{v}^T \mathbf{L}^{(\mathcal{G})} \mathbf{v} \\ &\stackrel{(3.8)}{=} \min_{\substack{\mathbf{v} \in \mathbb{R}^n \\ \mathbf{v}^T \mathbf{1} = 0 \\ \|\mathbf{v}\|=1}} \sum_{\{i,i'\} \in \mathcal{E}} A_{i,i'} (v_i - v_{i'})^2. \end{aligned} \quad (3.12)$$

By (3.8), we can compute the total variation of a collection of model parameters via the quadratic form $\mathbf{w}^T (\mathbf{L}^{(\mathcal{G})} \otimes \mathbf{I}_{d \times d}) \mathbf{w}$. This quadratic form involves the vector $\mathbf{w} \in \mathbb{R}^{nd}$ which is obtained by stacking the local model parameters $\mathbf{w}^{(i)}$, for $i = 1, \dots, n$. Another consequence of (3.8) is that any collection of identical local model parameters, stacked into the vector

$$\mathbf{w} = \text{stack}\{\mathbf{c}\} = (\mathbf{c}^T, \dots, \mathbf{c}^T)^T, \text{ with some } \mathbf{c} \in \mathbb{R}^d, \quad (3.13)$$

is an eigenvector of $\mathbf{L}^{(\mathcal{G})} \otimes \mathbf{I}$ with corresponding eigenvalue $\lambda_1 = 0$ (see (3.10)). Thus, the Laplacian matrix of any FL network is singular (non-invertible).

⁶Consider the scenario where the list (3.10) contains repeated entries, i.e., several ordered eigenvalues are identical.

The second eigenvalue λ_2 of the Laplacian matrix provides a great deal of information about the connectivity structure of \mathcal{G} .⁷

- Consider the case $\lambda_2 = 0$: Here, beside the eigenvector (3.13), we can find at least one additional eigenvector

$$\tilde{\mathbf{w}} = \text{stack}\{\mathbf{w}^{(i)}\}_{i=1}^n \text{ with } \mathbf{w}^{(i)} \neq \mathbf{w}^{(i')} \text{ for some } i, i' \in \mathcal{V}, \quad (3.14)$$

of $\mathbf{L}^{(\mathcal{G})} \otimes \mathbf{I}$ with eigenvalue equal to 0. In this case, the graph \mathcal{G} is not connected, i.e., we can find two subsets (components) of nodes that do not have any edge between them (see Fig. 3.4). For each connected component \mathcal{C} , we can construct the eigenvector by assigning the same (non-zero) vector $\mathbf{c} \in \mathbb{R}^d \setminus \{\mathbf{0}\}$ to all nodes $i \in \mathcal{C}$ and the zero vector $\mathbf{0}$ to the remaining nodes $i \in \mathcal{V} \setminus \mathcal{C}$.

- On the other hand, if $\lambda_2 > 0$ then \mathcal{G} is connected. Moreover, the larger the value of λ_2 , the stronger the connectivity between the nodes in \mathcal{G} . Indeed, adding edges to \mathcal{G} can only increase the objective in (3.12) and, in turn, λ_2 .

In what follows, we will make use of the lower bound [33]

$$\sum_{\{i,i'\} \in \mathcal{E}} A_{i,i'} \left\| \mathbf{w}^{(i)} - \mathbf{w}^{(i')} \right\|_2^2 \geq \lambda_2 \sum_{i=1}^n \left\| \mathbf{w}^{(i)} - \text{avg}\{\mathbf{w}^{(i)}\} \right\|_2^2. \quad (3.15)$$

Here, $\text{avg}\{\mathbf{w}^{(i)}\} := (1/n) \sum_{i=1}^n \mathbf{w}^{(i)}$ is the average of all local model parameters. The bound (3.15) follows from (3.8) and the CFW for the eigenvalues of the matrix $\mathbf{L}^{(\mathcal{G})} \otimes \mathbf{I}$.

⁷Much of spectral graph theory is devoted to the analysis of λ_2 for different graph constructions [31, 32].




Fig. 3.4: FL network with graph \mathcal{G} consisting of $n=6$ nodes that form two connected components $\mathcal{C}^{(1)}, \mathcal{C}^{(2)}$.

The quantity $\sum_{i=1}^n \|\mathbf{w}^{(i)} - \text{avg}\{\mathbf{w}^{(i)}\}_{i=1}^n\|_2^2$ on the RHS of (3.15) has an interesting geometric interpretation: It is the squared Euclidean norm of the projection of the stacked local model parameters $\mathbf{w} := \left((\mathbf{w}^{(1)})^T, \dots, (\mathbf{w}^{(n)})^T \right)^T$ on the orthogonal complement of the subspace

$$\mathcal{S} := \left\{ \mathbf{1} \otimes \mathbf{a} : \mathbf{a} \in \mathbb{R}^d \right\} = \left\{ (\mathbf{a}^T, \dots, \mathbf{a}^T)^T, \text{ for some } \mathbf{a} \in \mathbb{R}^d \right\} \subseteq \mathbb{R}^{dn}. \quad (3.16)$$

Indeed, the projection $\mathbf{P}_{\mathcal{S}}\mathbf{w}$ of $\mathbf{w} \in \mathbb{R}^{nd}$ on \mathcal{S} given explicitly as

$$\mathbf{P}_{\mathcal{S}}\mathbf{w} = (\mathbf{a}^T, \dots, \mathbf{a}^T)^T, \text{ with } \mathbf{a} = \text{avg}\{\mathbf{w}^{(i)}\}_{i=1}^n. \quad (3.17)$$

The projection on the orthogonal complement \mathcal{S}^\perp , in turn, is given explicitly as

$$\mathbf{P}_{\mathcal{S}^\perp}\mathbf{w} = \mathbf{w} - \mathbf{P}_{\mathcal{S}}\mathbf{w} = \text{stack}\left\{ \mathbf{w}^{(i)} - \text{avg}\{\mathbf{w}^{(i)}\}_{i=1}^n \right\}_{i=1}^n. \quad (3.18)$$

3.4 Generalized Total Variation Minimization

Consider some FL network \mathcal{G} with nodes $i \in \mathcal{V}$ representing devices that learn personalized model parameters $\mathbf{w}^{(i)}$. The usefulness of a specific choice for the model parameters $\mathbf{w}^{(i)}$ is measured by a local loss function $L_i(\mathbf{w}^{(i)})$. We are mainly interested in FL applications where the local loss functions do

not provide enough information for learning accurate model parameters.⁸ We therefore require learnt model parameters to not only incur a small local loss but also to have a small total variation $\sum_{\{i,i'\} \in \mathcal{E}} A_{i,i'} \phi(\mathbf{w}^{(i)} - \mathbf{w}^{(i')})$.

GTV minimization (GTVMin) optimally balances the (average) local loss and total variation ,

$$\{\widehat{\mathbf{w}}^{(i)}\}_{i=1}^n \in \underset{\text{stack}\{\mathbf{w}^{(i)}\}_{i=1}^n}{\operatorname{argmin}} \sum_{i \in \mathcal{V}} L_i(\mathbf{w}^{(i)}) + \alpha \sum_{\{i,i'\} \in \mathcal{E}} A_{i,i'} \phi(\mathbf{w}^{(i)} - \mathbf{w}^{(i')}) \quad (\text{GTVMin}). \quad (3.19)$$

Our main focus will be on the special case of (3.19), obtained for $\phi(\cdot) := \|\cdot\|_2^2$,

$$\{\widehat{\mathbf{w}}^{(i)}\}_{i=1}^n \in \underset{\text{stack}\{\mathbf{w}^{(i)}\}_{i=1}^n}{\operatorname{argmin}} \sum_{i \in \mathcal{V}} L_i(\mathbf{w}^{(i)}) + \alpha \sum_{\{i,i'\} \in \mathcal{E}} A_{i,i'} \left\| \mathbf{w}^{(i)} - \mathbf{w}^{(i')} \right\|_2^2. \quad (3.20)$$

Note that GTVMin is an instance of regularized empirical risk minimization (RERM): The regularizer is the generalized total variation (GTV) of local model parameters over the weighted edges $A_{i,i'}$ of the FL network. Clearly, the FL network is an important design choice for GTVMin-based methods. This choice can be guided by computational aspects and statistical aspects of GTVMin-based FL systems.

Some application domains allow to leverage domain expertise to guess a useful choice for the FL network. If local datasets are generated at different geographic locations, we might use nearest-neighbour graphs based on geodesic distances between data generators (e.g., FMI weather stations). Chapter 7 will also discuss graph learning methods that determine the edge weights $A_{i,i'}$ in a data-driven fashion, i.e., directly from the local datasets $\mathcal{D}^{(i)}, \mathcal{D}^{(i')}$.

⁸For example, the local loss function can be obtained from the training error on a local dataset that is much too small relative to the effective dimension of the local model (see Sec. 2.6).

Let us now consider the special case of GTVMin with local models being a linear model. For each node $i \in \mathcal{V}$ of the FL network, we want to learn the parameters $\mathbf{w}^{(i)}$ of a linear hypothesis $h^{(i)}(\mathbf{x}) := (\mathbf{w}^{(i)})^T \mathbf{x}$. We measure the quality of the parameters via the average squared error loss

$$\begin{aligned} L_i(\mathbf{w}^{(i)}) &:= (1/m_i) \sum_{r=1}^{m_i} \left(y^{(i,r)} - (\mathbf{w}^{(i)})^T \mathbf{x}^{(i,r)} \right)^2 \\ &\stackrel{(3.2)}{=} (1/m_i) \|\mathbf{y}^{(i)} - \mathbf{X}^{(i)} \mathbf{w}^{(i)}\|_2^2. \end{aligned} \quad (3.21)$$

Inserting (3.21) into (3.20), yields the following instance of GTVMin to train local linear models,

$$\{\widehat{\mathbf{w}}^{(i)}\} \in \underset{\{\mathbf{w}^{(i)}\}_{i=1}^n}{\operatorname{argmin}} \sum_{i \in \mathcal{V}} (1/m_i) \|\mathbf{y}^{(i)} - \mathbf{X}^{(i)} \mathbf{w}^{(i)}\|_2^2 + \alpha \sum_{\{i,i'\} \in \mathcal{E}} A_{i,i'} \|\mathbf{w}^{(i)} - \mathbf{w}^{(i')}\|_2^2. \quad (3.22)$$

The identity (3.8) allows to rewrite (3.22) using the Laplacian matrix $\mathbf{L}^{(\mathcal{G})}$ as

$$\widehat{\mathbf{w}}^{(i)} \in \underset{\mathbf{w} = \text{stack}\{\mathbf{w}^{(i)}\}}{\operatorname{argmin}} \sum_{i \in \mathcal{V}} (1/m_i) \|\mathbf{y}^{(i)} - \mathbf{X}^{(i)} \mathbf{w}^{(i)}\|_2^2 + \alpha \mathbf{w}^T (\mathbf{L}^{(\mathcal{G})} \otimes \mathbf{I}_d) \mathbf{w}. \quad (3.23)$$

Let us rewrite the objective function in (3.23) as

$$\mathbf{w}^T \left(\begin{pmatrix} \mathbf{Q}^{(1)} & 0 & \cdots & 0 \\ 0 & \mathbf{Q}^{(2)} & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & \mathbf{Q}^{(n)} \end{pmatrix} + \alpha \mathbf{L}^{(\mathcal{G})} \otimes \mathbf{I} \right) \mathbf{w} + ((\mathbf{q}^{(1)})^T, \dots, (\mathbf{q}^{(n)})^T) \mathbf{w} \quad (3.24)$$

with $\mathbf{Q}^{(i)} = (1/m_i) (\mathbf{X}^{(i)})^T \mathbf{X}^{(i)}$, and $\mathbf{q}^{(i)} := (-2/m_i) (\mathbf{X}^{(i)})^T \mathbf{y}^{(i)}$.

Thus, like linear regression (2.7) and ridge regression (2.27), also GTVMin (3.23) (for local linear models $\mathcal{H}^{(i)}$) minimizes a convex quadratic function,

$$\{\widehat{\mathbf{w}}^{(i)}\}_{i=1}^n \in \underset{\mathbf{w} = \text{stack}\{\mathbf{w}^{(i)}\}_{i=1}^n}{\operatorname{argmin}} \mathbf{w}^T \mathbf{Q} \mathbf{w} + \mathbf{q}^T \mathbf{w}. \quad (3.25)$$

Here, we used the psd matrix

$$\mathbf{Q} := \begin{pmatrix} \mathbf{Q}^{(1)} & \mathbf{0} & \dots & \mathbf{0} \\ \mathbf{0} & \mathbf{Q}^{(2)} & \dots & \mathbf{0} \\ \vdots & \vdots & \ddots & \vdots \\ \mathbf{0} & \mathbf{0} & \dots & \mathbf{Q}^{(n)} \end{pmatrix} + \alpha \mathbf{L}^{(\mathcal{G})} \otimes \mathbf{I} \text{ with } \mathbf{Q}^{(i)} := (1/m_i) (\mathbf{X}^{(i)})^T \mathbf{X}^{(i)} \quad (3.26)$$

and the vector

$$\mathbf{q} := ((\mathbf{q}^{(1)})^T, \dots, (\mathbf{q}^{(n)})^T)^T, \text{ with } \mathbf{q}^{(i)} := (-2/m_i) (\mathbf{X}^{(i)})^T \mathbf{y}^{(i)}. \quad (3.27)$$

3.4.1 Computational Aspects of GTVMin

Chapter 5 will apply optimization methods to solve GTVMin, resulting in practical FL algorithms. Different instances of GTVMin favour different classes of optimization methods. For example, using a differentiable loss function allows to apply gradient-based methods (see Chapter 4) to solve GTVMin.

Another important class of loss functions are those for which we can efficiently compute the proximal operator

$$\mathbf{prox}_{L,\rho}(\mathbf{w}) := \underset{\mathbf{w}' \in \mathbb{R}^d}{\operatorname{argmin}} L(\mathbf{w}') + (\rho/2) \|\mathbf{w} - \mathbf{w}'\|_2^2 \text{ for some } \rho > 0.$$

Some authors refer to functions L for which $\mathbf{prox}_{L,\rho}(\mathbf{w})$ can be computed easily as *simple* or proximable [34]. GTVMin with proximable loss functions can be solved quite efficiently via proximal algorithms [35].

Besides influencing the choice of optimization method, the design choices underlying GTVMin also determine the amount of computation that is required by a given optimization method. For example, using an FL network

with relatively few edges (“sparse graphs”) typically results in a smaller computational complexity. Indeed, Chapter 5 discusses GTVMin-based algorithms requiring an amount of computation that is proportional to the number of edges in the FL network.

Let us now consider the computational aspects of GTVMin (3.22) to train local linear models. As discussed above, this instance is equivalent to solving (3.25). Any solution $\widehat{\mathbf{w}}$ of (3.25) (and, in turn, (3.22)) is characterized by the zero-gradient condition

$$\mathbf{Q}\widehat{\mathbf{w}} = -(1/2)\mathbf{q}, \quad (3.28)$$

with \mathbf{Q}, \mathbf{q} as defined in (3.26) and (3.27). If the matrix \mathbf{Q} in (3.28) is invertible, the solution to (3.28) and, in turn, to the GTVMin instance (3.22) is unique and given by $\widehat{\mathbf{w}} = (-1/2)\mathbf{Q}^{-1}\mathbf{q}$.

The size of the matrix \mathbf{Q} (see (3.26)) is proportional to the number of nodes in the FL network \mathcal{G} , which might be on the order of millions (or even billions) for internet-scale applications. For such large systems, we typically cannot use direct matrix inversion methods (e.g., based on Gaussian elimination) to compute \mathbf{Q}^{-1} .⁹ Instead, we typically need to resort to iterative methods [36, 37].

One important family of such iterative methods are the gradient-based methods which we will discuss in Chapter 4. Starting from an initial choice for the local model parameters $\widehat{\mathbf{w}}_0 = (\widehat{\mathbf{w}}_0^{(1)}, \dots, \widehat{\mathbf{w}}_0^{(n)})$, these methods repeat (variants of) the gradient step,

$$\widehat{\mathbf{w}}_{k+1} := \widehat{\mathbf{w}}_k - \eta(2\mathbf{Q}\widehat{\mathbf{w}}_k + \mathbf{q}) \text{ for } k = 0, 1, \dots$$

⁹How many arithmetic operations (addition, multiplication) do you think are required to invert an arbitrary matrix $\mathbf{Q} \in \mathbb{R}^{d \times d}$?

The gradient step results in the updated local model parameters $\widehat{\mathbf{w}}^{(i)}$ which we stacked into

$$\widehat{\mathbf{w}}_{k+1} := \left((\widehat{\mathbf{w}}^{(1)})^T, \dots, (\widehat{\mathbf{w}}^{(n)})^T \right)^T.$$

We repeat the gradient step for a sufficient number of times, according to some stopping criterion (see Chapter 4).

3.4.2 Statistical Aspects of GTVMin

How useful are the solutions of GTVMin (3.20) as a choice for the local model parameters? To answer this question, we use - as for the statistical analysis of ERM in Chapter 2 - a probabilistic model for the local datasets. In particular, we use a variant of an i.i.d. assumption: Each local dataset $\mathcal{D}^{(i)}$, consists of data points whose features and labels are realizations of i.i.d. RVs

$$\mathbf{y}^{(i)} = \underbrace{(\mathbf{x}^{(i,1)}, \dots, \mathbf{x}^{(i,m_i)})^T}_{\text{local feature matrix } \mathbf{X}^{(i)}} \overline{\mathbf{w}}^{(i)} + \boldsymbol{\epsilon}^{(i)} \text{ with } \mathbf{x}^{(i,r)} \stackrel{\text{i.i.d.}}{\sim} \mathcal{N}(\mathbf{0}, \mathbf{I}), \boldsymbol{\epsilon}^{(i)} \sim \mathcal{N}(\mathbf{0}, \sigma^2 \mathbf{I}). \quad (3.29)$$

In contrast to the probabilistic model (2.13) (which we used for the analysis of ERM), the probabilistic model (3.29) allows for different node-specific parameters $\overline{\mathbf{w}}^{(i)}$, for $i \in \mathcal{V}$. In particular, the entire dataset obtained from pooling all local datasets does not conform to an i.i.d. assumption.

In what follows, we focus on the GTVMin instance (3.22) to learn the parameters $\mathbf{w}^{(i)}$ of a local linear model for each node $i \in \mathcal{V}$. For a reasonable choice for FL network, the parameters $\overline{\mathbf{w}}^{(i)}, \overline{\mathbf{w}}^{(i')}$ at connected nodes $\{i, i'\} \in \mathcal{E}$ should be similar. However, we cannot choose the edge weights based on parameters $\overline{\mathbf{w}}^{(i)}$ as they are unknown. We can only use (noisy) estimates for $\overline{\mathbf{w}}^{(i)}$ the features and labels of the data points in the local datasets (see

Chapter 7).

Consider an FL network with nodes carrying local datasets generated from the probabilistic model (3.29) with true model parameters $\bar{\mathbf{w}}^{(i)}$. For ease of exposition, we assume that

$$\bar{\mathbf{w}}^{(i)} = \bar{\mathbf{c}}, \text{ for some } \bar{\mathbf{c}} \in \mathbb{R}^d \text{ and all } i \in \mathcal{V}. \quad (3.30)$$

To study the deviation between the solutions $\hat{\mathbf{w}}^{(i)}$ of (3.22) and the true underlying parameters $\bar{\mathbf{w}}^{(i)}$, we decompose it as

$$\hat{\mathbf{w}}^{(i)} =: \tilde{\mathbf{w}}^{(i)} + \hat{\mathbf{c}}, \text{ with } \hat{\mathbf{c}} := (1/n) \sum_{i'=1}^n \hat{\mathbf{w}}^{(i')}. \quad (3.31)$$

The component $\hat{\mathbf{c}}$ is identical at all nodes $i \in \mathcal{V}$ and obtained as the orthogonal projection of $\hat{\mathbf{w}} = \text{stack}\{\hat{\mathbf{w}}^{(i)}\}_{i=1}^n$ on the subspace (3.16). The component $\tilde{\mathbf{w}}^{(i)} := \hat{\mathbf{w}}^{(i)} - (1/n) \sum_{i'=1}^n \hat{\mathbf{w}}^{(i')}$ consists of the deviations, for each node i , between the GTVMin solution $\hat{\mathbf{w}}^{(i)}$ and their average over all nodes. Trivially, the average of the deviations $\tilde{\mathbf{w}}^{(i)}$ across all nodes is the zero vector, $(1/n) \sum_{i=1}^n \tilde{\mathbf{w}}^{(i)} = \mathbf{0}$.

The decomposition (3.31) entails an analogous (orthogonal) decomposition of the error $\hat{\mathbf{w}}^{(i)} - \bar{\mathbf{w}}^{(i)}$. Indeed, for identical true underlying model parameters (3.30) (which makes $\bar{\mathbf{w}}$ an element of the subspace (3.16)), we have

$$\sum_{i=1}^n \|\hat{\mathbf{w}}^{(i)} - \bar{\mathbf{w}}^{(i)}\|_2^2 \stackrel{(3.30), (3.31)}{=} \underbrace{\sum_{i=1}^n \|\mathbf{c} - \hat{\mathbf{c}}\|_2^2}_{n\|\mathbf{c} - \hat{\mathbf{c}}\|_2^2} + \sum_{i=1}^n \|\tilde{\mathbf{w}}^{(i)}\|_2^2. \quad (3.32)$$

The following proposition provides an upper bound on the second error component in (3.32).

Proposition 3.1. *Consider a connected FL network, i.e., $\lambda_2 > 0$ (see (3.10)), and the solution (3.31) to GTVMin (3.22) for the local datasets (3.29). If the true local model parameters in (3.29) are identical (see (3.30)), we can upper bound the deviation $\tilde{\mathbf{w}}^{(i)} := \hat{\mathbf{w}}^{(i)} - (1/n) \sum_{i=1}^n \hat{\mathbf{w}}^{(i)}$ of learnt model parameters $\hat{\mathbf{w}}^{(i)}$ from their average, as*

$$\sum_{i=1}^n \|\tilde{\mathbf{w}}^{(i)}\|_2^2 \leq \frac{1}{\lambda_2 \alpha} \sum_{i=1}^n (1/m_i) \|\boldsymbol{\epsilon}^{(i)}\|_2^2. \quad (3.33)$$

Proof. See Section 3.8.1. □

Note that Prop. 3.1 only applies to GTVMin over a FL network with a connected graph \mathcal{G} . A necessary and sufficient condition for \mathcal{G} to be connected is that the second smallest eigenvalue is positive, $\lambda_2 > 0$. However, for an FL network with a graph \mathcal{G} that is not connected, we can still apply Prop. 3.1 separately to each connected component of \mathcal{G} .

The upper bound (3.33) involves three components:

- the properties of local datasets, via the noise terms $\boldsymbol{\epsilon}^{(i)}$ in (3.29),
- the FL network via the eigenvalue $\lambda_2(\mathbf{L}^{(\mathcal{G})})$ (see (3.10)),
- the GTVMin parameter α .

According to (3.33), we can ensure a small error component $\tilde{\mathbf{w}}^{(i)}$ of the GTVMin solution by choosing a large value α . Thus, by (3.32), for sufficiently large α , the local model parameters $\hat{\mathbf{w}}^{(i)}$ delivered by GTVMin are approximately identical for all nodes $i \in \mathcal{V}$ of a connected FL network (where $\lambda_2(\mathbf{L}^{(\mathcal{G})}) > 0$).

Enforcing identical local model parameters at all nodes is desirable for FL applications that require to learn a common (global) model for all nodes [14].

However, some FL applications involve heterogeneous nodes that carry local datasets with significantly different statistics. For such applications it is detrimental to enforce a common model at all nodes (see Chapter 6). Rather,

3.5 How to Handle Non-Parametric Models

In its basic form (3.20), GTVMin can only be applied to parametric local models with model parameters belonging to the same Euclidean space \mathbb{R}^d . Some FL application involve local models that are non-parametric (such as decision trees) or having different parametrizations (such as deep nets of different shape and size). For such settings we cannot use model parameters to measure the discrepancy between $h^{(i)}$ and $h^{(i')}$ across an edge $\{i, i'\}$.

One simple approach to measure the discrepancy between hypothesis maps $h^{(i)}, h^{(i')}$ is to compare their predictions on a dataset

$$\mathcal{D}^{\{i, i'\}} = \left\{ \mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m')} \right\}. \quad (3.34)$$

For each edge $\{i, i'\}$, the connected nodes need to agree on dataset $\mathcal{D}^{\{i, i'\}}$. This dataset is typically different for different edges and can be obtained by i.i.d. realizations of some probability distribution or via using subsets of $\mathcal{D}^{(i)}, \mathcal{D}^{(i')}$.

We compare the predictions delivered by $h^{(i)}$ and $h^{(i')}$ on $\mathcal{D}^{\{i, i'\}}$ using some loss function $L(\cdot, \cdot)$,

$$\begin{aligned} D(h^{(i)}, h^{(i')}) := (1/m') \sum_{\mathbf{x} \in \mathcal{D}^{\{i, i'\}}} (1/2) [& L\left(\left(\mathbf{x}, h^{(i)}(\mathbf{x})\right), h^{(i')}\right) \\ & + L\left(\left(\mathbf{x}, h^{(i')}(\mathbf{x})\right), h^{(i)}\right)]. \end{aligned} \quad (3.35)$$

Different choices for the loss function in (3.35) result in different computational and statistical properties of the resulting FL algorithms (see Sec. 5.6). For real-valued predictions we could use the squared error loss,

$$D(h^{(i)}, h^{(i')}) := (1/m') \sum_{\mathbf{x} \in \mathcal{D}^{i,i'}} \left(h^{(i)}(\mathbf{x}) - h^{(i')}(\mathbf{x}) \right)^2. \quad (3.36)$$

We can generalize GTVMin by replacing $\|\mathbf{w}^{(i)} - \mathbf{w}^{(i')}\|_2^2$ in (3.20) with the discrepancy $D(h^{(i)}, h^{(i')})$ given by (3.35) (or the special case (3.36)),

$$\{\hat{h}^{(i)}\}_{i=1}^n \in \operatorname{argmin}_{\substack{h^{(i)} \in \mathcal{H}^{(i)} \\ i \in \mathcal{V}}} \sum_{i \in \mathcal{V}} L_i(h^{(i)}) + \alpha \sum_{\{i,i'\} \in \mathcal{E}} A_{i,i'} D(h^{(i)}, h^{(i')}). \quad (3.37)$$

3.6 Interpretations

We next discuss some interpretations of GTVMin (3.19).

Generalized Convex Clustering. One important special case of GTVMin (3.19) is convex clustering [38, 39]. Indeed, convex clustering is obtained from (3.19) using the local loss function

$$L_i(\mathbf{w}^{(i)}) = \|\mathbf{w}^{(i)} - \mathbf{a}^{(i)}\|^2, \text{ for all nodes } i \in \mathcal{V} \quad (3.38)$$

and GTV penalty function $\phi(\mathbf{u}) = \|\mathbf{u}\|_p$ being a p -norm $\|\mathbf{u}\|_p := \left(\sum_{j=1}^d |u_j|^p \right)^{1/p}$ with some $p \geq 1$. The vectors $\mathbf{a}^{(i)}$ in (3.38) are the observations that we wish to cluster. Thus, we can interpret GTVMin as a generalization of convex clustering: we replace the terms $\|\mathbf{w}^{(i)} - \mathbf{a}^{(i)}\|^2$ with a more general local loss function.

Dual of Minimum-Cost Flow Problem. The optimization variables of GTVMin (3.19) are the local model parameters $\mathbf{w}^{(i)}$, for each node $i \in \mathcal{V}$ of an FL network \mathcal{G} . The optimization of variables $\mathbf{w}^{(i)}$ assigned to nodes

\mathcal{V} of a graph \mathcal{G} is naturally associated with a dual problem [40]. This dual problem optimizes variables $\mathbf{u}^{\{i,i'\}}$ assigned to each edge $\{i,i'\} \in \mathcal{E}$ of \mathcal{G} . It can be formulated as

$$\max_{\mathbf{u}^{(e)}, e \in \mathcal{E}} - \sum_{i \in \mathcal{V}} L_i^* (\mathbf{w}^{(i)}) - \alpha \sum_{e \in \mathcal{E}} A_e \phi^* (\mathbf{u}^{(e)} / (\alpha A_e)) \quad (3.39)$$

$$\text{subject to } -\mathbf{w}^{(i)} = \sum_{e \in \mathcal{E}} \sum_{i=e_+} \mathbf{u}^{(e)} - \sum_{i=e_-} \mathbf{u}^{(e)} \quad (3.40)$$

$$\text{for all nodes } i \in \mathcal{V}. \quad (3.41)$$

Here, we used the convention, for any edge $e = \{i,i'\}$, $e_- = \min\{i,i'\}$ and $\max\{i,i'\}$. Moreover, we used the convex conjugates $L_i^* (\cdot)$, ϕ^* of the local loss function $L_i (\cdot)$ and GTV penalty function ϕ . The convex conjugate of a function $f : \mathbb{R}^d \rightarrow \mathbb{R}$ is defined as [41]

$$f^*(\mathbf{x}) := \sup_{\mathbf{z} \in \mathbb{R}^d} \mathbf{x}^T \mathbf{z} - f(\mathbf{z}). \quad (3.42)$$

The dual problem (3.39) generalizes the optimal flow problem [42, Sec. 1J] to vector-valued flows. The special case of (3.39), obtained when the GTV penalty function ϕ is a norm, is equivalent to a generalized minimum-cost flow problem [43, Sec. 1.2.1]. Indeed, the maximization problem (3.39) is equivalent to the minimization

$$\begin{aligned} \min_{\mathbf{u} \in \mathcal{U}} \sum_{i \in \mathcal{V}} L_i^* (\mathbf{w}^{(i)}) \\ \text{subject to } -\mathbf{w}^{(i)} &= \sum_{e \in \mathcal{E}} \sum_{i=e_+} \mathbf{u}^{(e)} - \sum_{i=e_-} \mathbf{u}^{(e)} \text{ for all nodes } i \in \mathcal{V} \\ \|\mathbf{u}^{(e)}\|_* &\leq \alpha A_e \text{ for all edges } e \in \mathcal{E}. \end{aligned} \quad (3.43)$$

The optimization problem (3.43) reduces to the minimum-cost flow problem

[43, Eq. (1.3) - (1.5)] for local model parameters of length $d = 1$ (i.e., local models parametrized by a scalar).

Locally Weighted Learning. For many relevant settings, the solutions of GTVMin tend to be clustered: Each node i belongs to a cluster \mathcal{C} (subset of nodes approximately identical local model parameters. $\widehat{\mathbf{w}}^{(i)} \approx \overline{\mathbf{w}}^{(\mathcal{C})}$ [30]. The cluster-wise model parameters $\overline{\mathbf{w}}^{(\mathcal{C})}$ are the solutions of

$$\min_{\mathbf{w}} \sum_{i' \in \mathcal{C}} L_{i'}(\mathbf{w}), \quad (3.44)$$

which, in turn, is an instance of a locally weighted learning problem [44, Sec. 3.1.2]

$$\overline{\mathbf{w}}^{(\mathcal{C})} = \operatorname{argmin}_{\mathbf{w} \in \mathbb{R}^d} \sum_{i' \in \mathcal{V}} \rho_{i'} L_{i'}(\mathbf{w}). \quad (3.45)$$

Indeed, we obtain (3.44) from (3.45) by setting the weights $\rho_{i'}$ equal to 1 if $i' \in \mathcal{C}$ and 0 otherwise.

3.7 Problems

3.1. Spectral Radius of Laplacian Matrix. The spectral radius $\rho(\mathbf{Q})$ of a square matrix \mathbf{Q} is the largest magnitude of an eigenvalue,

$$\rho(\mathbf{Q}) := \max\{|\lambda| : \lambda \text{ is an eigenvalue of } \mathbf{Q}\}.$$

Consider the Laplacian matrix $\mathbf{L}^{(\mathcal{G})}$ of a FL network with undirected graph \mathcal{G} . Show that $\rho(\mathbf{L}^{(\mathcal{G})}) = \lambda_n(\mathbf{L}^{(\mathcal{G})})$ and verify the upper bound $\lambda_n(\mathbf{L}^{(\mathcal{G})}) \leq 2d_{\max}^{(\mathcal{G})}$. Try to find a graph \mathcal{G} such that $\lambda_n(\mathbf{L}^{(\mathcal{G})}) \approx 2d_{\max}^{(\mathcal{G})}$.

3.2. Kernel of the Laplacian matrix. Consider an undirected weighted graph \mathcal{G} with Laplacian matrix $\mathbf{L}^{(\mathcal{G})}$. A component of \mathcal{G} is a subset $\mathcal{C} \subseteq \mathcal{V}$ of




Fig. 3.5: A FL network whose nodes $i = 1, \dots, 5$ represent devices that hold local datasets whose size is indicated next to each node.

nodes that are connected but there is no edge between \mathcal{C} and the rest $\mathcal{V} \setminus \mathcal{C}$. The *nullspace* (or *kernel*) of $\mathbf{L}^{(\mathcal{G})}$ is the subspace $\mathcal{K} \subseteq \mathbb{R}^n$ constituted by all vectors $\mathbf{v} \in \mathbb{R}^n$ such that $\mathbf{L}^{(\mathcal{G})}\mathbf{v} = \mathbf{0}$. Show that the dimension of \mathcal{K} coincides with the number of components in \mathcal{G} .

3.3. Adding an Edge Increases Connectivity. Consider an undirected weighted graph \mathcal{G} with Laplacian matrix $\mathbf{L}^{(\mathcal{G})}$. We construct a new graph \mathcal{G}' , with Laplacian matrix $\mathbf{L}^{(\mathcal{G})'}$, by adding a new edge to \mathcal{G} . Show that $\lambda_2(\mathcal{G}') \geq \lambda_2(\mathcal{G})$, i.e., the second smallest eigenvalue of $\mathbf{L}^{(\mathcal{G})'}$ is at least as large as the second smallest eigenvalue of $\mathbf{L}^{(\mathcal{G})}$.

3.4. Capacity of FL Network. Consider the FL network depicted in Fig. 3.5. Each node carries a local dataset whose size is indicated by the numbers next it. The devices represented by the nodes can communicate via bi-directional communication links with a capacity indicated by the numbers placed next to the edges. How long does it take at least for the left-most node to collect the local datasets from all other nodes.

3.5. Discrepancy Measures. Consider a FL network with nodes carrying parametric local models, each having model parameters $\mathbf{w}^{(i)} \in \mathbb{R}^d$. Is it possible to construct a dataset $\mathcal{D}^{\{i,i'\}}$ such that (3.36) coincides with $\mathbf{w}^{(i)} - \mathbf{w}^{(i')}$.

3.6. Computing the Average. Consider a FL network with nodes carrying parametric local models, each having a single parameter $w^{(i)}$. Each node also carries a local dataset that consists of a single number $y^{(i)}$. Construct an instance of GTVMin such that its solutions are given by $\hat{w}^{(i)} \approx (1/n) \sum_{i=1}^n y^{(i)}$ for all $i = 1, \dots, n$.

3.7. Computing the Average over Star. Consider the FL network depicted in Fig. 3.6, which consists of a centre node i_0 which is connected to $n - 1$ peripheral nodes $\mathcal{P} := \mathcal{V} \setminus \{i_0\}$. Each peripheral node $i \in \mathcal{P}$ carries a local dataset that consists of a single real-valued observation $y^{(i)} \in \mathbb{R}$. Construct an instance of GTVMin, using real-valued local model parameters $w^{(i)} \in \mathbb{R}$, such that the solution satisfies $\hat{w}^{(i_0)} \approx (1/(n - 1)) \sum_{i \in \mathcal{P}} y^{(i)}$.




Fig. 3.6: A FL network that consists of a centre node i_0 that is connected to several peripheral nodes $\mathcal{P} := \mathcal{V} \setminus \{i_0\}$.

3.8. Fundamental Limits. Consider a FL network with nodes carrying

parametric local models, each having a single parameter $w^{(i)}$. Each node also carries a local dataset that consists of a single number $y^{(i)}$. We use a probabilistic model for the local datasets: $y^{(i)} = \bar{w} + n^{(i)}$. Here, \bar{w} is some fixed but unknown number and $n^{(i)} \sim \mathcal{N}(0, 1)$ are i.i.d. Gaussian RVs. We use a message-passing FL algorithm to estimate c based on the local datasets. What is a fundamental limit on the accuracy of the estimate $\hat{c}^{(i)}$ delivered at some fixed node i by such an algorithm after two iterations? Compare this limit with the risk $\mathbb{E}\{(\hat{w}^{(i)} - \bar{w})^2\}$ incurred by the estimate $\hat{w}^{(i)}$ delivered by running Algorithm 5.1 for two iterations.

3.9. Counting Number of Paths. Consider an undirected graph \mathcal{G} with each edge $\{i, i'\} \in \mathcal{E}$ having unit edge-weight $A_{i,i'} = 1$. For $k \in \{1, 2, \dots\}$, a k -hop path between two nodes $i, i' \in \mathcal{V}$ is a node sequence $i^{(1)}, \dots, i^{(k+1)}$ such that $i^{(1)} = i$, $i^{(k+1)} = i'$, and $\{i^{(r)}, i^{(r+1)}\} \in \mathcal{E}$ for each $r = 1, \dots, k$. Show that the number of k -hop paths between two nodes $i, i' \in \mathcal{V}$ is given by $(\mathbf{A}^k)_{i,i'}$.




Fig. 3.7: Some FL network with each node i having degree $d^{(i)} = 3$. We use a FL algorithm to learn a local model parameters $w^{(i)}$. If the algorithm uses message passing, executing the first iteration only provides access to the local datasets of the neighbours in $\mathcal{N}^{(i)}$ (located along the inner dashed circle). The second iteration provides access also to the local datasets of the neighbours $\mathcal{N}^{(i')}$ of each $i' \in \mathcal{N}^{(i)}$. These *second-hop* neighbours are located along the outer dashed circle.

3.8 Proofs

3.8.1 Proof of Proposition 3.1

Let us introduce the shorthand $f(\mathbf{w}^{(i)})$ for the objective function of the GTVMin instance (3.22). We verify the bound (3.33) by showing that if it does not hold, the choice for the local model parameters $\mathbf{w}^{(i)} := \bar{\mathbf{w}}^{(i)}$ (see (3.29)) results in a smaller objective function value, $f(\bar{\mathbf{w}}^{(i)}) < f(\hat{\mathbf{w}}^{(i)})$. This would contradict the fact that $\hat{\mathbf{w}}^{(i)}$ is a solution to (3.22).

First, note that

$$\begin{aligned}
f(\bar{\mathbf{w}}^{(i)}) &= \sum_{i \in \mathcal{V}} (1/m_i) \|\mathbf{y}^{(i)} - \mathbf{X}^{(i)} \bar{\mathbf{w}}^{(i)}\|_2^2 + \alpha \sum_{\{i, i'\} \in \mathcal{E}} A_{i, i'} \|\bar{\mathbf{w}}^{(i)} - \bar{\mathbf{w}}^{(i')}\|_2^2 \\
&\stackrel{(3.30)}{=} \sum_{i \in \mathcal{V}} (1/m_i) \|\mathbf{y}^{(i)} - \mathbf{X}^{(i)} \bar{\mathbf{w}}^{(i)}\|_2^2 \\
&\stackrel{(3.29)}{=} \sum_{i \in \mathcal{V}} (1/m_i) \|\mathbf{X}^{(i)} \bar{\mathbf{w}}^{(i)} + \boldsymbol{\varepsilon}^{(i)} - \mathbf{X}^{(i)} \bar{\mathbf{w}}^{(i)}\|_2^2 \\
&= \sum_{i \in \mathcal{V}} (1/m_i) \|\boldsymbol{\varepsilon}^{(i)}\|_2^2.
\end{aligned} \tag{3.46}$$

Inserting (3.31) into (3.22),

$$\begin{aligned}
f(\hat{\mathbf{w}}^{(i)}) &= \underbrace{\sum_{i \in \mathcal{V}} (1/m_i) \|\mathbf{y}^{(i)} - \mathbf{X}^{(i)} \hat{\mathbf{w}}^{(i)}\|_2^2}_{\geq 0} + \alpha \sum_{\{i, i'\} \in \mathcal{E}} A_{i, i'} \underbrace{\|\hat{\mathbf{w}}^{(i)} - \hat{\mathbf{w}}^{(i')}\|_2^2}_{\stackrel{(3.31)}{=} \|\tilde{\mathbf{w}}^{(i)} - \tilde{\mathbf{w}}^{(i')}\|_2^2} \\
&\geq \alpha \sum_{\{i, i'\} \in \mathcal{E}} A_{i, i'} \|\tilde{\mathbf{w}}^{(i)} - \tilde{\mathbf{w}}^{(i')}\|_2^2 \\
&\stackrel{(3.15)}{\geq} \alpha \lambda_2 \sum_{i=1}^n \|\tilde{\mathbf{w}}^{(i)}\|_2^2.
\end{aligned} \tag{3.47}$$

If the bound (3.33) would not hold, then by (3.47) and (3.46) we would obtain $f(\hat{\mathbf{w}}^{(i)}) > f(\bar{\mathbf{w}}^{(i)})$. This is a contradiction to the fact that $\hat{\mathbf{w}}^{(i)}$ solves (3.22).

4 Gradient Methods

Chapter 3 introduced GTVMin as a central design principle for FL methods. Many significant instances of GTVMin minimize a smooth objective function $f(\mathbf{w})$ over the parameter space (typically a subset of \mathbb{R}^d). This chapter explores gradient-based methods, a widely used family of iterative algorithms to find the minimum of a smooth function. These methods approximate the objective function locally using its gradient at the current choice for the model parameters. Chapter 5 focuses on FL algorithms obtained from applying gradient-based methods to solve GTVMin.

4.1 Learning Goals

After completing this chapter, you will

- have some intuition about the effect of a gradient step,
- understand the role of the step size (or learning rate),
- know some examples of a stopping criterion,
- be able to analyze the effect of perturbations in the gradient step,
- know about projected gradient descent (GD) to cope with constraints on model parameters.

4.2 Gradient Descent

Gradient-based methods are iterative algorithms for finding the minimum of a differentiable objective function $f(\mathbf{w})$ of a vector-valued argument \mathbf{w} (e.g., the model parameters in a ML method). Unless stated otherwise, we consider an objective function of form

$$f(\mathbf{w}) = \mathbf{w}^T \mathbf{Q} \mathbf{w} + \mathbf{q}^T \mathbf{w}. \quad (4.1)$$

Note that (4.1) defines an entire family of convex quadratic functions $f(\mathbf{w})$. Each member of this family is specified by a psd matrix $\mathbf{Q} \in \mathbb{R}^{d \times d}$ and a vector $\mathbf{q} \in \mathbb{R}^d$.

We have already encountered some ML and FL methods that minimize an objective function of the form (4.1): Linear regression (2.3) and ridge regression (2.27) in Chapter 2 as well as GTVMin (3.22) for local linear models in Chapter 3. Moreover, (4.1) might serve as a useful approximation for the objective functions arising from larger classes of ML methods [45–47].

Given model parameters $\mathbf{w}^{(k)}$, we want to update them towards a minimum of (4.1). To this end, we use the gradient $\nabla f(\mathbf{w}^{(k)})$ to locally approximate $f(\mathbf{w})$ (see Figure 4.1). The gradient $\nabla f(\mathbf{w}^{(k)})$ indicates the direction in which the function $f(\mathbf{w})$ maximally increases. Therefore, it seems reasonable to update $\mathbf{w}^{(k)}$ in the opposite direction of $\nabla f(\mathbf{w}^{(k)})$,

$$\begin{aligned} \mathbf{w}^{(k+1)} &:= \mathbf{w}^{(k)} - \eta \nabla f(\mathbf{w}^{(k)}) \\ &\stackrel{(4.1)}{=} \mathbf{w}^{(k)} - \eta (2\mathbf{Q}\mathbf{w}^{(k)} + \mathbf{q}). \end{aligned} \quad (4.2)$$

The gradient step (4.2) involves the factor η which we refer to as step-size or learning rate. Algorithm 4.1 summarizes the most basic instance of gradient-

based methods which simply repeats (iterates) (4.2) until some stopping criterion is met.




Fig. 4.1: We can approximate a differentiable function $f(\mathbf{w})$ locally around a point $\mathbf{w}^{(k)} \in \mathbb{R}^d$ using the linear function $f(\mathbf{w}^{(k)}) + (\mathbf{w} - \mathbf{w}^{(k)})^T \nabla f(\mathbf{w}^{(k)})$. Geometrically, we approximate the graph of $f(\mathbf{w})$ by a hyperplane with normal vector $\mathbf{n} = (\nabla f(\mathbf{w}^{(k)}), -1)^T \in \mathbb{R}^{d+1}$ of this approximating hyperplane is determined by the gradient $\nabla f(\mathbf{w}^{(k)})$ [2].

The usefulness of gradient-based methods crucially depends on the computational complexity of evaluating the gradient $\nabla f(\mathbf{w})$. Modern software libraries for automatic differentiation enable the efficient evaluation of the gradients arising in widely-used ERM-based methods [48]. Besides the actual computation of the gradient, it might be already challenging to gather the required data points which define the objective function $f(\mathbf{w})$ (e.g., being the average loss over a large training set). Indeed, the matrix \mathbf{Q} and vector \mathbf{q} in (4.1) are constructed from the features and labels of data points in the training set. For example, the gradient of the objective function in ridge

regression (2.27) is

$$\nabla f(\mathbf{w}) = -(2/m) \sum_{r=1}^m \mathbf{x}^{(r)} (y^{(r)} - \mathbf{w}^T \mathbf{x}^{(r)}) + 2\alpha \mathbf{w}.$$

Evaluating this gradient requires roughly $d \times m$ arithmetic operations (sums and multiplications).

Algorithm 4.1 A blueprint for gradient-based methods

Input: some objective function $f(\mathbf{w})$ (e.g., the average loss of a hypothesis $h(\mathbf{w})$ on a training set); learning rate $\eta > 0$; some stopping criterion;

Initialize: set $\mathbf{w}^{(0)} := \mathbf{0}$; set iteration counter $k := 0$

1: **repeat**

2: $k := k + 1$ (increase iteration counter)

3: $\mathbf{w}^{(k)} := \mathbf{w}^{(k-1)} - \eta \nabla f(\mathbf{w}^{(k-1)})$ (do a gradient step (4.2))

4: **until** stopping criterion is met

Output: learnt model parameters $\hat{\mathbf{w}} := \mathbf{w}^{(k)}$ (hopefully $f(\hat{\mathbf{w}}) \approx \min_{\mathbf{w}} f(\mathbf{w})$)

Note that Algorithm 4.1, like most other gradient-based methods, involves two hyper-parameters: (i) the learning rate η used for the gradient step and (ii) a stopping criterion to decide when to stop repeating the gradient step. We next discuss how to choose these hyper-parameters.

4.3 Learning Rate

The learning rate must not be too large to avoid moving away from the optimum by *overshooting* (see Figure 4.2-(a)). On the other hand, if the learning rate is chosen too small, the gradient step makes too little progress towards the solutions of (4.1) (see Figure 4.2-(b)). Note that in practice we can only afford to repeat the gradient step for a finite number of iterations.




Fig. 4.2: Effect of inadequate learning rates η in the gradient step (4.2). (a) If η is too large, the gradient steps might “overshoot” such that the iterates $\mathbf{w}^{(k)}$ might diverge from the optimum, i.e., $f(\mathbf{w}^{(k+1)}) > f(\mathbf{w}^{(k)})$! (b) If η is too small, the gradient steps make very little progress towards the optimum or even fail to reach the optimum at all.

One approach to choose the learning rate is to start with some initial value (first guess) and monitor the decrease of the objective function. If this decrease does not agree with the decrease predicted by the (local linear approximation using the) gradient, we decrease the learning rate by a constant factor. After we decrease the learning rate, we re-consider the decrease of the objective function. We repeat this procedure until a sufficient decrease of the objective function is achieved [49, Sec 6.1].

Alternatively, we can use a prescribed sequence (schedule) η_k , for $k = 1, 2, \dots$, of learning rates that vary across successive gradient steps [50]. We can guarantee convergence of (4.2) (under mild assumptions on the objective function $f(\mathbf{w})$) by using a learning rate schedule that satisfies [49, Sec. 6.1]

$$\lim_{k \rightarrow \infty} \eta_k = 0, \quad \sum_{k=1}^{\infty} \eta_k = \infty, \quad \text{and} \quad \sum_{k=1}^{\infty} \eta_k^2 < \infty. \quad (4.3)$$

The first condition (4.3) requires that the learning rate eventually become

sufficiently small to avoid overshooting. The third condition (4.3) ensures that this required decay of the learning rate does not take “forever”. Note that the first and third condition in (4.3) could be satisfied by the trivial learning rate schedule $\eta_k = 0$ which is clearly not useful as the gradient step has no effect.

The trivial schedule $\eta_k = 0$ is ruled out by the middle condition of (4.3). This middle condition ensures that the learning rate η_k is large enough such that the gradient steps make sufficient progress towards a minimizer of the objective function.

We emphasize that the conditions (4.3) do not involve any properties of the matrix \mathbf{Q} in (4.1). Note that the matrix \mathbf{Q} is determined by data points (see, e.g., (2.3)) whose statistical properties typically can be controlled only to a limited extend (e.g., via data normalization).

4.4 When to Stop?

For the stopping criterion we might use a fixed number k_{\max} of iterations. This hyper-parameter k_{\max} might be dictated by limited resources (such as computational time) for implementing GD or tuned via validation techniques.

We can obtain another stopping criterion from monitoring the decrease in the objective function $f(\mathbf{w}^{(k)})$: we stop repeating the gradient step (4.2) whenever $|f(\mathbf{w}^{(k)}) - f(\mathbf{w}^{(k+1)})| \leq \varepsilon^{(\text{tol})}$ for a given tolerance $\varepsilon^{(\text{tol})}$. The tolerance level $\varepsilon^{(\text{tol})}$ is a hyper-parameter of the resulting ML method which could be chosen via validation techniques (see Chapter 2).

The above technique for deciding when to stop the gradient steps are convenient as they do not require any information about the psd matrix \mathbf{Q}

in (4.1). However, they might result in sub-optimal use of computational resources by implementing “useless” additional gradient steps. We use information about the psd matrix \mathbf{Q} in (4.1) to avoid unnecessary computations.¹⁰ Indeed, the choice for the learning rate η and the stopping criterion can be guided by the eigenvalues

$$0 \leq \lambda_1(\mathbf{Q}) \leq \dots \leq \lambda_d(\mathbf{Q}). \quad (4.4)$$

Even if we do not know these eigenvalues precisely, we might know (or be able to ensure via feature learning) some upper and lower bounds,

$$0 \leq L \leq \lambda_1(\mathbf{Q}) \leq \dots \leq \lambda_d(\mathbf{Q}) \leq U. \quad (4.5)$$

In what follows, we assume that \mathbf{Q} is invertible and we know some positive lower bound $L > 0$ (see (4.5)).¹¹ The objective function (4.1) has then a unique solution $\widehat{\mathbf{w}}$. It turns out that a gradient step (4.2) reduces the distance $\|\mathbf{w}^{(k)} - \widehat{\mathbf{w}}\|_2$ to $\widehat{\mathbf{w}}$ by a constant factor [49, Ch. 6],

$$\|\mathbf{w}^{(k+1)} - \widehat{\mathbf{w}}\|_2 \leq \kappa^{(\eta_k)}(\mathbf{Q}) \|\mathbf{w}^{(k)} - \widehat{\mathbf{w}}\|_2. \quad (4.6)$$

Here, we used the contraction factor

$$\kappa^{(\eta)}(\mathbf{Q}) := \max\{|1 - \eta 2\lambda_1|, |1 - \eta 2\lambda_d|\}. \quad (4.7)$$

¹⁰For linear regression (2.7), the matrix \mathbf{Q} is determined by the features of the data points in the training set. We can influence the properties of \mathbf{Q} to some extent by feature transformation methods. One important example for such a transformation is the normalization of features.

¹¹What are sufficient conditions for the local datasets and the edge weights used in GTVMin such that \mathbf{Q} in (3.26) is invertible?

The contraction factor depends on the learning rate η which is a hyperparameter of gradient-based methods that we can control. However, the contraction factor also depends on the eigenvalues of the matrix \mathbf{Q} in (4.1). In ML and FL applications, this matrix typically depends on data and can be controlled only to some extent, e.g., using feature transformation [6, Ch. 5]. We can ensure $\kappa^{(\eta)}(\mathbf{Q}) < 1$ if we use a positive learning rate $\eta_k < 1/U$.

For a contraction factor $\kappa^{(\eta)}(\mathbf{Q}) < 1$, a sufficient condition on the number $k^{(\varepsilon^{(\text{tol})})}$ of gradient steps required to ensure an optimization error $\|\mathbf{w}^{(k+1)} - \widehat{\mathbf{w}}\|_2 \leq \varepsilon^{(\text{tol})}$ is (see (4.6))

$$\text{modelparameters}^{(\varepsilon^{(\text{tol})})} \geq \frac{\log(\|\mathbf{w}^{(0)} - \widehat{\mathbf{w}}\|_2 / \varepsilon^{(\text{tol})})}{\log(1/\kappa^{(\eta)}(\mathbf{Q}))}. \quad (4.8)$$

According to (4.6), smaller values of the contraction factor $\kappa^{(\eta)}(\mathbf{Q})$ guarantee a faster convergence of (4.2) towards the solution of (4.1). Figure 4.3 illustrates the dependence of $\kappa^{(\eta)}(\mathbf{Q})$ on the learning rate η . Thus, choosing a small η (close to 0) will typically result in a larger $\kappa^{(\eta)}(\mathbf{Q})$ and, in turn, require more iterations to ensure optimization error level $\varepsilon^{(\text{tol})}$ via (4.6).

We can minimize this contraction factor by choosing the learning rate (see Figure 4.3)

$$\eta^{(*)} := \frac{1}{\lambda_1 + \lambda_d}. \quad (4.9)$$

[Note that evaluating (4.9) requires to know the extremal eigenvalues λ_1, λ_d of \mathbf{Q} .] Inserting the optimal learning rate (4.9) into (4.6),

$$\|\mathbf{w}^{(k+1)} - \widehat{\mathbf{w}}\|_2 \leq \underbrace{\frac{(\lambda_d/\lambda_1) - 1}{(\lambda_d/\lambda_1) + 1}}_{:=\kappa^{(*)}(\mathbf{Q})} \|\mathbf{w}^{(k)} - \widehat{\mathbf{w}}\|_2. \quad (4.10)$$

Carefully note that the formula (4.10) is valid only if the matrix \mathbf{Q} in (4.1) is invertible, i.e., if $\lambda_1 > 0$. If the matrix \mathbf{Q} is singular ($\lambda_1 = 0$), the




Fig. 4.3: The contraction factor $\kappa^{(\eta)}(\mathbf{Q})$ (4.7), used in the upper bound (4.6), as a function of the learning rate η . Note that $\kappa^{(\eta)}(\mathbf{Q})$ also depends on the eigenvalues of the matrix \mathbf{Q} in (4.1).

convergence of (4.2) towards a solution of (4.1) is much slower than the decrease of the bound (4.10). However, we can still ensure the convergence of gradient steps $\mathbf{w}^{(k)}$ by using a fixed learning rate $\eta_k = \eta$ that satisfies [51, Thm. 2.1.14]

$$0 < \eta < 1/\lambda_d(\mathbf{Q}). \quad (4.11)$$

It is interesting to note that for linear regression, the matrix \mathbf{Q} depends only on the features $\mathbf{x}^{(r)}$ of the data points in the training set (see (2.17)) but not on their labels $y^{(r)}$. Thus, the convergence of gradient steps is only affected by the features, whereas the labels are irrelevant. The same is true for ridge

regression and GTVMin (using local linear models).

Note that both, the optimal learning rate (4.9) and the optimal contraction factor

$$\kappa^*(\mathbf{Q}) := \frac{(\lambda_d/\lambda_1) - 1}{(\lambda_d/\lambda_1) + 1} \quad (4.12)$$

depend on the eigenvalues of the matrix \mathbf{Q} in (4.1). According to (4.10), the ideal case is when all eigenvalues are identical which leads, in turn, to a contraction factor $\kappa^*(\mathbf{Q}) = 0$. Here, a single gradient step arrives at the unique solution of (4.1).

In general, we do not have full control over the matrix \mathbf{Q} and its eigenvalues. For example, the matrix \mathbf{Q} arising in linear regression (2.7) is determined by the features of data points in the training set. These features might be obtained from sensing devices and therefore beyond our control. However, other applications might allow for some design freedom in the choice of feature vectors. We might also use feature transformations that nudge the resulting \mathbf{Q} in (2.7) more towards a scaled identity matrix.

4.5 Perturbed Gradient Step

Consider the gradient step (4.2) used to find a minimum of (4.1). We again assume that the matrix \mathbf{Q} in (4.1) is invertible ($\lambda_1(\mathbf{Q}) > 0$) and, in turn, (4.1) has a unique solution $\hat{\mathbf{w}}$.

In some applications, it might be difficult to evaluate the gradient $\nabla f(\mathbf{w}) = 2\mathbf{Q}\mathbf{w} + \mathbf{q}$ of (4.1) exactly. For example, this evaluation might require to gather data points from distributed storage locations. These storage locations might become unavailable during the computation of $\nabla f(\mathbf{w})$ due to software or hardware failures (e.g., limited connectivity).

We can model imperfections during the computation of (4.2) as the perturbed gradient step

$$\begin{aligned}\mathbf{w}^{(k+1)} &:= \mathbf{w}^{(k)} - \eta \nabla f(\mathbf{w}^{(k)}) + \boldsymbol{\varepsilon}^{(k)} \\ &\stackrel{(4.1)}{=} \mathbf{w}^{(k)} - \eta(2\mathbf{Q}\mathbf{w}^{(k)} + \mathbf{q}) + \boldsymbol{\varepsilon}^{(k)}.\end{aligned}\tag{4.13}$$

We can use the contraction factor (4.7) to upper bound the deviation between $\mathbf{w}^{(k)}$ and the optimum $\widehat{\mathbf{w}}$ as (see (4.6))

$$\|\mathbf{w}^{(k)} - \widehat{\mathbf{w}}\|_2 \leq (\kappa^{(\eta)}(\mathbf{Q}))^{(k)} \|\mathbf{w}^{(0)} - \widehat{\mathbf{w}}\|_2 + \sum_{k'=1}^{(k)} (\kappa^{(\eta)}(\mathbf{Q}))^{(k')} \|\boldsymbol{\varepsilon}^{(k-k')}\|_2.\tag{4.14}$$

This bound applies for any number of iterations $k = 1, 2, \dots$ of the perturbed gradient step (4.13).

Finally, note that the perturbed gradient step (4.13) could also be used as a tool to analyze the (exact) gradient step for an objective function $\tilde{f}(\mathbf{w})$ which does not belong to the family (4.1) of convex quadratic functions. Indeed, we can write the gradient step for minimizing $\tilde{f}(\mathbf{w})$ as

$$\begin{aligned}\mathbf{w}^{(k+1)} &:= \mathbf{w}^{(k)} - \eta \nabla \tilde{f}(\mathbf{w}) \\ &= \mathbf{w}^{(k)} - \eta \nabla f(\mathbf{w}) + \underbrace{\eta(\nabla f(\mathbf{w}) - \nabla \tilde{f}(\mathbf{w}))}_{:= \boldsymbol{\varepsilon}^{(k)}}.\end{aligned}$$

The last identity is valid for any choice of surrogate function $f(\mathbf{w})$. In particular, we can choose $f(\mathbf{w})$ as a convex quadratic function (4.1) that approximates $\tilde{f}(\mathbf{w})$. Note that the perturbation term $\boldsymbol{\varepsilon}^{(k)}$ is scaled by the learning rate η .

4.6 Handling Constraints - Projected Gradient Descent

Many important ML and FL methods amount to the minimization of an objective function of the form (4.1). The optimization variable \mathbf{w} in (4.1) represents some model parameters.

Sometimes we might require the parameters \mathbf{w} to belong to a subset $\mathcal{C} \subset \mathbb{R}^d$. One example is regularization via model pruning (see Chapter 2). Another example are FL methods that learn identical local model parameters $\mathbf{w}^{(i)}$ at all nodes $i \in \mathcal{V}$ of an FL network. This can be implemented by requiring the stacked local model parameters $\mathbf{w} = (\mathbf{w}^{(1)}, \dots, \mathbf{w}^{(n)})^T$ to belong to the subset

$$\mathcal{C} = \left\{ (\mathbf{w}^{(1)}, \dots, \mathbf{w}^{(n)})^T : \mathbf{w}^{(1)} = \dots = \mathbf{w}^{(n)} \right\}.$$

Let us now show how to adapt the gradient step (4.2) to solve the constrained problem

$$f^* = \min_{\mathbf{w} \in \mathcal{C}} \mathbf{w}^T \mathbf{Q} \mathbf{w} + \mathbf{q}^T \mathbf{w}. \quad (4.15)$$

We assume that the constraint set $\mathcal{C} \subseteq \mathbb{R}^d$ is such that we can efficiently compute the projection

$$P_{\mathcal{C}}(\mathbf{w}) = \operatorname{argmin}_{\mathbf{w}' \in \mathcal{C}} \|\mathbf{w} - \mathbf{w}'\|_2 \text{ for any } \mathbf{w} \in \mathbb{R}^d. \quad (4.16)$$

A suitable modification of the gradient step (4.2) to solve the constrained variant (4.15) is [49]

$$\begin{aligned} \mathbf{w}^{(k+1)} &:= P_{\mathcal{C}}(\mathbf{w}^{(k)} - \eta \nabla f(\mathbf{w}^{(k)})) \\ &\stackrel{(4.1)}{=} P_{\mathcal{C}}(\mathbf{w}^{(k)} - \eta(2\mathbf{Q}\mathbf{w}^{(k)} + \mathbf{q})). \end{aligned} \quad (4.17)$$

The projected GD step (4.17) amounts to

1. compute ordinary gradient step $\mathbf{w}^{(k)} \mapsto \mathbf{w}^{(k)} - \eta \nabla f(\mathbf{w}^{(k)})$ and
2. project back to the constraint set \mathcal{C} .

Note that we re-obtain the basic gradient step (4.2) from the projected gradient step (4.17) for the specific constraint set $\mathcal{C} = \mathbb{R}^d$.




Fig. 4.4: Projected GD augments a basic gradient step with a projection back onto the constraint set \mathcal{C} .

The approaches for choosing the learning rate η and stopping criterion for basic gradient step (4.2) explained in Sections 4.3 and 4.4 work also for the projected gradient step (4.17). In particular, the convergence speed of the projected gradient step is also characterized by (4.6) [49, Ch. 6]. This follows from the fact that the concatenation of a contraction (such as the gradient step (4.2) for sufficiently small η) and a projection (such as $P_{\mathcal{C}}(\cdot)$) results again in a contraction with the same contraction factor.

Thus, the convergence speed of projected GD, in terms of number of iterations required to ensure a given level of optimization error, is essentially the same as that of basic GD. However, the bound (4.6) is only telling about the number of projected gradient steps (4.17) required to achieve a guaranteed level of sub-optimality $|f(\mathbf{w}^{(k)}) - f^*|$. The iteration (4.17) of projected GD

might require significantly more computation than the basic gradient step, as it requires to compute the projection (4.16).

4.7 Generalizing the Gradient Step

The gradient-based methods discussed so far can be used to learn a hypothesis from a parametric model. Let us now sketch one possible generalization of the gradient step (4.2) for a model \mathcal{H} without a parametrization.

We start with rewriting the gradient step (4.2) as the optimization

$$\mathbf{w}^{(k+1)} = \underset{\mathbf{w} \in \mathbb{R}^d}{\operatorname{argmin}} \left[(1/(2\eta)) \|\mathbf{w} - \mathbf{w}^{(k)}\|_2^2 + \underbrace{f(\mathbf{w}^{(k)}) + (\mathbf{w} - \mathbf{w}^{(k)})^T \nabla f(\mathbf{w}^{(k)})}_{\approx f(\mathbf{w})} \right]. \quad (4.18)$$

The objective function in (4.18) includes the first-order approximation

$$f(\mathbf{w}) \approx f(\mathbf{w}^{(k)}) + (\mathbf{w} - \mathbf{w}^{(k)})^T \nabla f(\mathbf{w}^{(k)})$$

of the function $f(\mathbf{w})$ around the location $\mathbf{w} = \mathbf{w}^{(k)}$ (see Figure 4.1).

Let us modify (4.18) by using $f(\mathbf{w})$ itself (instead of an approximation),

$$\mathbf{w}^{(k+1)} = \underset{\mathbf{w} \in \mathbb{R}^d}{\operatorname{argmin}} (1/(2\eta)) \|\mathbf{w} - \mathbf{w}^{(k)}\|_2^2 + f(\mathbf{w}). \quad (4.19)$$

Like the gradient step, also (4.19) maps a given vector $\mathbf{w}^{(k)}$ to an updated vector $\mathbf{w}^{(k+1)}$. Note that (4.19) is nothing but the proximal operator of the function $f(\mathbf{w})$ [35]. Similar to the role of the gradient step as the main building block of gradient-based methods, the proximal operator (4.19) is the main building block of proximal algorithms [35].

To obtain a version of (4.19) for a non-parametric model, we need to be able to evaluate its objective function directly in terms of a hypothesis h

instead of its parameters \mathbf{w} . The objective function (4.19) consists of two components. The second component $f(\cdot)$, which is the function we want to minimize, is obtained from a training error incurred by a hypothesis, which might be parametric $h^{(\mathbf{w})}$. Thus, we can evaluate the function $f(h)$ by computing the training error for a given hypothesis.

The first component of the objective function in (4.19) uses $\|\mathbf{w} - \mathbf{w}^{(k)}\|_2^2$ to measure the difference between the hypothesis maps $h^{(\mathbf{w})}$ and $h^{(\mathbf{w}^{(k)})}$. Another measure for the difference between two hypothesis maps can be obtained by using some test dataset $\mathcal{D}' = \{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m')}\}$: The average squared difference between their predictions,

$$(1/m') \sum_{r=1}^{m'} \left(h(\mathbf{x}^{(r)}) - h^{(k)}(\mathbf{x}^{(r)}) \right)^2, \quad (4.20)$$

is a measure for the difference between h and $h^{(k)}$. Note that (4.20) only requires the predictions delivered by the hypothesis maps $h, h^{(k)}$ on \mathcal{D}' - no other information is needed about these maps.

It is interesting to note that (4.20) coincides with $\|\mathbf{w} - \mathbf{w}^{(k)}\|_2^2$ for the linear model $h^{(\mathbf{w})}(\mathbf{x}) := \mathbf{w}^T \mathbf{x}$ and a specific construction of the dataset \mathcal{D}' . This construction uses the realizations $\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots$ of i.i.d. RVs with a common probability distribution $\mathbf{x} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$. Indeed, by the law of large

numbers

$$\begin{aligned}
& \lim_{m' \rightarrow \infty} (1/m') \sum_{r=1}^{m'} \left(h^{(\mathbf{w})}(\mathbf{x}^{(r)}) - h^{(\mathbf{w}^{(k)})}(\mathbf{x}^{(r)}) \right)^2 \\
&= \lim_{m' \rightarrow \infty} (1/m') \sum_{r=1}^{m'} \left((\mathbf{w} - \mathbf{w}^{(k)})^T \mathbf{x}^{(r)} \right)^2 \\
&= \lim_{m' \rightarrow \infty} (1/m') \sum_{r=1}^{m'} (\mathbf{w} - \mathbf{w}^{(k)})^T \mathbf{x}^{(r)} (\mathbf{x}^{(r)})^T (\mathbf{w} - \mathbf{w}^{(k)}) \\
&= (\mathbf{w} - \mathbf{w}^{(k)})^T \underbrace{\left[\lim_{m' \rightarrow \infty} (1/m') \sum_{r=1}^{m'} \mathbf{x}^{(r)} (\mathbf{x}^{(r)})^T \right]}_{=\mathbf{I}} (\mathbf{w} - \mathbf{w}^{(k)}) \\
&= \|\mathbf{w} - \mathbf{w}^{(k)}\|_2^2. \tag{4.21}
\end{aligned}$$

Finally, we arrive at a generalized gradient step for the training of a non-parametric model \mathcal{H} by replacing $\|\mathbf{w} - \mathbf{w}^{(k)}\|_2^2$ in (4.19) with (4.20). In other words,

$$h^{(k+1)} = \operatorname{argmin}_{h \in \mathcal{H}} \left[(1/(2\eta m')) \sum_{r=1}^{m'} \left(h(\mathbf{x}^{(r)}) - h^{(k)}(\mathbf{x}^{(r)}) \right)^2 + f(h) \right]. \tag{4.22}$$

We can turn gradient-based methods for the training of parametric models into corresponding training methods for non-parametric models by replacing the gradient step with the update (4.22). For example, we obtain Algorithm 4.2 from Algorithm 4.1 by modifying step 3 suitably.

Algorithm 4.2 A blueprint for generalized gradient-based methods

Input: some objective function $f : \mathcal{H} \rightarrow \mathbb{R}$ (e.g., the average loss of a hypothesis $h \in \mathcal{H}$ on a training set); learning rate $\eta > 0$; some stopping criterion; test dataset $\mathcal{D}' = \{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m')}\}$

Initialize: set $h^{(0)} := \mathbf{0}$; set iteration counter $k := 0$

1: **repeat**

2: $k := k + 1$ (increase iteration counter)

3: do a generalized gradient step (4.22),

$$h^{(k)} = \operatorname{argmin}_{h \in \mathcal{H}} \left[\left(1/(2\eta m') \right) \sum_{r=1}^{m'} \left(h(\mathbf{x}^{(r)}) - h^{(k-1)}(\mathbf{x}^{(r)}) \right)^2 + f(h) \right]$$

4: **until** stopping criterion is met

Output: learnt hypothesis $\hat{h} := h^{(k)}$ (hopefully $f(\hat{h}) \approx \min_{h \in \mathcal{H}} f(h)$)

4.8 Gradient Methods as Fixed-Point Iterations

The iterative optimization methods discussed in the previous sections are all special cases of a fixed-point iteration,

$$\mathbf{w}^{(k)} = \mathcal{F}\mathbf{w}^{(k-1)}, \text{ for } k = 1, 2, \dots \quad (4.23)$$

In what follows, we tacitly assume that the operator $\mathcal{F} : \mathbb{R}^m \rightarrow \mathbb{R}^m$ is defined on a Euclidean space \mathbb{R}^m with some fixed dimension m . We will use fixed-point iterations of the form (4.23) to solve GTVMin (see Ch. 3). For parametrized local models, we can learn local model parameters by the iterations (4.23) using the stacking $\mathbf{w}^{(k)}$ according to $\mathbf{w}^{(k)} = (\mathbf{w}^{(1,k)}, \dots, \mathbf{w}^{(n,k)})^T$. For a suitable choice for \mathcal{F} , the sequence model parameters $\mathbf{w}^{(k,i)}$ converges towards the optimal local model parameters (according to GTVMin) at each node

$i = 1, \dots, n$.

The FL algorithms in Ch. 5 are implementations of (4.23) with an operator \mathcal{F} having a GTVMin-solution $\widehat{\mathbf{w}} \in \mathbb{R}^{d \cdot n}$ as a fixed point,

$$\mathcal{F}\widehat{\mathbf{w}} = \widehat{\mathbf{w}}. \quad (4.24)$$

Given an instance of GTVMin, there are (infinitely) many different operators \mathcal{F} that satisfy (4.24). We obtain different FL algorithms by using different choices for \mathcal{F} in (4.23). A useful choice for \mathcal{F} should reduce the distance to a solution,

$$\underbrace{\|\mathbf{w}^{(k+1)} - \widehat{\mathbf{w}}\|_2}_{\stackrel{(4.23), (4.24)}{=} \|\mathcal{F}\mathbf{w}^{(k)} - \mathcal{F}\widehat{\mathbf{w}}\|_2} \leq \|\mathbf{w}^{(k)} - \widehat{\mathbf{w}}\|_2. \quad (4.25)$$

Thus, we require \mathcal{F} to be at least non-expansive (i.e., the iteration (4.23) should not result in worse model parameters). Moreover, each iteration (4.23) should make progress towards (reduce the distance from) the solution. This requirement can be made precise using the notion of a contractive operator [52, 53].

The operator \mathcal{F} is contractive (or a contraction mapping) if, for some $\kappa \in [0, 1)$,

$$\|\mathcal{F}\mathbf{w} - \mathcal{F}\mathbf{w}'\|_2 \leq \kappa \|\mathbf{w} - \mathbf{w}'\|_2 \text{ holds for any } \mathbf{w}, \mathbf{w}' \in \mathbb{R}^{dn}. \quad (4.26)$$

For a contractive \mathcal{F} , the fixed-point iteration (4.23) generates a sequence $\mathbf{w}^{(k)}$ that converges to a GTVMin solution $\widehat{\mathbf{w}}$ quite rapidly. In particular [2, Theorem 9.23],

$$\|\mathbf{w}^{(k)} - \widehat{\mathbf{w}}\|_2 \leq \kappa^k \|\mathbf{w}^{(0)} - \widehat{\mathbf{w}}\|_2. \quad (4.27)$$

Here, $\|\mathbf{w}^{(0)} - \widehat{\mathbf{w}}\|_2$ is the distance between the initialization $\mathbf{w}^{(0)}$ and the solution $\widehat{\mathbf{w}}$.




Fig. 4.5: Examples for a non-expansive operator $\mathcal{F}^{(1)}$, a firmly non-expansive operator $\mathcal{F}^{(2)}$ and a contractive operator $\mathcal{F}^{(3)}$.

4.9 Problems

4.1. Online Gradient Descent. Linear regression aims at learning model parameters of a linear model with minimum risk $\mathbb{E}\{(y - \mathbf{w}^T \mathbf{x})^2\}$ where (\mathbf{x}, y) is a RV. In practice we do not observe the RV (\mathbf{x}, y) itself but a (realization of a) sequence of i.i.d. samples $(\mathbf{x}^{(t)}, y^{(t)})$, for $t = 1, 2, \dots$. Online GD is an online learning method that updates the current model parameters $\mathbf{w}^{(t)}$, after observing $(\mathbf{x}^{(t)}, y^{(t)})$,

$$\mathbf{w}^{(t+1)} := \mathbf{w}^{(t)} + 2\eta_t \mathbf{x}^{(t)} (y - (\mathbf{w}^{(t)})^T \mathbf{x}^{(t)}) \text{ at time } t = 1, 2, \dots \quad (4.28)$$

Starting with initialization $\mathbf{w}^{(1)} := \mathbf{0}$, we run online GD for M time steps, resulting in the learnt model parameters $\mathbf{w}^{(M+1)}$. Develop upper bounds on the risk $\mathbb{E}\{(y - (\mathbf{w}^{(M)})^T \mathbf{x})^2\}$ for two choices for the learning rate schedule: $\eta_t := 1/(t + 5)$ or $\eta_t := 1/\sqrt{t + 5}$.

4.2. Computing the Average - I. Consider a FL network with graph \mathcal{G} and its Laplacian matrix $\mathbf{L}^{(\mathcal{G})}$. Each node carries a local dataset which consists of a single measurement $y^{(i)} \in \mathbb{R}$. To compute their average $(1/n) \sum_{i=1}^n y^{(i)}$ we try an iterative method that, starting from the initialization $\mathbf{u}^{(0)} := (y^{(1)}, \dots, y^{(n)})^T \in \mathbb{R}^n$, repeats the update

$$\mathbf{u}^{(k+1)} = \mathbf{u}^{(k)} - \eta \mathbf{L}^{(\mathcal{G})} \mathbf{u}^{(k)} \text{ for } k = 1, 2, \dots \quad (4.29)$$

Can you find a choice for η such that (4.29) becomes a fixed-point iteration (4.23) with a contractive operator \mathcal{F} . Given such a choice for η , how is the limit $\lim_{k \rightarrow \infty} \mathbf{u}^{(k+1)}$ related to the average $(1/n) \sum_{i=1}^n y^{(i)}$?

4.3. Computing the Average - II. Consider the FL network from Problem 4.2. Try to construct an instance of GTVMin for learning scalar local model parameters $w^{(i)}$ which coincide, for each node $i = 1, \dots, n$ with the average $(1/n) \sum_{i'=1}^n y^{(i')}$. If you find such an instance of GTVMin, solve it using GD (see Section 4.2).

5 FL Algorithms

Chapter 3 introduced GTVMin as a flexible design principle for FL methods that arise from different design choices for the local models and edge weights of the FL network. The solutions of GTVMin are local model parameters that strike a balance between the loss incurred on local datasets and the total variation. This chapter applies the gradient-based methods from Chapter 4 to solve GTVMin. We obtain FL algorithms by implementing these optimization methods as *message* passing across the edges of the FL network. These messages contain intermediate results of the computations carried out by FL algorithms. The details of how this message passing is implemented physically (e.g., via short range wireless technology) is beyond the scope of this book.

Section 5.2 studies the gradient step for the GTVMin instance obtained for training local linear models. In particular, we show how the convergence rate of the gradient step can be characterized by the properties of the local datasets and their FL network. Section (5.3) spells out the gradient step in the form of message passing over the FL network. Section 5.4 presents a FL algorithm that is obtained by replacing the exact GD with a stochastic approximation. Section 5.5 discusses FL algorithms for the single-model setting where we want to train a single global model in a distributed fashion.

5.1 Learning Goals

After completing this chapter, you

- can apply gradient-based methods to GTVMin for local linear models,
- can implement a gradient step via message passing over FL networks,

- know some motivation for stochastic gradient descent (SGD),
- know how federated averaging (FedAvg) is obtained from projected GD.

5.2 Gradient Descent for GTVMin

Consider a collection of n local datasets represented by the nodes $\mathcal{V} = \{1, \dots, n\}$ of an FL network $\mathcal{G} = (\mathcal{V}, \mathcal{E})$. Each undirected edge $\{i, i'\} \in \mathcal{E}$ in FL network \mathcal{G} has a known edge weight $A_{i,i'}$. We want to learn local model parameters $\mathbf{w}^{(i)}$ of a personalized linear model for each node $i = 1, \dots, n$. To this end, we solve the GTVMin instance

$$\{\widehat{\mathbf{w}}^{(i)}\}_{i=1}^n \in \underset{\{\mathbf{w}^{(i)}\}}{\operatorname{argmin}} \underbrace{\sum_{i \in \mathcal{V}} \overbrace{(1/m_i) \|\mathbf{y}^{(i)} - \mathbf{X}^{(i)} \mathbf{w}^{(i)}\|_2^2}^{\text{local loss } L_i(\mathbf{w}^{(i)})} + \alpha \sum_{\{i,i'\} \in \mathcal{E}} A_{i,i'} \|\mathbf{w}^{(i)} - \mathbf{w}^{(i')}\|_2^2}_{=: f(\mathbf{w})}. \quad (5.1)$$

As discussed in Chapter 3, the objective function in (5.1) - viewed as a function of the stacked local model parameters $\mathbf{w} := \operatorname{stack}\{\mathbf{w}^{(i)}\}_{i=1}^n$ - is a quadratic function of the form (4.1),

$$\mathbf{w}^T \underbrace{\left(\begin{pmatrix} \mathbf{Q}^{(1)} & \mathbf{0} & \dots & \mathbf{0} \\ \mathbf{0} & \mathbf{Q}^{(2)} & \dots & \mathbf{0} \\ \vdots & \vdots & \ddots & \vdots \\ \mathbf{0} & \mathbf{0} & \dots & \mathbf{Q}^{(n)} \end{pmatrix} + \alpha \mathbf{L}^{(\mathcal{G})} \otimes \mathbf{I} \right)}_{:= \mathbf{Q}} \mathbf{w} + \underbrace{\left((\mathbf{q}^{(1)})^T, \dots, (\mathbf{q}^{(n)})^T \right)}_{:= \mathbf{q}^T} \mathbf{w} \quad (5.2)$$

with $\mathbf{Q}^{(i)} = (1/m_i) (\mathbf{X}^{(i)})^T \mathbf{X}^{(i)}$, and $\mathbf{q}^{(i)} := (-2/m_i) (\mathbf{X}^{(i)})^T \mathbf{y}^{(i)}$.

Therefore, the discussion and analysis of gradient-based methods from Chapter 4 also apply to GTVMin (5.1). In particular, we can use the gradient step

$$\begin{aligned}\mathbf{w}^{(k+1)} &:= \mathbf{w}^{(k)} - \eta \nabla f(\mathbf{w}^{(k)}) \\ &\stackrel{(5.2)}{=} \mathbf{w}^{(k)} - \eta(2\mathbf{Q}\mathbf{w}^{(k)} + \mathbf{q})\end{aligned}\tag{5.3}$$

to iteratively compute an approximate solution $\widehat{\mathbf{w}}$ to (5.1). This solution consists of learnt local model parameters $\widehat{\mathbf{w}}^{(i)}$, i.e., $\widehat{\mathbf{w}} = \text{stack}\{\widehat{\mathbf{w}}^{(i)}\}$. Section 5.3 will formulate the gradient step (5.3) directly in terms of local model parameters, resulting in a message passing over the FL network \mathcal{G} .

According to the convergence analysis in Chapter 4, the convergence rate of the iterations (5.3) is determined by the eigenvalues $\lambda_j(\mathbf{Q})$ of the matrix \mathbf{Q} in (5.2). Clearly, these eigenvalues are related to the eigenvalues $\lambda_j(\mathbf{Q}^{(i)})$ and to the eigenvalues $\lambda_j(\mathbf{L}^{(\mathcal{G})})$ of the Laplacian matrix of the FL network \mathcal{G} . In particular, we will use the following two summary parameters

$$\lambda_{\max} := \max_{i=1,\dots,n} \lambda_d(\mathbf{Q}^{(i)}), \text{ and } \bar{\lambda}_{\min} := \lambda_1\left((1/n) \sum_{i=1}^n \mathbf{Q}^{(i)}\right).\tag{5.4}$$

We first present an upper bound U (see (4.5)) on the eigenvalues of the matrix \mathbf{Q} in (5.2).

Proposition 5.1. *The eigenvalues of \mathbf{Q} in (5.2) are upper-bounded as*

$$\begin{aligned}\lambda_j(\mathbf{Q}) &\leq \lambda_{\max} + \alpha \lambda_n(\mathbf{L}^{(\mathcal{G})}) \\ &\leq \underbrace{\lambda_{\max} + \alpha 2d_{\max}^{(\mathcal{G})}}_{:=U}, \text{ for } j = 1, \dots, dn.\end{aligned}\tag{5.5}$$

Proof. See Section 5.9.1. □

The next result offers a lower bound on the eigenvalues $\lambda_j(\mathbf{Q})$.

Proposition 5.2. *Consider the matrix \mathbf{Q} in (5.2). If $\lambda_2(\mathbf{L}^{(\mathcal{G})}) > 0$ (i.e., the FL network in (5.1) is connected) and $\bar{\lambda}_{\min} > 0$ (i.e., the average of the matrices $\mathbf{Q}^{(i)}$ is non-singular), then the matrix \mathbf{Q} is invertible and its smallest eigenvalue is lower bounded as*

$$\lambda_1(\mathbf{Q}) \geq \frac{1}{1 + \rho^2} \min\{\lambda_2(\mathbf{L}^{(\mathcal{G})})\alpha\rho^2, \bar{\lambda}_{\min}/2\}. \quad (5.6)$$

Here, we used the shorthand $\rho := \bar{\lambda}_{\min}/(4\lambda_{\max})$ (see (5.4)).

Proof. See Section 5.9.2. □

Prop. 5.2 and Prop. 5.1 can provide some guidance for the design choices of GTVMin. According to the convergence analysis of gradient-based methods in Chapter 4, the eigenvalue $\lambda_1(\mathbf{Q})$ should be close to $\lambda_{dn}(\mathbf{Q})$ to ensure fast convergence. This favours FL networks \mathcal{G} , with corresponding eigenvalue $\lambda_2(\mathbf{L}^{(\mathcal{G})})$ and maximum node degree $d_{\max}^{(\mathcal{G})}$, having a small ratio between the upper bound (5.5) and the lower bound (5.6).

The bounds in (5.5) and (5.6) also depend on the GTVMin parameter α . While these bounds might provide some guidance for the choice of α , the exact dependence of the convergence speed of (5.3) on α is complicated. For a fixed value of learning rate in (5.3), using larger values for α might slow down the convergence of (5.3) for some collection of local datasets but speed up the convergence of (5.3) for another collection of local datasets (see Prob. 5.1).

5.3 Message Passing Implementation

We now discuss in more detail the implementation of gradient-based methods to solve the GTVMin instances with a differentiable objective function $f(\mathbf{w})$.

One such instance is GTVMin for local linear models (see (5.1)). The core of gradient-based methods is the gradient step

$$\mathbf{w}^{(k+1)} := \mathbf{w}^{(k)} - \eta \nabla f(\mathbf{w}^{(k)}). \quad (5.7)$$

The iterate $\mathbf{w}^{(k)}$ contains local model parameters $\mathbf{w}^{(i,k)}$,

$$\mathbf{w}^{(k)} =: \text{stack}\{\mathbf{w}^{(i,k)}\}_{i=1}^n. \quad (5.8)$$

Inserting (5.1) into (5.7), we obtain the gradient step

$$\begin{aligned} \mathbf{w}^{(i,k+1)} := \mathbf{w}^{(i,k)} + \eta \left[\underbrace{(2/m_i)(\mathbf{X}^{(i)})^T (\mathbf{y}^{(i)} - \mathbf{X}^{(i)} \mathbf{w}^{(i,k)})}_{\text{(I)}} \right. \\ \left. + 2\alpha \underbrace{\sum_{i' \in \mathcal{N}^{(i)}} A_{i,i'} (\mathbf{w}^{(i',k)} - \mathbf{w}^{(i,k)})}_{\text{(II)}} \right]. \quad (5.9) \end{aligned}$$

The update (5.9) consists of two components, denoted (I) and (II). The component (I) is nothing but the negative gradient $-\nabla L_i(\mathbf{w}^{(i,k)})$ of the local loss $L_i(\mathbf{w}^{(i)}) := (1/m_i) \|\mathbf{y}^{(i)} - \mathbf{X}^{(i)} \mathbf{w}^{(i)}\|_2^2$. Component (I) drives the updated local model parameters $\mathbf{w}^{(i,k+1)}$ towards the minimum of $L_i(\cdot)$, i.e., having a small deviation between labels $y^{(i,r)}$ and the predictions $(\mathbf{w}^{(i,k+1)})^T \mathbf{x}^{(i,r)}$. Note that we can rewrite the component (I) in (5.9), as

$$(2/m_i) \sum_{r=1}^{m_i} \mathbf{x}^{(i,r)} (y^{(i,r)} - (\mathbf{x}^{(i,r)})^T \mathbf{w}^{(i,k)}). \quad (5.10)$$

The purpose of component (II) in (5.9) is to force the local model parameters to be similar across an edge $\{i, i'\}$ with large weight $A_{i,i'}$. We control the relative importance of (II) and (I) using the GTVMin parameter α : Choosing a large value for α puts more emphasis on enforcing similar local

model parameters across the edges. Using a smaller α puts more emphasis on learning local model parameters delivering accurate predictions (incurring a small loss) on the local dataset.




Fig. 5.1: At the beginning of iteration k , node $i = 1$ collects the current local model parameters $\mathbf{w}^{(2,k)}$ and $\mathbf{w}^{(3,k)}$ from its neighbours. Then, it computes the gradient step (5.9) to obtain the new local model parameters $\mathbf{w}^{(1,k+1)}$. These updated parameters are then used in the next iteration for the local updates at the neighbours $i = 2, 3$.

The execution of the gradient step (5.9) requires only local information at node i . Indeed, the update (5.9) node i depends only on its current model parameters $\mathbf{w}^{(i,k)}$, the local loss function $L_i(\cdot)$, the neighbours' model parameters $\mathbf{w}^{(i',k)}$, for $i' \in \mathcal{N}^{(i)}$, and the corresponding edge weights $A_{i,i'}$ (see Figure 5.1). In particular, the update (5.9) does not depend on any properties (nor edge weights) of the FL network beyond the neighbours $\mathcal{N}^{(i)}$.

We obtain Algorithm 5.1 by repeating the gradient step (5.9), simultaneously for each node $i \in \mathcal{V}$, until a stopping criterion is met. Algorithm 5.1 allows for potentially different learning rates $\eta_{k,i}$ at different nodes i and iterations k (see Section 4.3). It is important to note that Algorithm 5.1 requires a synchronous (simultaneous) execution of the updates (5.9) at all nodes $i \in \mathcal{V}$ [54, 55]. Loosely speaking, all nodes i relies on a single global

Algorithm 5.1 FedGD for Local Linear Models

Input: FL network \mathcal{G} ; GTV parameter α ; learning rate $\eta_{k,i}$

local dataset $\mathcal{D}^{(i)} = \{(\mathbf{x}^{(i,1)}, y^{(i,1)}), \dots, (\mathbf{x}^{(i,m_i)}, y^{(i,m_i)})\}$ for each i ; some stopping criterion.

Output: linear model parameters $\widehat{\mathbf{w}}^{(i)}$ for each node $i \in \mathcal{V}$

Initialize: $k := 0$; $\mathbf{w}^{(i,0)} := \mathbf{0}$

- 1: **while** stopping criterion is not satisfied **do**
 - 2: **for** all nodes $i \in \mathcal{V}$ (simultaneously) **do**
 - 3: share local model parameters $\mathbf{w}^{(i,k)}$ with neighbours $i' \in \mathcal{N}^{(i)}$
 - 4: update local model parameters via (5.9)
 - 5: **end for**
 - 6: increment iteration counter: $k := k + 1$
 - 7: **end while**
 - 8: $\widehat{\mathbf{w}}^{(i)} := \mathbf{w}^{(i,k)}$ for all nodes $i \in \mathcal{V}$
-




Fig. 5.2: Algorithm 5.1 alternates between message passing across the edges of the FL network (left and right) and updates of local model parameters (centre).

clock that maintains the current iteration counter k [56].

At the beginning of iteration k , each node $i \in \mathcal{V}$ sends its current model parameters $\mathbf{w}^{(i,k)}$ to their neighbours $i' \in \mathcal{N}^{(i)}$. Then, each node $i \in \mathcal{V}$ updates their model parameters according to (5.9), resulting in the updated model parameters $\mathbf{w}^{(i,k+1)}$. As soon as these local updates are completed, the global clock increments the counter $k \mapsto k + 1$ and triggers the next iteration to be executed by all nodes.

The implementation of Algorithm 5.1 in real-world computational infrastructures might incur deviations from the exact synchronous execution of (5.9) [57, Sec. 10]. This deviation can be modelled as a perturbation of the gradient step (5.7) and therefore analyzed using the concepts of Section 4.5 on perturbed GD. Section 8.3 will also discuss the effect of imperfect computation in the context of key requirements for trustworthy FL.

We close this section by generalizing Algorithm 5.1 which is limited FL networks using local linear models. This generalization, summarized in Algorithm 5.2, can be used to train parametric local models $\mathcal{H}^{(i)}$ with a differentiable loss function $L_i(\mathbf{w}^{(i)})$, for $i = 1, \dots, n$.

Algorithm 5.2 FedGD for Parametric Local Models

Input: FL network \mathcal{G} ; GTV parameter α ; learning rate $\eta_{k,i}$

local loss function $L_i(\mathbf{w}^{(i)})$ for each $i = 1, \dots, n$; some stopping criterion.

Output: linear model parameters $\hat{\mathbf{w}}^{(i)}$ for each node $i \in \mathcal{V}$

Initialize: $k := 0$; $\mathbf{w}^{(i,0)} := \mathbf{0}$

- 1: **while** stopping criterion is not satisfied **do**
- 2: **for** all nodes $i \in \mathcal{V}$ (simultaneously) **do**
- 3: share local model parameters $\mathbf{w}^{(i,k)}$ with neighbours $i' \in \mathcal{N}^{(i)}$
- 4: update local model parameters via

$$\mathbf{w}^{(i,k+1)} := \mathbf{w}^{(i,k)} - \eta_{k,i} \left[\nabla L_i(\mathbf{w}^{(i,k)}) + 2\alpha \sum_{i' \in \mathcal{N}^{(i)}} A_{i,i'} (\mathbf{w}^{(i,k)} - \mathbf{w}^{(i',k)}) \right].$$

- 5: **end for**
 - 6: increment iteration counter: $k := k + 1$
 - 7: **end while**
 - 8: $\hat{\mathbf{w}}^{(i)} := \mathbf{w}^{(i,k)}$ for all nodes $i \in \mathcal{V}$
-

5.4 FedSGD

Consider Algorithm 5.1 for training local linear models $h^{(i)}(\mathbf{x}) = \mathbf{x}^T \mathbf{w}^{(i)}$ for each node $i = 1, \dots, n$ of a FL network. Note that step 4 of Algorithm 5.1 requires to compute the sum (5.10). It might be infeasible to compute this

sum exactly, e.g., when local datasets are generated by remote devices with limited connectivity. It is then useful to approximate the sum by

$$\underbrace{(2/B) \sum_{r \in \mathcal{B}} \mathbf{x}^{(i,r)} (y^{(i,r)} - (\mathbf{x}^{(i,r)})^T \mathbf{w}^{(i,k)})}_{\approx (5.10)}. \quad (5.11)$$

The approximation (5.11) uses a subset (so-called *batch*)

$$\mathcal{B} = \{(\mathbf{x}^{(r_1)}, y^{(r_1)}), \dots, (\mathbf{x}^{(r_B)}, y^{(r_B)})\}$$

of B randomly chosen data points from $\mathcal{D}^{(i)}$. While (5.10) requires summing over m data points, the approximation requires to sum over B (typically $B \ll m$) data points.

Inserting the approximation (5.11) into the gradient step (5.9) yields the approximate gradient step

$$\begin{aligned} \mathbf{w}^{(i,k+1)} := & \mathbf{w}^{(i,k)} + \eta_{k,i} \left[\underbrace{(2/B) \sum_{r \in \mathcal{B}} \mathbf{x}^{(i,r)} (y^{(i,r)} - (\mathbf{x}^{(i,r)})^T \mathbf{w}^{(i,k)})}_{\approx (5.10)} \right. \\ & \left. + 2\alpha \sum_{i' \in \mathcal{N}^{(i)}} A_{i,i'} (\mathbf{w}^{(i',k)} - \mathbf{w}^{(i,k)}) \right]. \end{aligned} \quad (5.12)$$

We obtain Algorithm 5.3 from Algorithm 5.1 by replacing the gradient step (5.9) with the approximation (5.12).

We close this section by generalizing Algorithm 5.3 which is limited FL networks using local linear models. This generalization, summarized in Algorithm 5.4, can be used to train parametric local models $\mathcal{H}^{(i)}$ with a differentiable loss function $L_i(\mathbf{w}^{(i)})$, for $i = 1, \dots, n$. Algorithm 5.4 does not require these local loss function themselves, but only an oracle $\mathbf{g}^{(i)}(\cdot)$ for each node $i = 1, \dots, n$. For a given vector $\mathbf{w}^{(i)}$, the oracle at node i delivers an

Algorithm 5.3 FedSGD for Local Linear Models

Input: FL network \mathcal{G} ; GTV parameter α ; learning rate $\eta_{k,i}$;

local datasets $\mathcal{D}^{(i)} = \{(\mathbf{x}^{(i,1)}, y^{(i,1)}), \dots, (\mathbf{x}^{(i,m_i)}, y^{(i,m_i)})\}$ for each node i ;

batch size B ; some stopping criterion.

Output: linear model parameters $\widehat{\mathbf{w}}^{(i)}$ at each node $i \in \mathcal{V}$

Initialize: $k := 0$; $\mathbf{w}^{(i,0)} := \mathbf{0}$

- 1: **while** stopping criterion is not satisfied **do**
 - 2: **for** all nodes $i \in \mathcal{V}$ (simultaneously) **do**
 - 3: share local model parameters $\mathbf{w}^{(i,k)}$ with all neighbours $i' \in \mathcal{N}^{(i)}$
 - 4: draw fresh batch $\mathcal{B}^{(i)} := \{r_1, \dots, r_B\}$
 - 5: update local model parameters via (5.12)
 - 6: **end for**
 - 7: increment iteration counter $k := k + 1$
 - 8: **end while**
 - 9: $\widehat{\mathbf{w}}^{(i)} := \mathbf{w}^{(i,k)}$ for all nodes $i \in \mathcal{V}$
-

approximate gradient (or estimate) $\mathbf{g}^{(i)}(\mathbf{w}^{(i)}) \approx \nabla L_i(\mathbf{w}^{(i)})$. The analysis of Algorithm 5.4 can be facilitated by a probabilistic model which interprets the oracle output $\mathbf{g}^{(i)}(\mathbf{w}^{(i)})$ as the realization of a RV. Under such a probabilistic model, we refer to an oracle as unbiased if $\mathbb{E}\{\mathbf{g}^{(i)}(\mathbf{w}^{(i)})\} = \nabla L_i(\mathbf{w}^{(i)})$.

Algorithm 5.4 FedSGD for Parametric Local Models

Input: FL network \mathcal{G} ; GTV parameter α ; learning rate $\eta_{k,i}$

gradient oracle $\mathbf{g}^{(i)}(\cdot)$ for each node $i = 1, \dots, n$; some stopping criterion.

Output: linear model parameters $\widehat{\mathbf{w}}^{(i)}$ for each node $i \in \mathcal{V}$

Initialize: $k := 0$; $\mathbf{w}^{(i,0)} := \mathbf{0}$

- 1: **while** stopping criterion is not satisfied **do**
- 2: **for** all nodes $i \in \mathcal{V}$ (simultaneously) **do**
- 3: share local model parameters $\mathbf{w}^{(i,k)}$ with neighbours $i' \in \mathcal{N}^{(i)}$
- 4: update local model parameters via

$$\mathbf{w}^{(i,k+1)} := \mathbf{w}^{(i,k)} - \eta_{k,i} \left[\mathbf{g}^{(i)}(\mathbf{w}^{(i,k)}) + 2\alpha \sum_{i' \in \mathcal{N}^{(i)}} A_{i,i'} (\mathbf{w}^{(i,k)} - \mathbf{w}^{(i',k)}) \right].$$

- 5: **end for**
 - 6: increment iteration counter: $k := k + 1$
 - 7: **end while**
 - 8: $\widehat{\mathbf{w}}^{(i)} := \mathbf{w}^{(i,k)}$ for all nodes $i \in \mathcal{V}$
-

5.5 FedAvg

Consider a FL application that requires to learn model parameters $\hat{\mathbf{w}} \in \mathbb{R}^d$ for a single (global) linear model from de-centralized collection local datasets $\mathcal{D}^{(i)}$, $i = 1, \dots, n$.¹² How can we learn $\hat{\mathbf{w}}$ without exchanging local datasets but only some model parameters (updates)?

One approach is to apply Algorithm 5.1 to GTVMin (5.1) with a sufficiently large α . According to our analysis in Chapter 3 (in particular, Prop. 3.1), for sufficiently large α , the GTVMin solutions $\hat{\mathbf{w}}^{(i)}$ are almost identical for all nodes $i \in \mathcal{V}$. We can interpret the local model parameters delivered by GTVMin as a local copy of the global model parameters.

Note that the bound in Prop. 3.1 only applies if the FL network (used in GTVMin) is connected. One example for a connected FL network is the star as depicted in Figure 5.3. Here, we choose one node $i = 1$ as a centre node that is connected by an edge with weight $A_{1,i}$ to the remaining nodes $i = 2, \dots, n$. The star graph is distinct in the sense of using the minimum number of edges required to connect n nodes [58].

Instead of using GTVMin with a connected FL network and a large value of α , we can also enforce identical local copies $\hat{\mathbf{w}}^{(i)}$ via a constraint:

$$\begin{aligned} \hat{\mathbf{w}} \in \arg \min_{\mathbf{w} \in \mathcal{C}} \sum_{i \in \mathcal{V}} (1/m_i) \|\mathbf{y}^{(i)} - \mathbf{X}^{(i)} \mathbf{w}^{(i)}\|_2^2 \\ \text{with } \mathcal{C} = \left\{ \mathbf{w} = \text{stack}\{\mathbf{w}^{(i)}\}_{i=1}^n : \mathbf{w}^{(i)} = \mathbf{w}^{(i')} \text{ for any } i, i' \in \mathcal{V} \right\}. \end{aligned} \quad (5.13)$$

Note that the constraint set \mathcal{C} is nothing but the subspace defined in (3.16). The projection of a given collection of local model parameters $\mathbf{w} = \text{stack}\{\mathbf{w}^{(i)}\}$

¹²This setting is a special case of horizontal FL which we discuss in Sec. 6.4.




Fig. 5.3: Star graph $\mathcal{G}^{(\text{star})}$ with a centre node $i = 1$ representing a server that trains a (global) model which is shared with peripheral nodes. These peripheral nodes represent *clients* generating local datasets. The training process at the server is facilitated by receiving updates on the model parameters from the clients.

on \mathcal{C} is given by

$$P_{\mathcal{C}}(\mathbf{w}) = (\mathbf{v}^T, \dots, \mathbf{v}^T)^T \text{ with } \mathbf{v} := (1/n) \sum_{i \in \mathcal{V}} \mathbf{w}^{(i)}. \quad (5.14)$$

We can solve (5.13) using projected GD from Chapter 4 (see Section 4.6). The resulting projected gradient step for solving (5.13) is

$$\widehat{\mathbf{w}}_k^{(i)} := \underbrace{\mathbf{w}^{(i,k)} + \eta_{i,k} (2/m_i) (\mathbf{X}^{(i)})^T (\mathbf{y}^{(i)} - \mathbf{X}^{(i)} \mathbf{w}^{(i,k)})}_{\text{(local gradient step)}} \quad (5.15)$$

$$\mathbf{w}^{(i,k+1)} := (1/n) \sum_{i' \in \mathcal{V}} \widehat{\mathbf{w}}_k^{(i')} \quad (\text{projection}). \quad (5.16)$$

We can implement (5.16) conveniently in a server-client system with each node i being a client:

- First, each node computes the update (5.15), i.e., a gradient step towards a minimum of the local loss $\|\mathbf{y}^{(i)} - \mathbf{X}^{(i)} \mathbf{w}\|_2^2$.

- Second, each node i sends the result $\widehat{\mathbf{w}}_k^{(i)}$ of its local gradient step to a server.
- Finally, after receiving the updates $\widehat{\mathbf{w}}_k^{(i)}$ from all nodes $i \in \mathcal{V}$, the server computes the projection step (5.16). This projection results in the new local model parameters $\mathbf{w}^{(i,k+1)}$ that are sent back to each client i .

The averaging step (5.16) might take much longer to execute than the local update step (5.15). Indeed, (5.16) typically requires transmission of local model parameters from every client $i \in \mathcal{V}$ to a server or central computing unit. Thus, after the client $i \in \mathcal{V}$ has computed the local gradient step (5.15), it must wait until the server (i) has collected the updates $\widehat{\mathbf{w}}_k^{(i)}$ from all clients and (ii) sent back their average $\mathbf{w}^{(i,k+1)}$ to $i \in \mathcal{V}$.

Instead of using a computationally inexpensive gradient step (5.15),¹³ and then being forced to wait for receiving $\mathbf{w}^{(i,k+1)}$ back from the server, a client might make better use of its resources. For example, the device i could execute several local gradient steps (5.15) in order to make more progress towards the optimum. Another option is to use a local minimization of $L_i(\mathbf{v}) := (1/m_i) \|\mathbf{y}^{(i)} - \mathbf{X}^{(i)}\mathbf{v}\|_2^2$ around $\mathbf{w}^{(i,k)}$,

$$\widehat{\mathbf{w}}_k^{(i)} := \operatorname{argmin}_{\mathbf{v} \in \mathbb{R}^d} \left[\underbrace{(1/m_i) \|\mathbf{y}^{(i)} - \mathbf{X}^{(i)}\mathbf{v}\|_2^2}_{=L_i(\mathbf{v})} + (1/\eta) \|\mathbf{v} - \mathbf{w}^{(i,k)}\|_2^2 \right]. \quad (5.17)$$

Note that (5.17) is nothing but the proximal operator of $L_i(\mathbf{v})$ (see (4.19)). We obtain Algorithm 5.5 from (5.15)-(5.16) by replacing the gradient step (5.15) with the local minimization (5.17).

¹³For a large local dataset, the local gradient step (5.15) might actually be computationally expensive and should be replaced by an approximation, e.g., based on the stochastic gradient approximation (5.11).

The local minimization step (5.17) is closely related to ridge regression. Indeed, (5.17) is obtained from ridge regression (2.25) by replacing the regularizer $\mathcal{R}\{\mathbf{v}\} := \|\mathbf{v}\|_2^2$ with the regularizer $\mathcal{R}\{\mathbf{v}\} := \|\mathbf{v} - \mathbf{w}^{(i,k)}\|_2^2$.

As the notation in (5.17) indicates, the parameter η plays a role similar to the learning rate of a gradient step (4.2). It controls the size of the neighbourhood of $\mathbf{w}^{(i,k)}$ over which (5.17) optimizes the local loss function $L_i(\cdot)$. Choosing a small η forces the update (5.17) to not move too far from the given model parameters $\mathbf{w}^{(i,k)}$.

We can interpret the update (5.17) as a form of ERM (2.3) for linear regression. Indeed,

$$\begin{aligned} \min_{\mathbf{v} \in \mathbb{R}^d} & \left[(1/m_i) \|\mathbf{y}^{(i)} - \mathbf{X}^{(i)} \mathbf{v}\|_2^2 + (1/\eta) \|\mathbf{w}^{(i,k)} - \mathbf{v}\|_2^2 \right] = \\ \min_{\mathbf{v} \in \mathbb{R}^d} & \left[(1/m_i) \sum_{r=1}^{m_i} (y^{(i,r)} - (\mathbf{x}^{(i,r)})^T \mathbf{v})^2 + (1/\eta) \|\mathbf{w}^{(i,k)} - \mathbf{v}\|_2^2 \right] = \\ \min_{\mathbf{v} \in \mathbb{R}^d} & (1/m_i) \left[\sum_{r=1}^{m_i} (y^{(i,r)} - (\mathbf{x}^{(i,r)})^T \mathbf{v})^2 + \sum_{r=1}^d (m_i/\eta) (w_r^{(i,k)} - (\mathbf{e}^{(r)})^T \mathbf{v})^2 \right]. \quad (5.18) \end{aligned}$$

Here, $\mathbf{e}^{(r)}$ denotes the vector obtained from the r th column of the identity matrix \mathbf{I} . The only difference between (5.18) and (2.3) is the presence of the sample weights (m_i/η) .¹⁴

Note that Algorithm 5.5 provides only an abstract description (or mathematical model) of a practical FL system. The details of an actual implementation Algorithm 5.5, such as providing means for synchronous communication between the server and all clients (see steps 2 and 4 on the server side) is beyond the scope of this book. Instead, we refer the reader to relevant

¹⁴The “`.fit()`” method of the Python class `sklearn.linear_model.LinearRegression` allows to specify sample weights via the parameter `sample_weight` [59].

Algorithm 5.5 Server and Client-Side of FedAvg to train a linear model

The Server. Initialize. $k := 0$

- 1: **while** stopping criterion is not satisfied **do**
- 2: receive local model parameters $\mathbf{w}^{(i)}$ from all clients $i = 1, \dots, n$
- 3: update global model parameters

$$\bar{\mathbf{w}}[k] := (1/n) \sum_{i=1}^n \mathbf{w}^{(i)}.$$

- 4: send updated global model parameters $\bar{\mathbf{w}}[k]$ to all clients $i = 1, \dots, n$
- 5: $k := k + 1$
- 6: **end while**

Client i .

- 1: **while** stopping criterion is not satisfied **do**
- 2: receive global model parameters $\bar{\mathbf{w}}[k]$ from server
- 3: update local model parameters by RERM (see (5.17))

$$\mathbf{w}^{(i)} := \operatorname{argmin}_{\mathbf{v} \in \mathbb{R}^d} \left[(1/m_i) \sum_{r=1}^{m_i} (\mathbf{v}^T \mathbf{x}^{(i,r)} - y^{(i,r)})^2 + (1/\eta) \|\mathbf{v} - \bar{\mathbf{w}}[k]\|_2^2 \right].$$

- 4: send $\mathbf{w}^{(i)}$ back to server
 - 5: **end while**
-

literature on the implementation of distributed computing systems [55, 60].

5.6 FedRelax

We now apply a simple block-coordinate minimization method [54] to solve GTVMin (3.20). To this end, we rewrite (3.20) as

$$\begin{aligned} \hat{\mathbf{w}} &\in \arg \min_{\mathbf{w} \in \mathbb{R}^{d \cdot n}} \underbrace{\sum_{i \in \mathcal{V}} f^{(i)}(\mathbf{w})}_{:= f^{(\text{GTV})}(\mathbf{w})} \\ \text{with } f^{(i)}(\mathbf{w}) &:= L_i(\mathbf{w}^{(i)}) + (\alpha/2) \sum_{i' \in \mathcal{N}^{(i)}} A_{i,i'} \left\| \mathbf{w}^{(i)} - \mathbf{w}^{(i')} \right\|_2^2. \end{aligned} \quad (5.19)$$

According to (5.19), the objective function of (3.20) decomposes into components, one for each node \mathcal{V} of the FL network. Moreover, the local model parameters $\mathbf{w}^{(i)}$ influence the objective function only via the components at the nodes $i \cup \mathcal{N}^{(i)}$. We exploit this structure of (5.19) to decouple the optimization of the local model parameters $\{\hat{\mathbf{w}}^{(i)}\}_{i \in \mathcal{V}}$ as described next.

Consider some local model parameters $\hat{\mathbf{w}}_k^{(i)}$, for $i = 1, \dots, n$, at time k . We then update (in parallel) each $\hat{\mathbf{w}}_k^{(i)}$ by minimizing $f^{(\text{GTV})}(\cdot)$ along $\mathbf{w}^{(i)}$ with the other local model parameters $\mathbf{w}^{(i')} := \hat{\mathbf{w}}_k^{(i')}$ held fixed for all $i' \neq i$,

$$\begin{aligned} \hat{\mathbf{w}}_{k+1}^{(i)} &\in \underset{\mathbf{w}^{(i)} \in \mathbb{R}^d}{\operatorname{argmin}} f^{(\text{GTV})} \left(\hat{\mathbf{w}}_k^{(1)}, \dots, \hat{\mathbf{w}}_k^{(i-1)}, \mathbf{w}^{(i)}, \hat{\mathbf{w}}_k^{(i+1)}, \dots \right) \\ &\stackrel{(5.19)}{=} \underset{\mathbf{w}^{(i)} \in \mathbb{R}^d}{\operatorname{argmin}} f^{(i)} \left(\hat{\mathbf{w}}_k^{(1)}, \dots, \hat{\mathbf{w}}_k^{(i-1)}, \mathbf{w}^{(i)}, \hat{\mathbf{w}}_k^{(i+1)}, \dots \right) \\ &\stackrel{(5.19)}{=} \underset{\mathbf{w}^{(i)} \in \mathbb{R}^d}{\operatorname{argmin}} L_i(\mathbf{w}^{(i)}) + \alpha \sum_{i' \in \mathcal{N}^{(i)}} A_{i,i'} \left\| \mathbf{w}^{(i)} - \hat{\mathbf{w}}_k^{(i')} \right\|_2^2. \end{aligned} \quad (5.20)$$

The update (5.20) is an instance of the *non-linear Jacobi algorithm* (applied to (5.19)) [54, Sec. 3.2.4.]. Another interpretation of (5.20) is as a variant of

block-coordinate optimization [61]. We obtain Algorithm 5.6 by repeating the update (5.20) for a sufficient number of iterations.

Algorithm 5.6 FedRelax for Parametric Models

Input: FL network \mathcal{G} with local loss functions $L_i(\cdot)$, GTV parameter α

Initialize: $k := 0$; $\widehat{\mathbf{w}}_0^{(i)} := \mathbf{0}$

```

1: while stopping criterion is not satisfied do
2:   for all nodes  $i \in \mathcal{V}$  in parallel do
3:     compute  $\widehat{\mathbf{w}}_{k+1}^{(i)}$  via (5.20)
4:     share  $\widehat{\mathbf{w}}_{k+1}^{(i)}$  with neighbours  $\mathcal{N}^{(i)}$ 
5:   end for
6:    $k := k + 1$ 
7: end while

```

A Model Agnostic Method. The applicability of Algorithm 5.6 is limited to FL networks with parametric local models (such as linear regression or ANNs with a common structure). Note that Algorithm 5.6 results from the application of the non-linear Jacobi method to solve GTVMin (3.20) for parametric local models. We can generalize Algorithm 5.6 to non-parametric local models by applying the non-linear Jacobi algorithm to the GTVMin variant (3.37). This results in the update

$$\widehat{h}_{k+1}^{(i)} \in \operatorname{argmin}_{h^{(i)} \in \mathcal{H}^{(i)}} L_i(h^{(i)}) + \alpha \sum_{i' \in \mathcal{N}^{(i)}} A_{i,i'} \underbrace{D(h^{(i)}, \widehat{h}_k^{(i)})}_{\text{see (3.35)}}. \quad (5.21)$$

We obtain Algorithm 5.7 by replacing the update (5.20) with (5.21) in step 3 of Algorithm 5.6. Algorithm 5.7 is model-agnostic in the sense of allowing the devices of an FL network to train parametric as well as non-parametric

local models. The only restriction for the local models is that they allow for efficient (approximate) solution of (5.21).

Algorithm 5.7 Model Agnostic FedRelax

Input: FL network with \mathcal{G} , local models $\mathcal{H}^{(i)}$, loss functions $L_i(\cdot)$, GTV parameter α

Initialize: $k := 0; \widehat{h}_0^{(i)} := \mathbf{0}$

```

1: while stopping criterion is not satisfied do
2:   for all nodes  $i \in \mathcal{V}$  in parallel do
3:     compute  $\widehat{h}_{k+1}^{(i)}$  via (5.21)
4:   end for
5:    $k := k + 1$ 
6: end while

```

5.7 Asynchronous FL Algorithms

The FL algorithms presented so far require a synchronization between the devices $i = 1, \dots, n$ of an FL network [55, Ch. 6]. Indeed, only when all participating device have completed their local updates (e.g., a variant of a gradient step) and shared the results with its neighbours, the algorithm can proceed to the next iteration [62, Sec. 10], [54, Sec. 1.4]. Such a synchronous operation can be detrimental for several reasons. First, the entire algorithm fails as soon as a single device stops its operation. Second, the computational and communication capabilities of devices might be very different, resulting in some devices wasting time by waiting for the results of slow devices (so called *stragglers*).

Let us now discuss a generic FL algorithm that does not require synchronization among the devices $i = 1, \dots, n$ [54, Sec. 6]. We formulate this algorithm for parametric local models, each having its own vector $\mathbf{w}^{(i)}$ of local model parameters. The generalization of this algorithm to non-parametric local models will not pose a significant challenge.

This generic asynchronous algorithm consists of a sequence of *update events*, indexed by $k = 1, 2, \dots$. During event k , some nodes $i \in \mathcal{V}$ update their local model parameters to

$$\mathbf{w}^{(i,k+1)} := \mathcal{F}^{(i)}(\mathbf{w}^{(1,k_{i,1})}, \dots, \mathbf{w}^{(n,k_{i,n})}). \quad (5.22)$$

The nodes which execute the update (5.22) during event k constitute the active set $\mathcal{A}^{(k)} \subseteq \mathcal{V}$.

The update (5.22) involves some operator $\mathcal{F}^{(i)} : \mathbb{R}^{dn} \rightarrow \mathbb{R}^d$ that determines the resulting FL algorithm. We obtain different asynchronous FL algorithms by using different operators $\mathcal{F}^{(i)}$ in (5.22). For example, to obtain an asynchronous variant of Algorithm 5.1 we can use the choice

$$\mathcal{F}^{(i)}(\mathbf{w}^{(1)}, \dots, \mathbf{w}^{(n)}) = \mathbf{w}^{(i)} - \eta \left(\nabla L_i(\mathbf{w}^{(i)}) + \sum_{i' \in \mathcal{N}^{(i)}} 2A_{i,i'}(\mathbf{w}^{(i)} - \mathbf{w}^{(i')}) \right). \quad (5.23)$$

Note that the choice (5.23) involves the local loss functions and the weighted edges of a FL network.

The update (5.22), at an active node $i \in \mathcal{A}^{(k)}$, involves potentially outdated local model parameters $\mathbf{w}^{(i',k_{i,i'})}$, with $k_{i,i'} \leq k$, for $i' = 1, \dots, n$. We can interpret the differences $k - k_{i,i'}$ as a measure of the communication delay between node i' and node i .

Depending on the extent of the delays $k - k_{i,i'}$ in the update (5.22), we distinguish between [54]

- **totally asynchronous algorithms** where delays $k - k_{i,i'}$ can be arbitrarily large and
- **partially asynchronous algorithms** with delays $k - k_{i,i'} \leq B$ bounded by some constant B .

For some choices of $\mathcal{F}^{(i)}$ in (5.22), a partially asynchronous algorithm can converge for any value of B . However, there also choices of $\mathcal{F}^{(i)}$, for which a partially asynchronous algorithm will only converge if B is sufficiently small [54, Ch. 7].




Fig. 5.4: Illustration of an asynchronous FL algorithm that results in a sequence of events, indexed by $k = 1, 2, \dots$. During event k , each active node $i \in \mathcal{A}^{(i)} \subseteq \mathcal{V}$ of an FL network update their local model parameters $\mathbf{w}^{(i)}$ by computing (5.22). We indicate the active nodes by filled circles.

5.8 Problems

5.1. The converge speed of Gradient-based methods. Study the convergence speed of (5.3) for two different collections of local datasets assigned to the nodes of the FL network \mathcal{G} with nodes $\mathcal{V} = \{1, 2\}$ and (unit weight) edges $\mathcal{E} = \{\{1, 2\}\}$. The first collection of local datasets results in the local loss functions $L_1(w) := (w + 5)^2$ and $L_2(w) := 1000(w + 5)^2$. The second collection of local datasets results in the local loss functions $L_1(w) := 1000(w + 5)^2$ and $L_2(w) := 1000(w - 5)^2$. Use a fixed learning rate $\eta := 0.5 \cdot 10^{-3}$ for the iteration (5.3).

5.2. Converge Speed for Homogeneous Data. Study the convergence speed of (5.3) when applied to GTVMin (5.1) with the following FL network \mathcal{G} : Each node $i = 1, \dots, n$ carries a simple local model with single parameter $w^{(i)}$ and the local loss function $L_i(w) := (y^{(i)} - x^{(i)}w^{(i)})^2$. The local dataset consists of a constant $x^{(i)} := 1$ and some $y^{(i)} \in \mathbb{R}$. The edges \mathcal{E} are obtained by connecting each node i with 4 randomly chosen other nodes. We learn model parameters $\hat{w}^{(i)}$ by repeating (5.3), starting with the initializations $w^{(i,0)} := y^{(i)}$. Study the dependence of the convergence speed of (5.3) (towards a solution of (5.1)) on the value of α in (5.1).

5.9 Proofs

5.9.1 Proof of Proposition 5.1

The first inequality in (5.5) follows from well-known results on the eigenvalues of a sum of symmetric matrices (see, e.g., [3, Thm 8.1.5]). In particular,

$$\lambda_{\max}(\mathbf{Q}) \leq \max \left\{ \underbrace{\max_{i=1,\dots,n} \lambda_d(\mathbf{Q}^{(i)})}_{\stackrel{(5.4)}{=} \lambda_{\max}}, \lambda_{\max}(\alpha \mathbf{L}^{(\mathcal{G})} \otimes \mathbf{I}) \right\}. \quad (5.24)$$

The second inequality in (5.5) uses the following upper bound on the maximum eigenvalue $\lambda_n(\mathbf{L}^{(\mathcal{G})})$ of the Laplacian matrix:

$$\begin{aligned} \lambda_n(\mathbf{L}^{(\mathcal{G})}) &\stackrel{(a)}{=} \max_{\mathbf{v} \in \mathbb{S}^{(n-1)}} \mathbf{v}^T \mathbf{L}^{(\mathcal{G})} \mathbf{v} \\ &\stackrel{(3.8)}{=} \max_{\mathbf{v} \in \mathbb{S}^{(n-1)}} \sum_{\{i,i'\} \in \mathcal{E}} A_{i,i'} (v_i - v_{i'})^2 \\ &\stackrel{(b)}{\leq} \max_{\mathbf{v} \in \mathbb{S}^{(n-1)}} \sum_{\{i,i'\} \in \mathcal{E}} A_{i,i'} 2(v_i^2 + v_{i'}^2) \\ &\stackrel{(c)}{=} \max_{\mathbf{v} \in \mathbb{S}^{(n-1)}} \sum_{i \in \mathcal{V}} 2v_i^2 \sum_{i' \in \mathcal{N}^{(i)}} A_{i,i'} \\ &\stackrel{(3.6)}{\leq} \max_{\mathbf{v} \in \mathbb{S}^{(n-1)}} \sum_{i \in \mathcal{V}} 2v_i^2 d_{\max}^{(\mathcal{G})} \\ &= 2d_{\max}^{(\mathcal{G})}. \end{aligned} \quad (5.25)$$

Here, step (a) uses the CFW of eigenvalues [3, Thm. 8.1.2.] and step (b) uses the inequality $(u+v)^2 \leq 2(u^2+v^2)$ for any $u, v \in \mathbb{R}$. For step (c) we use the identity $\sum_{i \in \mathcal{V}} \sum_{i' \in \mathcal{N}^{(i)}} f(i, i') = \sum_{\{i,i'\}} (f(i, i') + f(i', i))$ (see Figure 5.25). The bound (5.25) is essentially tight.¹⁵

¹⁵Consider a FL network being a chain (or path).




Fig. 5.5: Illustration of step (c) in (5.25).

5.9.2 Proof of Proposition 5.2

Similar to the upper bound (5.25) we also start with the CFW for the eigenvalues of \mathbf{Q} in (5.2). In particular,

$$\lambda_1 = \min_{\|\mathbf{w}\|_2^2=1} \mathbf{w}^T \mathbf{Q} \mathbf{w}. \quad (5.26)$$

We next analyze the RHS of (5.26) by partitioning the constraint set $\{\mathbf{w} : \|\mathbf{w}\|_2^2 = 1\}$ of (5.26) into two complementary regimes for the optimization variable $\mathbf{w} = \text{stack}\{\mathbf{w}^{(i)}\}$. To define these two regimes, we use the orthogonal decomposition

$$\mathbf{w} = \underbrace{\mathbf{P}_{\mathcal{S}} \mathbf{w}}_{=:\bar{\mathbf{w}}} + \underbrace{\mathbf{P}_{\mathcal{S}^\perp} \mathbf{w}}_{=:\tilde{\mathbf{w}}} \text{ for subspace } \mathcal{S} \text{ in (3.16)}. \quad (5.27)$$

Explicit expressions for the orthogonal components $\bar{\mathbf{w}}$, $\tilde{\mathbf{w}}$ are given by (3.17) and (3.18). In particular, the component $\bar{\mathbf{w}}$ satisfies

$$\bar{\mathbf{w}} = ((\mathbf{c})^T, \dots, (\mathbf{c})^T)^T \text{ with } \mathbf{c} := \text{avg}\{\mathbf{w}^{(i)}\}_{i=1}^n. \quad (5.28)$$

Note that

$$\|\mathbf{w}\|_2^2 = \|\bar{\mathbf{w}}\|_2^2 + \|\tilde{\mathbf{w}}\|_2^2. \quad (5.29)$$

Regime I. This regime is obtained for $\|\tilde{\mathbf{w}}\|_2 \geq \rho \|\bar{\mathbf{w}}\|_2$. Since $\|\mathbf{w}\|_2^2 = 1$, and due to (5.29), we have

$$\|\tilde{\mathbf{w}}\|_2^2 \geq \rho^2 / (1 + \rho^2). \quad (5.30)$$

This implies, in turn, via (3.15) that

$$\begin{aligned} \mathbf{w}^T \mathbf{Q} \mathbf{w} &\stackrel{(5.2)}{\geq} \alpha \mathbf{w}^T (\mathbf{L}^{(\mathcal{G})} \otimes \mathbf{I}) \mathbf{w} \\ &\stackrel{(3.8), (3.15)}{\geq} \alpha \lambda_2(\mathbf{L}^{(\mathcal{G})}) \|\tilde{\mathbf{w}}\|_2^2 \\ &\stackrel{(5.30)}{\geq} \alpha \lambda_2(\mathbf{L}^{(\mathcal{G})}) \rho^2 / (1 + \rho^2). \end{aligned} \quad (5.31)$$

Regime II. This regime is obtained for $\|\tilde{\mathbf{w}}\|_2 < \rho \|\bar{\mathbf{w}}\|_2$. Here we have $\|\bar{\mathbf{w}}\|_2^2 > (1/\rho^2)(1 - \|\bar{\mathbf{w}}\|_2^2)$ and, in turn,

$$n \|\mathbf{c}\|_2^2 = \|\bar{\mathbf{w}}\|_2^2 > 1/(1 + \rho^2). \quad (5.32)$$

We next develop the RHS of (5.26) according to

$$\begin{aligned} \mathbf{w}^T \mathbf{Q} \mathbf{w} &\stackrel{(5.2)}{\geq} \sum_{i=1}^n (\mathbf{w}^{(i)})^T \mathbf{Q}^{(i)} \mathbf{w}^{(i)} \\ &\stackrel{(5.27)}{\geq} \sum_{i=1}^n (\mathbf{c} + \tilde{\mathbf{w}}^{(i)})^T \mathbf{Q}^{(i)} (\mathbf{c} + \tilde{\mathbf{w}}^{(i)}) \\ &\stackrel{(5.32)}{\geq} \underbrace{\|\bar{\mathbf{w}}\|_2^2 \lambda_1 \left((1/n) \sum_{i=1}^n \mathbf{Q}^{(i)} \right)}_{\bar{\lambda}_{\min}} + \sum_{i=1}^n \left[2(\tilde{\mathbf{w}}^{(i)})^T \mathbf{Q}^{(i)} \mathbf{c} + \underbrace{(\tilde{\mathbf{w}}^{(i)})^T \mathbf{Q}^{(i)} \tilde{\mathbf{w}}^{(i)}}_{\geq 0} \right] \\ &\geq \|\bar{\mathbf{w}}\|_2^2 \bar{\lambda}_{\min} + \sum_{i=1}^n 2(\tilde{\mathbf{w}}^{(i)})^T \mathbf{Q}^{(i)} \mathbf{c}. \end{aligned} \quad (5.33)$$

To develop (5.33) further, we note that

$$\begin{aligned} \left| \sum_{i=1}^n 2(\tilde{\mathbf{w}}^{(i)})^T \mathbf{Q}^{(i)} \mathbf{c} \right| &\stackrel{(a)}{\leq} 2\lambda_{\max} \|\tilde{\mathbf{w}}\|_2 \|\bar{\mathbf{w}}\|_2 \\ &\stackrel{\|\tilde{\mathbf{w}}\|_2 < \rho \|\bar{\mathbf{w}}\|_2}{\leq} 2\lambda_{\max} \rho \|\bar{\mathbf{w}}\|_2^2. \end{aligned} \quad (5.34)$$

Here, step (a) follows from $\max_{\|\mathbf{y}\|_2=1, \|\mathbf{x}\|_2=1} \mathbf{y}^T \mathbf{Q} \mathbf{x} = \lambda_{\max}$. Inserting (5.34) into (5.33) for $\rho = \bar{\lambda}_{\min}/(4\lambda_{\max})$,

$$\mathbf{w}^T \mathbf{Q} \mathbf{w} \geq \|\bar{\mathbf{w}}\|_2^2 \bar{\lambda}_{\min}/2 \stackrel{(5.32)}{\geq} (1/(1+\rho^2)) \bar{\lambda}_{\min}/2 \quad (5.35)$$

For each \mathbf{w} with $\|\mathbf{w}\|_2^2 = 1$, either (5.31) or (5.35) must hold.

6 Main Flavours of FL

Chapter 3 discussed GTVMin as a main design principle for FL algorithms. GTVMin learns local model parameters that optimally balance the individual local loss with their variation across the edges of an FL network. Chapter 5 discussed how to obtain practical FL algorithms. These algorithms solve GTVMin using distributed optimization methods, such as those from Chapter 4.

This chapter discusses important special cases of GTVMin, known as “main flavours”. These flavours arise for specific construction of local datasets, choices of local models, measures for their variation and, last but not least, the weighted edges in the FL network. We next briefly summarize the resulting main flavours of FL discussed in the following sections.

Section 6.2 discusses single-model FL that learns model parameters of a single (global) model from local datasets. This single-model flavour can be obtained from GTVMin using a connected FL network with large edge weights or, equivalently, a sufficient large value for the GTVMin parameter.

Section 6.3 discusses how clustered FL is obtained from GTVMin over FL networks with a clustering structure. CFL exploits the presence of clusters (subsets of local datasets) which can be approximated using an i.i.d. assumption. GTVMin captures these clusters if they are well-connected by many (large weight) edges in the FL network.

Section 6.4 discusses horizontal FL which is obtained from GTVMin over an FL network whose nodes carry different subsets of a single underlying global dataset. Loosely speaking, horizontal FL involves local datasets characterized by the same set of features but obtained from different data points from an

underlying dataset.

Section 6.5 discusses vertical FL which is obtained from GTVMin over an FL network whose nodes the same data points but using different features. As an example, consider the local datasets at different public institutions (tax authority, social insurance institute, supermarkets) which contain different informations about the same underlying population (anybody who has a Finnish social security number).

Section 6.6 shows how personalized FL can be obtained from GTVMin by using specific measures for the total variation of local model parameters. For example, using deep neural networks as local models, we might only use the model parameters corresponding to the first few input layers to define the total variation.

6.1 Learning Goals

After this chapter, you will know particular design choices for GTVMin corresponding to some main flavours of FL:

- single-model FL
- CFL (generalization of clustering methods)
- horizontal FL (relation to semi-supervised learning)
- personalized FL/multi-task learning
- vertical FL.

6.2 Single-Model FL

Some FL use cases require to train a single (global) model \mathcal{H} from a decentralized collection of local datasets $\mathcal{D}^{(i)}$, $i = 1, \dots, n$ [15, 63]. In what follows we assume that the model \mathcal{H} is parametrized by a vector $\mathbf{w} \in \mathbb{R}^d$. Figure 6.1 depicts a server-client architecture for an iterative FL algorithm that generates a sequence of (global) model parameters $\mathbf{w}^{(k)}$, $k = 1, \dots$. After computing the new model parameters $\mathbf{w}^{(k+1)}$, the server broadcasts it to the clients i . During the next iteration, each client i uses the current global model parameters $\mathbf{w}^{(k)}$ to compute a local update $\mathbf{w}^{(i,k)}$ based on its local dataset $\mathcal{D}^{(i)}$. The precise implementation of this local update step depends on the choice for the global model \mathcal{H} (trained by the server). One example for such a local update has been discussed in Chapter 5 (see (5.17)).

Chapter 5 already hinted at an alternative to the server-based system in Figure 6.1. Indeed, we might learn local model parameters $\mathbf{w}^{(i)}$ for each client i using a distributed optimization of GTVMin. We can force the resulting model parameters $\mathbf{w}^{(i)}$ to be (approximately) identical by using a connected FL network and a sufficiently large GTVMin parameter α .

To minimize the computational complexity of the resulting single-model FL system, we prefer FL networks with small number of edges such as the star graph in Figure 5.3 [58]. However, to increase the robustness against node/link failures we might prefer using an FL network that has more edges. This redundancy helps to ensure that the FL network is connected even after removing some of its edges.

Much like the server-based system from Figure 6.1, GTVMin-based methods using a star graph offers a single point of failure (the server in Figure 6.1

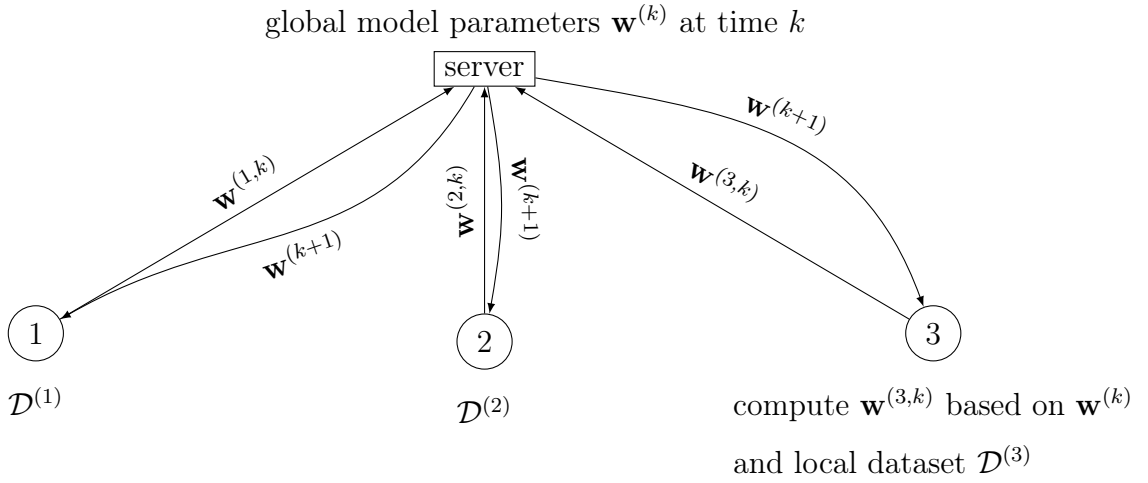


Fig. 6.1: The operation of a server-based (or centralized) FL system during iteration k . First, the server broadcasts the current global model parameters $\mathbf{w}^{(k)}$ to each client $i \in \mathcal{V}$. Each client i then computes the update $\mathbf{w}^{(i,k)}$ by combining the previous model parameters $\mathbf{w}^{(k)}$ (received from the server) and its local dataset $\mathcal{D}^{(i)}$. The updates $\mathbf{w}^{(i,k)}$ are then sent back to the server who aggregates them to obtain the updated global model parameters $\mathbf{w}^{(k+1)}$.

or the centre node in Figure 5.3). Chapter 8 will discuss the robustness of GTVMin-based FL systems in slightly more detail (see Section 8.3).

6.3 Clustered FL

Single-model FL systems require the local datasets to be well approximated as i.i.d. realizations from a common underlying probability distribution. However, requiring homogeneous local datasets, generated from the same probability distribution, might be overly restrictive. Indeed, the local datasets might be heterogeneous and need to be modelled using different probability distribution [18, 30].

CFL relaxes the requirement of a common probability distribution underlying all local datasets. Instead, we approximate subsets of local datasets as i.i.d. realizations from a common probability distribution. In other words, CFL assumes that local datasets form clusters. Each cluster $\mathcal{C} \subseteq \mathcal{V}$ has a cluster-specific probability distribution $p^{(\mathcal{C})}$.

The idea of CFL is to pool the local datasets $\mathcal{D}^{(i)}$ in the same cluster \mathcal{C} to obtain a training set to learn cluster-specific $\hat{\mathbf{w}}^{(\mathcal{C})}$. Each node $i \in \mathcal{C}$ then uses these learnt model parameters $\hat{\mathbf{w}}^{(\mathcal{C})}$. A main challenge in CFL is that the cluster assignments of the local datasets are unknown in general.

To determine a cluster \mathcal{C} , we could apply basic clustering methods, such as k -means or Gaussian mixture model (GMM) to vector representations for local datasets [6, Ch. 5]. We can obtain a vector representation for local dataset $\mathcal{D}^{(i)}$ via the learnt model parameters $\hat{\mathbf{w}}$ of some parametric ML model that is trained on $\mathcal{D}^{(i)}$.

We can also implement CFL via GTVMin with a suitably chosen FL

network. In particular, the FL network should contain many edges (with large weight) between nodes in the same cluster and few edges (with small weight) between nodes in different clusters. To fix ideas, consider the FL network in Figure 6.3, which contains a cluster $\mathcal{C} = \{1, 2, 3\}$.




Fig. 6.2: The solutions of GTVMin (3.20) are local model parameters that are approximately identical for all nodes in a tight-knit cluster \mathcal{C} .

Chapter 3 discussed how the eigenvalues of the Laplacian matrix can be used to measure the connectivity of \mathcal{G} . Similarly, we can measure the connectivity of a cluster \mathcal{C} via the eigenvalue $\lambda_2(\mathbf{L}^{(\mathcal{C})})$ of the Laplacian matrix $\mathbf{L}^{(\mathcal{C})}$ of the induced sub-graph $\mathcal{G}^{(\mathcal{C})}$.¹⁶ The larger $\lambda_2(\mathbf{L}^{(\mathcal{C})})$, the better the connectivity among the nodes in \mathcal{C} . While $\lambda_2(\mathbf{L}^{(\mathcal{C})})$ describes the intrinsic connectivity of a cluster \mathcal{C} , we also need to characterize its connectivity with the other nodes in the FL network. To this end, we use the cluster boundary

$$|\partial[\mathcal{C}]| := \sum_{\{i, i'\} \in \partial\mathcal{C}} A_{i, i'} \text{ with } \partial\mathcal{C} := \{\{i, i'\} \in \mathcal{E} : i \in \mathcal{C}, i' \notin \mathcal{C}\}. \quad (6.1)$$

Note that for a single-node cluster $\mathcal{C} = \{i\}$, the cluster boundary coincides with the node degree, $|\partial[\mathcal{C}]| = d^{(i)}$ (see (3.5)).

¹⁶The graph $\mathcal{G}^{(\mathcal{C})}$ consists of the nodes in \mathcal{C} and the edges $\{i, i'\} \in \mathcal{E}$ for $i, i' \in \mathcal{C}$.

Intuitively, GTVMin tends to deliver (approximately) identical model parameters $\mathbf{w}^{(i)}$ for nodes $i \in \mathcal{C}$ if $\lambda_2(\mathbf{L}^{(\mathcal{C})})$ is large and the cluster boundary $|\partial[\mathcal{C}]|$ is small. The following result makes this intuition more precise for the special case of GTVMin (5.1) for local linear models.

Proposition 6.1. *Consider an FL network \mathcal{G} which contains a cluster \mathcal{C} of local datasets with labels $\mathbf{y}^{(i)}$ and feature matrix $\mathbf{X}^{(i)}$ related via*

$$\mathbf{y}^{(i)} = \mathbf{X}^{(i)} \bar{\mathbf{w}}^{(\mathcal{C})} + \boldsymbol{\varepsilon}^{(i)}, \text{ for all } i \in \mathcal{C}. \quad (6.2)$$

We learn local model parameters $\hat{\mathbf{w}}^{(i)}$ via solving GTVMin (5.1). If the cluster is connected, the error component

$$\tilde{\mathbf{w}}^{(i)} := \hat{\mathbf{w}}^{(i)} - (1/|\mathcal{C}|) \sum_{i \in \mathcal{C}} \hat{\mathbf{w}}^{(i)} \quad (6.3)$$

is upper bounded as

$$\sum_{i \in \mathcal{C}} \|\tilde{\mathbf{w}}^{(i)}\|_2^2 \leq \frac{1}{\alpha \lambda_2(\mathbf{L}^{(\mathcal{C})})} \left[\sum_{i \in \mathcal{C}} \frac{1}{m_i} \|\boldsymbol{\varepsilon}^{(i)}\|_2^2 + \alpha |\partial[\mathcal{C}]| 2 \left(\|\bar{\mathbf{w}}^{(\mathcal{C})}\|_2^2 + R^2 \right) \right]. \quad (6.4)$$

Here, we used $R := \max_{i' \in \mathcal{V} \setminus \mathcal{C}} \|\hat{\mathbf{w}}^{(i')}\|_2$.

Proof. See Sec. 6.9.1. □

The bound (6.4) depends on the cluster \mathcal{C} (via the eigenvalue $\lambda_2(\mathbf{L}^{(\mathcal{C})})$ and the boundary $|\partial[\mathcal{C}]|$) and the GTVMin parameter α . Using a larger \mathcal{C} might result in a decreased eigenvalue $\lambda_2(\mathbf{L}^{(\mathcal{C})})$.¹⁷ According to (6.4), we

¹⁷Consider an FL network (with uniform edge weights) that contains a fully connected cluster \mathcal{C} which is connected via a single edge with another node $i' \in \mathcal{V} \setminus \mathcal{C}$ (see Figure 6.3). Compare the corresponding eigenvalues $\lambda_2(\mathbf{L}^{(\mathcal{C})})$ and $\lambda_2(\mathbf{L}^{(\mathcal{C}')})$ of \mathcal{C} and the enlarged cluster $\mathcal{C}' := \mathcal{C} \cup \{i'\}$.

should then increase α to maintain a small deviation $\widetilde{\mathbf{w}}^{(i)}$ of the learnt local model parameters from their cluster-wise average. Thus, increasing α in (3.20) enforces its solutions to be approximately constant over increasingly larger subsets (clusters) of nodes (see Figure 6.3).

For a connected FL network \mathcal{G} , using a sufficiently large α for GTVMin results in learnt model parameters that are approximately identical for all nodes \mathcal{V} in \mathcal{G} . The resulting approximation error is quantified by Prop. 6.1 for the extreme case where the entire FL network forms a single cluster, i.e., $\mathcal{C} = \mathcal{V}$. Trivially, the cluster boundary is then equal to 0 and the bound (6.4) specializes to (3.33).

We hasten to add that the bound (6.4) only applies for local datasets that conform with the probabilistic model (6.2). In particular, it assumes that all cluster nodes $i \in \mathcal{C}$ have identical model parameters $\overline{\mathbf{w}}^{(\mathcal{C})}$. Trivially, this is no restriction if we allow for arbitrary error terms $\boldsymbol{\varepsilon}^{(i)}$ in the probabilistic model (6.4). However, as soon as we place additional assumptions on these error terms (such as being realizations of i.i.d. Gaussian RVs) we should verify their validity using principled statistical tests [29, 64]. Finally, we might replace $\|\overline{\mathbf{w}}^{(\mathcal{C})}\|_2^2$ in (6.4) with an upper bound for this quantity.




Fig. 6.3: The solutions of GTVMin (3.20) become increasingly clustered for increasing α .

6.4 Horizontal FL

Horizontal FL uses local datasets $\mathcal{D}^{(i)}$, for $i \in \mathcal{V}$, that contain data points characterized by the same features [65]. As illustrated in Figure 6.4, we can think of each local dataset $\mathcal{D}^{(i)}$ as being a subset (or batch) of an underlying global dataset

$$\mathcal{D}^{(\text{global})} := \{(\mathbf{x}^{(1)}, y^{(1)}), \dots, (\mathbf{x}^{(m)}, y^{(m)})\}.$$

In particular, local dataset $\mathcal{D}^{(i)}$ is constituted by the data points of $\mathcal{D}^{(\text{global})}$ with indices in $\{r_1, \dots, r_{m_i}\}$,

$$\mathcal{D}^{(i)} := \{(\mathbf{x}^{(r_1)}, y^{(r_1)}), \dots, (\mathbf{x}^{(r_{m_i})}, y^{(r_{m_i})})\}.$$

We can interpret horizontal FL as a generalization of semi-supervised learning (SSL) [66]: For some local datasets $i \in \mathcal{U}$ we might not have access to the label values of data points. Still, we can use the features of the data points to construct (the weighted edges of) the FL network. To implement SSL, we can solve GTVMin using a trivial loss function $L_i(\mathbf{w}^{(i)}) = 0$ for each unlabelled node $i \in \mathcal{U}$. Solving GTVMin delivers model parameters $\mathbf{w}^{(i)}$




Fig. 6.4: Horizontal FL uses the same features to characterize data points in different local datasets. Different local datasets are constituted by different subsets of an underlying global dataset.

for all nodes i (including the unlabelled ones \mathcal{U}). GTVMin-based methods combine the information in the labelled local datasets $\mathcal{D}^{(i)}$, for $i \in \mathcal{V} \setminus \mathcal{U}$ and their connections (via the edges in \mathcal{G}) with nodes in \mathcal{U} (see Figure 6.4).




Fig. 6.5: Horizontal FL includes SSL as a special case: There is a subset of nodes \mathcal{U} , for which the local datasets do not contain labels. We can take this into account by using the trivial loss function $L_i(\cdot) = 0$ for each node $i \in \mathcal{U}$. However, we can still use the features in $\mathcal{D}^{(i)}$ to construct an FL network \mathcal{G} .

6.5 Vertical FL

Vertical FL uses local datasets that are constituted by the same (identical!) data points. However, each local dataset uses a different choice of features to characterize these data points [67]. Formally, vertical FL applications revolve around an underlying global dataset

$$\mathcal{D}^{(\text{global})} := \{(\mathbf{x}^{(1)}, y^{(1)}), \dots, (\mathbf{x}^{(m)}, y^{(m)})\}.$$

Each data point in the global dataset is characterized by d' features $\mathbf{x}^{(r)} = (x_1^{(r)}, \dots, x_{d'}^{(r)})^T$. The global dataset can only be accessed indirectly via local datasets that use different subsets of the feature vectors $\mathbf{x}^{(r)}$ (see Figure 6.6).

Formally, the local dataset $\mathcal{D}^{(i)}$ contains the pairs $(\mathbf{x}^{(i,r)}, y^{(i,r)})$, for $r = 1, \dots, m$. The labels $y^{(i,r)}$ are identical to the labels in the global dataset, $y^{(i,r)} = y^{(r)}$. The feature vectors $\mathbf{x}^{(i,r)}$ are obtained by a subset $\mathcal{F}^{(i)} :=$

$\{j_1, \dots, j_d\}$ of the original d' features in $\mathbf{x}^{(r)}$,

$$\mathbf{x}^{(i,r)} = (x_{j_1}^{(r)}, \dots, x_{j_d}^{(r)})^T.$$

$$\underbrace{\begin{array}{c} \mathcal{D}^{(1)} \quad \mathcal{D}^{(i)} \\ \left[\begin{array}{cccc|c} x_1^{(1)} & x_2^{(1)} & \cdots & x_d^{(1)} & y^{(1)} \\ x_1^{(2)} & x_2^{(2)} & \cdots & x_d^{(2)} & y^{(2)} \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ x_1^{(m)} & x_2^{(m)} & \cdots & x_d^{(m)} & y^{(m)} \end{array} \right] \end{array}}_{\mathcal{D}^{(\text{global})}}$$

Fig. 6.6: Vertical FL uses local datasets that are derived from the same data points. The local datasets differ in the choice of features used to characterize the common data points.

6.6 Personalized Federated Learning

Consider GTVMin (3.20) for learning local model parameters $\hat{\mathbf{w}}^{(i)}$ for each local dataset $\mathcal{D}^{(i)}$. If the value of α in (3.20) is not too large, the local model parameters $\hat{\mathbf{w}}^{(i)}$ can be different for each $i \in \mathcal{V}$. However, the local model parameters are still coupled via the GTV term in (3.20).

For some FL use-cases we should use different coupling strengths for different components of the local model parameters. For example, if local models are deep ANNs we might enforce the parameters of input layers to be

identical while the parameters of the deeper layers might be different for each local dataset.

The partial parameter sharing for local models can be implemented in many different ways [68, Sec. 4.3.]. One way is to use a choice for the GTV penalty that is different from $\phi = \|\mathbf{w}^{(i)} - \mathbf{w}^{(i')}\|_2^2$. In particular, we could construct the penalty function as a combination of two terms,

$$\phi(\mathbf{w}^{(i)} - \mathbf{w}^{(i')}) := \alpha^{(1)}\phi^{(1)}(\mathbf{w}^{(i)} - \mathbf{w}^{(i')}) + \alpha^{(2)}\phi^{(2)}(\mathbf{w}^{(i)} - \mathbf{w}^{(i')}). \quad (6.5)$$

Each component $\phi^{(1)}, \phi^{(2)}$ measures different components of the variation $\mathbf{w}^{(i)} - \mathbf{w}^{(i')}$ of local model parameters at connected nodes $\{i, i'\} \in \mathcal{E}$.

Moreover, we might use different regularization strengths $\alpha^{(1)}$ and $\alpha^{(2)}$ for different penalty components in (6.5) to enforce different subsets of the model parameters to be clustered with different granularity (cluster size). For local models being deep ANNs, we might want to enforce the low-level layers (closer to the input) to have same model parameters (weights and bias terms), while deeper (closer to the output) layers can have different model parameters. Figure 6.7 illustrates this setting for local models constituted by ANNs with a single hidden layer. Yet another technique for partial sharing of model parameters is to train a hyper-model which, in turn, is used to initialize the training of local models [69].




Fig. 6.7: Personalized FL with local models being ANNs with one hidden layer. The ANN $h^{(i)}$ is parametrized by the vector $\mathbf{w}^{(i)} = \left((\mathbf{u}^{(i)})^T, (\mathbf{v}^{(i)})^T \right)^T$, with parameters $\mathbf{u}^{(i)}$ of hidden layer and the parameters $\mathbf{v}^{(i)}$ of the output layer. We couple the training of $\mathbf{u}^{(i)}$ via GTVMin using the discrepancy measure $\phi = \|\mathbf{u}^{(i)} - \mathbf{u}^{(i')}\|_2^2$.

6.7 Few-Shot Learning

Some maximum likelihood applications involve data points belonging to a large number of different categories. A prime example is the detection of a specific object in a given image [70, 71]. Here, the object category is the label $y \in \mathcal{Y}$ of a data point (image). The label space \mathcal{Y} is constituted by the possible object categories and, in turn, can be quite large. Moreover, for some categories we might have only few example images in the training set.

Few-shot learning leverages similarities between object categories in order to accurately detect objects for which only very few (or even no) training examples are available. One principled approach to few shot learning is via GTVMin. To this end, we define a FL network \mathcal{G} whose nodes $i \in \mathcal{I}$ represent the elements of the label space \mathcal{Y} . The edge weights of \mathcal{G} encode known similarities between different object categories.

Each node i of \mathcal{G} represents a specific object category. Solving GTVMin delivers, for each node i , model parameters $\hat{\mathbf{w}}^{(i)}$ for an object detector which is tailored for the i -th object category.

6.8 Problems

6.1. Horizontal FL of a Linear Model [62, Sec. 8.2] Linear regression learns the model parameters of a linear model by minimizing the average squared error loss on a given dataset \mathcal{D} . Consider an application where the data points are gathered by different devices. We can model such an application using a FL network with nodes i carrying different subsets of \mathcal{D} . Construct an instance of GTVMin such that its solutions coincide (approximately) with

the solution of plain vanilla linear regression.

6.2. Vertical FL of a Linear Model [62, Sec. 8.3] Linear regression learns the model parameters of a linear model by minimizing the average squared error loss on a given dataset \mathcal{D} . Consider an application where the features of a data point are measured by different devices. We can model such an application using a FL network with nodes i carrying different features of the same dataset \mathcal{D} . In particular, node i carries the features x_j with $j \in \mathcal{F}^{(i)}$. Construct an instance of GTVMin such that its solutions coincide (approximately) with the solution of plain vanilla linear regression.

6.9 Proofs

6.9.1 Proof of Proposition 6.1

To verify (6.4), we follow a similar argument as used in the proof (see Section 3.8.1) of Prop. 3.1.

First we decompose the objective function $f(\mathbf{w})$ in (5.1) as follows:

$$\begin{aligned}
 f(\mathbf{w}) = & \underbrace{\sum_{i \in \mathcal{C}} (1/m_i) \|\mathbf{y}^{(i)} - \mathbf{X}^{(i)} \mathbf{w}^{(i)}\|_2^2 + \alpha \left[\sum_{i, i' \in \mathcal{C}} A_{i, i'} \|\mathbf{w}^{(i)} - \mathbf{w}^{(i')}\|_2^2 + \sum_{\{i, i'\} \in \partial \mathcal{C}} A_{i, i'} \|\mathbf{w}^{(i)} - \mathbf{w}^{(i')}\|_2^2 \right]}_{=: f'(\mathbf{w})} \\
 & + f''(\mathbf{w}). \tag{6.6}
 \end{aligned}$$

Note that only the first component f' depends on the local model parameters $\mathbf{w}^{(i)}$ of cluster nodes $i \in \mathcal{C}$. Let us introduce the shorthand $f'(\mathbf{w}^{(i)})$ for the function obtained from $f'(\mathbf{w})$ for varying $\mathbf{w}^{(i)}$, $i \in \mathcal{C}$, but fixing $\mathbf{w}^{(i')} := \widehat{\mathbf{w}}^{(i')}$ for $i' \notin \mathcal{C}$.

We obtain the bound (6.4) via a proof by contradiction: If (6.4) does not hold, the local model parameters $\overline{\mathbf{w}}^{(i)} := \overline{\mathbf{w}}^{(\mathcal{C})}$, for $i \in \mathcal{C}$, result in a smaller value $f'(\overline{\mathbf{w}}^{(i)}) < f'(\widehat{\mathbf{w}}^{(i)})$ than the choice $\widehat{\mathbf{w}}^{(i)}$, for $i \in \mathcal{C}$. This would contradict the fact that $\widehat{\mathbf{w}}^{(i)}$ is a solution to (5.1).

First, note that

$$\begin{aligned}
f'(\bar{\mathbf{w}}^{(i)}) &= \sum_{i \in \mathcal{C}} (1/m_i) \|\mathbf{y}^{(i)} - \mathbf{X}^{(i)} \bar{\mathbf{w}}^{(\mathcal{C})}\|_2^2 \\
&\quad + \alpha \left[\sum_{\substack{\{i, i'\} \in \mathcal{E} \\ i, i' \in \mathcal{C}}} A_{i, i'} \|\bar{\mathbf{w}}^{(\mathcal{C})} - \bar{\mathbf{w}}^{(\mathcal{C})}\|_2^2 + \sum_{\substack{\{i, i'\} \in \mathcal{E} \\ i \in \mathcal{C}, i' \notin \mathcal{C}}} A_{i, i'} \|\bar{\mathbf{w}}^{(\mathcal{C})} - \hat{\mathbf{w}}^{(i')}\|_2^2 \right] \\
&\stackrel{(6.2)}{=} \sum_{i \in \mathcal{C}} (1/m_i) \|\boldsymbol{\varepsilon}^{(i)}\|_2^2 + \alpha \sum_{\substack{\{i, i'\} \in \mathcal{E} \\ i \in \mathcal{C}, i' \notin \mathcal{C}}} A_{i, i'} \|\bar{\mathbf{w}}^{(\mathcal{C})} - \hat{\mathbf{w}}^{(i')}\|_2^2 \\
&\stackrel{(a)}{\leq} \sum_{i \in \mathcal{C}} (1/m_i) \|\boldsymbol{\varepsilon}^{(i)}\|_2^2 + \alpha \sum_{\substack{\{i, i'\} \in \mathcal{E} \\ i \in \mathcal{C}, i' \notin \mathcal{C}}} A_{i, i'} 2 \left(\|\bar{\mathbf{w}}^{(\mathcal{C})}\|_2^2 + \|\hat{\mathbf{w}}^{(i')}\|_2^2 \right) \\
&\leq \sum_{i \in \mathcal{C}} (1/m_i) \|\boldsymbol{\varepsilon}^{(i)}\|_2^2 + \alpha |\partial[\mathcal{C}]| 2 \left(\|\bar{\mathbf{w}}^{(\mathcal{C})}\|_2^2 + R^2 \right). \tag{6.7}
\end{aligned}$$

Step (a) uses the inequality $\|\mathbf{u} + \mathbf{v}\|_2^2 \leq 2(\|\mathbf{u}\|_2^2 + \|\mathbf{v}\|_2^2)$ which is valid for any two vectors $\mathbf{u}, \mathbf{v} \in \mathbb{R}^d$.

On the other hand,

$$\begin{aligned}
f'(\hat{\mathbf{w}}^{(i)}) &\geq \alpha \sum_{i, i' \in \mathcal{C}} A_{i, i'} \underbrace{\|\hat{\mathbf{w}}^{(i)} - \hat{\mathbf{w}}^{(i')}\|_2^2}_{\stackrel{(6.3)}{=} \|\tilde{\mathbf{w}}^{(i)} - \tilde{\mathbf{w}}^{(i')}\|_2^2} \\
&\stackrel{(3.15)}{\geq} \alpha \lambda_2(\mathbf{L}^{(\mathcal{C})}) \sum_{i \in \mathcal{C}} \|\tilde{\mathbf{w}}^{(i)}\|_2^2. \tag{6.8}
\end{aligned}$$

If the bound (6.4) would not hold, then by (6.8) and (6.7) we would obtain $f'(\hat{\mathbf{w}}^{(i)}) > f'(\bar{\mathbf{w}}^{(i)})$, which contradicts the fact that $\hat{\mathbf{w}}^{(i)}$ solves (5.1).

7 Graph Learning

Chapter 3 discussed GTVMin as a main design principle for FL algorithms. In particular, Chapter 5 discusses FL algorithms that arise from the application of optimization methods, such as the gradient-based methods from Chapter 4, to solve GTVMin.

The computational and statistical properties of these algorithms depend crucially on the properties of the underlying FL network. For example, the amount of computation and communication required by FL systems typically grows with the number of edges in the FL network. Moreover, the connectivity of the FL network steers the pooling of local datasets into clusters that share common model parameters.

In some applications, domain expertise can guide the choice for the FL network. However, it might also be useful to learn the FL network in a more data-driven fashion. This chapter discusses methods that learn an FL network solely from a given collection of local datasets and corresponding local loss functions.

The outline of this chapter is as follows: Section 7.2 discusses how the computational and statistical properties of Algorithm 5.1 from Chapter 5 can guide the construction of the FL network. Section 7.3 presents some ideas for how to measure the discrepancy (lack of similarity) between two local datasets.

The discrepancy measure is an important design choice for the graph learning methods discussed in Section 7.4. We formulate these methods as the optimization of edge weights given the discrepancy measure for any pair of local datasets. The formulation as an optimization problem allows to include

connectivity constraints such as a minimum value for each node degree.

7.1 Learning Goals

After completing this chapter, you will

- have intuition for how the computational and statistical properties of GTVMin-based methods depend on the structure of the FL network,
- know some measures for (dis-)similarity (or discrepancy) between local datasets,
- be able to learn a graph from given pairwise similarities and structural constraints, such a prescribed (maximum) node degree.

7.2 Edges as Design Choice

Consider the GTVMin instance (3.22) for learning the model parameters of a local linear model for each local dataset $\mathcal{D}^{(i)}$. To solve (3.22), we use Algorithm 5.1 as a message passing implementation of the basic gradient step (5.3). Note that GTVMin (3.22) is defined for a given FL network \mathcal{G} . Therefore, the choice for \mathcal{G} is critical for the statistical and computational properties of Algorithm 5.1.

Statistical Properties. The statistical properties of Algorithm 5.1 can be assessed via a probabilistic model for the local datasets. One important example for such a probabilistic model is the clustering assumption (6.2) of CFL (see Section 3.4.1). For CFL, we would like to learn similar model parameters for nodes in the same cluster.

According to Prop. 6.1, the GTVMin solutions will be approximately constant over \mathcal{C} , if $\lambda_2(\mathbf{L}^{(\mathcal{C})})$ is large and the cluster boundary $|\partial[\mathcal{C}]|$ is small. Here, $\lambda_2(\mathbf{L}^{(\mathcal{C})})$ denotes the smallest non-zero eigenvalue of the Laplacian matrix associated with the induced sub-graph $\mathcal{G}^{(\mathcal{C})}$.

Roughly speaking, $\lambda_2(\mathbf{L}^{(\mathcal{C})})$ will be larger if there are more edges connecting the nodes in \mathcal{C} . This informal statement can be made precise using a celebrated result from spectral graph theory, known as Cheeger’s inequality [32, Ch. 21]. Alternatively, we can analyse $\lambda_2(\mathbf{L}^{(\mathcal{C})})$ by interpreting (or approximating) the induced sub-graph $\mathcal{G}^{(\mathcal{C})}$ as a typical realization of an Erdős-Rényi (ER) random graph.¹⁸

The ER model for $\mathcal{G}^{(\mathcal{C})}$ postulates that two nodes $i, i' \in \mathcal{C}$ are connected by an edge with probability p_e . The presence of an edge between a given pair of nodes does not depend on the presence of an edge between any other pair of nodes (“edges are i.i.d.”). The absence of dependencies between different edges facilitates the analysis of the ER model [72].

Within the ER model for a graph \mathcal{G} , also the node degrees $d^{(i)}$ as well as the eigenvalue $\lambda_2(\mathbf{L}^{(\mathcal{C})})$ become realizations of RVs. The expected node degree is given as $\mathbb{E}\{d^{(i)}\} = p_e(|\mathcal{C}| - 1)$. With high probability, a realization of an ER graph has maximum node degree

$$d_{\max}^{(\mathcal{G})} \approx \mathbb{E}\{d^{(i)}\} = p_e(|\mathcal{C}| - 1). \quad (7.1)$$

Thus, increasing the ER parameter p_e results in a larger node degree (i.e., a higher connectivity) of $\mathcal{G}^{(\mathcal{C})}$.

¹⁸This approximation is particularly useful if the FL network \mathcal{G} itself is (close to) a typical realization of an ER graph.

We can approximate $\lambda_2(\mathbf{L}^{(c)})$ with the second smallest eigenvalue $\lambda_2(\bar{\mathbf{L}})$ of the expected Laplacian matrix $\bar{\mathbf{L}} := \mathbb{E}\{\mathbf{L}^{(c)}\} = |\mathcal{C}|p_e\mathbf{I} - p_e\mathbf{1}\mathbf{1}^T$. A simple calculation reveals that $\lambda_2(\bar{\mathbf{L}}) = |\mathcal{C}|p_e$. Thus, we have the approximation

$$\lambda_2(\mathbf{L}^{(c)}) \approx \lambda_2(\bar{\mathbf{L}}) = |\mathcal{C}|p_e \stackrel{(7.1)}{\approx} d_{\max}^{(\mathcal{G})}. \quad (7.2)$$

The precise quantification of the approximation error in (7.2) is beyond the scope of this book. We refer the reader to relevant literature on the theory of random graphs [72, 73].

Computational Properties. The computational complexity of Algorithm 5.1 depends on the amount of computation required by a single iteration of its steps (3) - (4). Clearly, this “per-iteration” complexity of Algorithm 5.1 increases with increasing node degrees $d^{(i)}$. Indeed, the step (3) requires to communicate local model parameters across each edge of the FL network. This can be implemented by different physical channels such as a short-range wireless link or an optical fibre cable [74, 75].

To summarize, using an FL network with smaller $d^{(i)}$ translates into a smaller amount of computation and communication needed during a single iteration of Algorithm 5.1. Trivially, the per-iteration complexity of Algorithm 5.1 is minimized by $d^{(i)} = 0$, i.e., an empty FL network without any edges ($\mathcal{E} = \emptyset$). However, the overall computational complexity of Algorithm 5.1 also depends on the number of iterations required to achieve an approximate solution to GTVMin (3.22).

According to (4.6), the convergence speed of the gradient steps (5.9) used in Algorithm 5.1 depends on the condition number $\underbrace{\lambda_{nd}(\mathbf{Q})}_{=\lambda_{\max}(\mathbf{Q})} / \lambda_1(\mathbf{Q})$ of the matrix \mathbf{Q} in (5.2). Algorithm 5.1 tends to require fewer iterations when the

condition number of \mathbf{Q} is small (closer to 1). This condition number, which is the ratio between the largest and smallest eigenvalues, of \mathbf{Q} tends to be smaller for a smaller ratio between the maximum node degree $d_{\max}^{(\mathcal{G})}$ and the eigenvalue $\lambda_2(\mathbf{L}^{(\mathcal{G})})$ (see (5.5) and (5.6)). Thus, for a given maximum node degree $d_{\max}^{(\mathcal{G})}$, we should place the edges in an FL network such that $\lambda_2(\mathbf{L}^{(\mathcal{G})})$ is large and, in turn, Algorithm 5.1 converges faster.

Spectral graph theory also provides upper bounds on $\lambda_2(\mathbf{L}^{(\mathcal{G})})$ in terms of the node degrees [32, 33, 76]. These upper bounds can be used as a baseline for practical constructions of the FL network: If some construction results in a value $\lambda_2(\mathbf{L}^{(\mathcal{G})})$ close to the upper bound, there is little benefit in trying to further improve the construction (in the sense of achieving higher $\lambda_2(\mathbf{L}^{(\mathcal{G})})$). The next result provides one example for such an upper bound.

Proposition 7.1. *Consider an FL network \mathcal{G} with $n > 1$ nodes and associated Laplacian matrix $\mathbf{L}^{(\mathcal{G})}$. Then, $\lambda_2(\mathbf{L}^{(\mathcal{G})})$ cannot exceed the node degree $d^{(i)}$ of any node by more than a factor $n/(n-1)$. In other words,*

$$\lambda_2(\mathbf{L}^{(\mathcal{G})}) \leq \frac{n}{n-1} d^{(i)}, \text{ for every } i = 1, \dots, n. \quad (7.3)$$

Proof. The bound (7.3) follows from (3.12) and evaluating the quadratic form $\mathbf{w}^T \mathbf{L}^{(\mathcal{G})} \mathbf{w}$ for the specific vector

$$\tilde{\mathbf{w}} = \sqrt{\frac{n}{n-1}} \left(- (1/n), \dots, \underbrace{1 - (1/n)}_{\tilde{w}^{(i)}}, \dots, - (1/n) \right)^T.$$

Note that the vector $\tilde{\mathbf{w}}$ is tailored to a specific node $i \in \mathcal{V}$: its only positive entry is $\tilde{w}^{(i)} = 1 - (1/n)$. A basic calculation reveals that $\|\tilde{\mathbf{w}}\| = 1$ and $\tilde{\mathbf{w}}^T \mathbf{1} = 0$, i.e., it is feasible for the optimization in (3.12). \square

Alternative (and potentially tighter) upper bounds on $\lambda_2(\mathbf{L}^{(\mathcal{G})})$ can be found in the graph theory literature [31, 32, 72, 77].

To summarize, the per-iteration complexity of FL algorithms increases with increasing node degrees $d^{(i)}$ (and, in turn, total number of edges) of the FL network \mathcal{G} . On the other hand, the number of iterations required by Algorithm 5.1 will typically decrease with increasing $\lambda_2(\mathbf{L}^{(\mathcal{G})})$. By (7.3), the eigenvalue $\lambda_2(\mathbf{L}^{(\mathcal{G})})$ can only be large if the node degrees $d^{(i)}$ (and, in turn, the total number of edges) are sufficiently large.¹⁹ Fig. 7.1 illustrates the typical dependency of the per-iteration complexity and number of iterations required by FL algorithms.




Fig. 7.1: The per-iteration complexity and number of iterations required by Algorithm 5.1 depends on the number of edges in the underlying FL network in different manners.

¹⁹Some recent work studies graph constructions that maximize $\lambda_2(\mathbf{L}^{(\mathcal{G})})$ for a given (prescribed) maximum node degree $d_{\max}^{(\mathcal{G})} = \max_{i \in \mathcal{V}} d^{(i)}$ [78, 79].

7.3 Measuring (Dis-)Similarity Between Datasets

The main idea behind GTVMin is to enforce similar model parameters at two nodes i, i' that are connected by an edge $\{i, i'\}$ with (relatively) large edge weight $A_{i,i'}$. In general, the edges (and their weights) of the FL network are a design choice. Placing an edge between two nodes i, i' is typically only useful if the local datasets $\mathcal{D}^{(i)}, \mathcal{D}^{(i')}$ (generated by devices i, i') have similar statistical properties. We next discuss different approaches to measuring the similarity or, equivalently, the discrepancy (the lack of similarity) between two local datasets.

The first approach is based on a probabilistic model, i.e., we interpret the local dataset $\mathcal{D}^{(i)}$ as realizations of RVs with some parametrized probability distribution $p^{(i)}(\mathcal{D}^{(i)}; \mathbf{w}^{(i)})$. We can then measure the discrepancy between $\mathcal{D}^{(i)}$ and $\mathcal{D}^{(i')}$ via the Euclidean distance $\|\mathbf{w}^{(i)} - \mathbf{w}^{(i')}\|_2$ between the parameters $\mathbf{w}^{(i)}, \mathbf{w}^{(i')}$ of the probability distributions.

In most FL applications, we do not know the parameters of the probability distribution $p^{(i)}(\mathcal{D}^{(i)}; \mathbf{w}^{(i)})$ underlying a local dataset.²⁰ We might still be able to estimate these parameters using established statistical techniques such as maximum likelihood [6, Ch. 3]. Given the estimates $\hat{\mathbf{w}}^{(i)}, \hat{\mathbf{w}}^{(i')}$ for the model parameters, we can then compute the discrepancy measure $D(i, i') := \|\hat{\mathbf{w}}^{(i)} - \hat{\mathbf{w}}^{(i')}\|_2$.

Example. Consider local datasets being a single number $y^{(i)}$ which is modelled as a noisy observation $y^{(i)} = w^{(i)} + n^{(i)}$ with $n^{(i)} \sim \mathcal{N}(0, 1)$. The maximum likelihood estimator for $w^{(i)}$ is then obtained as $\hat{w}^{(i)} = y^{(i)}$ [27, 80]

²⁰One exception is when we generate the local dataset by drawing i.i.d. realizations from $p^{(i)}(\mathcal{D}^{(i)}; \mathbf{w}^{(i)})$.

and, in turn, the resulting discrepancy measure $D(i, i') := |y^{(i)} - y^{(i')}|$ [81].

Example. Consider an FL network with nodes $i \in \mathcal{V}$ that carry local datasets $\mathcal{D}^{(i)}$. Each $\mathcal{D}^{(i)}$ consists of data points with labels from the label space $\mathcal{Y}^{(i)}$. We can measure the similarity between i, i' by the fraction of data points in $\mathcal{D}^{(i)} \cup \mathcal{D}^{(i')}$ with label values in $\mathcal{Y}^{(i)} \cap \mathcal{Y}^{(i')}$ [82].

Example. Consider local datasets $\mathcal{D}^{(i)}$ constituted by images of handwritten digits $0, 1, \dots, 9$. We model a local dataset using a hierarchical probabilistic model: Each node $i \in \mathcal{V}$ is assigned a deterministic but unknown distribution $\boldsymbol{\alpha}^{(i)} = (\alpha_1^{(i)}, \dots, \alpha_9^{(i)})$. The entry $\alpha_j^{(i)}$ is the fraction of images at node i that show digit j . We interpret the labels $y^{(i,1)}, \dots, y^{(i,m_i)}$ as realizations of i.i.d. RVs, with values in $\{0, 1, \dots, 9\}$ and distributed according to $\boldsymbol{\alpha}^{(i)}$. The features are interpreted as realizations of RVs with conditional distribution $p(\mathbf{x}|y)$ which is the same for all nodes $i \in \mathcal{V}$. We can then estimate the dis-similarity between nodes i, i' via the distance between (estimations of) the parameters $\boldsymbol{\alpha}^{(i)}$ and $\boldsymbol{\alpha}^{(i')}$.

The above discrepancy measure construction (using estimates for the parameters of a probabilistic model) is a special case of a more generic two-step approach:

- First, we determine a vector representation $\mathbf{z}^{(i)} \in \mathbb{R}^{m'}$ for each node $i \in \mathcal{V}$ [6, 83].
- Second, we construct the discrepancy $D(i, i')$ between nodes i, i' via the distance between the representation vectors $\mathbf{z}^{(i)}, \mathbf{z}^{(i')} \in \mathbb{R}^{m'}$, e.g., using $D(i, i') := \|\mathbf{z}^{(i)} - \mathbf{z}^{(i')}\|$.

Let us next discuss three implementations of the first step to obtain the vector

for each node i .

Parametrized Probabilistic Models. If we use a parametrized probabilistic model $p(\mathcal{D}^{(i)}; \mathbf{w}^{(i)})$ for the local dataset $\mathcal{D}^{(i)}$, we can use an estimator $\widehat{\mathbf{w}}^{(i)}$ as $\mathbf{z}^{(i)}$. One popular approach for estimating the model parameters of a probabilistic model is the maximum likelihood principle [6].

Gradients. Let us next discuss a construction for the vector representation $\mathbf{z}^{(i)} \in \mathbb{R}^{m'}$ that is motivated by SGD (see Section 5.4). In particular, we could define the discrepancy between two local datasets by interpreting them as two potential batches used by SGD to train a model. If these two batches have similar statistical properties, then their corresponding gradient approximations (5.11) should be aligned. This suggests to use the gradient $\nabla f(\mathbf{w}')$ of the average loss $f(\mathbf{w}) := (1/|\mathcal{D}^{(i)}|) \sum_{(\mathbf{x}, y) \in \mathcal{D}^{(i)}} L((\mathbf{x}, y), h^{(\mathbf{w})})$ as a vector representation $\mathbf{z}^{(i)}$ for $\mathcal{D}^{(i)}$. We can generalize this construction, for parametrized local models $\mathcal{H}^{(i)}$, by using the gradient of the local loss function,

$$\mathbf{z}^{(i)} := \nabla L_i(\mathbf{v}). \quad (7.4)$$

Note that the construction (7.4) requires a choice for the model parameters \mathbf{v} at which the gradient is evaluated.

Feature Learning. We can also use an autoencoder [83, Ch. 14] to learn a vector representation for a local dataset. In particular, we fed it into an encoder network which has been trained jointly with a decoder network on some learning task. Figure 7.2 illustrates a generic autoencoder setup.




Fig. 7.2: An autoencoder consists of an encoder, mapping the input to a latent vector, and a decoder which tries to reconstruct the input as accurately as possible. The encoder and the decoder are trained jointly by minimizing some quantitative measure (a loss) of the reconstruction error (see [6, Ch. 9]). When using a local dataset as input, we can use the latent vector as its vector representation.

7.4 Graph Learning Methods

Assume we have constructed a useful measure $D(i, i') \in \mathbb{R}_+$ for the discrepancy between any two local datasets $\mathcal{D}^{(i)}, \mathcal{D}^{(i')}$. We could then construct an FL network by connecting each node i with its nearest neighbours. The nearest neighbours of i are those other nodes $i' \in \mathcal{V} \setminus \{i\}$ with smallest discrepancy $D(i, i')$. We next discuss an alternative to the nearest-neighbour graph construction. This alternative approach formulates graph learning as a constrained linear optimization problem.

Let us measure the usefulness of a given choice for the edge weights $A_{i,i'} \in \mathbb{R}_+$ via

$$\sum_{i,i' \in \mathcal{V}} A_{i,i'} D(i, i'). \quad (7.5)$$

The objective function (7.5) penalizes having a large edge weight $A_{i,i'}$ between two nodes i, i' with a large discrepancy $D(i, i')$. Note that the objective function (7.5) is minimized by the trivial choice $A_{i,i'} = 0$, i.e., the empty

graph without any edges $\mathcal{E} = \emptyset$.

As discussed in Section 7.2, the FL network should have a sufficient amount of edges in order to ensure that the GTVMin solutions are useful model parameters. Indeed, the desired pooling effect of GTVMin requires the eigenvalue $\lambda_2(\mathbf{L}^{(\mathcal{G})})$ to be sufficiently large. Ensuring a large $\lambda_2(\mathbf{L}^{(\mathcal{G})})$ requires, in turn, that the FL network \mathcal{G} contains a sufficiently large number of edges and corresponding node degrees (see (7.2)).

We can enforce the presence of edges (with positive weight) by adding constraints to (7.5). For example, we might require

$$A_{i,i} = 0, \quad \sum_{i' \neq i} A_{i,i'} = d_{\max}^{(\mathcal{G})} \text{ for all } i \in \mathcal{V}, \quad A_{i,i'} \in [0, 1] \text{ for all } i, i' \in \mathcal{V}. \quad (7.6)$$

The constraints (7.6) require that each node i is connected with other nodes using total edge weight (weighted node degree) $\sum_{i' \neq i} A_{i,i'} = d_{\max}^{(\mathcal{G})}$.

Combining the constraints (7.6) with the objective function (7.5) results in the following graph learning principle,

$$\begin{aligned} \hat{A}_{i,i'} &\in \operatorname{argmin}_{A_{i,i'} = A_{i',i}, i, i' \in \mathcal{V}} \sum A_{i,i'} D(i, i') \\ A_{i,i'} &\in [0, 1] \text{ for all } i, i' \in \mathcal{V}, \\ A_{i,i} &= 0 \text{ for all } i \in \mathcal{V}, \\ \sum_{i' \neq i} A_{i,i'} &= d_{\max}^{(\mathcal{G})} \text{ for all } i \in \mathcal{V}. \end{aligned} \quad (7.7)$$

Note that (7.7) is a special case of the constrained quadratic minimization problem (4.15). In fact, (7.7) is equivalent to a linear program [41, Sec. 4.3]. We can therefore use projected GD from Section 4.6 to compute (approximate) solutions to (7.7).

The first constraint in (7.7) requires each edge weight to belong to the interval $[0, 1]$. The second constraint prohibits any self-loops in the resulting FL network. Note that adding self-loops to an FL network has no effect on the resulting GTVMin-based method (see (3.20)). The last constraint of the learning principle (7.7) enforces a regular FL network: Each node i having the same weighted node degree $d^{(i)} = \sum_{i' \neq i} A_{i,i'} = d_{\max}^{(g)}$.

For some FL applications it might be detrimental to insist on identical node degrees of the FL network. Instead, we might prefer other structural properties such as a small total number of edges or presence of few hub nodes (with exceptionally large node degree) [13, 81].

We can enforce an upper bound on the total number E_{\max} of edges by modifying the last constraint in (7.7),

$$\begin{aligned} \hat{A}_{i,i'} &\in \operatorname{argmin}_{A_{i,i'} = A_{i',i}} \sum_{i,i' \in \mathcal{V}} A_{i,i'} D(i, i') \\ A_{i,i'} &\in [0, 1] \text{ for all } i, i' \in \mathcal{V}, \\ A_{i,i} &= 0 \text{ for all } i \in \mathcal{V}, \\ \sum_{i', i \in \mathcal{V}} A_{i,i'} &= E_{\max}. \end{aligned} \tag{7.8}$$

The problem has a closed-form solution as explained in [81]: It is obtained by placing the edges between those pairs $i, i' \in \mathcal{V}$ that result in the smallest discrepancy $D(i, i')$. However, it might still be useful to solve (7.8) via iterative optimization methods such as the gradient-based methods discussed in Chapter 4. These methods can be implemented in a fully distributed fashion as message passing over an underlying communication network [62]. This communication network might be significantly different from the learnt

FL network.²¹

7.5 Problems

7.1. A Simple Ranking Approach. Consider a collection of devices $i = 1, \dots, n = 100$, each carrying a local dataset that consists of a single vector $\mathbf{x} \in \mathbb{R}^{(m_i)}$. The vectors $\mathbf{x} \in \mathbb{R}^{m_i}$ can be modelled as statistically independent (across nodes) RVs. Moreover, the vector $\mathbf{x} \in \mathbb{R}^{m_i}$ is a realization of a multivariate normal distribution $\mathcal{N}(c_i \mathbf{1}, \mathbf{I})$ with given (fixed) quantities $c_i \in \{-1, 1\}$. We construct an FL network by determining for each node i its neighbours $\mathcal{N}^{(i)}$ as follows

- we randomly select a fraction $\mathcal{B}^{(i)}$ of 10 percent from all other nodes
- we define $\mathcal{N}^{(i)}$ as those $i' \in \mathcal{B}^{(i)}$ whose corresponding values $|(1/m_i)\mathbf{1}^T \mathbf{x}^{(i)} - (1/m_{i'})\mathbf{1}^T \mathbf{x}^{(i')}|$ are among the 3 smallest.

Analyze the probability that some neighbourhood $\mathcal{N}^{(i)}$ contains a node i' such that $c_i \neq c_{i'}$.

²¹Can you think of FL application domains where the connectivity (e.g., via short-range wireless links) of two clients $i, i' \in \mathcal{V}$ might also reflect the pair-wise similarities between probability distributions of local datasets $\mathcal{D}^{(i)}, \mathcal{D}^{(i')}$?

8 Trustworthy FL

The Story So Far. We have introduced GTVMin as a main design principle for FL in Chapter 3. Chapter 5 applied the gradient-based methods from Chapter 4 to solve GTVMin, resulting in practical FL systems. Our focus has been on computational and statistical properties of these FL systems. In this and the following chapters, we shift the focus from technical properties to the trustworthiness of FL systems.

Section 8.2 reviews key requirements for trustworthy AI, which includes FL systems, that have been put forward by the European Union [84, 85]. We will also discuss how these requirements guide the design choices for GTVMin-based methods. Our focus will be on the three design criteria for trustworthy FL: privacy, robustness and explainability. This chapter covers robustness and explainability, the leakage and protection of privacy in FL systems is the subject of Chapter 9.

Section 8.3 discusses the robustness of FL systems against perturbations of local datasets and computations. A special type of perturbation is the intentional modification or poisoning of local datasets (see Chapter 10). Section 8.4 introduces a measure for the (subjective) explainability of the personalized models trained by FL systems.

8.1 Learning Goals

After completing this chapter, you will

- know some key requirements for trustworthiness of FL
- be familiar with quantitative measures of robustness and explainability

- have some intuition about how robustness, privacy and transparency guides design choices for local models, loss functions and FL network in GTVMin

8.2 Seven Key Requirements by the EU

As part of their artificial intelligence (AI) strategy, the European Commission set up the High-Level Expert Group on Artificial Intelligence (AI HLEG) in 2018. This group put forward seven key requirements for trustworthy AI [84, 85]. We next discuss in step-by-step fashion how these requirements guide the design choices for GTVMin-based FL systems.

8.2.1 KR1 - Human Agency and Oversight.

“..AI systems should support human autonomy and decision-making, as prescribed by the principle of respect for human autonomy. This requires that AI systems should both act as enablers to a democratic, flourishing and equitable society by supporting the user’s agency and foster fundamental rights, and allow for human oversight...” [85, p.15]

Human Dignity. Learning personalized model parameters for recommender systems allows to boost addiction or in widespread emotional manipulation resulting in genocide [86–88]. KR1 might rule out certain design choices for the labels of data points. In particular, we might not use the mental and psychological characteristics of a user as the label. We might also avoid loss functions that can be used to train predictors of psychological characteristics. Using personalized ML models to predict user preferences for products or susceptibility towards propaganda is also referred to as “micro-targeting” [89].

Simple is Good. Human oversight can be facilitated by relying on simple local models. Examples include linear models with few features or decision trees with small depth. However, we are unaware of a widely accepted definition of when a model is simple. Loosely speaking, a simple model results in a learn hypothesis that allows humans to understand how features of a data point relate to the prediction $h(\mathbf{x})$. This notion of simplicity is closely related to the concept of explainability which we discuss in more detail in Section 8.4.

Continuous Monitoring. In its simplest form, GTVMin-based methods involve a single training phase, i.e., learning local model parameters by solving GTVMin. However, this approach is only useful if the data can be well approximated by an i.i.d. assumption. In particular, this approach works only if the statistical properties of local datasets do not change over time. For many FL applications, this assumption is unrealistic (consider a social network which is exposed to constant change of memberships and user behaviour). It is then important to continuously compute a validation error on a timely validation set which is then used, in turn, to diagnose the overall FL system (see [6, Sec. 6.6]).

8.2.2 KR2 - Technical Robustness and Safety.

“...Technical robustness requires that AI systems be developed with a preventative approach to risks and in a manner such that they reliably behave as intended while minimising unintentional and unexpected harm, and preventing unacceptable harm. ...” [85, p.16].

Practical FL systems are obtained by implementing FL algorithms in phys-

ical distributed computers [54, 55]. One example for a distributed computer is a collection of smartphones that are connected either by short-range wireless links or by a cellular network. Most distributed computers will incur imperfections, such as a temporary lack of connectivity or mobile devices becoming inactive due to running out of battery. Moreover, the data generation processes can be subject to perturbations such as statistical anomalies (outliers) or intentional modifications (see Chapter 10). Section 8.3 studies in some detail the robustness of GTVMin-based systems against different perturbations of data sources and imperfections of computational infrastructure.

8.2.3 KR3 - Privacy and Data Governance.

“..privacy, a fundamental right particularly affected by AI systems. Prevention of harm to privacy also necessitates adequate data governance that covers the quality and integrity of the data used...” [85, p.17].

So far, our discussion of FL system used a high level of abstraction using mathematical concepts such as GTVMin and its building blocks. However, to obtain actual FL systems we need implement these mathematical concepts in a given physical hardware. These implementations might incur deviations from the (idealized) GTVMin formulation (3.20) and the gradient-based methods (such as Algorithm 5.1) used to solve it. For example, we might have access only to quantized label values resulting in a quantization error. Moreover, the local datasets might deviate significantly from a typical realizations of i.i.d. RVs (such a situation is referred to as statistical bias [90, Sec. 3.3.]

Also, we might not be able to access all features of a data point due to data processing regulations [91–93]. For example, the general data protection

regulation (GDPR) includes a data minimization principle which requires to only access those features of data points (being users) that are relevant for predicting the label.

Data Governance. FL systems might use local datasets that are generated by human users, i.e., personal data. Whenever personal data is used, special care must be dedicated towards data protection regulations [93]. It might then be useful (or even compulsory) to designate a data protection officer and conduct a data protection impact assessment [85].

Privacy. The operation of a FL system must not violate the fundamental human right for privacy [94]. Chapter 9 discusses quantitative measures and methods for privacy protection in GTVMin-based FL systems.

8.2.4 KR4 - Transparency.

Traceability. This key requirement includes the documentation of design choices (and underlying business models) for a GTVMin-based FL system. This includes the source for the local datasets, the local models, the local loss function as well as the construction of the FL network. Moreover, the documentation should also cover the details of the implemented optimization method used to solve GTVMin. This documentation might also require to periodically store the current model parameters along with a time-stamp (“logging”).

Communication. Depending on the use case, FL systems need to communicate the capabilities and limitations to their end users (e.g., of a digital health app running on a smartphone). For example, we can indicate a measure of uncertainty about the predictions delivered by the trained

local models. Such an uncertainty measure can be obtained naturally from probabilistic model for the data, e.g., the conditional variance of the label y , given the features \mathbf{x} of a random data point. Another example for an uncertainty measure is the validation error of a trained local model.

Explainability. The transparency of a GTVMin-based FL system also includes the explainability of the trained local models. Section 8.4 discusses quantitative measures for the subjective explainability of a learnt hypothesis. We will also use this measure as a regularizer to obtain GTVMin-based systems that guarantee subjective explainability “by design”.

8.2.5 KR5 - Diversity, Non-Discrimination and Fairness.

“...we must enable inclusion and diversity throughout the entire AI system’s life cycle...this also entails ensuring equal access through inclusive design processes as well as equal treatment.” [85, p.18].

The local datasets used for the training of local models should be carefully selected to not enforce existing discrimination. In a health-care application, there might be significantly more training data for patients of a specific gender, resulting in models that perform best for that specific gender at the cost of worse performance for the minority [90, Sec. 3.3.]. Fairness is also important for ML methods used to determine credit score and, in turn, if a loan should be granted or not [95]. Here, we must ensure that ML methods do not discriminate customers based on ethnicity or race. To this end, we could augment data points via modifying any features that mainly reflect the ethnicity or race of a customer (see Figure 8.1).




Fig. 8.1: We can improve fairness of a ML method by augmenting the training set using perturbations of an irrelevant feature such as the gender of a person for which we want to predict the adequate compensation as the label.

8.2.6 KR6 - Societal and Environmental Well-Being.

“...Sustainability and ecological responsibility of AI systems should be encouraged, and research should be fostered into AI solutions addressing areas of global concern, such as for instance the Sustainable Development Goals.” [85, p.19].

Society. FL systems might be used to deliver personalized recommendations to users within a social media application (social network). These recommendations might be (fake) news used to boost polarization and, in the extreme case, social unrest [96].

Environment. Chapter 5 discussed FL algorithms that were obtained by applying gradient-based methods to solve GTVMin. These methods require computational resources to compute local updates for model parameters and to share them across the edges of the FL network. Computation and communi-

cation always requires energy should be generated in an environmental-friendly fashion [97].

8.2.7 KR7 - Accountability.

“...mechanisms be put in place to ensure responsibility and accountability for AI systems and their outcomes, both before and after their development, deployment and use.” [85, p. 19].

8.3 Technical Robustness of FL Systems

Consider a GTVMin-based FL system that aims at training a single (global) linear model in a distributed fashion from a collection of local datasets $\mathcal{D}^{(i)}$, for $i = 1, \dots, n$. As discussed Section 6.2, this single-model FL setting amounts to using GTVMin (3.22) over a connected FL network with a sufficiently large choice for α .

To ensure **KR2** we need to understand the effect of perturbations on a GTVMin-based FL system. These perturbations might be intentional or non-intentional and affect the local datasets used to evaluate the loss of local model parameters or the computational infrastructure used to implement a GTVMin-based method (see Chapter 5). We next explain how to use some of the theoretic tools from previous chapters to quantify the robustness of GTVMin-based FL systems.

8.3.1 Sensitivity Analysis

As pointed out in Chapter 3, GTVMin (3.22) can be rewritten as the minimization of a quadratic function,

$$\min_{\mathbf{w}=\text{stack}\{\mathbf{w}^{(i)}\}} \mathbf{w}^T \mathbf{Q} \mathbf{w} + \mathbf{q}^T \mathbf{w}. \quad (8.1)$$

The matrix \mathbf{Q} and vector \mathbf{q} are determined by the feature matrices $\mathbf{X}^{(i)}$ and label vector $\mathbf{y}^{(i)}$ at nodes $i \in \mathcal{V}$ (see (2.27)). We next study the sensitivity of (the solutions of) (8.1) towards external perturbations of the label vector.²²

Consider an additive perturbation $\tilde{\mathbf{y}}^{(i)} := \mathbf{y}^{(i)} + \boldsymbol{\varepsilon}^{(i)}$ of the label vector $\mathbf{y}^{(i)}$. Using the perturbed label vector $\tilde{\mathbf{y}}^{(i)}$ results also in a “perturbation” of GTVMin (8.1),

$$\min_{\mathbf{w}=\text{stack}\{\mathbf{w}^{(i)}\}} \mathbf{w}^T \mathbf{Q} \mathbf{w} + \mathbf{q}^T \mathbf{w} + \mathbf{n}^T \mathbf{w} + c. \quad (8.2)$$

An inspection of (2.27) yields that $\mathbf{n} = \left((\boldsymbol{\varepsilon}^{(1)})^T \mathbf{X}^{(1)}, \dots, (\boldsymbol{\varepsilon}^{(n)})^T \mathbf{X}^{(n)} \right)^T$. The next result provides an upper bound on the deviation between the solutions of (8.1) and (8.2).

Proposition 8.1. *Consider the GTVMin instance (8.1) for learning local model parameters of a linear model for each node $i \in \mathcal{V}$ of an FL network \mathcal{G} . We assume that the FL network is connected, i.e., $\lambda_2(\mathbf{L}^{(\mathcal{G})}) > 0$ and the local datasets are such that $\bar{\lambda}_{\min} > 0$ (see (5.4)). Then, the deviation between $\hat{\mathbf{w}}^{(i)}$ and the solution $\tilde{\mathbf{w}}^{(i)}$ to the perturbed problem (8.2) is upper bounded as*

$$\sum_{i=1}^n \|\hat{\mathbf{w}}^{(i)} - \tilde{\mathbf{w}}^{(i)}\|_2^2 \leq \frac{\lambda_{\max}(1 + \rho^2)^2}{[\min\{\lambda_2(\mathbf{L}^{(\mathcal{G})})\alpha\rho^2, \bar{\lambda}_{\min}/2\}]^2} \sum_{i=1}^n \|\boldsymbol{\varepsilon}^{(i)}\|_2^2. \quad (8.3)$$

²²Our study can be generalized to also take into account perturbations of the feature matrices $\mathbf{X}^{(i)}$, for $i = 1, \dots, n$.

Here, we used the shorthand $\rho := \bar{\lambda}_{\min}/(4\lambda_{\max})$ (see (5.4)).

Proof. Left as an exercise to the reader. \square

8.3.2 Estimation Error Analysis

Prop. 8.1 characterizes the sensitivity of GTVMin solutions against *external* perturbations of the local datasets. While this notion of robustness is important, it might not suffice for a comprehensive assessment of a FL system. For example, we can trivially achieve perfect robustness (in the sense of minimum sensitivity) by delivering constant model parameters, e.g., $\widehat{\mathbf{w}}^{(i)} = \mathbf{0}$.

Another form of robustness is to ensure a small estimation error of (3.22). To study this form of robustness, we use a variant of the probabilistic model (3.29): We assume that the labels and features of data points of each local dataset $\mathcal{D}^{(i)}$, for $i = 1, \dots, n$, are related via

$$\mathbf{y}^{(i)} = \mathbf{X}^{(i)}\bar{\mathbf{w}} + \boldsymbol{\varepsilon}^{(i)}. \quad (8.4)$$

In contrast to Ch. 3, we assume that all components of (8.4) are deterministic. In particular, the noise term $\boldsymbol{\varepsilon}^{(i)}$ is a deterministic but unknown quantity. This term accommodates any perturbation that might arise from technical imperfections or intrinsic label noise due to random fluctuations in the labelling process.

In the (unlikely) ideal case of no perturbation, we would have $\boldsymbol{\varepsilon}^{(i)} = \mathbf{0}$. However, in general might only know some upper bound measure for the size of the perturbation, e.g., $\|\boldsymbol{\varepsilon}^{(i)}\|_2^2$. We next present upper bounds on the estimation error $\widehat{\mathbf{w}}^{(i)} - \bar{\mathbf{w}}$ incurred by the GTVMin solutions $\widehat{\mathbf{w}}^{(i)}$.

This estimation error consists of two components, the first component being $\text{avg}\{\widehat{\mathbf{w}}^{(i')}\} - \bar{\mathbf{w}}$ for each node $i \in \mathcal{V}$. Note that this error component is identical for all nodes $i \in \mathcal{V}$. The second component of the estimation error is the deviation $\widetilde{\mathbf{w}}^{(i)} := \widehat{\mathbf{w}}^{(i)} - \text{avg}\{\widehat{\mathbf{w}}^{(i')}\}$ of the learnt local model parameters $\widehat{\mathbf{w}}^{(i')}$, for $i' = 1, \dots, n$, from their average $\text{avg}\{\widehat{\mathbf{w}}^{(i')}\} = (1/n) \sum_{i'=1}^n \widehat{\mathbf{w}}^{(i')}$. As discussed in Section 3.4.2, these two components correspond to two orthogonal subspaces of $\mathbb{R}^{d \cdot n}$.

According to Prop. 3.1, the second error component is upper bounded as

$$\sum_{i=1}^n \|\widetilde{\mathbf{w}}^{(i)}\|_2^2 \leq \frac{1}{\lambda_2 \alpha} \sum_{i=1}^n (1/m_i) \|\boldsymbol{\varepsilon}^{(i)}\|_2^2. \quad (8.5)$$

To bound the first error component $\bar{\mathbf{c}} - \bar{\mathbf{w}}$, using the shorthand $\bar{\mathbf{c}} := \text{avg}\{\widehat{\mathbf{w}}^{(i)}\}$, we first note that (see (3.22))

$$\bar{\mathbf{c}} = \underset{\mathbf{w} \in \mathbb{R}^d}{\text{argmin}} \sum_{i \in \mathcal{V}} (1/m_i) \|\mathbf{y}^{(i)} - \mathbf{X}^{(i)}(\mathbf{w} - \widetilde{\mathbf{w}}^{(i)})\|_2^2 + \alpha \sum_{\{i, i'\} \in \mathcal{E}} A_{i, i'} \|\widetilde{\mathbf{w}}^{(i)} - \widetilde{\mathbf{w}}^{(i')}\|_2^2. \quad (8.6)$$

Using a similar argument as in the proof for Prop. 2.1, we obtain

$$\|\bar{\mathbf{c}} - \bar{\mathbf{w}}\|_2^2 \leq \left\| \sum_{i=1}^n (1/m_i) (\mathbf{X}^{(i)})^T (\boldsymbol{\varepsilon}^{(i)} + \mathbf{X}^{(i)} \widetilde{\mathbf{w}}^{(i)}) \right\|_2^2 / (n \bar{\lambda}_{\min})^2. \quad (8.7)$$

Here, $\bar{\lambda}_{\min}$ is the smallest eigenvalue of $(1/n) \sum_{i=1}^n \mathbf{Q}^{(i)}$, i.e., the average of the matrices $\mathbf{Q}^{(i)} = (1/m_i) (\mathbf{X}^{(i)})^T \mathbf{X}^{(i)}$ over all nodes $i \in \mathcal{V}$.²³ Note that the bound (8.7) is only valid if $\bar{\lambda}_{\min} > 0$ which, in turn, implies that the solution to (8.6) is unique.

²³We encountered the quantity $\bar{\lambda}_{\min}$ already during our discussion of gradient-based methods for solving the GTVMin instance (3.22) (see (5.4)).

We can develop (8.7) further using

$$\begin{aligned}
& \left\| \sum_{i=1}^n (1/m_i) (\mathbf{X}^{(i)})^T (\boldsymbol{\varepsilon}^{(i)} + \mathbf{X}^{(i)} \tilde{\mathbf{w}}^{(i)}) \right\|_2 \\
& \stackrel{(a)}{\leq} \sum_{i=1}^n \left\| (1/m_i) (\mathbf{X}^{(i)})^T (\boldsymbol{\varepsilon}^{(i)} + \mathbf{X}^{(i)} \tilde{\mathbf{w}}^{(i)}) \right\|_2 \\
& \stackrel{(b)}{\leq} \sqrt{n} \sqrt{\sum_{i=1}^n \left\| (1/m_i) (\mathbf{X}^{(i)})^T (\boldsymbol{\varepsilon}^{(i)} + \mathbf{X}^{(i)} \tilde{\mathbf{w}}^{(i)}) \right\|_2^2} \\
& \stackrel{(c)}{\leq} \sqrt{n} \sqrt{\sum_{i=1}^n 2 \left\| (1/m_i) (\mathbf{X}^{(i)})^T \boldsymbol{\varepsilon}^{(i)} \right\|_2^2 + 2 \left\| (1/m_i) (\mathbf{X}^{(i)})^T \mathbf{X}^{(i)} \tilde{\mathbf{w}}^{(i)} \right\|_2^2} \\
& \stackrel{(d)}{\leq} \sqrt{n} \sqrt{\sum_{i=1}^n (2/m_i) \lambda_{\max} \left\| \boldsymbol{\varepsilon}^{(i)} \right\|_2^2 + 2 \lambda_{\max}^2 \left\| \tilde{\mathbf{w}}^{(i)} \right\|_2^2}. \tag{8.8}
\end{aligned}$$

Here, step (a) uses the triangle inequality form the norm $\|\cdot\|_2$, step (b) uses the Cauchy-Schwarz inequality, step (c) uses the inequality $\|\mathbf{a} + \mathbf{b}\|_2^2 \leq 2 \left(\|\mathbf{a}\|_2^2 + \|\mathbf{b}\|_2^2 \right)$ and step (d) uses the maximum eigenvalue $\lambda_{\max} := \max_{i \in \mathcal{V}} \lambda_d(\mathbf{Q}^{(i)})$ of the matrices $\mathbf{Q}^{(i)} = (1/m_i) (\mathbf{X}^{(i)})^T \mathbf{X}^{(i)}$ (see (5.4)).

Inserting (8.8) into (8.7) results in the upper bound

$$\begin{aligned}
\|\bar{\mathbf{c}} - \bar{\mathbf{w}}\|_2^2 & \leq 2 \sum_{i=1}^n \left[(1/m_i) \lambda_{\max} \left\| \boldsymbol{\varepsilon}^{(i)} \right\|_2^2 + \lambda_{\max}^2 \left\| \tilde{\mathbf{w}}^{(i)} \right\|_2^2 \right] / (n \bar{\lambda}_{\min}^2) \\
& \stackrel{(8.5)}{\leq} 2 (\lambda_{\max} + (\lambda_{\max}^2 / (\lambda_2 \alpha))) \sum_{i=1}^n (1/m_i) \left\| \boldsymbol{\varepsilon}^{(i)} \right\|_2^2 / (n \bar{\lambda}_{\min}^2). \tag{8.9}
\end{aligned}$$

The upper bound (8.9) on the estimation error of GTVMin-based methods depends on both, the FL network \mathcal{G} via the eigenvalue λ_2 of $\mathbf{L}^{(\mathcal{G})}$, and the feature matrices $\mathbf{X}^{(i)}$ of the local datasets (via the quantities λ_{\max} and $\bar{\lambda}_{\min}$ as defined in (5.4)). Let us next discuss how the upper bound (8.9) might

guide the choice for the FL network \mathcal{G} and the features of data points in the local datasets.

According to (8.9), we should use an FL network \mathcal{G} with large $\lambda_2(\mathbf{L}^{(\mathcal{G})})$ to ensure a small estimation error for GTVMin-based methods. Note that we came across the same design criterion already when discussing graph learning methods in Chapter 7. In particular, using an FL network with large $\lambda_2(\mathbf{L}^{(\mathcal{G})})$ tends also to speed up the convergence of gradient-based methods for solving GTVMin (such as Algorithm 5.1).

The upper bound (8.9) suggests to use features that result in a small ratio $\lambda_{\max}/\bar{\lambda}_{\min}$ between the quantities λ_{\max} and $\bar{\lambda}_{\min}$ (see (5.4)). We might use feature learning methods to optimize this ratio.

8.3.3 Network Resilience

The previous sections studied the robustness of GTVMin-based methods against perturbations of local datasets (see Prop. 8.1) and in terms of ensuring a small estimation error (see (8.9)). We also need to ensure that FL systems are robust against imperfections of the computational infrastructure used to solve GTVMin. These imperfections include hardware failures, running out of battery or lack of wireless connectivity.

Chapter 5 showed how to design FL algorithms by applying gradient-based methods to solve GTVMin (3.22). We obtain practical FL systems by implementing these algorithms, such as Algorithm 5.1, in a particular computational infrastructure. Two important examples for such an infrastructure are mobile networks and wireless sensor networks.

The effect of imperfections in the implementation of the GD based Algo-

rithm 5.1 can be modelled as perturbed GD (4.13) from Chapter 4. We can then analyze the robustness of the resulting FL system via the convergence analysis of perturbed GD (see Section 4.5).

According to (4.14), the performance of the decentralized Algorithm 5.1 degrades gracefully in the presence of imperfections such as missing or faulty communication links. In contrast, the client-server based implementation of FedAvg Algorithm 5.5 offers a single point of failure (the server).

8.3.4 Stragglers

Using perturbed GD to model imperfections in the actual implementation of Algorithm 5.1 is quite flexible. The flexibility in allowing for a wide range of imperfections might come at the cost of a coarse-grained analysis. In particular, the upper bound (4.14) might be too loose (pessimistic) to be of any use. We might then need to take into account the specific nature of the imperfection and the resulting perturbation $\epsilon^{(i)}$. Let us next focus on a particular type of imperfection that arises from the asynchronous implementation of Algorithm 5.1 at different nodes $i \in \mathcal{V}$.

Note that Algorithm 5.1 requires the nodes to operate synchronously: during each iteration the nodes need to exchange their current model parameters $\mathbf{w}^{(i,k)}$ simultaneously with their neighbours in the FL network. This requires, in turn, that the update in step 4 of Algorithm 5.1 is completed at every node i before the global clock ticks (triggering the next iteration).

Some of the nodes i might have limited computational resources and therefore require much longer for the update step 4 of Algorithm 5.1. The literature refers to such slower nodes sometimes as stragglers [98]. Instead of

forcing the other (faster) nodes to wait until also the slower ones are ready, we could instead let them continue with their local updates. This results in an asynchronous variant of Algorithm 5.1 which we summarize in Algorithm 8.1.

Note that, much like the synchronous Algorithm 5.1, also the asynchronous Algorithm 8.1 uses an iteration counter k . However, the practical meaning of k in the asynchronous variant is fundamentally different from the synchronous variant. Instead of indexing a global iteration that corresponds to the synchronous execution of local updates at all nodes \mathcal{V} , the counter k now denotes an event during which some nodes are active, i.e., compute local updates. We denote the set of active nodes (or devices) during event k by $\mathcal{A}^{(k)} \subseteq \mathcal{V}$. The remaining *inactive* nodes $i \notin \mathcal{A}^{(k)}$ leave their current model parameters unchanged, $\mathbf{w}^{(i,k+1)} := \mathbf{w}^{(i,k)}$.

For each *active* node $i \in \mathcal{A}^{(k)}$, the local update (8.10) uses potentially *outdated* model parameters $\mathbf{w}^{(i',k_{i,i'})}$ from its neighbours $i' \in \mathcal{N}^{(i)}$. Indeed, some of the neighbours might have not been in the active sets $\mathcal{A}^{(k-1)}, \mathcal{A}^{(k-2)}, \dots$ of the most recent iterations. In this case, the update (8.10) does not have access to $\mathbf{w}^{(i',k)}$ at *time instant* k . Instead, we can only use $\mathbf{w}^{(i',k_{i,i'})}$ that has been produced obtained during some previous iteration $k_{i,i'} < k$.

We can interpret the quantity $k_{i,i'}$ as the most recent time instant during which node i' has shared its updated local model parameters with node i . We can interpret the difference $k - k_{i,i'}$ as a measure for the communication delay from node i' to node i . The robustness of (the convergence of) Algorithm 8.1 against these communication delays is studied in-depth in [54, Ch. 6 and 7].

Algorithm 8.1 Asynchronous FedGD for Local Linear Models

Input: FL network \mathcal{G} ; GTV parameter α ; learning rate η

local dataset $\mathcal{D}^{(i)} = \{(\mathbf{x}^{(i,1)}, y^{(i,1)}), \dots, (\mathbf{x}^{(i,m_i)}, y^{(i,m_i)})\}$ for each i ; some stopping criterion.

Output: linear model parameters $\widehat{\mathbf{w}}^{(i)}$ for each node $i \in \mathcal{V}$

Initialize: $k := 0$; $\mathbf{w}^{(i,0)} := \mathbf{0}$ for all nodes $i \in \mathcal{V}$.

1: **while** stopping criterion is not satisfied **do**

2: **for** all active nodes $i \in \mathcal{A}^{(k)}$: **do**

3: update local model parameters via

$$\begin{aligned} \mathbf{w}^{(i,k+1)} := & \mathbf{w}^{(i,k)} + \eta \left[(2/m_i) (\mathbf{X}^{(i)})^T (\mathbf{y}^{(i)} - \mathbf{X}^{(i)} \mathbf{w}^{(i,k)}) \right. \\ & \left. + 2\alpha \sum_{i' \in \mathcal{N}^{(i)}} A_{i,i'} (\mathbf{w}^{(i',k_{i,i'})} - \mathbf{w}^{(i,k)}) \right]. \end{aligned} \quad (8.10)$$

4: share local model parameters $\mathbf{w}^{(i,k+1)}$ with neighbours $i' \in \mathcal{N}^{(i)}$

5: **end for**

6: **for** all inactive nodes $i \notin \mathcal{A}^{(k)}$: **do**

7: keep model parameters unchanged $\mathbf{w}^{(i,k+1)} := \mathbf{w}^{(i,k)}$

8: **end for**

9: increment iteration counter: $k := k + 1$

10: **end while**

11: $\widehat{\mathbf{w}}^{(i)} := \mathbf{w}^{(i,k)}$ for all nodes $i \in \mathcal{V}$

8.4 Subjective Explainability of FL Systems

Let us now discuss how to ensure key requirement **KR4 - Transparency** in GTVMin-based FL systems. This key requirement includes also the explainability of a trained personalized model $\hat{h}^{(i)} \in \mathcal{H}^{(i)}$ (and their predictions). It is important to note that the explainability of $\hat{h}^{(i)}$ is subjective: A given learnt hypothesis $\hat{h}^{(i)}$ might offer high degree of explainability to one user (a graduate student at a university) but a low degree of explainability to another user (a high-school student). We must ensure explainability for the specific user, which we will also denote by i , that “consumes” the predictions at node $i \in \mathcal{V}$ of the FL network.

The explainability of trained ML models is closely related to its simulatability [99–101]: How well can a user anticipate (or guess) the prediction $\hat{y} = \hat{h}^{(i)}(\mathbf{x})$ delivered by $\hat{h}^{(i)}$ for a data point with features \mathbf{x} . We can then measure the explainability of $\hat{h}^{(i)}(\mathbf{x})$ to the user at node i by comparing the prediction $\hat{h}^{(i)}(\mathbf{x})$ with the corresponding *guess* (or *simulation*) $u^{(i)}(\mathbf{x})$.

We can enforce (subjective) explainability of FL systems by modifying the local loss functions in GTVMin. For ease of exposition we will focus on the GTVMin instance (5.1) for training local (personalized) linear models. For each node $i \in \mathcal{V}$, we construct a test-set $\mathcal{D}_t^{(i)}$ and ask user i to deliver a guess $u^{(i)}(\mathbf{x})$ for each data point in $\mathcal{D}_t^{(i)}$.²⁴

We measure the (subjective) explainability of a linear hypothesis with

²⁴We only use the features of the data points in $\mathcal{D}_t^{(i)}$, i.e., this dataset can be constructed from unlabeled data.

model parameters $\mathbf{w}^{(i)}$ by

$$(1/|\mathcal{D}_t^{(i)}|) \sum_{\mathbf{x} \in \mathcal{D}_t^{(i)}} \left(u^{(i)}(\mathbf{x}) - \mathbf{x}^T \mathbf{w}^{(i)} \right)^2. \quad (8.11)$$

It seems natural to add this measure as a penalty term to the local loss function in (5.1), resulting in the new loss function

$$L_i(\mathbf{w}^{(i)}) := \underbrace{(1/m_i) \|\mathbf{y}^{(i)} - \mathbf{X}^{(i)} \mathbf{w}^{(i)}\|_2^2}_{\text{training error}} + \underbrace{\rho (1/|\mathcal{D}_t^{(i)}|) \sum_{\mathbf{x} \in \mathcal{D}_t^{(i)}} (u^{(i)}(\mathbf{x}) - \mathbf{x}^T \mathbf{w}^{(i)})^2}_{\text{subjective explainability}}. \quad (8.12)$$

The regularization parameter ρ controls the preference for a high subjective explainability of the hypothesis $h^{(i)}(\mathbf{x}) = (\mathbf{w}^{(i)})^T \mathbf{x}$ over a small training error [101]. It can be shown that (8.12) is the average weighted squared error loss of $h^{(i)}(\mathbf{x})$ on an augmented version of $\mathcal{D}^{(i)}$. This augmented version includes the data point $(\mathbf{x}, u^{(i)}(\mathbf{x}))$ for each data point \mathbf{x} in the test-set $\mathcal{D}_t^{(i)}$.

So far, we have focused on the problem of explaining a trained personalized model to some user. The general idea is to provide partial information, in the form of some explanation, about the learnt hypothesis map \hat{h} . Explanations should help the user to anticipate the prediction $\hat{h}(\mathbf{x})$ for any given data point. Instead of explaining a trained model, i.e., a learnt hypothesis \hat{h} , we might also want to explain an entire FL algorithm.

Mathematically, we can interpret an FL algorithm as a map \mathcal{A} that reads in local datasets and delivers learnt hypothesis maps $\hat{h}^{(i)}$. Explaining an FL algorithm amounts to providing partial information about this map \mathcal{A} . Thus, mathematically speaking, the problem of explaining a learnt hypothesis is essentially the same as explaining an entire FL algorithm: Provide partial

information (explanation) about a map such that user can anticipate the results of applying the map to arbitrary arguments. However, the map \mathcal{A} might be much more complicated compared to a learnt hypothesis (which could be a linear map for linear models). The different level of complexity of these two families of maps requires to use different forms of explanation. For example, we might explain a FL algorithm using a pseudo-code such as Algorithm 5.1. Another form of explanation could be a Python code snippet that illustrates a potential implementation of the algorithm.

```

1 from sklearn.datasets import load_iris
2 from sklearn.model_selection import train_test_split
3 from sklearn.tree import DecisionTreeClassifier
4 from sklearn.metrics import accuracy_score
5
6 # Load the Iris dataset
7 data = load_iris()
8 X = data.data
9 y = data.target
10
11 # Split the dataset into training and test sets
12 X_train, X_test, y_train, y_test = train_test_split(X, y,
13                                                    test_size=0.3, random_state=42)
14
15 # Create a Decision Tree classifier
16 clf = DecisionTreeClassifier(random_state=42)
17
18 # Train the classifier
19 clf.fit(X_train, y_train)
20
21 # Make predictions on the test data
22 y_pred = clf.predict(X_test)
23
24 # Calculate accuracy
25 accuracy = accuracy_score(y_test, y_pred)
26 accuracy

```

Fig. 8.2: Python code for fitting a decision tree model on the Iris dataset.

8.5 Overview of Coding Assignment

Python Notebook. `TrustworthyFL_CodingAssignment.ipynb`

Data File. `Assignment_MLBasicsData.csv`

This coding assignment continues from the coding assignment “FL Algorithms” (see Section ??). This previous assignment required you to implement Algorithm 5.1 and apply it to an FL network whose nodes are FMI weather stations. We now study the effect of adding perturbations $\boldsymbol{\epsilon}^{(i)}$ to the label vector $\mathbf{y}^{(i)}$, consisting of temperature measurements taken at the FMI station i . In particular, we replace the label vector $\mathbf{y}^{(i)}$ with the perturbed label vector $\mathbf{y}^{(i)} + \boldsymbol{\epsilon}^{(i)}$. How much do the local model parameters delivered by Algorithm 5.1 change when adding perturbations of different strength $\|\boldsymbol{\epsilon}^{(i)}\|_2$.

9 Privacy-Protection in FL

The core idea of FL is to share information contained in collections of local datasets in order to improve the training of (personalized) ML models. Chapter 5 discussed FL algorithms that share information in the form of model parameters that are computed from the local loss function. Each node $i \in \mathcal{V}$ receives current model parameters from other nodes and, after executing a local update, shares its new model parameters with other nodes.

Depending on the design choices for GTVMin-based methods, sharing model parameters allows to reconstruct local loss functions and, in turn, to estimate private information about individual data points such as health-care customers (“patients”) [102]. Thus, the bad news is that FL systems will almost inevitably incur some leakage of private information. The good news is, however, that the extent of privacy leakage can be controlled by (i) careful design choices for GTVMin and (ii) applying slight modifications of basic FL algorithms (such as those from Chapter 5).

This chapter revolves around two main questions:

- How can we measure privacy leakage of a FL system?
- How can we control (minimize) privacy leakage of a FL system?

Section 9.2 addresses the first question while Sections 9.3 and 9.4 address the second question.

9.1 Learning Goals

After completing this chapter, you will

- be aware of threats to privacy in FL and the need to protect it,
- know some quantitative measures for privacy leakage,
- understand how GTVMin can facilitate privacy protection,
- be able to implement FL algorithms with guaranteed levels of privacy protection.

9.2 Measuring Privacy Leakage

Consider a FL system that trains a personalized model for the users, indexed by $i = 1, \dots, n$, of heart rate sensors. Each user i generates a local dataset $\mathcal{D}^{(i)}$ that consists of time-stamped heart rate measurements. We define a single data point as a single continuous activity, e.g. as a 50-minute long run. The features of such a data point (activity) might include the trajectory in the form of a time-series of GPS coordinates (e.g., measured every 30 seconds). The label of a data point (activity) could be the average heart rate during the activity. Let us assume that this average heart rate is private information that should not be shared with anybody.²⁵

Our FL system also exploits the information provided by a fitness expert that determines pair-wise similarities $A_{i,i'}$ between users i, i' (e.g., due to body weight and height). We then use Algorithm 5.1 to learn, for each user i , the model parameters $\mathbf{w}^{(i)}$ for some AI health-care assistant [103]. In what follows, we interpret Algorithm 5.1 as a map $\mathcal{A}(\cdot)$ (see Figure 9.1). The map \mathcal{A}

²⁵In particular, we might not want to share our heart rate profiles with a potential future employer who prefers candidates with a long life expectation.




Fig. 9.1: Algorithm 5.1 maps the collection \mathcal{D} of local datasets $\mathcal{D}^{(i)}$ to the learnt model parameters $\hat{\mathbf{w}}^{(i)}$, for each node $i = 1, \dots, n$. These learnt model parameters are (approximate) solutions to the GTVMin instance (3.22).

reads in the dataset $\mathcal{D} := \{\mathcal{D}^{(i)}\}_{i=1}^n$ (constituted by the local datasets $\mathcal{D}^{(i)}$ for $i = 1, \dots, n$) and delivers the local model parameters $\mathcal{A}(\mathcal{D}) := \underbrace{\text{stack}\{\hat{\mathbf{w}}^{(i)}\}_{i=1}^n}_{\hat{\mathbf{w}}}$.

A privacy-preserving FL system should not allow to infer, solely from the learnt model parameters, the average heart rate $y^{(i,r)}$ during a specific single activity r of a specific user i . Mathematically, we must ensure that the map \mathcal{A} is not invertible: The learnt model parameters should not change if we would apply the FL algorithm to a perturbed dataset that includes a different value for the average heart rate $y^{(i,r)}$. Figure 9.2 illustrates an algorithm that is partially invertible in the sense of allowing to infer the label of some data points used in the training set.

The sole requirement for a FL algorithm \mathcal{A} to be not invertible is not useful in general. Indeed, we can easily make any algorithm \mathcal{A} by simple pre- or post-processing techniques whose effect is limited to irrelevant regions of the input space (which is the space of all possible datasets). The level of privacy-protection offered by \mathcal{A} can be characterized by a measure of its “non-invertibility”.




Fig. 9.2: Scatterplot of data points $r = 1, 2, \dots$, each characterized by features $\mathbf{x}^{(r)} = (x_1^{(r)}, x_2^{(r)})^2$ and a binary label $y^{(r)} \in \{\circ, \times\}$. The plot also indicates the decision regions of a hypothesis \hat{h} that has been learnt via ERM. Would you be able to infer the label of data point $r = 2$ if you knew the decision regions?




Fig. 9.3: Two probability distributions of the learnt model parameters $\hat{\mathbf{w}} = \text{stack}\{\hat{\mathbf{w}}^{(i)}\}_{i=1}^n$ delivered by some FL algorithm (such as Algorithm 5.1). These two probability distributions correspond to two different choices for the input dataset, denoted by \mathcal{D}' and \mathcal{D} . For example \mathcal{D}' might be obtained from \mathcal{D} by changing the value of a private feature of some data point in \mathcal{D} . We also indicate an “acceptance region” \mathcal{T} that is used to detect if \mathcal{D} (or a neighbouring dataset \mathcal{D}') has been fed into the algorithm.

A simple measure of non-invertibility is the sensitivity of the output $\mathcal{A}(\mathcal{D})$ against varying the heart rate value $y^{(i,r)}$,

$$\frac{\|\mathcal{A}(\mathcal{D}) - \mathcal{A}(\mathcal{D}')\|_2}{\varepsilon}. \quad (9.1)$$

Here, \mathcal{D} denotes some given collection of local datasets and \mathcal{D}' is a modified dataset. In particular, \mathcal{D}' is obtained by replacing the actual average heart rate $y^{(i,r)}$ with the modified value $y^{(i,r)} + \varepsilon$. The privacy-protection of \mathcal{A} is higher for smaller values (9.1), i.e., the output changes only little when varying the value of the average heart rate.

Another measure for the non-invertibility of \mathcal{A} is referred to as differential privacy (DP). This measure is particularly useful for stochastic algorithms that use some random mechanism. One example for such a random mechanism is the selection of a random subset of data points (a batch) within Algorithm 5.1. Section 9.3 discusses another example of a random mechanism: add the realization of a RV to the (intermediate) results of an algorithm.

A stochastic algorithm \mathcal{A} can be described by a probability distribution $p(\hat{\mathbf{w}}; \mathcal{D})$ over a measurable space that is constituted by the possible values of the learnt model parameters $\hat{\mathbf{w}}$ (see Figure 9.3).²⁶ This probability distribution is parametrized by the dataset \mathcal{D} that is fed as input to the algorithm \mathcal{A} .

DP measures the non-invertibility of a stochastic algorithm \mathcal{A} via the similarity of the probability distributions obtained for two datasets $\mathcal{D}, \mathcal{D}'$ that are considered neighbouring or adjacent [90, 107]. Typically, we consider \mathcal{D}' as adjacent to \mathcal{D} if it is obtained by modifying the features or label of a single data point in \mathcal{D} . As a case in point, consider data points representing

²⁶For more details about the concept of a measurable space, we refer to the literature on probability and measure theory [104–106].

physical activities which are characterized by a binary feature $x_j \in \{0, 1\}$ that indicates an excessively high average heart rate during the activity. We could then define neighbouring datasets via flipping the feature x_j of a single data point. In general, the notion of neighbouring datasets is a design choice used in the formal definition of DP.

Definition 1. (from [107]) *A stochastic algorithm \mathcal{A} is (ε, δ) -DP if for any measurable set \mathcal{S} and any two neighbouring datasets $\mathcal{D}, \mathcal{D}'$,*

$$\text{Prob}\{\mathcal{A}(\mathcal{D}) \in \mathcal{S}\} \leq \exp(\varepsilon)\text{Prob}\{\mathcal{A}(\mathcal{D}') \in \mathcal{S}\} + \delta. \quad (9.2)$$

It appears that Definition 1 is the current de-facto standard for measuring the (lack of) privacy-protection in FL systems [90, 107]. Nevertheless, there are also other measures for the similarity between probability distributions $p(\hat{\mathbf{w}}; \mathcal{D})$ and $p(\hat{\mathbf{w}}; \mathcal{D}')$ that might be more useful for practical or theoretical reasons [108]. One such alternative measure is the Rényi divergence of order $\alpha > 1$,

$$D_\alpha\left(p(\hat{\mathbf{w}}; \mathcal{D}) \parallel p(\hat{\mathbf{w}}; \mathcal{D}')\right) := \frac{1}{\alpha - 1} \mathbb{E}_{p(\hat{\mathbf{w}}; \mathcal{D}')} \left[\left(\frac{dp(\hat{\mathbf{w}}; \mathcal{D})}{dp(\hat{\mathbf{w}}; \mathcal{D}')} \right)^\alpha \right]. \quad (9.3)$$

The Rényi divergence allows to define the following variant of DP [108, 109].

Definition 2. (from [107]) *A stochastic algorithm \mathcal{A} is (α, γ) -RDP if, for any two neighbouring datasets $\mathcal{D}, \mathcal{D}'$,*

$$D_\alpha\left(p(\hat{\mathbf{w}}; \mathcal{D}) \parallel p(\hat{\mathbf{w}}; \mathcal{D}')\right) \leq \gamma. \quad (9.4)$$

A recent use-case of (α, γ) -RDP is the analysis of DP guarantees offered by variants of SGD [108]. This analysis uses the fact that (α, γ) -RDP implies (ε, δ) -DP for suitable choices of ε, δ [108].

One important property of the DP notions in Definition 1 and Definition 2 is that they are preserved by post-processing:

Proposition 9.1. *Consider a FL system \mathcal{A} that is applied to some dataset \mathcal{D} and some (possibly stochastic) map \mathcal{B} that does not depend on \mathcal{D} . If \mathcal{A} is (ε, δ) -DP (or (α, γ) -RDP), then so is also the composition $\mathcal{B} \circ \mathcal{A}$.*

Proof. See, e.g., [107, Sec. 2.3]. □

By Prop. (9.1), we cannot reduce the (level of) DP of \mathcal{A} by any post-processing method \mathcal{B} that has no access to the raw data itself. It seems almost natural to make this “post-processing immunity” a defining property of any useful notion of DP [109]. However, due to Prop. (9.1), this property is already “built-in” into the Definition 1 (and the Definition 2).

Operational Meaning of DP. The above (mathematically precise) definitions of DP might seem somewhat abstract. It is instructive to interpret them from the perspective of statistical testing: We could use the output of an algorithm \mathcal{A} to test (or detect) if the underlying dataset fed into \mathcal{A} was \mathcal{D} or if it actually was a neighbouring dataset \mathcal{D}' [110]. Such a statistical test amounts to specifying a region \mathcal{T} and to declare either

- “dataset \mathcal{D} seems to be used”, if $\mathcal{A} \in \mathcal{T}$, or
- “dataset \mathcal{D}' seems to be used”, if $\mathcal{A} \notin \mathcal{T}$.

The performance of a test \mathcal{T} is characterized by two error probabilities:

- The probability of declaring \mathcal{D}' but actually \mathcal{D} was fed into \mathcal{A} , which is $P_{\mathcal{D} \rightarrow \mathcal{D}'} := 1 - \int_{\mathcal{T}} p(\hat{\mathbf{w}}; \mathcal{D})$.

- The probability of declaring \mathcal{D} but actually \mathcal{D}' was fed into \mathcal{A} , which is

$$P_{\mathcal{D}' \rightarrow \mathcal{D}} := \int_{\mathcal{T}} p(\hat{\mathbf{w}}; \mathcal{D}').$$

For a privacy-preserving algorithm \mathcal{A} , there should be no test \mathcal{T} for which both $P_{\mathcal{D} \rightarrow \mathcal{D}'}$ and $P_{\mathcal{D}' \rightarrow \mathcal{D}}$ are small (close to 0). This intuition can be made precise as follows (see, e.g., [111, Thm. 2.1.] or [112]): If an algorithm \mathcal{A} is (ε, δ) -DP, then

$$\exp(\varepsilon)P_{\mathcal{D} \rightarrow \mathcal{D}'} + P_{\mathcal{D}' \rightarrow \mathcal{D}} \geq 1 - \delta. \quad (9.5)$$

Thus, if \mathcal{A} is (ε, δ) -DP with small ε, δ (close to 0), then (9.5) implies $P_{\mathcal{D} \rightarrow \mathcal{D}'} + P_{\mathcal{D}' \rightarrow \mathcal{D}} \approx 1$.

9.3 Ensuring Differential Privacy

Depending on the underlying design choices (for data, model and optimization method), a GTVMin-based method \mathcal{A} might already ensure DP by design. However, for some design choices the resulting GTVMin-based method \mathcal{A} might not ensure DP. However, according to Prop. 9.1, we might then still be able to ensure DP by applying pre- and/or post-processing techniques to the input (local datasets) and output (learnt model parameters) of \mathcal{A} . Formally, this means to compose the map \mathcal{A} with two (possibly stochastic) maps \mathcal{I} and \mathcal{O} , resulting in a new algorithm with map $\mathcal{A}' := \mathcal{O} \circ \mathcal{A} \circ \mathcal{I}$. The output of \mathcal{A}' for a given dataset \mathcal{D} is obtained by

- first applying the pre-processing $\mathcal{I}(\mathcal{D})$,
- then the original algorithm $\mathcal{A}(\mathcal{I}(\mathcal{D}))$,
- and the final post-processing $\mathcal{O}(\mathcal{A}(\mathcal{I}(\mathcal{D}))) =: \mathcal{A}'(\mathcal{D})$.

Post-Processing. Maybe the most widely used post-processing to ensure DP is simply to add *some noise* [107],

$$\mathcal{O}(\mathcal{A}) := \mathcal{A} + \mathbf{n}, \text{ with noise } \mathbf{n} = (n_1, \dots, n_{nd})^T, n_1, \dots, n_{nd} \stackrel{i.i.d.}{\sim} p(n). \quad (9.6)$$

Note that the post-processing (9.6) is parametrized by the choice for the probability distribution $p(n)$ of the noise entries. Two important choices are the Laplacian distribution $p(n) := \frac{1}{2b} \exp(-\frac{|n|}{b})$ and the norm distribution $p(n) := \frac{1}{\sqrt{2\pi}\sigma^2} \exp(-\frac{n^2}{2\sigma^2})$ (i.e., using Gaussian noise $n \sim \mathcal{N}(0, \sigma^2)$).

When using Gaussian noise $n \sim \mathcal{N}(0, \sigma^2)$ in (9.6), the variance σ^2 can be chosen based on the sensitivity

$$\Delta_2(\mathcal{A}) := \max_{\mathcal{D}, \mathcal{D}'} \|\mathcal{A}(\mathcal{D}) - \mathcal{A}(\mathcal{D}')\|_2. \quad (9.7)$$

Here, the maximum is over all pairs of neighbouring datasets $\mathcal{D}, \mathcal{D}'$. Adding Gaussian noise with variance $\sigma^2 > \sqrt{2 \ln(1.25/\delta)} \cdot \Delta_2(\mathcal{A})/\varepsilon$ ensures that \mathcal{A} is (ε, δ) -DP [107, Thm. 3.22]. It might be difficult to evaluate the sensitivity (9.7) for a given FL algorithm \mathcal{A} [113]. For a GTVMin-based method, i.e., $\mathcal{A}(\mathcal{D})$ is a solution to (3.20), we might obtain upper bounds on $\Delta_2(\mathcal{A})$ by a perturbation analysis similar in spirit to the proof of Prop. 8.1.

Pre-Processing. Instead of ensuring DP via post-processing the output of a FL algorithm \mathcal{A} , we can ensure DP by applying a pre-processing map $\mathcal{I}(\mathcal{D})$ to the dataset \mathcal{D} . The result of the pre-processing is a new dataset $\widehat{\mathcal{D}} = \mathcal{I}(\mathcal{D})$ which can be made available (publicly!) to any algorithm \mathcal{A} that has no direct access to \mathcal{D} . According to Prop. 9.1, as long as the pre-processing map \mathcal{I} is (ε, δ) -DP (see Definition 1), so will be the composition $\mathcal{A} \circ \mathcal{I}$.

As for post-processing, one important approach to pre-processing is to “add” or “inject” noise. This results in a stochastic pre-processing map $\widehat{\mathcal{D}} = \mathcal{I}(\mathcal{D})$

that is characterized by a probability distribution. The noise mechanisms used for pre-processing might be different from just adding the realization of a RV (see (9.6)): ²⁷

- For a classification method with discrete label space $\mathcal{Y} = \{1, \dots, K\}$, we can inject noise by replacing the true label of a data point with a randomly selected element of \mathcal{Y} [114, Mechanism 1]. The noise injection might also include the replacement of the features of a data point by a realization of a RV whose probability distribution is somehow matched to the dataset \mathcal{D} [114, Mechanism 2].
- Another form of noise injection is to construct $\mathcal{I}(\mathcal{D})$ by randomly selecting data points from the original (private) dataset \mathcal{D} [115]. Note that such a form of noise injection is naturally provided by SGD methods (see, e.g., step 4 of Algorithm 5.3).

How To Be Sure? Consider some algorithm \mathcal{A} , possibly obtained by pre- and post-processing techniques, that is claimed to be (ε, δ) -DP. In practice, we might not know the detailed implementation of the algorithm. For example, we might not have access to the noise generation mechanism used in the pre- or post-processing steps. How can we verify a claim about DP of an algorithm \mathcal{A} without having access to the detailed implementation of \mathcal{A} ? One approach could be to apply the algorithm to synthetic datasets $\mathcal{D}_{\text{syn}}^{(1)}, \dots, \mathcal{D}_{\text{syn}}^{(L)}$ that differ only in some private attribute of a single data point. We can then try to predict the private attribute $s^{(r)}$ of the dataset $\mathcal{D}_{\text{syn}}^{(r)}$ by applying a

²⁷Can you think of a simple pre-processing map that is deterministic and guarantees maximum DP?

learnt hypothesis \hat{h} to the output $\mathcal{A}(\mathcal{D}_{\text{syn}}^{(r)})$ delivered by the “algorithm under test” \mathcal{A} . The hypothesis \hat{h} might be learnt by a ERM-based method (see Algorithm 2.1) using a training set consisting of pairs $(\mathcal{A}(\mathcal{D}_{\text{syn}}^{(r)}), s^{(r)})$ for some $r \in \{1, \dots, L\}$.

9.4 Private Feature Learning

Section 9.3 discussed pre-processing techniques that ensure DP of a FL algorithm. We next discuss pre-processing techniques that are not directly motivated from a DP perspective. Instead, we cast privacy-friendly pre-processing of a dataset as a feature learning problem [6, Ch. 9].

Consider a data point characterized by a feature vector $\mathbf{x} \in \mathbb{R}^d$ and a label $y \in \mathbb{R}$. Moreover, each data point is characterized by a private attribute s . We want to learn a (potentially stochastic) feature map $\Phi : \mathbb{R}^d \rightarrow \mathbb{R}^{d'}$ such that the new features $\mathbf{z} = \Phi(\mathbf{x}) \in \mathbb{R}^{d'}$ do not allow to accurately predict the private attribute s . Trivially, we can make prediction of s from $\Phi(\mathbf{x})$ impossible by using a constant map, e.g., $\Phi(\mathbf{x}) = 0$. However, we still want the new features $\mathbf{z} = \Phi(\mathbf{x})$ to allow an accurate prediction (using a suitable hypothesis) for the label y of a data point.

Privacy Funnel. To quantify the predictability of the private attribute s solely from the transformed features $\mathbf{z} = \phi(\mathbf{x})$ we can use the i.i.d. assumption as a simple but useful probabilistic model. Indeed, we can then use the mutual information (MI) $I(s; \Phi(\mathbf{x}))$ as a measure for the predicability of s from $\Phi(\mathbf{x})$. A small value of $I(s; \Phi(\mathbf{x}))$ indicates that it is difficult to predict the private attribute s solely from $\Phi(\mathbf{x})$, i.e., a high level of privacy protection.²⁸ Similarly,

²⁸The relation between MI-based privacy measures and DP has been studied in some

we can use the MI $I(y; \Phi(\mathbf{x}))$ to measure the predicability of the label y from $\Phi(\mathbf{x})$. A large value $I(y; \Phi(\mathbf{x}))$ indicates that $\Phi(\mathbf{x})$ allows to accurately predict y (which is of course preferable).

It seems natural to use a feature map $\Phi(\mathbf{x})$ that optimally balances a small $I(s; \Phi(\mathbf{x}))$ (privacy protection) with a sufficiently large $I(y; \Phi(\mathbf{x}))$ (allowing to accurately predict y). The mathematically precise formulation of this plan is known as the privacy funnel [117, Eq. (2)],

$$\min_{\Phi(\cdot)} I(s; \Phi(\mathbf{x})) \text{ such that } I(y; \Phi(\mathbf{x})) \geq R. \quad (9.8)$$

Figure 9.4 illustrates the solution of (9.8) for varying R , i.e., minimum value of $I(y; \Phi(\mathbf{x}))$.

Optimal Private Linear Transformation. The privacy funnel (9.8) uses the MI $I(s; \Phi(\mathbf{x}))$ to quantify the privacy leakage of a feature map $\Phi(\mathbf{x})$. An alternative measure for the privacy leakage is the minimum reconstruction error $s - \hat{s}$. The reconstruction \hat{s} is obtained by applying a reconstruction map $r(\cdot)$ to the transformed features $\Phi(\mathbf{x})$. If the joint probability distribution $p(s, \mathbf{x})$ is a multivariate normal and the $\Phi(\cdot)$ is a linear map (of the form $\Phi(\mathbf{x}) := \mathbf{F}\mathbf{x}$ with some matrix \mathbf{F}), then the optimal reconstruction map is again linear [27].

We would like to find the linear feature map $\Phi(\mathbf{x}) := \mathbf{F}\mathbf{x}$ such that for any linear reconstruction map \mathbf{r} (resulting in $\hat{s} := \mathbf{r}^T \mathbf{F}\mathbf{x}$) the expected squared error $\mathbb{E}\{(s - \hat{s})^2\}$ is large. The smallest possible expected squared error loss

$$\varepsilon(\mathbf{F}) := \min_{\mathbf{r} \in \mathbb{R}^{d'}} \mathbb{E}\{(s - \mathbf{r}^T \mathbf{F}\mathbf{x})^2\} \quad (9.9)$$

detail recently [116].




Fig. 9.4: The solutions of the privacy funnel (9.8) trace out (for varying constraint R) a curve in the plane spanned by the values of $I(s; \Phi(\mathbf{x}))$ (measuring the privacy leakage) and $I(y; \Phi(\mathbf{x}))$ (measuring the usefulness of the transformed features for predicting the label).

measures the level of privacy protection offered by the new features $\mathbf{z} = \mathbf{F}\mathbf{x}$. The larger the value $\varepsilon(\mathbf{F})$, the more privacy protection is offered. It can be shown that $\varepsilon(\mathbf{F})$ is maximized by any \mathbf{F} that is orthogonal to the cross-covariance vector $\mathbf{c}_{\mathbf{x},s} := \mathbb{E}\{\mathbf{x}s\}$, i.e., whenever $\mathbf{F}\mathbf{c}_{\mathbf{x},s} = \mathbf{0}$. One specific choice for \mathbf{F} that satisfies this orthogonality condition is

$$\mathbf{F} = \mathbf{I} - (1/\|\mathbf{c}_{\mathbf{x},s}\|_2^2)\mathbf{c}_{\mathbf{x},s}\mathbf{c}_{\mathbf{x},s}^T. \quad (9.10)$$

Figure 9.5 illustrates a dataset for which we want to find a linear feature map \mathbf{F} such that the new features $\mathbf{z} = \mathbf{F}\mathbf{x}$ do not allow to accurately predict a private attribute.



Fig. 9.5: A toy dataset \mathcal{D} whose data points represent customers, each characterized by features $\mathbf{x} = (x_1, x_2)^T$. These raw features carry information about a private attribute s (gender) and the label y (food preference) of a person. The scatter-plot suggests that we can find a linear feature transformation $\mathbf{F} := \mathbf{f}^T \in \mathbb{R}^{1 \times 2}$ resulting in a new feature $z := \mathbf{F}\mathbf{x}$ that does not allow to predict s , while still allowing to predict y .

9.5 Problems

9.1. Where is *Alice*? Consider a device, named *Alice*, that implements the asynchronous Algorithm 8.1. The local dataset of the device are temperature measurements from some FMI weather station. Assuming that no other device interacts with *Alice* except your device, named *Bob*. Develop a software for *Bob* that interacts with *Alice* according to Algorithm 8.1 in order to determine at which FMI station we can find *Alice*.

9.2. Linear Discriminant Analysis with Privacy Protection. Consider binary classification problem with data points characterized by a feature vector $\mathbf{x} \in \mathbb{R}^d$ and binary label $y \in \{-1, 1\}$. Each data point has a sensitive attribute $s = \mathbf{F}\mathbf{x}$, obtained by applying a fixed matrix \mathbf{F} to the feature vector \mathbf{x} . We use a probabilistic model - interpreting data points (\mathbf{x}, y) as i.i.d. realizations of a RV - with the feature vector having multivariate normal distribution $\mathcal{N}(\mu^{(y)}, \mathbf{C}^{(y)})$ conditioned on y . The label is uniformly distributed over the label space $\{-1, 1\}$. Try to find a vector \mathbf{a} such that the transformed feature vector $z' := \mathbf{a}^T \mathbf{x}$ optimally balances the privacy leakage (information carried by z' about s) with the information carried by z' about the label y .

9.6 Overview of Coding Assignment

Python Notebook. `PrivacyProtectionFL_CodingAssignment.ipynb`

Data File. `Assignment_MLBasicsData.csv`

This coding assignment builds on the coding assignment “ML Basics” (see Section ??). Similar to this previous coding assignment, we consider data points, indexed by $r = 1, \dots, m$, that represent weather measurements recorded by the FMI. The r -th data point is characterized by (normalized) coordinates, a time stamp and a temperature measurement.

The overall idea of the assignment is to illustrate a model inversion attack: An attacker tries to determine the private attribute of an individual by evaluating a trained model (learnt hypothesis). Here, the private attribute is the latitude and longitude of a person that shares a snapshot of a thermometer during a Finland journey. The attacker might then use a trained ML model to predict the latitude and longitude from the disclosed (public) attributes \mathbf{x} consisting of the time-stamp and temperature value.

9.6.1 Where Are You?

Consider a social media post of a friend that is travelling across Finland. This post includes a snapshot of a temperature measurement and a clock. Can you guess the latitude and longitude of the location where your friend took this snapshot? We can use ERM to do this: Use Algorithm 2.1 to learn a vector-valued hypothesis \hat{h} for predicting latitude and longitude from the time and value of a temperature measurement. For the training set and validation set, we use the FMI recordings stored in the data file.

9.6.2 Ensuring Privacy with Pre-Processing

Repeat the privacy attack described in Section 9.6.1 but this time using a pre-processed version of the raw data. In particular, try out combinations of randomly selecting a subset of the data points in the data file and also adding noise to their features and label. How well can you predict the latitude and longitude from the time and value of a temperature measurement using a hypothesis $\hat{\mathbf{h}}$ learnt from the perturbed data.

9.6.3 Ensuring Privacy with Post-Processing

Repeat the privacy attack described in Section (9.6.1) but this time using a post-processing of the learnt hypothesis $\hat{\mathbf{h}}$ (obtained from Algorithm 2.1 applied to the data file). In particular, study how well you can predict the latitude and longitude from the time and value of a temperature measurement using a noisy prediction hypothesis $\hat{\mathbf{h}}(\mathbf{x}) + \mathbf{n}$. Here, \mathbf{n} is a realization drawn from a multivariate normal distribution $\mathcal{N}(\mathbf{0}, \sigma^2 \mathbf{I})$.

9.6.4 Private Feature Learning

We now use the original features $\mathbf{x} \in \mathbb{R}^7$ of a FMI weather recording as we used in the “ML Basics” assignment (see Section ??) The goal of this task is to learn a linear feature transformation $\mathbf{z} = \mathbf{F}\mathbf{x}$ such that the new features do not allow to recover the latitude x_1 (which is considered a private attribute s of the data point). In particular, we construct the matrix \mathbf{F} according to (9.10) by replacing the exact cross-covariance vector $\mathbf{c}_{\mathbf{x},s}$ with an estimate (or approximation) $\hat{\mathbf{c}}_{\mathbf{x},s}$. This estimate is computed as follows:

1. read in all data points stored in the data file and construct a feature matrix $\mathbf{X} \in \mathbb{R}^{m \times 7}$ with m being the total number of data points
2. remove the sample means from each feature, resulting in the centered feature matrix

$$\hat{\mathbf{X}} := \mathbf{X} - (1/m)\mathbf{1}\mathbf{1}^T\mathbf{X}, \mathbf{1} := (1, \dots, 1)^T \in \mathbb{R}^m. \quad (9.11)$$

3. extract the private attribute (centred normalized latitude) for each data point and store in the vector

$$\mathbf{s} := (\hat{x}_1^{(1)}, \hat{x}_1^{(2)}, \dots, \hat{x}_1^{(m)})^T. \quad (9.12)$$

4. compute the approximate cross-covariance vector

$$\hat{\mathbf{c}}_{\mathbf{x},s} := (1/m)(\hat{\mathbf{X}})^T\mathbf{s} \quad (9.13)$$

The matrix \mathbf{F} obtained from (9.10) by replacing $\mathbf{c}_{\mathbf{x},s}$ with $\hat{\mathbf{c}}_{\mathbf{x},s}$, is then used to compute the privacy preserving features $\mathbf{z}^{(r)} = \mathbf{F}\mathbf{x}^{(r)}$ for $r = 1, \dots, m$. To verify if these new features are indeed privacy preserving, we use linear regression (as implemented by the `LinearRegression` class of the Python package `scikit-learn`) to learn the model parameters of a linear model to predict the private attribute $s^{(r)} = x_1^{(r)}$ (the latitude of the FMI station at which the r -th measurement has been taken) from the features $\mathbf{z}^{(r)}$. (Note that you have already trained and validated a linear model within the “ML Basics” assignment.) We can use the validation error of the trained linear model as a measure for the privacy protection (larger validation error means higher privacy protection). Beside the privacy protection, the new features

$\mathbf{z}^{(r)}$ should still allow to accurately predict the label $y^{(r)}$ of a data point. To this end, we learn the model parameters of a linear model to predict the private property $y^{(r)}$ solely from the features $\mathbf{z}^{(r)}$.

10 Data and Model Poisoning in FL

Every ML method, including ERM or GTVMin, is to some extent at the mercy of the data generator. Indeed, the model parameters learnt via ERM (for basic ML) or via GTVMin (for FL) are determined by the statistical properties of the training set. We must hope (e.g., via an i.i.d. assumption) that the data points in the training set truthfully reflect the statistical properties of the underlying data generation process. However, these data points might have been intentionally perturbed (or *poisoned*).

In general, it is impossible to perfectly detect if data points have been poisoned. The perfect detection of those perturbed data points requires complete knowledge of the underlying probability distribution. However, we typically do not know this probability distribution but can only estimate it from (possibly perturbed) data points. We can then use this estimate to detect perturbed data points via outlier detection techniques.

Instead of trying to identify and remove poisoned data points we can also try to make GTVMin-based FL systems more robust against data poisoning. As discussed in Section 8.3, the level of robustness crucially depends on the design choices for the local models, local loss functions and FL network in GTVMin.

Besides data poisoning, FL systems can also be subject to model poisoning attacks. FL systems involve message passing between devices in order to implement the training of local models (see Section 5). For some applications it is not feasible (or desirable) to enforce sophisticated authentication techniques to ensure only benign devices participate in the message passing. Some messages might then be manipulated (poisoned) in order to disturb the

training of some local models.

10.1 Learning Goals

This chapter discusses the robustness of FL systems against data poisoning and model poisoning attacks. After completing this chapter, you will

1. be aware of data poisoning and model poisoning attacks at FL systems
2. know how design choices for GTVMin-based methods affect the vulnerability of FL systems against these attacks.

10.2 Attack Types

Consider a FL system that implements Algorithm 5.1 over a computer network. Each computer corresponds to one node of the FL network, and implements the local update in step 4 of Algorithm 5.1. The model parameters sharing in step 3 of Algorithm 5.1 is implemented over some communication channel (e.g., short-range wireless links).

An attack on such a FL system could be carried out in different forms, depending on the level of control that the attacker has over the implementation of Algorithm 5.1. If the attacker has some control over the communication links between the nodes, it can directly manipulate the model parameters shared between nodes (model poisoning). The attacker might instead have only access to the local datasets of some (vulnerable) nodes $\mathcal{W} \subset \mathcal{V}$. It could then manipulate (poison) the local datasets at these vulnerable nodes to perturb the corresponding local updates (see step 4 of Algorithm 5.1). The perturbations at the nodes $i' \in \mathcal{W}$ then propagate over the edges of the FL

network (via the model parameters sharing step 3 in Algorithm 5.1) and result in perturbed model parameters at other nodes (whose local datasets have not been poisoned).

Based on the objective of an attack, we distinguish between the following attack types:

- Denial-of-Service Attacks.** The goal of a denial-of-service attack is to make the learnt hypothesis $\bar{h}^{(i)}$, at some target node i , useless in the sense of having unacceptable large prediction errors. We can detect a denial-of-service attack by continuously monitoring the performance of the learnt model parameters $\mathbf{w}^{(i)}$. Denial-of-service attacks on GTVMin-based FL systems can be launched by manipulating some of the local datasets at few nodes in the FL network. These manipulated (poisoned) local datasets influence the model parameters at the target node i indirectly via the sharing of model parameters across the edges of the FL network. Instead of poisoning local datasets, denial-of-service attacks might manipulate the sharing (communication) of model parameters across the edges in the FL network. Such model poisoning is possible when the sharing of model parameters is carried out over in-secure communication links. Figure 10.1 illustrates, for some (target) node i , the result of a denial-of-service attack on a FL system.
- Backdoor Attacks.** A backdoor attack tries to make a target node i learn a hypothesis $\tilde{h}^{(i)}$ that behaves well on the local dataset $\mathcal{D}^{(i)}$ but highly irregular for a certain range of feature values. The goal of the attacker is to exploit this irregular behaviour by preparing a data point with feature values falling in this range such that the hypothesis $\tilde{h}^{(i)}$

delivers a specific prediction (e.g., a prediction that results in granting access to a restricted area within a building). Figure 10.1 illustrates, for some (target) node i , the result of a backdoor attack on a FL system.

- **Privacy Attacks (see Chapter 9).** The goal of a privacy attack is to determine private attributes of the data points in the local dataset at some target node i . To this end, an attacker might try to enforce some other (vulnerable) node i' to learn a copy of the model parameters $\mathbf{w}^{(i)}$ at node i . For GTVMin-based methods, this could be achieved by using a trivial local loss function $L_{i'}()$ (e.g., being identically equal to zero) and having edges between i' and nodes that are in the same cluster as node i (see Section 6.3). After the attacker has obtained a copy of the learned model parameters $\widehat{\mathbf{w}}^{(i)}$, they can try to probe the resulting hypothesis in order to infer private attributes of the data points in $\mathcal{D}^{(i)}$ (see Figure 9.2).




Fig. 10.1: A local dataset $\mathcal{D}^{(i)}$ along with three hypothesis maps learnt via some GTVMin-based method such as Algorithm 5.1. These three maps are obtained for different attacks to the FL system: The map $\hat{h}^{(i)}$ is obtained when no manipulation is applied (no attacks). A backdoor attack aims at nudging the learnt hypothesis $\tilde{h}^{(i)}$ to behave similar to $\hat{h}^{(i)}$ when applied to data points in $\mathcal{D}^{(i)}$. However, it behaves very different for a certain range of feature values outside of $\mathcal{D}^{(i)}$. This value range can be interpreted as a *backdoor* that is used to trigger a malicious prediction. In contrast to backdoor attacks, a denial-of-service attack aims at enforcing a learnt hypothesis $\bar{h}^{(i)}$ that delivers poor predictions on the local dataset $\mathcal{D}^{(i)}$.

10.3 Data Poisoning

Consider a GTVMin-based FL system that trains local models for device $i = 1, \dots, n$ based on their local datasets. A data poisoning attack on such a FL system consists of the manipulation (*poisoning*) of (some of) the local datasets.

The manipulation (poisoning) of data points can consist of adding the realization of RVs to the features and label of a data point: We poison a data point by replacing its features \mathbf{x} and label y with $\tilde{\mathbf{x}} := \mathbf{x} + \Delta\mathbf{x}$ and $\tilde{y} = y + \Delta y$.

For classification problems, with discrete label spaces, we distinguish between the following data poisoning strategies [118]:

- Label Poisoning: The attacker manipulates the labels of data points in the training set.
- Clean-Label Attack: The attacker leaves the labels untouched and only manipulates the features of data points in the training set.

From a GTVMin-perspective, the effect of a data poisoning attack is that the original local loss functions $L_i(\cdot)$ in GTVMin (3.22) are replaced by perturbed local loss functions $\tilde{L}_i(\cdot)$. The extend of perturbation depends on the fraction of data points that are poisoned as well as on the loss function used to measure the prediction errors.

Different choices for the underlying loss function offer different levels of robustness against poisoning. For example, using the absolute error loss yields higher robustness against perturbations of the label values of few data points

compared to the squared error loss. Another class of robust loss functions is obtained by including a penalty term (as in regularization).

10.4 Model Poisoning

A model poisoning attack on Algorithm 5.1 manipulates the model parameters sharing step such that a target node i receives perturbed model parameters of its neighbours. Note that Algorithm 5.1 is nothing but a message passing implementation of plain GD (see Section 4.2). Thus, the effect of a model poisoning attack on Algorithm 5.1 is that it becomes an instance of perturbed GD. We can use the analysis of perturbed GD (see Section 4.5) to study the impact of model poisoning on the learnt model parameters $\mathbf{w}^{(i)}$.

10.5 Assignment

Python Notebook. `PoisoningFL_CodingAssignment.ipynb`

Data File. `Assignment_MLBasicsData.csv`

This coding assignment builds on the coding assignments “ML Basics” (see Section ??) and “FL Algorithm” (see Section ??). As in these previous assignments, we consider data points, indexed by $r = 1, \dots, m$, that represent weather measurements recorded by the FMI. The r -th data point is characterized by (normalized) coordinates, a time stamp and a temperature measurement.

The overall idea of the assignment is to illustrate two different types of data poisoning attacks: a denial-of-service attack and a backdoor attack.

10.5.1 Denial-of-Service Attack

Construct an FL network of FMI stations and store it as a `networkx.Graph()` object. Implement Algorithm 5.1 to learn, for each node $i = 1, \dots, n$, the model parameters of local linear model. Then, implement a denial-of-service attack by poisoning the local datasets as increasingly many nodes $i' \neq 1$. The goal of the attack is to increase the validation error of the learnt model parameters $\mathbf{w}^{(1)}$ (at target node $i = 1$) by 20 %.

10.5.2 Backdoor Attack

We now use a different collection of features for a data point (= temperature recording). In particular, we replace the numeric feature representing the hour of the measurement with 24 new features. These new features are the one-hot encoding of the hour. For example if the hour was 0 then $x'_1 = 1, x'_2 = 0, \dots$

Glossary

k -means The k -means algorithm is a hard clustering method which assigns each data point of a dataset to precisely one of k different clusters. The method alternates between updating the cluster assignments (to the cluster with nearest cluster mean) and, given the updated cluster assignments, re-calculating the cluster means [6, Ch. 8].. 7, 24, 108

accuracy Consider data points characterized by features $\mathbf{x} \in \mathcal{X}$ and a categorical label y which takes on values from a finite label space \mathcal{Y} . The accuracy of a hypothesis $h : \mathcal{X} \rightarrow \mathcal{Y}$, when applied to the data points in a dataset $\mathcal{D} = \{(\mathbf{x}^{(1)}, y^{(1)}), \dots, (\mathbf{x}^{(m)}, y^{(m)})\}$ is then defined as $1 - (1/m) \sum_{r=1}^m L((\mathbf{x}^{(r)}, y^{(r)}), h)$ using the 0/1 loss.. 23

activation function Each artificial neuron within an ANN is assigned an activation function $g(\cdot)$ that maps a weighted combination of the neuron inputs x_1, \dots, x_d to a single output value $a = g(w_1x_1 + \dots + w_dx_d)$. Note that each neuron is parameterized by the weights w_1, \dots, w_d .. 17

artificial intelligence Artificial intelligence refers to systems that behave rational in the sense of maximizing a long-term reward. The ML-based approach to AI is to train a model that allows to predict optimal actions for a given observed state of the environment. What sets AI applications apart from more basic ML applications is the choice of loss function. AI systems rarely have access to a labeled training set that allows to measure the average loss for a given choice of model parameters. Rather, AI systems typically use a loss function that can only be estimated from

observed reward signals.. 8, 136

artificial neural network An artificial neural network is a graphical (signal-flow) representation of a map from features of a data point at its input to a prediction for the label as its output.. 1, 11, 13, 17, 23, 24, 36, 46, 49, 95, 115–117

autoencoder An autoencoder is a ML method that jointly learns an encoder map $h(\cdot) \in \mathcal{H}$ and a decoder map $h^*(\cdot) \in \mathcal{H}^*$. It is an instance of ERM using a loss computed from the reconstruction error $\mathbf{x} - h^*(h(\mathbf{x}))$.. 130, 131

backdoor A backdoor attack refers to the intentional manipulation of the training process underlying a ML method. This manipulation can be implemented by perturbing the training set (data poisoning) or the optimization algorithm used by an ERM-based method. The goal of a backdoor attack is to nudge the learnt hypothesis \hat{h} towards specific predictions for a certain range of feature values. This range of feature values serves as a key (or trigger) to unlock a *backdoor* in the sense of delivering anomalous predictions. The key \mathbf{x} and the corresponding anomalous prediction $\hat{h}(\mathbf{x})$ are only known to the attacker.. 5

baseline Consider some ML method that delivers a learnt hypothesis (or trained model) $\hat{h} \in \mathcal{H}$. We evaluate the quality of a trained model by computing the average loss on a test set. But how do we know that the resulting test set performance is good (enough)? How can we determine if the trained model performs close to optimal and there is little point in

investing more resources (for data collection or computation) to improve it? To this end, it is useful to have a reference (or baseline) level against which we can compare the performance of the trained model. Such a reference value might be obtained from human performance, e.g., the misclassification rate of dermatologists who diagnose cancer from visual inspection of skin. Another source for a baseline is an existing, but for some reason unsuitable, ML method. For example, the existing ML method might be computationally too expensive for the intended ML application. However, we might still use its test set error as a baseline. Another, somewhat more principled, approach to constructing a baseline is via a probabilistic model. For a wide range of probabilistic models $p(\mathbf{x}, y)$ we can precisely determine the minimum achievable risk among *any* hypothesis (not even required to belong to the hypothesis space \mathcal{H}) [27]. This minimum achievable risk (referred to as the Bayes risk) is the risk of the Bayes estimator for the label y of a data point, given its features \mathbf{x} . Note that, for a given choice of loss function, the Bayes estimator (if it exists) is completely determined by the probability distribution $p(\mathbf{x}, y)$ [27, Chapter 4]. However, there are two challenges to computing the Bayes estimator and Bayes risk: i) the probability distribution $p(\mathbf{x}, y)$ is unknown and needs to be estimated but (ii) even if we know $p(\mathbf{x}, y)$ it might be computationally too expensive to compute the Bayes risk exactly. A widely used probabilistic model is the multivariate normal distribution $(\mathbf{x}, y) \sim \mathcal{N}(\boldsymbol{\mu}, \boldsymbol{\Sigma})$ for data points characterised by numeric features and labels. Here, for the squared error loss, the Bayes estimator is given by the posterior mean $\mu_{y|\mathbf{x}}$ of

the label y given the features \mathbf{x} [27, 119]. The corresponding Bayes risk is given by the posterior variance $\sigma_{y|\mathbf{x}}^2$ (see Figure 10.2). . 15, 21, 23,




Fig. 10.2: If features and label of a data point are drawn from a multivariate normal distribution, we can achieve minimum risk (under squared error loss) by using the Bayes estimator $\mu_{y|\mathbf{x}}$ to predict the label y of a data point with features \mathbf{x} . The corresponding minimum risk is given by the posterior variance $\sigma_{y|\mathbf{x}}^2$. We can use this quantity as a baseline for the average loss of a trained model \hat{h} .

24, 126

batch In the context of SGD, a batch refers to a randomly chosen subset of the overall training set. We use the data points in this subset to estimate the gradient of training error and, in turn, to update the model parameters.. 6, 86, 87, 112, 130

Bayes estimator Consider a probabilistic model with joint probability distribution $p(\mathbf{x}, y)$ for the features \mathbf{x} and label y of a data point. For a given loss function $L(\cdot, \cdot)$, we refer to a hypothesis h as a Bayes estimator if its risk $\mathbb{E}\{L((\mathbf{x}, y), h)\}$ is minimum [27]. Note that the property

of a hypothesis being a Bayes estimator depends on the underlying probability distribution and the choice for the loss function $L(\cdot, \cdot)$. 3–5, 21

Bayes risk Consider a probabilistic model with joint probability distribution $p(\mathbf{x}, y)$ for the features \mathbf{x} and label y of a data point. The Bayes risk is the minimum possible risk that can be achieved by any hypothesis $h : \mathcal{X} \rightarrow \mathcal{Y}$. Any hypothesis that achieves the Bayes risk is referred to as a Bayes estimator [27]. 3, 4, 15

bias Consider a ML method using a parameterized hypothesis space \mathcal{H} . It learns the model parameters $\mathbf{w} \in \mathbb{R}^d$ using the dataset $\mathcal{D} = \{(\mathbf{x}^{(r)}, y^{(r)})\}_{r=1}^m$. To analyze the properties of the ML method, we typically interpret the data points as realizations of i.i.d. RVs,

$$y^{(r)} = h(\bar{\mathbf{w}})(\mathbf{x}^{(r)}) + \epsilon^{(r)}, r = 1, \dots, m.$$

We can then interpret the ML method as an estimator $\hat{\mathbf{w}}$, computed from \mathcal{D} (e.g., by solving ERM). The (squared) bias incurred by the estimate $\hat{\mathbf{w}}$ is then defined as $B^2 := \|\mathbb{E}\{\hat{\mathbf{w}}\} - \bar{\mathbf{w}}\|_2^2$. 16

classification Classification is the task of determining a discrete-valued label y of a data point based solely on its features \mathbf{x} . The label y belongs to a finite set, such as $y \in \{-1, 1\}$, or $y \in \{1, \dots, 19\}$ and represents a category to which the corresponding data point belongs to. Some classification problems involve a countably infinite label space. 6, 11, 15, 16, 23, 26, 27

classifier A classifier is a hypothesis (map) $h(\mathbf{x})$ used to predict a label taking values from a finite label space. We might use the function value $h(\mathbf{x})$ itself as a prediction \hat{y} for the label. However, it is customary to use a map $h(\cdot)$ that delivers a numeric quantity. The prediction is then obtained by a simple thresholding step. For example, in a binary classification problem with $\mathcal{Y} \in \{-1, 1\}$, we might use real-valued hypothesis map $h(\mathbf{x}) \in \mathbb{R}$ as classifier. A prediction \hat{y} can then be obtained via thresholding,

$$\hat{y} = 1 \text{ for } h(\mathbf{x}) \geq 0, \text{ and } \hat{y} = -1 \text{ otherwise.} \quad (10.1)$$

We can characterize a classifier by its decision regions \mathcal{R}_a , for every possible label value $a \in \mathcal{Y}$. 30, 50

cluster A cluster is a subset of data points that are more similar to each other than to the data points outside the cluster. The quantitative measure of similarity between data points is a design choice. If data points are characterized by Euclidean feature vectors $\mathbf{x} \in \mathbb{R}^d$, we can define the similarity between two data points via the Euclidean distance between their feature vectors.. 1, 6, 7, 11, 24, 45, 108–111, 116, 122, 124

clustered federated learning (CFL) Clustered FL (CFL) assumes that local datasets form clusters. The local datasets belonging to the same cluster have similar statistical properties. CFL pools local datasets in the same cluster to obtain a training set for training a cluster-specific model. GTVMin implements this pooling implicitly by forcing the local

model parameters to be approximately identical over well-connected subsets of the FL network.. 104, 105, 108, 123

clustering Clustering methods decompose a given set of data points into few subsets, which are referred to as clusters. Each cluster consists of data points that are more similar to each other than to data points outside the cluster. Different clustering methods use different measures for the similarity between data points and different forms of cluster representations. The clustering method k -means uses the average feature vector (*cluster mean*) of a cluster as its representative. A popular soft clustering method based on GMM represents a cluster by a multivariate normal distribution.. 1, 18, 45, 105, 108

clustering assumption The clustering assumption postulates that data points in a dataset form a (small) number of groups or clusters. Data points in the same cluster are more similar with each other than with those outside the cluster [66]. We obtain different clustering methods by using different notions of similarity between data points.. 123

computational aspects By computational aspects of a ML method, we mainly refer to the computational resources required for its implementation. For example, if a ML method uses iterative optimization techniques to solve ERM, then its computational aspects include (i) how many arithmetic operations are needed to implement a single iteration (gradient step) and (ii) how many iterations are needed to obtain useful model parameters. One important example of an iterative optimization technique is GD.. 14, 27, 30, 41, 44, 48

condition number The condition number $\kappa(\mathbf{Q}) \geq 1$ of a positive definite matrix $\mathbf{Q} \in \mathbb{R}^{d \times d}$ is the ratio $\lambda_{\max}/\lambda_{\min}$ between the largest λ_{\max} and the smallest λ_{\min} eigenvalue of \mathbf{Q} . The condition number is useful for the analysis of ML methods. The computational complexity of gradient-based methods for linear regression crucially depends on the condition number of the matrix $\mathbf{Q} = \mathbf{X}\mathbf{X}^T$, with the feature matrix \mathbf{X} of the training set. Thus, from a computational perspective, we prefer features of data points such that \mathbf{Q} has a condition number close to 1.. 125, 126

connected graph A undirected graph $\mathcal{G} = (\mathcal{V}, \mathcal{E})$ is connected if it does not contain a (non-empty) subset $\mathcal{V}' \subset \mathcal{V}$ with no edges leaving \mathcal{V}' .. 39

convex A subset $\mathcal{C} \subseteq \mathbb{R}^d$ of the Euclidean space \mathbb{R}^d is referred to as convex if it contains the line segment between any two points of that set. We define a function as convex if its epigraph is a convex set [41].. 14, 15, 18, 19, 27, 36, 38, 39, 42, 49, 58

Courant–Fischer–Weyl min-max characterization Consider a psd matrix $\mathbf{Q} \in \mathbb{R}^{d \times d}$ with EVD (or *spectral decomposition*),

$$\mathbf{Q} = \sum_{j=1}^d \lambda_j \mathbf{u}^{(j)} (\mathbf{u}^{(j)})^T.$$

Here, we used the ordered (in increasing fashion) eigenvalues

$$\lambda_1 \leq \dots \leq \lambda_n.$$

. The Courant–Fischer–Weyl min-max characterization [3, Thm. 8.1.2.] amounts to representing the eigenvalues as solutions of optimization problems.. 38, 39, 100, 101

covariance matrix The covariance matrix of a RV $\mathbf{x} \in \mathbb{R}^d$ is defined as

$$\mathbb{E}\left\{(\mathbf{x} - \mathbb{E}\{\mathbf{x}\})(\mathbf{x} - \mathbb{E}\{\mathbf{x}\})^T\right\}.. 13, 18, 19, 27, 34$$

data See dataset.. 1–3, 6, 9, 11, 15, 26, 41, 64, 151

data augmentation Data augmentation methods add synthetic data points to an existing set of data points. These synthetic data points are obtained by perturbations (e.g., adding noise to physical measurements) or transformations (e.g., rotations of images) of the original data points. These perturbations and transformations are such that the resulting synthetic data points should still have the same label. As a case in point, a rotated cat image is still a cat image even if their feature vectors (obtained by stacking pixel color intensities) are very different. Data augmentation can be an efficient form of regularization.. 25–28, 41, 42

data minimization principle European data protection regulation includes a data minimization principle. This principle requires a data controller to limit the collection of personal information to what is directly relevant and necessary to accomplish a specified purpose. The data should be retained only for as long as necessary to fulfill that purpose [93, Article 5(1)(c)] [?]. 20

data point A data point is any object that conveys information [120]. Data points might be students, radio signals, trees, forests, images, RVs, real numbers or proteins. We characterize data points using two types of properties. One type of property is referred to as a feature. Features are properties of a data point that can be measured or computed in

an automated fashion. A different kind of property is referred to as labels. The label of a data point represents some higher-level fact (or quantity of interest). In contrast to features, determining the label of a data point typically requires human experts (domain experts). Roughly speaking, ML aims to predict the label of a data point based solely on its features.. 1–21, 24–33, 36–38, 40–50, 59, 62, 63, 65, 66, 86, 104, 105, 112–115, 118, 119, 129, 136–140, 144, 147, 151, 152

data poisoning Data poisoning refers to the intentional manipulation (or fabrication) of data points to steer the training of a ML model [121,122]. The protection against data poisoning is particularly important in distributed ML applications where datasets are de-centralized.. 1, 2, 6, 12

dataset With a slight abuse of language we use the term *dataset* or *set of data points* to refer to an indexed list of data points $\mathbf{z}^{(1)}, \mathbf{z}^{(2)}, \dots$. Thus, there is a first data point $\mathbf{z}^{(1)}$, a second data point $\mathbf{z}^{(2)}$ and so on. Strictly speaking, a dataset is a list and not a set [123]. We need to keep track of the order of data points in order to cope with several data points having the same features and labels. Database theory studies formal languages for defining, structuring, and reasoning about datasets [?]. 1, 3–12, 14–19, 21, 22, 25, 29–31, 33, 37, 45, 47, 48, 53, 71, 73, 83, 87, 104, 105, 112–114, 116, 118, 119, 150, 151

decision boundary Consider a hypothesis map h that reads in a feature vector $\mathbf{x} \in \mathbb{R}^d$ and delivers a value from a finite set \mathcal{Y} . The decision boundary of h is the set of vectors $\mathbf{x} \in \mathbb{R}^d$ that lie between different

decision regions. More precisely, a vector \mathbf{x} belongs to the decision boundary if and only if each neighbourhood $\{\mathbf{x}' : \|\mathbf{x} - \mathbf{x}'\| \leq \varepsilon\}$, for any $\varepsilon > 0$, contains at least two vectors with different function values.. 27

decision region Consider a hypothesis map h that delivers values from a finite set \mathcal{Y} . For each label value (category) $a \in \mathcal{Y}$, the hypothesis h determines a subset of feature values $\mathbf{x} \in \mathcal{X}$ that result in the same output $h(\mathbf{x}) = a$. We refer to this subset as a decision region of the hypothesis h .. 4, 6, 11, 17, 30

decision tree A decision tree is a flow-chart-like representation of a hypothesis map h . More formally, a decision tree is a directed graph containing a root node that reads in the feature vector \mathbf{x} of a data point. The root node then forwards the data point to one of its children nodes based on some elementary test on the features \mathbf{x} . If the receiving children node is not a leaf node, i.e., it has itself children nodes, it represents another test. Based on the test result, the data point is further pushed to one of its descendants. This testing and forwarding of the data point is continued until the data point ends up in a leaf node (having no children nodes). . 23, 24, 48, 137

deep net A deep net is a ANN with a (relatively) large number of hidden layers. Deep learning is an umbrella term for ML methods that use a deep net as their model [83].. 16, 24, 28, 48

degree of belonging A number that indicates the extent by which a data point belongs to a cluster [6, Ch. 8]. The degree of belonging can be interpreted as a soft cluster assignment. Soft clustering methods can

encode the degree of belonging by a real number in the interval $[0, 1]$. Hard clustering is obtained as the extreme case when the degree of belonging only takes on values 0 or 1.. 45

denial-of-service attack A denial-of-service attack aims (e.g., via data poisoning) to steer the training of a model such that it performs poorly for typical data points. 5

device Any physical system that is can be used to store and process data. In the context of ML, we typically mean a computer that is able to read in data points from different sources and, in turn, to train a ML model using these data points.. 1, 3, 6, 16, 134

differentiable A function real-valued function $f : \mathbb{R}^d \rightarrow \mathbb{R}$ is differentiable if it can, at any point, be approximated locally by a linear function. The local linear approximation at the point \mathbf{x} is determined by the gradient $\nabla f(\mathbf{x})$ [2].. 10, 20, 21, 23, 43, 44, 58, 59, 85, 86

differential privacy Consider some ML method \mathcal{A} that reads in a dataset (e.g., the training set used for ERM) and delivers some output $\mathcal{A}(\mathcal{D})$. The output could be either the learnt model parameters or the predictions for specific data points. Differential privacy is a precise measure of privacy leakage incurred by revealing the output. Roughly speaking, a ML method is differentially private if the probability distribution of the output $\mathcal{A}(\mathcal{D})$ does not change too much if the sensitive attribute of one data point in the training set is changed. Note that differential privacy builds on a probabilistic model for a ML method, i.e., we interpret its output $\mathcal{A}(\mathcal{D})$ as the realization of a RV. The randomness in the output

can be ensured by intentionally adding the realization of an auxiliary RV (*noise*) to the output of the ML method.. 6–12, 37

discrepancy Consider a FL application with networked data represented by a FL network. FL methods use a discrepancy measure to compare hypothesis maps from local models at nodes i, i' connected by an edge in the FL network.. 35, 117, 122, 123, 128–131, 133

edge weight Each edge $\{i, i'\}$ of a FL network is assigned a non-negative edge weight $A_{i,i'} \geq 0$. A zero edge weight $A_{i,i'} = 0$ indicates the absence of an edge between nodes $i, i' \in \mathcal{V}$.. 13, 28, 32–35, 41, 45, 63, 77, 78, 82, 104, 118, 122, 128, 131, 133

effective dimension The effective dimension $d_{\text{eff}}(\mathcal{H})$ of an infinite hypothesis space \mathcal{H} is a measure of its size. Loosely speaking, the effective dimension is equal to the effective number of independent tunable parameters of the model. These parameters might be the coefficients used in a linear map or the weights and bias terms of an ANN.. 16, 23, 40, 41

eigenvalue We refer to a number $\lambda \in \mathbb{R}$ as eigenvalue of a square matrix $\mathbf{A} \in \mathbb{R}^{d \times d}$ if there is a non-zero vector $\mathbf{x} \in \mathbb{R}^d \setminus \{\mathbf{0}\}$ such that $\mathbf{A}\mathbf{x} = \lambda\mathbf{x}$.. 8, 10, 13, 14, 17–19, 36–39, 47, 51, 52, 63–66, 79, 80, 100, 101, 109, 110, 124–127, 132, 145, 146

eigenvalue decomposition The eigenvalue decomposition for a square matrix $\mathbf{A} \in \mathbb{R}^{d \times d}$ is a factorization of the form

$$\mathbf{A} = \mathbf{V}\mathbf{\Lambda}\mathbf{V}^{-1}.$$

The columns of the matrix $\mathbf{V} = (\mathbf{v}^{(1)}, \dots, \mathbf{v}^{(d)})$ are the eigenvectors of the matrix \mathbf{V} . The diagonal matrix $\mathbf{\Lambda} = \text{diag}\{\lambda_1, \dots, \lambda_d\}$ contains the eigenvalues λ_j corresponding to the eigenvectors $\mathbf{v}^{(j)}$. Note that the above decomposition exists only if the matrix \mathbf{A} is diagonalizable.. 8, 14, 19, 37

eigenvector An eigenvector of a matrix $\mathbf{A} \in \mathbb{R}^{d \times d}$ is a non-zero vector $\mathbf{x} \in \mathbb{R}^d \setminus \{\mathbf{0}\}$ such that $\mathbf{Ax} = \lambda\mathbf{x}$ with some eigenvalue λ .. 14, 36–39

empirical risk The empirical risk $\widehat{L}(h|\mathcal{D})$ of a hypothesis on a dataset \mathcal{D} is the average loss incurred by h when applied to the data points in \mathcal{D} .. 13, 14, 16, 35, 46, 47

empirical risk minimization Empirical risk minimization is the optimization problem of finding a hypothesis (out of a model) with minimum average loss (or empirical risk) on a given dataset \mathcal{D} (the training set). Many ML methods are obtained from empirical risk via specific design choices for the dataset, model and loss [6, Ch. 3].. 1, 2, 4–7, 11–14, 16–18, 21, 22, 24–26, 28–30, 34, 35, 40, 41, 45–48, 59, 92

estimation error Consider data points with feature vectors \mathbf{x} and label y . In some applications we can model the relation between the features and the label of a data point as $y = \bar{h}(\mathbf{x}) + \varepsilon$. Here, we used some true underlying hypothesis \bar{h} and a noise term ε which summarized any modelling or labelling errors. The estimation error incurred by a ML method that learns a hypothesis \widehat{h} , e.g., using ERM, is defined as $\widehat{h}(\mathbf{x}) - \bar{h}(\mathbf{x})$, for some feature vector. For a parameterized hypothesis space, consisting of hypothesis maps that are determined by a model parameters \mathbf{w} , we

can define the estimation error as $\Delta \mathbf{w} = \hat{\mathbf{w}} - \overline{\mathbf{w}}$ [80, 124].. 18, 19, 144, 146, 147

Euclidean space The Euclidean space \mathbb{R}^d of dimension $d \in \mathbb{N}$ consists of vectors $\mathbf{x} = (x_1, \dots, x_d)$, with d real-valued entries $x_1, \dots, x_d \in \mathbb{R}$. Such an Euclidean space is equipped with a geometric structure defined by the inner product $\mathbf{x}^T \mathbf{x}' = \sum_{j=1}^d x_j x'_j$ between any two vectors $\mathbf{x}, \mathbf{x}' \in \mathbb{R}^d$ [2].. 4, 8, 17, 30, 35, 40, 48, 73

expert ML aims to learn a hypothesis h that accurately predicts the label of a data point based on its features. We measure the prediction error using some loss function. Ideally we want to find a hypothesis that incurs minimum loss on any data point. We can make this informal goal precise via the i.i.d. assumption and using the Bayes risk as the baseline for the (average) loss of a hypothesis. An alternative approach to obtain a baseline is to use the hypothesis h' learnt by an existing ML method. We refer to this hypothesis h' as an expert [125]. Regret minimization methods learn a hypothesis that incurs a loss comparable to the best expert [125, 126].. 23

explainability We define the (subjective) explainability of a ML method as the level of simulatability [99] of the predictions delivered by a ML system to a human user. Quantitative measures for the (subjective) explainability of a trained model can be constructed by comparing its predictions with the predictions provided by a user on a test-set [99, 101]. Alternatively, we can use probabilistic models for data and measure explainability of a trained ML model via the conditional (differential)

entropy of its predictions, given the user predictions [100, 127]. . 20, 135, 137, 140, 151, 152

feature A feature of a data point is one of its properties that can be measured or computed easily without the need for human supervision. For example, if a data point is a digital image (e.g., stored in as a jpeg file), then we could use the red-green-blue intensities of its pixels as features. Domain-specific synonyms for the term feature are *covariate*, *explanatory variable*, *independent variable*, *input (variable)*, *predictor (variable)* or *regressor* [128–130].. 1–21, 23–25, 27–33, 36, 37, 40, 41, 43–45, 49, 50, 59, 63–66, 104, 105, 112–115, 119, 129, 137–141, 144, 146, 147, 151

feature map A map that transforms the original features of a data point into new features. The so-obtained new features might be preferable over the original features for several reasons. For example, the arrangement of data points might become simpler (of *more linear*) in the new feature space, allowing to use linear models in the new features. This idea is a main driver for the development of kernel methods [131]. Moreover, the hidden layers of a deep net can be interpreted as a trainable feature map followed by a linear model in the form of the output layer. Another reason for learning a feature map could be that learning a small number of new features helps to avoid overfitting and ensure interpretability [132]. The special case of a feature map delivering two numeric features is particularly useful for data visualization. Indeed, we can depict data points in a scatterplot by using two features as the coordinates of a

data point.. 13, 14, 16, 17

feature matrix Consider a dataset \mathcal{D} with m data points with feature vectors $\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)} \in \mathbb{R}^d$. It is convenient to collect the individual feature vectors into a feature matrix $\mathbf{X} := (\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)})^T$ of size $m \times d$.. 8, 18–20, 33, 45, 110

feature space The feature space of a given ML application or method is constituted by all potential values that the feature vector of a data point can take on. A widely used choice for the feature space is the Euclidean space \mathbb{R}^d with dimension d being the number of individual features of a data point.. 14, 16, 17, 24–26

feature vector A vector $\mathbf{x} = (x_1, \dots, x_d)^T$ whose entries are individual features x_1, \dots, x_d . Many ML methods use feature vectors that belong to some finite-dimensional Euclidean space \mathbb{R}^d . However, for some ML methods it is more convenient to work with feature vectors that belong to an infinite-dimensional vector space (see, e.g., kernel method).. 8, 9, 16, 26, 27, 29, 41

federated averaging (FedAvg) Federated averaging is an iterative FL algorithm that alternates between local model trainings and averaging the resulting local model parameters. Different variants of this algorithm are obtained by different techniques for the model training. The authors of [14] consider federated averaging methods where the local model training is implemented by running several GD steps.. 78, 148

federated learning (FL) Federated learning is an umbrella term for ML

methods that train models in a collaborative fashion using decentralized data and computation.. 1–8, 10–13, 17, 18, 24, 30–32, 34–36, 40, 41, 43, 47–49, 54, 55, 57, 58, 64, 68, 74, 77, 84, 89, 92, 96–98, 104–108, 112–115, 117, 122, 127, 128, 133–142, 144, 147, 148, 151–153

federated learning (FL) network A federated network is an undirected weighted graph whose nodes represent data generators that aim to train a local (or personalized) model. Each node in a federated network represents some device, capable to collect a local dataset and, in turn, train a local model. FL methods learn a local hypothesis $h^{(i)}$, for each node $i \in \mathcal{V}$, such that it incurs small loss on the local datasets.. 1–3, 5–8, 13, 18, 19, 30–38, 40–47, 49, 51–53, 55, 68, 76–80, 82–89, 94–100, 104–106, 108–112, 114, 118, 119, 122–129, 131–134, 136, 139, 141–143, 146–148, 150, 151, 155

Finnish Meteorological Institute The Finnish Meteorological Institute is a government agency responsible for gathering and reporting weather data in Finland.. 1, 7–10, 12, 16–19, 41, 44, 155

Gaussian mixture model A Gaussian mixture model (GMM) is particular type of probabilistic models for a numeric vector \mathbf{x} (e.g., the features of a data point). Within a GMM, the vector \mathbf{x} is drawn from a randomly selected multivariate normal distribution $p^{(c)} = \mathcal{N}(\boldsymbol{\mu}^{(c)}, \mathbf{C}^{(c)})$ with $c = I$. The index $I \in \{1, \dots, k\}$ is a RV with probabilities $p(I = c) = p_c$. Note that a GMM is parameterized by the probability p_c , the mean vector $\boldsymbol{\mu}^{(c)}$ and covariance matrix $\boldsymbol{\Sigma}^{(c)}$ for each $c = 1, \dots, k$. GMMs are widely used for clustering, density estimation and as a generative

model.. 7, 45, 108

Gaussian random variable A standard Gaussian RV is a real-valued random variable x with probability density function (pdf) [5, 119, 133]

$$p(x) = \frac{1}{\sqrt{2\pi}} \exp^{-x^2/2}.$$

Given a standard Gaussian RV x , we can construct a general Gaussian RV x' with mean μ and variance σ^2 via $x' := \sigma(x + \mu)$. The probability distribution of a Gaussian RV is referred to as normal distribution, denoted $\mathcal{N}(\mu, \sigma)$.

A Gaussian random vector $\mathbf{x} \in \mathbb{R}^d$ with covariance matrix \mathbf{C} and mean $\boldsymbol{\mu}$ can be constructed via $\mathbf{x} := \mathbf{A}(\mathbf{z} + \boldsymbol{\mu})$. Here, \mathbf{A} is any matrix that satisfies $\mathbf{A}\mathbf{A}^T = \mathbf{C}$ and $\mathbf{z} := (z_1, \dots, z_d)^T$ is a vector whose entries are i.i.d. standard Gaussian RVs z_1, \dots, z_d . Gaussian random processes generalize Gaussian random vectors by applying linear transformations to infinite sequences of standard Gaussian RVs [?].

Gaussian RVs are widely used probabilistic models for the statistical analysis of machine learning methods. Their significance arises partly from the central limit theorem which states that the sum of many independent RVs (not necessarily Gaussian themselves) tends to a Gaussian RV [?]. 41

General Data Protection Regulation The General Data Protection Regulation (GDPR) was enacted by the European Union (EU), effective from May 25, 2018 [93]. It safeguards the privacy and data rights of individuals in the EU. The GDPR has significant implications for how

data is collected, stored, and used in ML applications. Key provisions include:

- Data minimization principle: ML systems should only use necessary amount of personal data for their purpose.
- Transparency and Explainability: ML systems should enable their users to understand how they make decisions that impact them.
- Data Subject Rights: Including the rights to access, rectify, and delete personal data, as well as to object to automated decision-making and profiling.
- Accountability: Organizations must ensure robust data security and demonstrate compliance through documentation and regular audits.

. 138

generalized total variation Generalized total variation measures the changes of vector-valued node attributes over a weighted undirected graph.. 24, 30, 41, 47, 49, 50, 83, 85, 87, 88, 95, 96, 115, 116, 150

gradient For a real-valued function $f : \mathbb{R}^d \rightarrow \mathbb{R} : \mathbf{w} \mapsto f(\mathbf{w})$, a vector \mathbf{g} such that $\lim_{\mathbf{w} \rightarrow \mathbf{w}'} \frac{f(\mathbf{w}) - (f(\mathbf{w}') + \mathbf{g}^T(\mathbf{w} - \mathbf{w}'))}{\|\mathbf{w} - \mathbf{w}'\|} = 0$ is referred to as the gradient of f at \mathbf{w}' . If such a vector exists it is denoted $\nabla f(\mathbf{w}')$ or $\nabla f(\mathbf{w})|_{\mathbf{w}'}$ [2].. 4, 8, 10, 12, 21, 23, 44–46, 49, 57–61, 66, 81, 88, 130

gradient descent (GD) Gradient descent is an iterative method for finding the minimum of a differentiable function $f(\mathbf{w})$ of a vector-valued argument $\mathbf{w} \in \mathbb{R}^d$. Consider a current guess or approximation $\mathbf{w}^{(k)}$

for minimum. We would like to find a new (better) vector $\mathbf{w}^{(k+1)}$ that has smaller objective value $f(\mathbf{w}^{(k+1)}) < f(\mathbf{w}^{(k)})$ than the current guess $\mathbf{w}^{(k)}$. We can achieve this typically by using a gradient step

$$\mathbf{w}^{(k+1)} = \mathbf{w}^{(k)} - \eta \nabla f(\mathbf{w}^{(k)}) \quad (10.2)$$

with a sufficiently small step size $\eta > 0$. Figure 10.3 illustrates the effect of a single GD step (10.2). . 4, 6, 7, 17, 21, 23, 29, 38, 45, 46, 57, 62,




Fig. 10.3: A single gradient step (10.2) towards the minimizer $\bar{\mathbf{w}}$ of $f(\mathbf{w})$.

68, 69, 75–78, 84, 90, 132, 147, 148

gradient step Given a differentiable real-valued function $f(\cdot) : \mathbb{R}^d \rightarrow \mathbb{R}$ and a vector $\mathbf{w} \in \mathbb{R}^d$, the gradient step updates \mathbf{w} by adding the scaled negative gradient $\nabla f(\mathbf{w})$ to obtain the new vector

$$\hat{\mathbf{w}} := \mathbf{w} - \eta \nabla f(\mathbf{w}). \quad (10.3)$$

Mathematically, the gradient step is a (typically non-linear) operator $\mathcal{T}^{(f,\eta)}$ that is parametrized by the function f and the step size η . Note




Fig. 10.4: The basic gradient step (10.3) maps a given vector \mathbf{w} to the updated vector \mathbf{w}' . It defines an operator $\mathcal{T}^{(f,\eta)}(\cdot) : \mathbb{R}^d \rightarrow \mathbb{R}^d : \mathbf{w} \mapsto \hat{\mathbf{w}}$.

that the gradient step (10.3) optimizes locally (confined to a neighbourhood defined by the step size η) a linear approximation to the function $f(\cdot)$. A natural generalization of (10.3) is to locally optimize the function itself (instead of its linear approximation),

$$\hat{\mathbf{w}} = \operatorname{argmin}_{\mathbf{w}' \in \mathbb{R}^d} f(\mathbf{w}') + (1/\eta) \|\mathbf{w} - \mathbf{w}'\|_2^2. \quad (10.4)$$

We intentionally use the same symbol η for the parameter in (10.4) as we used for the step-size in (10.3). The larger we choose η in (10.4), the more progress the update will make towards reducing the function value $f(\hat{\mathbf{w}})$. Note that, much like the gradient step (10.3), also the update (10.4) defines a (typically non-linear) operator that is parameterized by the function $f(\cdot)$ and the parameter η . For convex $f(\cdot)$, this operator is known as the proximal operator of $f(\cdot)$ [35]. . 4, 7, 16, 38, 44, 45, 57, 58, 60–70, 72, 73, 77, 79, 81, 82, 84, 86, 90–92, 96, 123, 125




Fig. 10.5: A generalized gradient step updates a vector \mathbf{w} by minimizing a penalized version of the function $f(\cdot)$. The penalty term is the squared Euclidean distance from the vector \mathbf{w} .

gradient-based method Gradient-based methods are iterative techniques for finding the minimum (or maximum) of a differentiable objective function of the model parameters. These methods construct a sequence of approximations to an optimal choice for model parameters that results in a minimum (or maximum) value of the objective function. As their name indicates, gradient-based methods use the gradients of the objective function evaluated during previous iterations to construct new (hopefully) improved model parameters. One important example for a gradient-based method is GD.. 4, 7, 8, 15, 16, 36, 43, 44, 57–60, 64, 70, 72, 73, 77, 79–81, 99, 122, 133, 135, 138, 141, 145, 147

graph A graph $\mathcal{G} = (\mathcal{V}, \mathcal{E})$ is a pair that consists of a node set \mathcal{V} and an edge set \mathcal{E} . In its most general form, a graph is specified by a map that assigns to each edge $e \in \mathcal{E}$ a pair of nodes [42]. One important family of graphs are simple undirected graphs. A simple undirected graph is

obtained by identifying each edge $e \in \mathcal{E}$ with two different nodes $\{i, i'\}$.
 Weighted graphs also specify numeric weights A_e for each edge $e \in \mathcal{E}$.
 18, 20, 27, 35, 36, 41, 50, 123, 132

GTV minimization GTV minimization is an instance of RERM using the
 GTV of local model parameters as a regularizer [30]. 1–7, 9, 10, 30,
 31, 41–49, 51, 53, 56–58, 63, 66, 73, 74, 76–81, 89, 94, 95, 99, 104–106,
 108–113, 115, 117–119, 122–125, 128, 132, 133, 135–147, 151

hard clustering Hard clustering refers to the task of partitioning a given
 set of data points into (few) non-overlapping clusters. The most widely
 used hard clustering method is k -means.. 12

horizontal FL Horizontal FL uses local datasets that are constituted by dif-
 ferent data points but using the same features to characterize them [65].
 For example, weather forecasting uses a network of spatially distributed
 weather (observation) stations. Each weather station measures the
 same quantities such as daily temperature, air pressure and precipita-
 tion. However, different weather stations measure the characteristics or
 features of different spatio-temporal regions (each such region being a
 separate data point).. 89, 104, 105, 112–114

hypothesis A map (or function) $h : \mathcal{X} \rightarrow \mathcal{Y}$ from the feature space \mathcal{X} to the
 label space \mathcal{Y} . Given a data point with features \mathbf{x} , we use a hypothesis
 map h to estimate (or approximate) the label y using the prediction
 $\hat{y} = h(\mathbf{x})$. ML is all about learning (or finding) a hypothesis map h
 such that $y \approx h(\mathbf{x})$ for any data point (having features \mathbf{x} and label y)..

1–6, 10–19, 21, 22, 24–26, 28, 29, 31–36, 40, 42–45, 47, 48, 60, 70, 71, 73, 137, 140, 151–153

hypothesis space Every practical ML method uses a hypothesis space (or model) \mathcal{H} . The hypothesis space of a ML method is a subset of all possible maps from the feature space to label space. The design choice of the hypothesis space should take into account available computational resources and statistical aspects. If the computational infrastructure allows for efficient matrix operations, and there is a (approximately) linear relation between features and label, a useful choice for the hypothesis space might be the linear model.. 3, 5, 11, 13, 14, 16, 23–26, 30, 31, 33, 36, 42, 47, 49

i.i.d. It can be useful to interpret data points $\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}$ as realizations of independent and identically distributed RVs with a common probability distribution. If these RVs are continuous-valued, their joint pdf is $p(\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}) = \prod_{r=1}^m p(\mathbf{z}^{(r)})$ with $p(\mathbf{z})$ being the common marginal pdf of the underlying RVs.. 3, 5, 10, 16, 17, 19, 20, 24–29, 32, 38, 43, 45, 48, 54, 71, 75, 108, 111, 124, 128, 129, 138

i.i.d. assumption The i.i.d. assumption interprets data points of a dataset as the realizations of i.i.d. RVs.. 1, 12, 15, 17, 25, 43, 45, 104, 137

interpretability A ML method is interpretable for a specific user if they can well anticipate the predictions delivered by the method. The notion of interpretability can be made precise using quantitative measures of the uncertainty about the predictions [100].. 31

kernel Consider data points characterized by a feature vector $\mathbf{x} \in \mathcal{X}$ with a generic feature space \mathcal{X} . A (real-valued) kernel $K : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ assigns each pair of feature vectors $\mathbf{x}, \mathbf{x}' \in \mathcal{X}$ a real number $K(\mathbf{x}, \mathbf{x}')$. The value $K(\mathbf{x}, \mathbf{x}')$ is often interpreted as a measure for the similarity between \mathbf{x} and \mathbf{x}' . Kernel methods use a kernel to transform the feature vector \mathbf{x} to a new feature vector $\mathbf{z} = K(\mathbf{x}, \cdot)$. This new feature vector belongs to a linear feature space \mathcal{X}' which is (in general) different from the original feature space \mathcal{X} . The feature space \mathcal{X}' has a specific mathematical structure, i.e., it is a reproducing kernel Hilbert space [131, 134].. 16, 26, 36

kernel method A kernel method is a ML method that uses a kernel K to map the original (raw) feature vector \mathbf{x} of a data point to a new (transformed) feature vector $\mathbf{z} = K(\mathbf{x}, \cdot)$ [131, 134]. The motivation for transforming the feature vectors is that, using a suitable kernel, the data points have a *more pleasant* geometry in the transformed feature space. For example, in a binary classification problem, using transformed feature vectors \mathbf{z} might allow to use linear models, even if the data points are not linearly separable in the original feature space (see Figure 10.6). . 17, 26

label A higher-level fact or quantity of interest associated with a data point. For example, if the data point is an image, the label could indicate whether the image contains a cat or not. Synonyms for label, commonly used in specific domains, include *response variable*, *output variable*, and *target* [128–130].. 2–6, 9–20, 24, 25, 27–33, 36, 37, 40, 43–45, 47, 50, 59,




Fig. 10.6: Five data points characterized by feature vectors $\mathbf{x}^{(r)}$ and labels $y^{(r)} \in \{\circ, \square\}$, for $r = 1, \dots, 5$. With these feature vectors, there is no way to separate the two classes by a straight line (representing the decision boundary of a linear classifier). In contrast, the transformed feature vectors $\mathbf{z}^{(r)} = K(\mathbf{x}^{(r)}, \cdot)$ allow to separate the data points using a linear classifier.

65, 81, 112, 114, 118, 129, 136, 138, 139, 141, 143, 144

label space Consider a ML application that involves data points characterized by features and labels. The label space is constituted by all potential values that the label of a data point can take on. Regression methods, aiming at predicting numeric labels, often use the label space $\mathcal{Y} = \mathbb{R}$. Binary classification methods use a label space that consists of two different elements, e.g., $\mathcal{Y} = \{-1, 1\}$, $\mathcal{Y} = \{0, 1\}$ or $\mathcal{Y} = \{\text{“cat image”}, \text{“no cat image”}\}$. 1, 5, 6, 11, 15, 16, 24, 25, 29, 30, 118, 129

Laplacian matrix The structure of a graph \mathcal{G} , with nodes $i = 1, \dots, n$, can be analyzed using the properties of special matrices that are associated with \mathcal{G} . One such matrix is the graph Laplacian matrix $\mathbf{L}^{(\mathcal{G})} \in \mathbb{R}^{n \times n}$ which is defined for an undirected and weighted graph [135, 136]. It is




Fig. 10.7: Left: Some undirected graph \mathcal{G} with three nodes $i = 1, 2, 3$. Right: Laplacian matrix $\mathbf{L}^{(\mathcal{G})} \in \mathbb{R}^{3 \times 3}$ of \mathcal{G} .

defined element-wise as (see Fig. 10.7)

$$L_{i,i'}^{(\mathcal{G})} := \begin{cases} -A_{i,i'} & \text{for } i \neq i', \{i, i'\} \in \mathcal{E} \\ \sum_{i'' \neq i} A_{i,i''} & \text{for } i = i' \\ 0 & \text{else.} \end{cases} \quad (10.5)$$

Here, $A_{i,i'}$ denotes the edge weight of an edge $\{i, i'\} \in \mathcal{E}$. . 18, 31, 36–39, 42, 51, 52, 76, 79, 100, 109, 124–126

Large Language Model Large Language Models (LLMs) is an umbrella term for ML methods that process and generate human-like text. These methods typically use deep nets with billions (or even trillions) of parameters. A widely used choice for the network architecture is referred to as Transformers [?]. The training of LLMs is often based on the task of predicting a few words that are intentionally removed from a large text corpus. Thus, we can construct labelled data points simply by selecting some words of a text as labels and the remaining words as features of data points. This construction requires very little human supervision and allows for generating sufficiently large training sets for LLMs.. 3

law of large numbers The law of large numbers refers to the convergence of the average of an increasing (large) number of i.i.d. RVs to the mean of their common probability distribution. Different instances of the law of large numbers are obtained using different notions of convergence [133].. 24, 71

learning rate Consider an iterative method for finding or learning a useful hypothesis $h \in \mathcal{H}$. Such an iterative method repeats similar computational (update) steps that adjust or modify the current hypothesis to obtain an improved hypothesis. A prime example of such an iterative learning method is GD and its variants such as SGD or projected GD. We refer by learning rate to a parameter of an iterative learning method that controls the extent by which the current hypothesis can be modified during a single iteration. A prime example of such a parameter is the step size used in GD [6, Ch. 5].. 15, 45, 57, 58, 60–67, 69, 73, 75, 80, 82, 83, 85, 87, 88, 92, 99, 150

learning task Consider a dataset \mathcal{D} constituted by several data points, each of them characterized by features \mathbf{x} . For example, the dataset \mathcal{D} might be constituted by the images of a particular database. Sometimes it might be useful to represent a dataset \mathcal{D} , along with the choice of features, by a probability distribution $p(\mathbf{x})$. A learning task associated with \mathcal{D} consists of a specific choice for the label of a data point and the corresponding label space. Given a choice for the loss function and model, a learning task gives rise to an instance of ERM. Thus, we could define a learning task also via an instance of ERM, i.e., via an objective

function. Note that, for the same dataset, we obtain different learning tasks by using different choices for the features and label of a data point. These learning tasks are related, as they are based on the same dataset, and solving them jointly (via multitask learning methods) is typically preferable over solving them separately [137–139].. 3, 17

least absolute shrinkage and selection operator (Lasso) The least absolute shrinkage and selection operator (Lasso) is an instance of structural risk minimization (SRM) to learn the weights \mathbf{w} of a linear map $h(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$ based on a training set. The Lasso is obtained from linear regression by adding the scaled ℓ_1 -norm $\alpha \|\mathbf{w}\|_1$ to the average squared error loss incurred on the training set.. 42

linear classifier Consider data points characterized by numeric features $\mathbf{x} \in \mathbb{R}^d$ and a label $y \in \mathcal{Y}$ from some finite label space \mathcal{Y} . A linear classifier characterized by having decision regions separated by hyperplanes in the Euclidean space \mathbb{R}^d [6, Ch. 2].. 27

linear model Consider data points, each characterized by a numeric feature vector $\mathbf{x} \in \mathbb{R}^d$. A linear model is a hypothesis space which consists of all linear maps,

$$\mathcal{H}^{(d)} := \{h(\mathbf{x}) = \mathbf{w}^T \mathbf{x} : \mathbf{w} \in \mathbb{R}^d\}. \quad (10.6)$$

Note that (10.6) defines an entire family of hypothesis spaces, which is parameterized by the number d of features that are linearly combined to form the prediction $h(\mathbf{x})$. The design choice of d is guided by computational aspects (smaller d means less computation), statistical

aspects (increasing d might reduce prediction error) and interpretability. A linear model using few carefully chosen features tends to be considered more interpretable [132, 140]. 8, 13, 16, 17, 19, 20, 23–26, 28, 42, 44, 45, 49, 58, 66, 71, 75, 77, 78, 81, 83, 85–89, 93, 110, 118, 119, 123, 137, 142, 143, 150, 151, 153

linear regression Linear regression aims to learn a linear hypothesis map to predict a numeric label based on numeric features of a data point. The quality of a linear hypothesis map is measured using the average squared error loss incurred on a set of labeled data points, which we refer to as the training set.. 8, 9, 13–15, 18, 19, 26–30, 41, 42, 58, 63, 65, 66, 75, 92, 95, 118, 119

local dataset The concept of a local dataset is in-between the concept of a data point and a dataset. A local dataset consists of several individual data points which are characterized by features and labels. In contrast to a single dataset used in basic ML methods, a local dataset is also related to other local datasets via different notions of similarities. These similarities might arise from probabilistic models or communication infrastructure and are encoded in the edges of a FL network.. 1–9, 16, 18, 24, 30–36, 41, 45–49, 52–55, 63, 76–78, 80, 82, 86, 89–91, 99, 104–108, 110–115, 122, 123, 128–131, 134, 135, 137–140, 142–144, 146, 147, 152

local model Consider a collections of local datasets that are assigned to the nodes of a FL network. A local model $\mathcal{H}^{(i)}$ is a hypothesis space assigned to a node $i \in \mathcal{V}$. Different nodes might be assigned different hypothesis

spaces, i.e., in general $\mathcal{H}^{(i)} \neq \mathcal{H}^{(i')}$ for different nodes $i, i' \in \mathcal{V}$. 1–3, 6, 18, 19, 30–35, 39–42, 44, 45, 48, 53, 54, 56, 68, 73, 77–79, 81–83, 85, 86, 88, 89, 91, 95–97, 99, 104–106, 109, 110, 115–117, 120, 130, 137, 140, 150

loss ML methods use a loss function $L(\mathbf{z}, h)$ to measure the error incurred by applying a specific hypothesis to a specific data point. With slight abuse of notation, we use the term *loss* for both, the loss function L itself and for its value $L(\mathbf{z}, h)$ for a specific data point \mathbf{z} and hypothesis h . 1, 2, 4, 6, 8, 11, 13–19, 21, 24–27, 33, 35, 41–43, 45, 47, 48, 50, 59, 60, 73, 77, 78, 81, 82, 90, 104, 130, 131, 142

loss function A loss function is a map

$$L : \mathcal{X} \times \mathcal{Y} \times \mathcal{H} \rightarrow \mathbb{R}_+ : ((\mathbf{x}, y), h) \mapsto L((\mathbf{x}, y), h)$$

which assigns a pair of a data point, with features \mathbf{x} and label y , and a hypothesis $h \in \mathcal{H}$ the non-negative real number $L((\mathbf{x}, y), h)$. The loss value $L((\mathbf{x}, y), h)$ quantifies the discrepancy between the true label y and the prediction $h(\mathbf{x})$. Lower (closer to zero) values $L((\mathbf{x}, y), h)$ indicate a smaller discrepancy between prediction $h(\mathbf{x})$ and label y . Figure 10.8 depicts a loss function for a given data point, with features \mathbf{x} and label y , as a function of the hypothesis $h \in \mathcal{H}$. 1, 3–7, 11–13, 15, 19, 21, 22, 29–34, 40, 41, 43, 48–50, 82, 85, 86, 92, 95–97, 99, 112, 114, 122, 130, 136, 139, 151, 152

maximum likelihood Consider data points $\mathcal{D} = \{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$ that are interpreted as realizations of i.i.d. RVs with a common probability




Fig. 10.8: Some loss function $L((\mathbf{x}, y), h)$ for a fixed data point, with feature vector \mathbf{x} and label y , and varying hypothesis h . ML methods try to find (learn) a hypothesis that incurs minimum loss.

distribution $p(\mathbf{z}; \mathbf{w})$ which depends on the model parameters $\mathbf{w} \in \mathcal{W} \subseteq \mathbb{R}^n$. Maximum likelihood methods learn model parameters \mathbf{w} by maximizing the probability (density) $p(\mathcal{D}; \mathbf{w}) = \prod_{r=1}^m p(\mathbf{z}^{(r)}; \mathbf{w})$ of observing the dataset is maximized. Thus, the maximum likelihood estimator is a solution to the optimization problem $\max_{\mathbf{w} \in \mathcal{W}} p(\mathcal{D}; \mathbf{w})$.

37, 118, 128, 130

mean The expectation $\mathbb{E}\{\mathbf{x}\}$ of a numeric RV \mathbf{x} .. 29

model In the context of ML methods, the term *model* typically refers to the hypothesis space used by a ML method [6, 141].. 1–7, 9, 11–18, 21–26, 28–31, 33, 35, 40, 42, 45, 47, 68, 70, 72, 90, 106, 136, 139, 152

model parameters Model parameters are quantities that are used to select a specific hypothesis map from a model. We can think of model parameters as a unique identifier for a hypothesis map, similar to how a social security number identifies a person in Finland.. 1–9, 12, 14, 17, 19, 20,

23, 24, 26, 28–30, 32, 33, 35, 37, 38, 40–42, 45–49, 51, 53, 55, 57, 58, 60, 64, 73–76, 82–85, 87–94, 97–99, 104–108, 110–112, 116, 118, 119, 122, 123, 125, 128, 130, 132, 136, 137, 139, 141–145, 148–150, 152, 155

multivariate normal distribution The multivariate normal distribution $\mathcal{N}(\mathbf{m}, \mathbf{C})$ is an important family of probability distributions for a continuous RV $\mathbf{x} \in \mathbb{R}^d$ [5, 119, 142]. This family is parameterized by the mean \mathbf{m} and covariance matrix \mathbf{C} of \mathbf{x} . If the covariance matrix is invertible, the probability distribution of \mathbf{x} is

$$p(\mathbf{x}) \propto \exp \left(- (1/2)(\mathbf{x} - \mathbf{m})^T \mathbf{C}^{-1}(\mathbf{x} - \mathbf{m}) \right).$$

. 3, 4, 7, 13, 16, 18, 134

mutual information The mutual information $I(\mathbf{x}; y)$ between two RVs \mathbf{x} , y defined on the same probability space is given by [120]

$$I(\mathbf{x}; y) := \mathbb{E} \left\{ \log \frac{p(\mathbf{x}, y)}{p(\mathbf{x})p(y)} \right\}.$$

It is a measure for how well we can estimate y based solely from \mathbf{x} . A large value of $I(\mathbf{x}; y)$ indicates that y can be well predicted solely from \mathbf{x} . This prediction could be obtained by a hypothesis learnt by a ERM-based ML method.. 12, 13, 37

neighbourhood The neighbourhood of a node $i \in \mathcal{V}$ is the subset of nodes constituted by the neighbours of i .. 18, 36, 134

neighbours The neighbours of a node $i \in \mathcal{V}$ within a FL network are those nodes $i' \in \mathcal{V} \setminus \{i\}$ that are connected (via an edge) to node i .. 34, 35, 55, 82–85, 87, 88, 95, 134, 150

networked data Networked data consists of local datasets that are related by some notion of pair-wise similarity. We can represent networked data using a graph whose nodes carry local datasets and edges encode pairwise similarities. One example for networked data arises in FL applications where local datasets are generated by spatially distributed devices.. 13, 35

node degree The degree $d^{(i)}$ of a node $i \in \mathcal{V}$ in an undirected graph is the number of its neighbours, $d^{(i)} := |\mathcal{N}^{(i)}|$.. 36, 80, 109, 123, 124, 126, 127, 132, 133

objective function An objective function is a map that assigns each value of an optimization variable, such as the model parameters \mathbf{w} of a hypothesis $h^{(\mathbf{w})}$, to an objective value $f(\mathbf{w})$. The objective value $f(\mathbf{w})$ could be the risk or the empirical risk of a hypothesis $h^{(\mathbf{w})}$.. 18, 23, 26, 29, 41, 42, 45, 49, 56–63, 67, 68, 70, 71, 73, 78, 80, 94, 120, 131, 132

overfitting Consider a ML method that uses ERM to learn a hypothesis with minimum empirical risk on a given training set. Such a method is *overfitting* the training set if it learns hypothesis with small empirical risk on the training set but significantly larger loss outside the training set.. 16, 23, 40

parameters The parameters of a ML model are tunable (learnable or adjustable) quantities that allow to choose between different hypothesis maps. For example, the linear model $\mathcal{H} := \{h^{(\mathbf{w})} : h^{(\mathbf{w})}(x) = w_1x + w_2\}$ consists of all hypothesis maps $h^{(\mathbf{w})}(x) = w_1x + w_2$ with a particular

choice for the parameters $\mathbf{w} = (w_1, w_2)^T \in \mathbb{R}^2$. Another example of parameters is the weights assigned to the connections between neurons of an ANN.. 13, 16, 18, 19, 26, 28, 39–42, 44–46, 56, 68, 71, 77–79, 81–83, 85, 87–89, 91, 106, 109, 110, 115–117, 120, 128, 129, 150

polynomial regression Polynomial regression aims at learning a polynomial hypothesis map to predict a numeric label based on numeric features of a data point. For data points characterized by a single numeric feature, polynomial regression uses the hypothesis space $\mathcal{H}_d^{(\text{poly})} := \{h(x) = \sum_{j=0}^{d-1} x^j w_j\}$. The quality of a polynomial hypothesis map is measured using the average squared error loss incurred on a set of labeled data points (which we refer to as training set).. 24

positive semi-definite A (real-valued) symmetric matrix $\mathbf{Q} = \mathbf{Q}^T \in \mathbb{R}^{d \times d}$ is referred to as positive semi-definite if $\mathbf{x}^T \mathbf{Q} \mathbf{x} \geq 0$ for every vector $\mathbf{x} \in \mathbb{R}^d$. The property of being psd can be extended from matrices to (real-valued) symmetric kernel maps $K : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ (with $K(\mathbf{x}, \mathbf{x}') = K(\mathbf{x}', \mathbf{x})$) as follows: For any finite set of feature vectors $\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}$, the resulting matrix $\mathbf{Q} \in \mathbb{R}^{m \times m}$ with entries $Q_{r,r'} = K(\mathbf{x}^{(r)}, \mathbf{x}^{(r')})$ is psd [131].. 4, 8–10, 12, 14, 17, 37, 43, 58, 62, 63

prediction A prediction is an estimate or approximation for some quantity of interest. ML revolves around learning or finding a hypothesis map h that reads in the features \mathbf{x} of a data point and delivers a prediction $\hat{y} := h(\mathbf{x})$ for its label y .. 2, 4–6, 10, 12, 13, 15–17, 24, 25, 30–32, 34–37, 43, 48–50, 71, 81, 82, 137, 139, 151, 152

privacy leakage Consider a (ML or FL) system that processes a local dataset

$\mathcal{D}^{(i)}$ and shares data, such as the predictions obtained for new data points, with other parties. Privacy leakage arises if the shared data carries information about a private (sensitive) feature of a data point (which might be a human) of $\mathcal{D}^{(i)}$. The amount of privacy leakage can be measured via MI using a probabilistic model for the local dataset. Another quantitative measure for privacy leakage is DP.. 1, 2, 16

privacy protection Consider some ML method \mathcal{A} that reads in a \mathcal{D} and delivers some output $\mathcal{A}(\mathcal{D})$. The output could be the learnt model parameters $\hat{\mathbf{w}}$ or the prediction $\hat{h}(\mathbf{x})$ obtained for a specific data point with features \mathbf{x} . Many important ML applications involve data points representing humans. Each data point is characterized by features \mathbf{x} , potentially a label y and a sensitive attribute s (e.g., a recent medical diagnosis). Roughly speaking, privacy protection means that it should be impossible to infer, from the output $\mathcal{A}(\mathcal{D})$, any of the sensitive attributes of data points in \mathcal{D} . Mathematically, privacy protection requires non-invertibility of the map $\mathcal{A}(\mathcal{D})$. In general, just making $\mathcal{A}(\mathcal{D})$ non-invertible is typically insufficient for privacy protection. We need to make $\mathcal{A}(\mathcal{D})$ sufficiently non-invertible.. 2, 37

probabilistic model A probabilistic model interprets data points as realizations of RVs with a joint probability distribution. This joint probability distribution typically involves parameters which have to be manually chosen or learnt via statistical inference methods such as maximum likelihood [27].. 3–5, 12, 15–21, 28, 31, 37, 45, 46, 54, 88, 111, 123, 128–130, 140, 144

probability density function (pdf) The probability density function (pdf) $p(x)$ of a real-valued RV $x \in \mathbb{R}$ is a particular representation of its probability distribution. If the pdf exists, it can be used to compute the probability that x takes on a value from a (measurable) set $\mathcal{B} \subseteq \mathbb{R}$ via $p(x \in \mathcal{B}) = \int_{\mathcal{B}} p(x') dx'$ [5, Ch. 3]. The pdf of a vector-valued RV $\mathbf{x} \in \mathbb{R}^d$ (if it exists) allows to compute the probability that \mathbf{x} falls into a (measurable) region \mathcal{R} via $p(\mathbf{x} \in \mathcal{R}) = \int_{\mathcal{R}} p(\mathbf{x}') dx'_1 \dots dx'_d$ [5, Ch. 3].
19, 25

probability distribution To analyze ML methods it can be useful to interpret data points as i.i.d. realizations of a RV. The typical properties of such data points are then governed by the probability distribution of this RV. The probability distribution of a binary RV $y \in \{0, 1\}$ is fully specified by the probabilities $p(y = 0)$ and $p(y = 1) = 1 - p(y = 0)$. The probability distribution of a real-valued RV $x \in \mathbb{R}$ might be specified by a probability density function $p(x)$ such that $p(x \in [a, b]) \approx p(a)|b - a|$. In the most general case, a probability distribution is defined by a probability measure [106, 119]. 1, 3–7, 10–13, 17, 19–21, 24, 25, 28, 29, 32, 34, 37–39, 43, 45, 48, 71, 108, 128, 134

projected gradient descent (projected GD) Projected GD extends basic GD for unconstrained optimization to handle constraints on the optimization variable (model parameters). A single iteration of projected GD consists of first taking a gradient step and then projecting the result back into a constrain set.. 29, 57, 69, 78, 90, 132

proximable A convex function for which the proximal operator can be

computed efficiently are sometimes referred to as *proximable* or *simple* [34].. 43

proximal operator Given a convex function and a vector \mathbf{w}' , we define its proximal operator as [35, 52]

$$\mathbf{prox}_{L_i(\cdot), 2\alpha}(\mathbf{w}') := \underset{\mathbf{w} \in \mathbb{R}^d}{\operatorname{argmin}} \left[f(\mathbf{w}) + (\rho/2) \|\mathbf{w} - \mathbf{w}'\|_2^2 \right] \text{ with } \rho > 0.$$

Convex functions for which the proximal operator can be computed efficiently are sometimes referred to as *proximable* or *simple* [34].. 22, 38, 43, 70, 91

quadratic function A quadratic function $f(\mathbf{w})$, reading in a vector $\mathbf{w} \in \mathbb{R}^d$ as its argument, is such that

$$f(\mathbf{w}) = \mathbf{w}^T \mathbf{Q} \mathbf{w} + \mathbf{q}^T \mathbf{w} + a,$$

with some matrix $\mathbf{Q} \in \mathbb{R}^{d \times d}$, vector $\mathbf{q} \in \mathbb{R}^d$ and scalar $a \in \mathbb{R}$.. 15, 18, 19, 27, 42, 58, 78, 143

Rényi divergence The Rényi divergence measures the (dis-)similarity between two probability distributions [?]. 7

random variable (RV) A random variable is a mapping from a probability space \mathcal{P} to a value space [106]. The probability space, whose elements are elementary events, is equipped with a probability measure that assigns a probability to subsets of \mathcal{P} . A binary random variable maps elementary events to a set containing two different values, e.g., $\{-1, 1\}$ or $\{\text{cat}, \text{no cat}\}$. A real-valued random variable maps elementary events

to real numbers \mathbb{R} . A vector-valued random variable maps elementary events to the Euclidean space \mathbb{R}^d . Probability theory uses the concept of measurable spaces to rigorously define and study the properties of (large) collections of random variables [106, 119].. 5, 6, 9, 11–13, 16–20, 24, 25, 28, 29, 32–34, 37, 38, 40, 41, 43, 45, 49, 54, 71, 75, 88, 111, 124, 128, 129, 134, 138

realization Consider a RV x which maps each element (outcome, or elementary event) $\omega \in \mathcal{P}$ of a probability space \mathcal{P} to an element a of a measurable space \mathcal{N} [2, 105, 106]. A realization of x is any element $a' \in \mathcal{N}$ such that there is an element $\omega' \in \mathcal{P}$ with $x(\omega') = a'$.. 5, 6, 11–13, 16–18, 20, 24–28, 37, 38, 41, 43, 45, 48, 71, 75, 88, 108, 111, 124, 128, 129, 134

regression Regression problems revolve around the problem of predicting a numeric label solely from the features of a data point [6, Ch. 2].. 15

regularization A key challenge of modern ML applications is that they often use large models, having an effective dimension in the order of billions. Using basic ERM-based methods to train a high-dimensional model is prone to overfitting: the learn hypothesis performs well on the training set but poorly outside the training set. Regularization refers to modifications of a given instance of ERM in order to avoid overfitting, i.e., to ensure the learnt hypothesis performs not much worse outside the training set. There are three routes for implementing regularization:

- **Model pruning.** We prune the original model \mathcal{H} to obtain a smaller model \mathcal{H}' . For a parametrized model, the pruning can

be implemented by including constraints on the model parameters (such as $w_1 \in [0.4, 0.6]$ for the weight of feature x_1 in linear regression).

- **Loss penalization.** We modify the objective function of ERM by adding a penalty term to the training error. The penalty term estimates how much larger the expected loss (risk) is compared to the average loss on the training set.
- **Data augmentation.** We can enlarge the training set \mathcal{D} by adding perturbed copies of the original data points in \mathcal{D} . One example for such a perturbation is to add the realization of a RV to the feature vector of a data point.

Figure 10.9 illustrates the above routes to regularization. These routes are often equivalent: data augmentation using Gaussian RVs to perturb the feature vectors in the training set of linear regression has the same effect as adding the penalty $\lambda \|\mathbf{w}\|_2^2$ to the training error (which is nothing but ridge regression).




Fig. 10.9: Three approaches to regularization: data augmentation, loss penalization and model pruning (via constraints on model parameters).

. 7, 9, 17, 23, 26, 28, 30, 43, 47, 68, 116, 152

regularized empirical risk minimization (RERM) Synonym for SRM..

7, 24, 41, 93

regularizer A regularizer assigns each hypothesis h from a hypothesis space \mathcal{H} a quantitative measure $\mathcal{R}\{h\}$ for how much its prediction error on a training set might differ from its prediction errors on data points outside the training set. Ridge regression uses the regularizer $\mathcal{R}\{h\} := \|\mathbf{w}\|_2^2$ for linear hypothesis maps $h^{(\mathbf{w})}(\mathbf{x}) := \mathbf{w}^T \mathbf{x}$ [6, Ch. 3]. The least absolute shrinkage and selection operator (Lasso) uses the regularizer $\mathcal{R}\{h\} := \|\mathbf{w}\|_1$ for linear hypothesis maps $h^{(\mathbf{w})}(\mathbf{x}) := \mathbf{w}^T \mathbf{x}$ [6, Ch. 3].

7, 24, 26, 41, 92, 140

ridge regression Ridge regression learns the weights \mathbf{w} of a linear hypoth-

esis map $h^{(\mathbf{w})}(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$. The quality of a particular choice for the parameter vector \mathbf{w} is measured by the sum of two components. The first component is the average squared error loss incurred by $h^{(\mathbf{w})}$ on a set of labeled data points (the training set). The second component is the scaled squared Euclidean norm $\alpha \|\mathbf{w}\|_2^2$ with a regularization parameter $\alpha > 0$. It can be shown that the effect of adding to $\alpha \|\mathbf{w}\|_2^2$ to the average squared error loss is equivalent to replacing the original data points by an ensemble of realizations of a RV centered around these data points.. 26–29, 41, 42, 58, 59, 65, 92

risk Consider a hypothesis h used to predict the label y of a data point based on its features \mathbf{x} . We measure the quality of a particular prediction using a loss function $L((\mathbf{x}, y), h)$. If we interpret data points as the realizations of i.i.d. RVs, also the $L((\mathbf{x}, y), h)$ becomes the realization of a RV. The i.i.d. assumption allows to define the risk of a hypothesis as the expected loss $\mathbb{E}\{L((\mathbf{x}, y), h)\}$. Note that the risk of h depends on both, the specific choice for the loss function and the probability distribution of the data points.. 3–5, 17, 21, 35, 41, 48

scatterplot A visualization technique that depicts data points by markers in a two-dimensional plane. . 4, 16

semi-supervised learning Semi-supervised learning methods use unlabeled data points to support the learning of a hypothesis from labeled data points [66]. This approach is particularly useful for ML applications that offer a large amount of unlabeled data points, but only a limited number of labeled data points.. 112, 114




Fig. 10.10: A scatterplot of data points that represent daily weather conditions in Finland. Each data point is characterized by its minimum daytime temperature x as feature and its maximum daytime temperature y as the label. The temperatures have been measured at the FMI weather station *Helsinki Kaisaniemi* during 1.9.2024 - 28.10.2024.

sensitive attribute ML revolves around learning a hypothesis map that allows to predict the label of a data point from its features. In some applications we must ensure that the output delivered by an ML system does not allow to infer sensitive attributes of a data point. Which parts of a data point is considered as a sensitive attribute is a design choice that varies across different application domains.. 12, 16, 37

smooth We refer to a real-valued function as smooth if it is differentiable and its gradient is continuous [51, 143]. A differentiable function $f(\mathbf{w})$ is referred to as β -smooth if the gradient $\nabla f(\mathbf{w})$ is Lipschitz continuous with Lipschitz constant β , i.e.,

$$\|\nabla f(\mathbf{w}) - \nabla f(\mathbf{w}')\| \leq \beta \|\mathbf{w} - \mathbf{w}'\|.$$

. 14, 15, 36, 49, 57

soft clustering Soft clustering refers to the task of partitioning a given set of data points into (few) overlapping clusters. Each data point is assigned

to several different clusters with varying degree of belonging. Soft clustering methods determine the degree of belonging (or soft cluster assignment) for each data point and each cluster. A principled approach to soft clustering is by interpreting data points as i.i.d. realizations of a GMM. We then obtain a natural choice for the degree of belonging as the conditional probability of a data point belonging to a specific mixture component.. 7, 11

squared error loss The squared error loss measures the prediction error of a hypothesis h when predicting a numeric label $y \in \mathbb{R}$ from the features \mathbf{x} of a data point. It is defined as

$$L((\mathbf{x}, y), h) := \left(y - \underbrace{h(\mathbf{x})}_{=\hat{y}} \right)^2.$$

. 3, 4, 7, 13, 17, 28–31, 36, 42, 43, 49, 118, 119, 152

statistical aspects By statistical aspects of a ML method, we refer to (properties of) the probability distribution of its output under a probabilistic model for the data fed into the method.. 15, 28, 30, 41, 48

step size See learning rate.. 21, 22, 57

stochastic gradient descent Stochastic GD is obtained from GD by replacing the gradient of the objective function with a stochastic approximation. A main application of stochastic gradient descent is to implement ERM for a parametrized model and a training set \mathcal{D} is either very large or not readily available (e.g., when data points are stored in a database distributed all over the planet). To evaluate the gradient



Fig. 10.11: Stochastic GD for ERM approximates the gradient $\sum_{r=1}^m \nabla_{\mathbf{w}} L(\mathbf{z}^{(r)}, \mathbf{w})$ by replacing the sum over all data points in the training set (indexed by $r = 1, \dots, m$) with a sum over a randomly chosen subset $\mathcal{B} \subseteq \{1, \dots, m\}$.

of the empirical risk (as a function of the model parameters \mathbf{w}), we need to compute a sum $\sum_{r=1}^m \nabla_{\mathbf{w}} L(\mathbf{z}^{(r)}, \mathbf{w})$ over all data points in the training set. We obtain a stochastic approximation to the gradient by replacing the sum $\sum_{r=1}^m \nabla_{\mathbf{w}} L(\mathbf{z}^{(r)}, \mathbf{w})$ with a sum $\sum_{r \in \mathcal{B}} \nabla_{\mathbf{w}} L(\mathbf{z}^{(r)}, \mathbf{w})$ over a randomly chosen subset $\mathcal{B} \subseteq \{1, \dots, m\}$ (see Figure 10.11). We often refer to these randomly chosen data points as a *batch*. The batch size $|\mathcal{B}|$ is an important parameter of stochastic GD. Stochastic GD with $|\mathcal{B}| > 1$ is referred to as mini-batch stochastic GD [?]. . 4, 7, 11, 29, 78, 130

stopping criterion Many ML methods use iterative algorithms that construct a sequence of model parameters (such as the weights of a linear map or the weights of an ANN) that (hopefully) converge to an optimal choice for the model parameters. In practice, given finite computational resources, we need to stop iterating after a finite number of times. A stopping criterion is any well-defined condition required for stopping

iterating.. 45, 57, 59, 60, 62, 63, 69, 73, 82, 83, 85, 87, 88, 93, 150

structural risk minimization Structural risk minimization is the problem of finding the hypothesis that optimally balances the average loss (or empirical risk) on a training set with a regularization term. The regularization term penalizes a hypothesis that is not robust against (small) perturbations of the data points in the training set.. 30, 42

test set A set of data points that have neither been used to train a model, e.g., via ERM, nor in a validation set to choose between different models.. 2, 3, 35

total variation See GTV.. 35, 38, 41, 77, 105

training error The average loss of a hypothesis when predicting the labels of data points in a training set. We sometimes refer by training error also the minimum average loss incurred on the training set by the optimal hypothesis from a hypothesis space.. 4, 11, 17, 21–25, 28, 29, 41, 47, 48, 71, 152

training set A dataset \mathcal{D} , constituted by some data points used in ERM to learn a hypothesis \hat{h} . The average loss of \hat{h} on the training set is referred to as the training error. The comparison of the training error with the validation error of \hat{h} allows to diagnose the ML method and informs how to improve them (e.g., using a different hypothesis space or collecting more data points) [6, Sec. 6.6].. 1–4, 6, 8, 12, 14–18, 21–31, 35, 36, 40–43, 45–48, 59, 60, 63, 65, 66, 73, 108, 118, 141

trustworthiness Beside the computational aspects and statistical aspects, a third main design aspect for ML methods is their trustworthiness [?]. The European Union has put forward seven key requirements (KRs) for trustworthy AI (that typically build on ML methods) [84]: **KR1 - Human Agency and Oversight**, **KR2 - Technical Robustness and Safety**, **KR3 - Privacy and Data Governance**, **KR4 - Transparency**, **KR5 - Diversity Non-Discrimination and Fairness**, **KR6 Societal and Environmental Well-Being**, **KR7 - Accountability**. . 135

validation Consider a hypothesis \hat{h} that has been learnt via some ML method, e.g., by solving ERM on a training set \mathcal{D} . Validation refers to the practice of evaluating the loss incurred by hypothesis \hat{h} on a validation set that consists of data points that are not contained in the training set \mathcal{D} . 11, 21, 62

validation error Consider a hypothesis \hat{h} which is obtained by some ML method, e.g., using ERM on a training set. The average loss of \hat{h} on a validation set, which is different from the training set, is referred to as the validation error.. 8, 11, 16, 19, 21–25, 28, 47, 48, 137, 140

validation set A set of data points used to estimate the risk of a hypothesis \hat{h} that has been learnt by some ML method (e.g., solving ERM). The average loss of \hat{h} on the validation set is referred to as the validation error and can be used to diagnose a ML method (see [6, Sec. 6.6.]). The comparison between training error and validation error can inform directions for improvements of the ML method (such as using a different

hypothesis space).. 16, 17, 21, 22, 24, 25, 47, 48, 137

variance The variance of a real-valued RV x is defined as the expectation $\mathbb{E}\{(x - \mathbb{E}\{x\})^2\}$ of the squared difference between x and its expectation $\mathbb{E}\{x\}$. We extend this definition to vector-valued RVs \mathbf{x} as $\mathbb{E}\{\|\mathbf{x} - \mathbb{E}\{\mathbf{x}\}\|_2^2\}$.. 4, 16

vertical FL Vertical FL uses local datasets that are constituted by the same data points but characterizing them with different features [67]. For example, different healthcare providers might all contain information about the same population of patients. However, different healthcare providers collect different measurements (blood values, electrocardiography, lung X-ray) for the same patients.. 105, 114, 115

weights Consider a parameterized hypothesis space \mathcal{H} . We use the term weights for numeric model parameters that are used to scale features or their transformations in order to compute $h^{(\mathbf{w})} \in \mathcal{H}$. A linear model uses weights $\mathbf{w} = (w_1, \dots, w_d)^T$ to compute the linear combination $h^{(\mathbf{w})}(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$. Weights are also used in ANNs to form linear combinations of features or the outputs of neurons in hidden layers.. 42

zero-gradient condition Consider the unconstrained optimization problem $\min_{\mathbf{w} \in \mathbb{R}^d} f(\mathbf{w})$ with a smooth and convex objective function $f(\mathbf{w})$. A necessary and sufficient condition for a vector $\hat{\mathbf{w}} \in \mathbb{R}^d$ to solve this problem is that the gradient $\nabla f(\hat{\mathbf{w}})$ is the zero-vector,

$$\nabla f(\hat{\mathbf{w}}) = \mathbf{0} \Leftrightarrow f(\hat{\mathbf{w}}) = \min_{\mathbf{w} \in \mathbb{R}^d} f(\mathbf{w}).$$

. 44

0/1 loss The 0/1 loss $L((\mathbf{x}, y), h)$ measures the quality of a classifier $h(\mathbf{x})$ that delivers a prediction \hat{y} (e.g., via thresholding (10.1)) for the label y of a data point with features \mathbf{x} . It is equal to 0 if the prediction is correct, i.e., $L((\mathbf{x}, y), h) = 0$ when $\hat{y} = y$. It is equal to 1 if the prediction is wrong, $L((\mathbf{x}, y), h) = 1$ when $\hat{y} \neq y$. 1

References

- [1] W. Rudin, *Real and Complex Analysis*, 3rd ed. New York: McGraw-Hill, 1987.
- [2] ———, *Principles of Mathematical Analysis*, 3rd ed. New York: McGraw-Hill, 1976.
- [3] G. H. Golub and C. F. Van Loan, *Matrix Computations*, 4th ed. Baltimore, MD: Johns Hopkins University Press, 2013.
- [4] G. Golub and C. van Loan, “An analysis of the total least squares problem,” *SIAM J. Numerical Analysis*, vol. 17, no. 6, pp. 883–893, Dec. 1980.
- [5] D. Bertsekas and J. Tsitsiklis, *Introduction to Probability*, 2nd ed. Athena Scientific, 2008.
- [6] A. Jung, *Machine Learning: The Basics*, 1st ed. Springer Singapore, Feb. 2022.
- [7] M. Wollschlaeger, T. Sauter, and J. Jasperneite, “The future of industrial communication: Automation networks in the era of the internet of things and industry 4.0,” *IEEE Industrial Electronics Magazine*, vol. 11, no. 1, pp. 17–27, 2017.
- [8] M. Satyanarayanan, “The emergence of edge computing,” *Computer*, vol. 50, no. 1, pp. 30–39, Jan. 2017. [Online]. Available: <https://doi.org/10.1109/MC.2017.9>

- [9] H. Ates, A. Yetisen, F. Güder, and C. Dincer, “Wearable devices for the detection of covid-19,” *Nature Electronics*, vol. 4, no. 1, pp. 13–14, 2021. [Online]. Available: <https://doi.org/10.1038/s41928-020-00533-1>
- [10] H. Boyes, B. Hallaq, J. Cunningham, and T. Watson, “The industrial internet of things (iiot): An analysis framework,” *Computers in Industry*, vol. 101, pp. 1–12, 2018. [Online]. Available: <https://www.sciencedirect.com/science/article/pii/S0166361517307285>
- [11] S. Cui, A. Hero, Z.-Q. Luo, and J. Moura, Eds., *Big Data over Networks*. Cambridge Univ. Press, 2016.
- [12] A. Barabási, N. Gulbahce, and J. Loscalzo, “Network medicine: a network-based approach to human disease,” *Nature Reviews Genetics*, vol. 12, no. 56, 2011.
- [13] M. E. J. Newman, *Networks: An Introduction*. Oxford Univ. Press, 2010.
- [14] B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. A. y. Arcas, “Communication-Efficient Learning of Deep Networks from Decentralized Data,” in *Proceedings of the 20th International Conference on Artificial Intelligence and Statistics*, ser. Proceedings of Machine Learning Research, A. Singh and J. Zhu, Eds., vol. 54. PMLR, 20–22 Apr 2017, pp. 1273–1282. [Online]. Available: <https://proceedings.mlr.press/v54/mcmahan17a.html>

- [15] T. Li, A. K. Sahu, A. Talwalkar, and V. Smith, “Federated learning: Challenges, methods, and future directions,” *IEEE Signal Processing Magazine*, vol. 37, no. 3, pp. 50–60, May 2020.
- [16] Y. Cheng, Y. Liu, T. Chen, and Q. Yang, “Federated learning for privacy-preserving ai,” *Communications of the ACM*, vol. 63, no. 12, pp. 33–36, Dec. 2020.
- [17] N. Agarwal, A. Suresh, F. Yu, S. Kumar, and H. McMahan, “cpSGD: Communication-efficient and differentially-private distributed sgd,” in *Proc. Neural Inf. Proc. Syst. (NIPS)*, 2018.
- [18] V. Smith, C.-K. Chiang, M. Sanjabi, and A. Talwalkar, “Federated Multi-Task Learning,” in *Advances in Neural Information Processing Systems*, vol. 30, 2017. [Online]. Available: <https://proceedings.neurips.cc/paper/2017/file/6211080fa89981f66b1a0c9d55c61d0f-Paper.pdf>
- [19] J. You, J. Wu, X. Jin, and M. Chowdhury, “Ship compute or ship data? why not both?” in *18th USENIX Symposium on Networked Systems Design and Implementation (NSDI 21)*. USENIX Association, April 2021, pp. 633–651. [Online]. Available: <https://www.usenix.org/conference/nsdi21/presentation/you>
- [20] D. Tse and P. Viswanath, *Fundamentals of Wireless Communication*. Cambridge University Press, 2005.
- [21] T. Yang, G. Andrew, H. Eichner, H. Sun, W. Li, N. Kong, D. Ramage, and F. Beaufays, “Applied federated learning: Improving

- google keyboard query suggestions,” 2018. [Online]. Available: <https://arxiv.org/abs/1812.02903>
- [22] A. Ghosh, J. Chung, D. Yin, and K. Ramchandran, “An efficient framework for clustered federated learning,” in *34th Conference on Neural Information Processing Systems (NeurIPS 2020)*, Vancouver, Canada, 2020.
 - [23] F. Sattler, K. Müller, and W. Samek, “Clustered federated learning: Model-agnostic distributed multitask optimization under privacy constraints,” *IEEE Transactions on Neural Networks and Learning Systems*, 2020.
 - [24] G. Strang, *Computational Science and Engineering*. Wellesley-Cambridge Press, MA, 2007.
 - [25] —, *Introduction to Linear Algebra*, 5th ed. Wellesley-Cambridge Press, MA, 2016.
 - [26] H. H. Bauschke and P. L. Combettes, *Convex Analysis and Monotone Operator Theory in Hilbert Spaces*. New York: Springer, 2011.
 - [27] E. L. Lehmann and G. Casella, *Theory of Point Estimation*, 2nd ed. New York: Springer, 1998.
 - [28] A. Esteva, B. Kuprel, R. A. Novoa, J. Ko, S. M. Swetter, H. M. Blau, and S. Thrun, “Dermatologist-level classification of skin cancer with deep neural networks,” *Nature*, vol. 542, 2017.

- [29] H. Lütkepohl, *New Introduction to Multiple Time Series Analysis*. New York: Springer, 2005.
- [30] Y. SarcheshmehPour, Y. Tian, L. Zhang, and A. Jung, “Clustered federated learning via generalized total variation minimization,” *IEEE Transactions on Signal Processing*, vol. 71, pp. 4240–4256, 2023.
- [31] F. Chung, “Spectral graph theory,” in *Regional Conference Series in Mathematics*, 1997, no. 92.
- [32] D. Spielman, “Spectral and algebraic graph theory,” 2019.
- [33] —, “Spectral graph theory,” in *Combinatorial Scientific Computing*, U. Naumann and O. Schenk, Eds. Chapman and Hall/CRC, 2012.
- [34] L. Condat, “A primal–dual splitting method for convex optimization involving lipschitzian, proximable and linear composite terms,” *Journal of Opt. Th. and App.*, vol. 158, no. 2, pp. 460–479, Aug. 2013.
- [35] N. Parikh and S. Boyd, “Proximal algorithms,” *Foundations and Trends in Optimization*, vol. 1, no. 3, pp. 123–231, 2013.
- [36] R. Peng and D. A. Spielman, “An efficient parallel solver for SDD linear systems,” in *Proc. ACM Symposium on Theory of Computing*, New York, NY, 2014, pp. 333–342.
- [37] N. K. Vishnoi, “ $Lx = b$ — Laplacian solvers and their algorithmic applications,” *Foundations and Trends in Theoretical Computer Science*, vol. 8, no. 1–2, pp. 1–141, 2012. [Online]. Available: <http://dx.doi.org/10.1561/04000000054>

- [38] D. Sun, K.-C. Toh, and Y. Yuan, “Convex clustering: Model, theoretical guarantee and efficient algorithm,” *Journal of Machine Learning Research*, vol. 22, no. 9, pp. 1–32, 2021. [Online]. Available: <http://jmlr.org/papers/v22/18-694.html>
- [39] K. Pelckmans, J. D. Brabanter, J. Suykens, and B. D. Moor, “Convex clustering shrinkage,” in *PASCAL Workshop on Statistics and Optimization of Clustering Workshop*, 2005.
- [40] R. T. Rockafellar, *Convex Analysis*. Princeton, NJ: Princeton Univ. Press, 1970.
- [41] S. Boyd and L. Vandenberghe, *Convex Optimization*. Cambridge, UK: Cambridge Univ. Press, 2004.
- [42] R. T. Rockafellar, *Network Flows and Monotropic Optimization*. Athena Scientific, Jul. 1998.
- [43] D. P. Bertsekas, *Network Optimization: Continuous and Discrete Models*. Athena Scientific, 1998.
- [44] C. G. Atkeson, A. W. Moore, and S. Schaal, “Locally weighted learning,” *Artif. Intell. Rev.*, vol. 11, no. 1–5, pp. 11–73, feb 1997. [Online]. Available: <https://doi.org/10.1023/A:1006559212014>
- [45] A. Jacot, F. Gabriel, and C. Hongler, “Neural tangent kernel: Convergence and generalization in neural networks,” in *Advances in Neural Information Processing Systems*, S. Bengio, H. Wallach, H. Larochelle, K. Grauman, N. Cesa-Bianchi, and

- R. Garnett, Eds., vol. 31. Curran Associates, Inc., 2018. [Online]. Available: https://proceedings.neurips.cc/paper_files/paper/2018/file/5a4be1fa34e62bb8a6ec6b91d2462f5a-Paper.pdf
- [46] S. S. Du, X. Zhai, B. Póczos, and A. Singh, “Gradient descent provably optimizes over-parameterized neural networks,” in *7th International Conference on Learning Representations, ICLR 2019, New Orleans, LA, USA, May 6-9, 2019*. OpenReview.net, 2019. [Online]. Available: <https://openreview.net/forum?id=S1eK3i09YQ>
- [47] W. E. C. Ma, and L. Wu, “A comparative analysis of optimization and generalization properties of two-layer neural network and random feature models under gradient descent dynamics,” *Science China Mathematics*, vol. 63, no. 7, pp. 1235–1258, 2020. [Online]. Available: <https://doi.org/10.1007/s11425-019-1628-5>
- [48] A. Paszke, S. Gross, F. Massa, A. Lerer, J. Bradbury, G. Chanan, T. Killeen, Z. Lin, N. Gimelshein, L. Antiga, A. Desmaison, A. Köpf, E. Yang, Z. DeVito, M. Raison, A. Tejani, S. Chilamkurthy, B. Steiner, L. Fang, J. Bai, and S. Chintala, *PyTorch: An Imperative Style, High-Performance Deep Learning Library*. Red Hook, NY, USA: Curran Associates Inc., 2019.
- [49] D. P. Bertsekas, *Convex Optimization Algorithms*. Athena Scientific, 2015.

- [50] Y. L. T. Schaul, X. Zhang, “No more pesky learning rates,” in *Proc. of the 30th International Conference on Machine Learning, PMLR 28(3)*, vol. 28, Atlanta, Georgia, June 2013, pp. 343–351.
- [51] Y. Nesterov, *Introductory lectures on convex optimization*, ser. Applied Optimization. Kluwer Academic Publishers, Boston, MA, 2004, vol. 87, a basic course. [Online]. Available: <http://dx.doi.org/10.1007/978-1-4419-8853-9>
- [52] H. Bauschke and P. Combettes, *Convex Analysis and Monotone Operator Theory in Hilbert Spaces*, 2nd ed. New York: Springer, 2017.
- [53] V. Istrăţescu, *Fixed point theory: An Introduction*, ser. Mathematics and its applications ; 7. Dordrecht: Reidel, 1981.
- [54] D. P. Bertsekas and J. Tsitsiklis, *Parallel and Distributed Computation: Numerical Methods*. Athena Scientific, 2015.
- [55] M. van Steen and A. Tanenbaum, *Distributed Systems*, 3rd ed., Feb. 2017, self-published, open publication.
- [56] D. Mills, “Internet time synchronization: the network time protocol,” *IEEE Transactions on Communications*, vol. 39, no. 10, pp. 1482–1493, 1991.
- [57] J. Hirvonen and J. Suomela. (2023) Distributed algorithms 2020.
- [58] R. Diestel, *Graph Theory*. Springer Berlin Heidelberg, 2005.

- [59] F. Pedregosa, “Scikit-learn: Machine learning in python,” *Journal of Machine Learning Research*, vol. 12, no. 85, pp. 2825–2830, 2011. [Online]. Available: <http://jmlr.org/papers/v12/pedregosa11a.html>
- [60] A. Tanenbaum and D. Wetherall, *Computer Networks*, 5th ed. USA: Prentice Hall Press, 2010.
- [61] P. Tseng, “Convergence of a block coordinate descent method for nondifferentiable minimization,” *Journal of Optimization Theory and Applications*, vol. 109, no. 3, pp. 475–494, 2001. [Online]. Available: <https://doi.org/10.1023/A:1017501703105>
- [62] S. Boyd, N. Parikh, E. Chu, B. Peleato, and J. Eckstein, *Distributed Optimization and Statistical Learning via the Alternating Direction Method of Multipliers*. Hanover, MA: Now Publishers, 2010, vol. 3, no. 1.
- [63] Q. Yang, Y. Liu, Y. Cheng, Y. Kang, T. Chen, and H. Yu, *Federated Learning*, 1st ed. Springer, 2022.
- [64] D. J. Spiegelhalter, “An omnibus test for normality for small samples,” *Biometrika*, vol. 67, no. 2, pp. 493–496, 2024/03/25/ 1980. [Online]. Available: <http://www.jstor.org/stable/2335498>
- [65] Q. Yang, Y. Liu, Y. Cheng, Y. Kang, T. Chen, and H. Yu, *Horizontal Federated Learning*. Cham: Springer International Publishing, 2020, pp. 49–67. [Online]. Available: https://doi.org/10.1007/978-3-031-01585-4_4

- [66] O. Chapelle, B. Schölkopf, and A. Zien, Eds., *Semi-Supervised Learning*. Cambridge, Massachusetts: The MIT Press, 2006.
- [67] Q. Yang, Y. Liu, Y. Cheng, Y. Kang, T. Chen, and H. Yu, *Vertical Federated Learning*. Cham: Springer International Publishing, 2020, pp. 69–81. [Online]. Available: https://doi.org/10.1007/978-3-031-01585-4_5
- [68] H. Ludwig and N. Baracaldo, Eds., *Federated Learning: A Comprehensive Overview of Methods and Applications*. Springer, 2022.
- [69] A. Shamsian, A. Navon, E. Fetaya, and G. Chechik, “Personalized federated learning using hypernetworks,” in *Proceedings of the 38th International Conference on Machine Learning*, ser. Proceedings of Machine Learning Research, M. Meila and T. Zhang, Eds., vol. 139. PMLR, 18–24 Jul 2021, pp. 9489–9502. [Online]. Available: <https://proceedings.mlr.press/v139/shamsian21a.html>
- [70] F. Sung, Y. Yang, L. Zhang, T. Xiang, P. H. Torr, and T. M. Hospedales, “Learning to compare: Relation network for few-shot learning,” in *2018 IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2018, pp. 1199–1208.
- [71] V. Satorras and J. Bruna, “Few-shot learning with graph neural networks.” in *ICLR (Poster)*. OpenReview.net, 2018. [Online]. Available: <http://dblp.uni-trier.de/db/conf/iclr/iclr2018.html#SatorrasE18>

- [72] B. Bollobas, W. Fulton, A. Katok, F. Kirwan, and P. Sarnak, *Random graphs*. Cambridge studies in advanced mathematics., 2001, vol. 73.
- [73] J. Tropp, “An introduction to matrix concentration inequalities,” *Found. Trends Mach. Learn.*, May 2015.
- [74] G. Keiser, *Optical Fiber Communication*, 4th ed. New Delhi: Mc-Graw Hill, 2011.
- [75] D. Tse and P. Viswanath, *Fundamentals of wireless communication*. USA: Cambridge University Press, 2005.
- [76] M. Fiedler, “Algebraic connectivity of graphs,” *Czechoslovak Mathematical Journal*, vol. 23, no. 2, pp. 298–305, 1973.
- [77] S. Hoory, N. Linial, and A. Wigderson, “Expander graphs and their applications,” *Bull. Amer. Math. Soc.*, vol. 43, no. 04, pp. 439–562, Aug. 2006.
- [78] B. Ying, K. Yuan, Y. Chen, H. Hu, P. PAN, and W. Yin, “Exponential graph is provably efficient for decentralized deep training,” in *Advances in Neural Information Processing Systems*, M. Ranzato, A. Beygelzimer, Y. Dauphin, P. Liang, and J. W. Vaughan, Eds., vol. 34. Curran Associates, Inc., 2021, pp. 13 975–13 987. [Online]. Available: https://proceedings.neurips.cc/paper_files/paper/2021/file/74e1ed8b55ea44fd7dbb685c412568a4-Paper.pdf
- [79] Y.-T. Chow, W. Shi, T. Wu, and W. Yin, “Expander graph and communication-efficient decentralized optimization,” in *2016 50th Asilo-*

- mar Conference on Signals, Systems and Computers*, 2016, pp. 1715–1720.
- [80] S. M. Kay, *Fundamentals of Statistical Signal Processing: Estimation Theory*. Englewood Cliffs, NJ: Prentice Hall, 1993.
 - [81] S. Chepuri, S. Liu, G. Leus, and A. Hero, “Learning sparse graphs under smoothness prior,” in *Proc. of the IEEE Int. Conf. on Acoustics, Speech and Signal Processing*, 2017, pp. 6508–6512.
 - [82] J. Tan, Y. Zhou, G. Liu, J. H. Wang, and S. Yu, “pFedSim: Similarity-Aware Model Aggregation Towards Personalized Federated Learning,” *arXiv e-prints*, p. arXiv:2305.15706, May 2023.
 - [83] I. Goodfellow, Y. Bengio, and A. Courville, *Deep Learning*. MIT Press, 2016.
 - [84] E. Commission, C. Directorate-General for Communications Networks, and Technology, *The Assessment List for Trustworthy Artificial Intelligence (ALTAI) for self assessment*. Publications Office, 2020.
 - [85] H.-L. E. G. on Artificial Intelligence, “Ethics guidelines for trustworthy AI,” European Commission, Tech. Rep., April 2019.
 - [86] D. Kuss and O. Lopez-Fernandez, “Internet addiction and problematic internet use: A systematic review of clinical research.” *World J Psychiatry*, vol. 6, no. 1, pp. 143–176, Mar 2016.
 - [87] L. Munn, “Angry by design: toxic communication and technical architectures,” *Humanities and Social Sciences Communications*, vol. 7,

- no. 1, p. 53, 2020. [Online]. Available: <https://doi.org/10.1057/s41599-020-00550-7>
- [88] P. Mozur, “A genocide incited on facebook, with posts from myanmar’s military,” *The New York Times*, 2018.
- [89] A. Simchon, M. Edwards, and S. Lewandowsky, “The persuasive effects of political microtargeting in the age of generative artificial intelligence.” *PNAS Nexus*, vol. 3, no. 2, p. pgae035, Feb 2024.
- [90] J. Near and D. Darais, “Guidelines for evaluating differential privacy guarantees,” National Institute of Standards and Technology, Gaithersburg, MD, Tech. Rep., 2023.
- [91] S. Wachter, “Data protection in the age of big data,” *Nature Electronics*, vol. 2, no. 1, pp. 6–7, 2019. [Online]. Available: <https://doi.org/10.1038/s41928-018-0193-y>
- [92] P. Samarati, “Protecting respondents identities in microdata release,” *IEEE Transactions on Knowledge and Data Engineering*, vol. 13, no. 6, pp. 1010–1027, 2001.
- [93] E. Comission, “Regulation (eu) 2016/679 of the european parliament and of the council of 27 april 2016 on the protection of natural persons with regard to the processing of personal data and on the free movement of such data, and repealing directive 95/46/ec (general data protection regulation) (text with eea relevance),” no. 119, pp. 1–88, May 2016.
- [94] U. N. G. Assembly, *The Universal Declaration of Human Rights (UDHR)*, New York, 1948.

- [95] N. Kozodoi, J. Jacob, and S. Lessmann, “Fairness in credit scoring: Assessment, implementation and profit implications,” *European Journal of Operational Research*, vol. 297, no. 3, pp. 1083–1094, 2022. [Online]. Available: <https://www.sciencedirect.com/science/article/pii/S0377221721005385>
- [96] J. Gonçalves-Sá and F. Pinheiro, *Societal Implications of Recommendation Systems: A Technical Perspective*. Cham: Springer International Publishing, 2024, pp. 47–63. [Online]. Available: https://doi.org/10.1007/978-3-031-41264-6_3
- [97] A. Abrol and R. Jha, “Power optimization in 5g networks: A step towards green communication,” *IEEE Access*, vol. 4, pp. 1355–1374, 2016.
- [98] C. Wang, Y. Yang, and P. Zhou, “Towards efficient scheduling of federated mobile devices under computational and statistical heterogeneity,” *IEEE Transactions on Parallel and Distributed Systems*, vol. 32, no. 2, pp. 394–410, 2021.
- [99] J. Colin, T. Fel, R. Cadène, and T. Serre, “What I Cannot Predict, I Do Not Understand: A Human-Centered Evaluation Framework for Explainability Methods.” *Advances in Neural Information Processing Systems*, vol. 35, pp. 2832–2845, 2022.
- [100] A. Jung and P. Nardelli, “An information-theoretic approach to personalized explainable machine learning,” *IEEE Sig. Proc. Lett.*, vol. 27, pp. 825–829, 2020.

- [101] L. Zhang, G. Karakasidis, A. Odnoblyudova, L. Dogruel, Y. Tian, and A. Jung, “Explainable empirical risk minimization,” *Neural Computing and Applications*, vol. 36, no. 8, pp. 3983–3996, 2024. [Online]. Available: <https://doi.org/10.1007/s00521-023-09269-3>
- [102] M. J. Sheller, B. Edwards, G. A. Reina, J. Martin, S. Pati, A. Kotrotsou, M. Milchenko, W. Xu, D. Marcus, R. R. Colen, and S. Bakas, “Federated learning in medicine: facilitating multi-institutional collaborations without sharing patient data,” *Scientific Reports*, vol. 10, no. 1, p. 12598, 2020. [Online]. Available: <https://doi.org/10.1038/s41598-020-69250-1>
- [103] P. Amin, N. R. Anikireddypally, S. Khurana, S. Vadakkemadathil, and W. Wu, “Personalized health monitoring using predictive analytics,” in *2019 IEEE Fifth International Conference on Big Data Computing Service and Applications (BigDataService)*, 2019, pp. 271–278.
- [104] R. B. Ash, *Probability and Measure Theory*, 2nd ed. New York: Academic Press, 2000.
- [105] P. R. Halmos, *Measure Theory*. New York: Springer, 1974.
- [106] P. Billingsley, *Probability and Measure*, 3rd ed. New York: Wiley, 1995.
- [107] C. Dwork and A. Roth, “The algorithmic foundations of differential privacy,” *Foundations and Trends® in Theoretical Computer Science*, vol. 9, no. 3–4, pp. 211–407, 2014. [Online]. Available: <http://dx.doi.org/10.1561/04000000042>

- [108] S. Asodeh, J. Liao, F. P. Calmon, O. Kosut, and L. Sankar, “A Better Bound Gives a Hundred Rounds: Enhanced Privacy Guarantees via f -Divergences,” *arXiv e-prints*, p. arXiv:2001.05990, Jan. 2020.
- [109] I. Mironov, “Rényi differential privacy,” in *2017 IEEE 30th Computer Security Foundations Symposium (CSF)*, 2017, pp. 263–275.
- [110] S. M. Kay, *Fundamentals of statistical signal processing. Vol. 2., Detection theory*, ser. Prentice-Hall signal processing series. Upper Saddle River, NJ: Prentice-Hall PTR, 1998.
- [111] P. Kairouz, S. Oh, and P. Viswanath, “The composition theorem for differential privacy,” in *Proceedings of the 32nd International Conference on Machine Learning*, ser. Proceedings of Machine Learning Research, F. Bach and D. Blei, Eds., vol. 37. Lille, France: PMLR, 07–09 Jul 2015, pp. 1376–1385. [Online]. Available: <https://proceedings.mlr.press/v37/kairouz15.html>
- [112] Q. Geng and P. Viswanath, “The optimal noise-adding mechanism in differential privacy,” *IEEE Transactions on Information Theory*, vol. 62, no. 2, pp. 925–951, 2016.
- [113] H. Shu and H. Zhu, “Sensitivity analysis of deep neural networks,” in *Proceedings of the Thirty-Third AAAI Conference on Artificial Intelligence*, ser. AAAI’19/IAAI’19/EAAI’19. AAAI Press, 2019. [Online]. Available: <https://doi.org/10.1609/aaai.v33i01.33014943>
- [114] R. Busa-Fekete, A. Munoz-Medina, U. Syed, and S. Vassilvitskii, “Label differential privacy and private training data release,” in

- Proceedings of the 40th International Conference on Machine Learning*, ser. Proceedings of Machine Learning Research, vol. 202. PMLR, 23–29 Jul 2023, pp. 3233–3251. [Online]. Available: <https://proceedings.mlr.press/v202/busa-fekete23a.html>
- [115] B. Balle, G. Barthe, and M. Gaboardi, “Privacy amplification by sub-sampling: tight analyses via couplings and divergences,” in *Proceedings of the 32nd International Conference on Neural Information Processing Systems*, ser. NIPS’18. Red Hook, NY, USA: Curran Associates Inc., 2018, pp. 6280–6290.
- [116] P. Cuff and L. Yu, “Differential privacy as a mutual information constraint,” in *Proceedings of the 2016 ACM SIGSAC Conference on Computer and Communications Security*, ser. CCS ’16. New York, NY, USA: Association for Computing Machinery, 2016, pp. 43–54. [Online]. Available: <https://doi.org/10.1145/2976749.2978308>
- [117] A. Makhdoumi, S. Salamatian, N. Fawaz, and M. Médard, “From the information bottleneck to the privacy funnel,” in *2014 IEEE Information Theory Workshop (ITW 2014)*, 2014, pp. 501–505.
- [118] A. Turner, D. Tsipras, and A. Madry, “Clean-label backdoor attacks,” 2019. [Online]. Available: <https://openreview.net/forum?id=HJg6e2CCK7>
- [119] R. Gray, *Probability, Random Processes, and Ergodic Properties*, 2nd ed. New York: Springer, 2009.

- [120] T. M. Cover and J. A. Thomas, *Elements of Information Theory*, 2nd ed. New Jersey: Wiley, 2006.
- [121] X. Liu, H. Li, G. Xu, Z. Chen, X. Huang, and R. Lu, “Privacy-enhanced federated learning against poisoning adversaries,” *IEEE Transactions on Information Forensics and Security*, vol. 16, pp. 4574–4588, 2021.
- [122] J. Zhang, B. Chen, X. Cheng, H. T. T. Binh, and S. Yu, “PoisonGAN: Generative poisoning attacks against federated learning in edge computing systems,” *IEEE Internet of Things Journal*, vol. 8, no. 5, pp. 3310–3322, 2021.
- [123] P. Halmos, *Naive set theory*. Springer-Verlag, 1974.
- [124] T. Hastie, R. Tibshirani, and J. Friedman, *The Elements of Statistical Learning*, ser. Springer Series in Statistics. New York, NY, USA: Springer, 2001.
- [125] N. Cesa-Bianchi and G. Lugosi, *Prediction, Learning, and Games*. New York, NY, USA: Cambridge University Press, 2006.
- [126] E. Hazan, *Introduction to Online Convex Optimization*. Now Publishers Inc., 2016.
- [127] J. Chen, L. Song, M. Wainwright, and M. Jordan, “Learning to explain: An information-theoretic perspective on model interpretation,” in *Proc. 35th Int. Conf. on Mach. Learning*, Stockholm, Sweden, 2018.
- [128] D. Gujarati and D. Porter, *Basic Econometrics*. McGraw Hill, 2009.

- [129] Y. Dodge, *The Oxford Dictionary of Statistical Terms*. Oxford University Press, 2003.
- [130] B. Everitt, *Cambridge Dictionary of Statistics*. Cambridge University Press, 2002.
- [131] B. Schölkopf and A. Smola, *Learning with Kernels: Support Vector Machines, Regularization, Optimization, and Beyond*. Cambridge, MA, USA: MIT Press, Dec. 2002.
- [132] M. Ribeiro, S. Singh, and C. Guestrin, ““Why should i trust you?”: Explaining the predictions of any classifier,” in *Proc. 22nd ACM SIGKDD*, Aug. 2016, pp. 1135–1144.
- [133] A. Papoulis and S. U. Pillai, *Probability, Random Variables, and Stochastic Processes*, 4th ed. New York: Mc-Graw Hill, 2002.
- [134] C. Lampert, “Kernel methods in computer vision,” *Foundations and Trends in Computer Graphics and Vision*, 2009.
- [135] U. von Luxburg, “A tutorial on spectral clustering,” *Statistics and Computing*, vol. 17, no. 4, pp. 395–416, Dec. 2007.
- [136] A. Y. Ng, M. I. Jordan, and Y. Weiss, “On spectral clustering: Analysis and an algorithm,” in *Adv. Neur. Inf. Proc. Syst.*, 2001.
- [137] R. Caruana, “Multitask learning,” *Machine Learning*, vol. 28, no. 1, pp. 41–75, 1997. [Online]. Available: <https://doi.org/10.1023/A:1007379606734>

- [138] A. Jung, G. Hannak, and N. Görtz, “Graphical LASSO Based Model Selection for Time Series,” *IEEE Sig. Proc. Letters*, vol. 22, no. 10, Oct. 2015.
- [139] A. Jung, “Learning the conditional independence structure of stationary time series: A multitask learning approach,” *IEEE Trans. Signal Processing*, vol. 63, no. 21, Nov. 2015.
- [140] C. Rudin, “Stop explaining black box machine learning models for high-stakes decisions and use interpretable models instead,” *Nature Machine Intelligence*, vol. 1, no. 5, pp. 206–215, 2019.
- [141] S. Shalev-Shwartz and S. Ben-David, *Understanding Machine Learning – from Theory to Algorithms*. Cambridge University Press, 2014.
- [142] A. Lapidoth, *A Foundation in Digital Communication*. New York: Cambridge University Press, 2009.
- [143] S. Bubeck, “Convex optimization. algorithms and complexity.” in *Foundations and Trends in Machine Learning*. Now Publishers, 2015, vol. 8.